Intro to Visualizations

Daniel Anderson

Week 2, Class 1

Agenda

- Quick note on projects and `here::here()`

Discuss different visualizations

- Visualizing distributions
  - histograms
  - density plots
  - Empirical cumulative density plots
  - QQ plots
- Visualizing amounts
  - bar plots
  - dot plots
  - heatmaps

Learning Objectives

- Understand various ways the same underlying data can be displayed
- Think through pros/cons of each
- Understand the basic structure of the code to produce the various plots

What type of data do you have?

We'll focus primarily on standard continuous/categorical data

What is your purpose?

Exploratory? Communication?

One continuous variable

Histogram

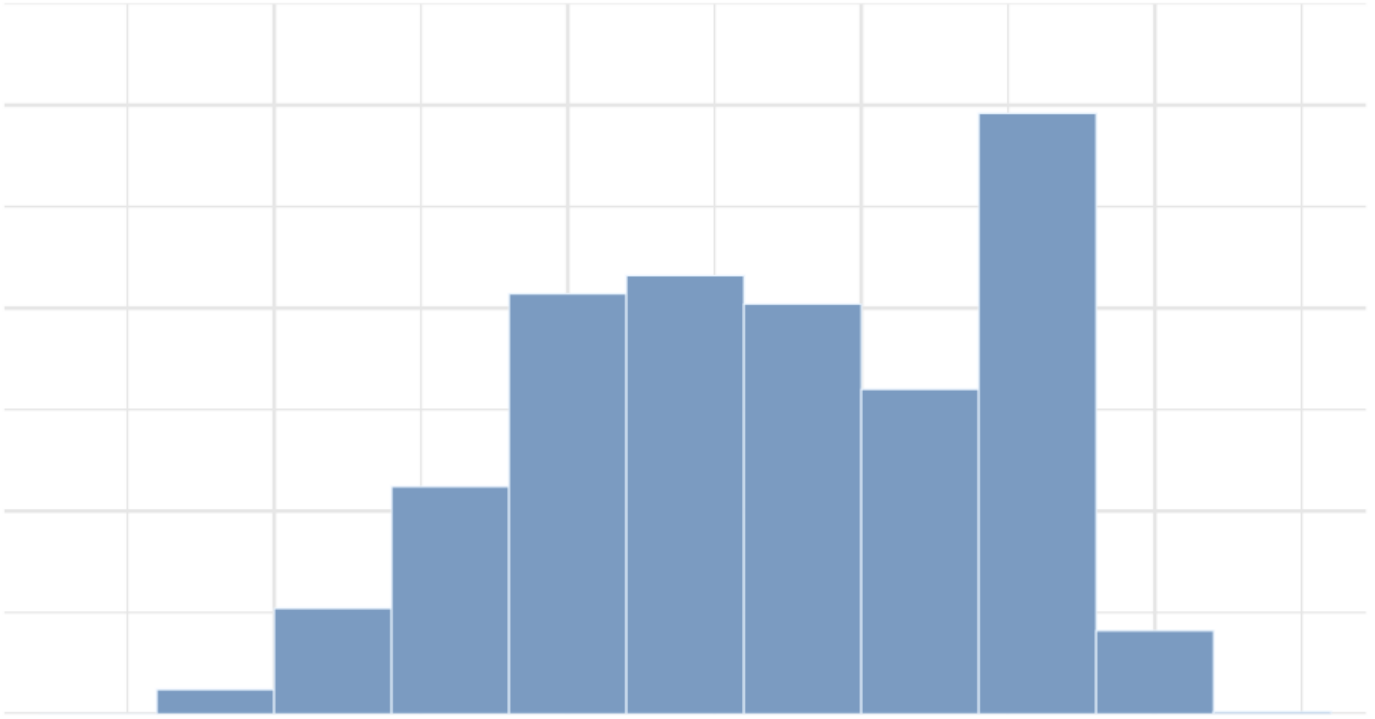

Density plot

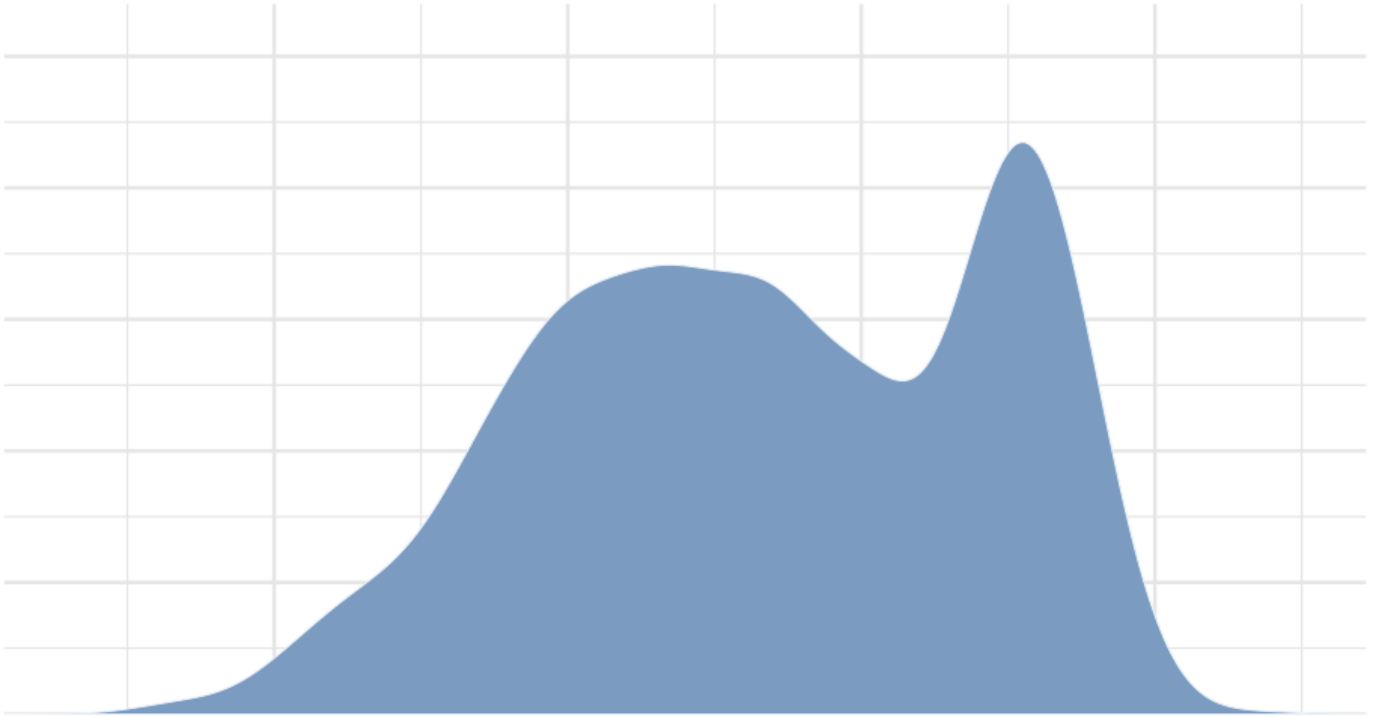

(Empirical) Cumulative Density

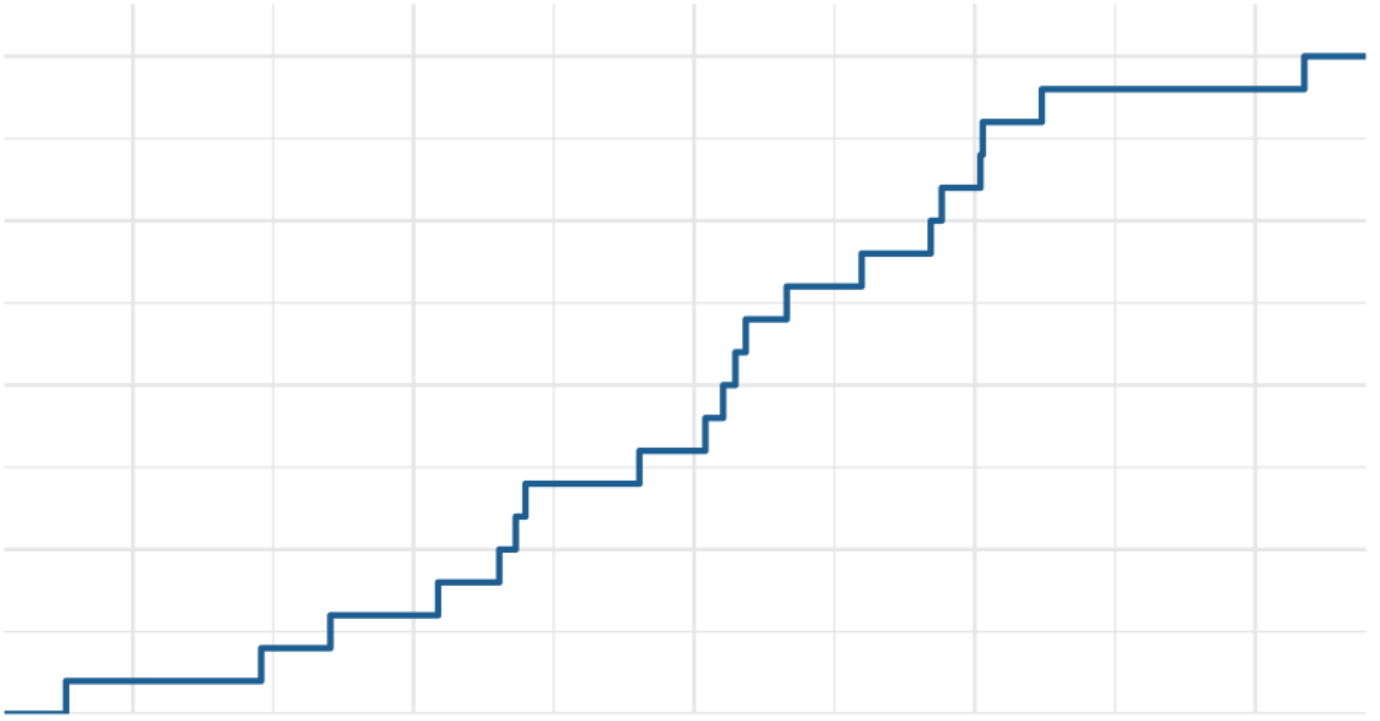

QQ Plot

Compare to theoretical quantiles (for normality)

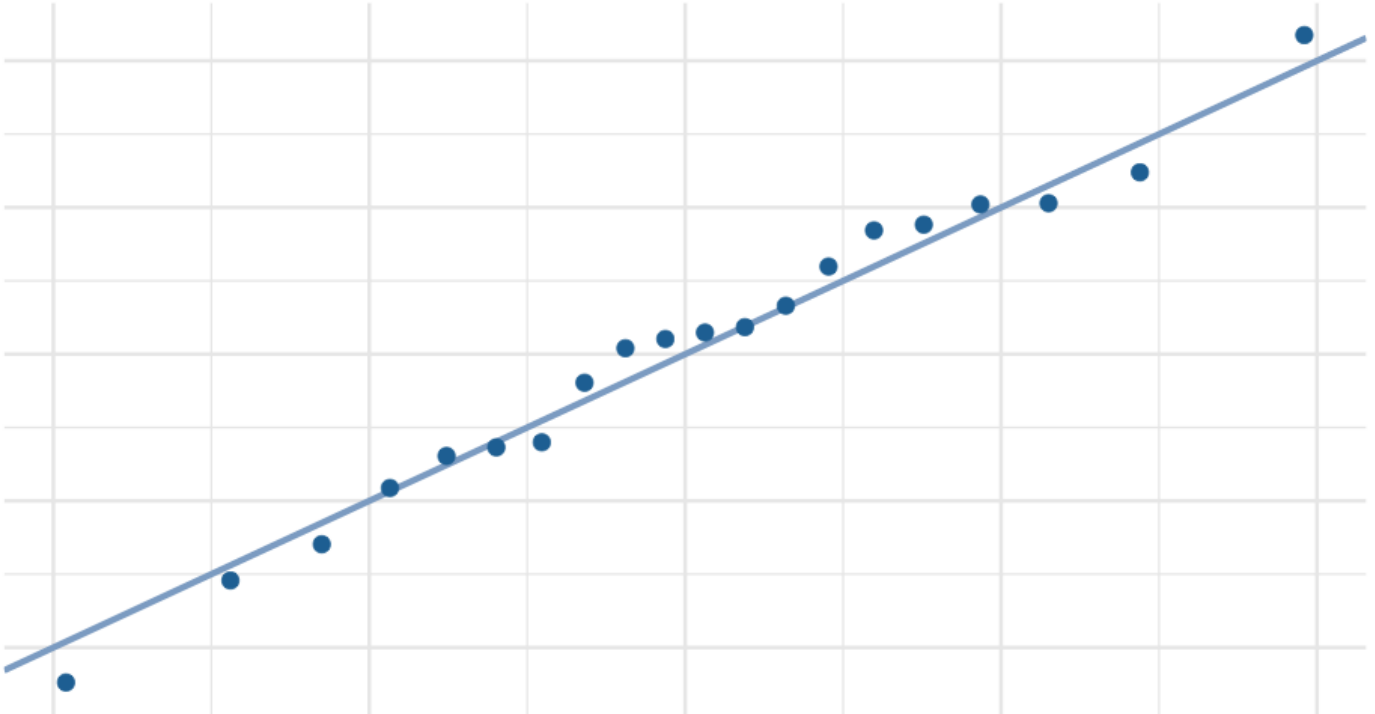

Empirical examples

I'll move fast, but if you want to (try to) follow along, or recreate anything here later, first run

```
remotes::install_github("clauswilke/dviz.supp")
```

Titanic data

```
head(titanic)
```

```
##                                                name class   age    sex survi
## 1                    Allen, Miss Elisabeth Walton   1st 29.00 female      
## 2                     Allison, Miss Helen Loraine   1st  2.00 female      
## 3              Allison, Mr Hudson Joshua Creighton   1st 30.00   male      
## 4 Allison, Mrs Hudson JC (Bessie Waldo Daniels)   1st 25.00 female      
## 5                   Allison, Master Hudson Trevor   1st  0.92   male      
## 6                              Anderson, Mr Harry   1st 47.00   male
```

Basic histogram

```
ggplot(titanic, aes(x = age)) +
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

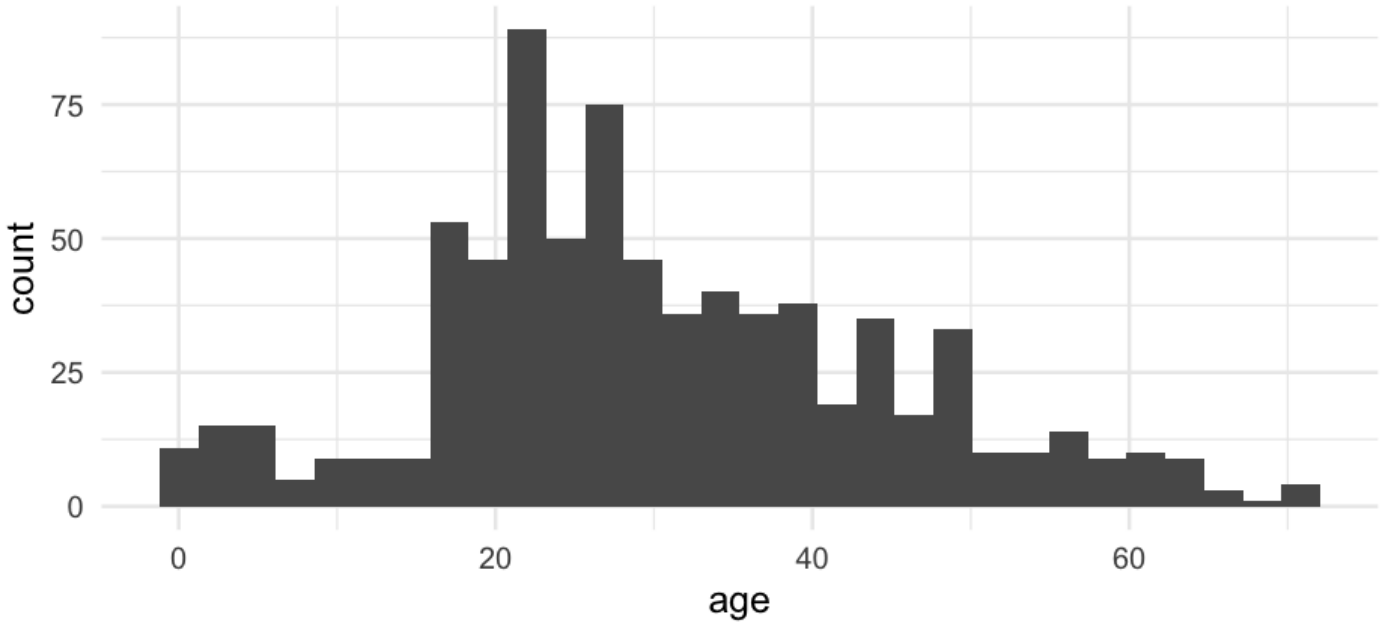

Make it a little prettier

```
ggplot(titanic, aes(x = age)) +
  geom_histogram(fill = "#56B4E9",
                 color = "white",
                 alpha = 0.9)
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

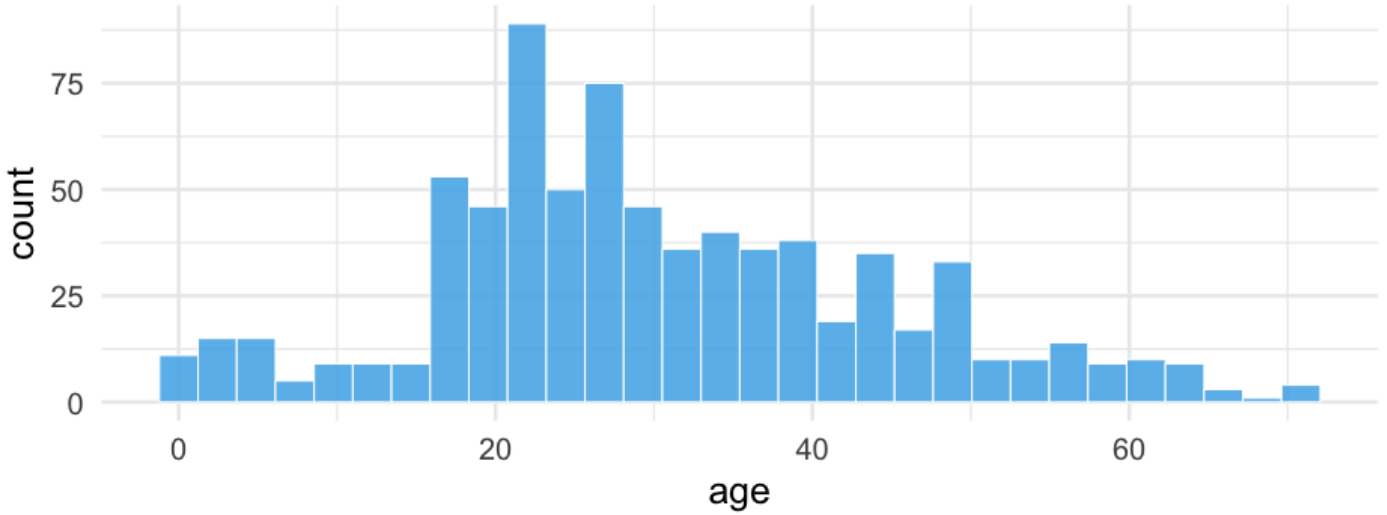

Change the number of bins

```
ggplot(titanic, aes(x = age)) +
  geom_histogram(fill = "#56B4E9",
                 color = "white",
                 alpha = 0.9,
                 bins = 50)
```

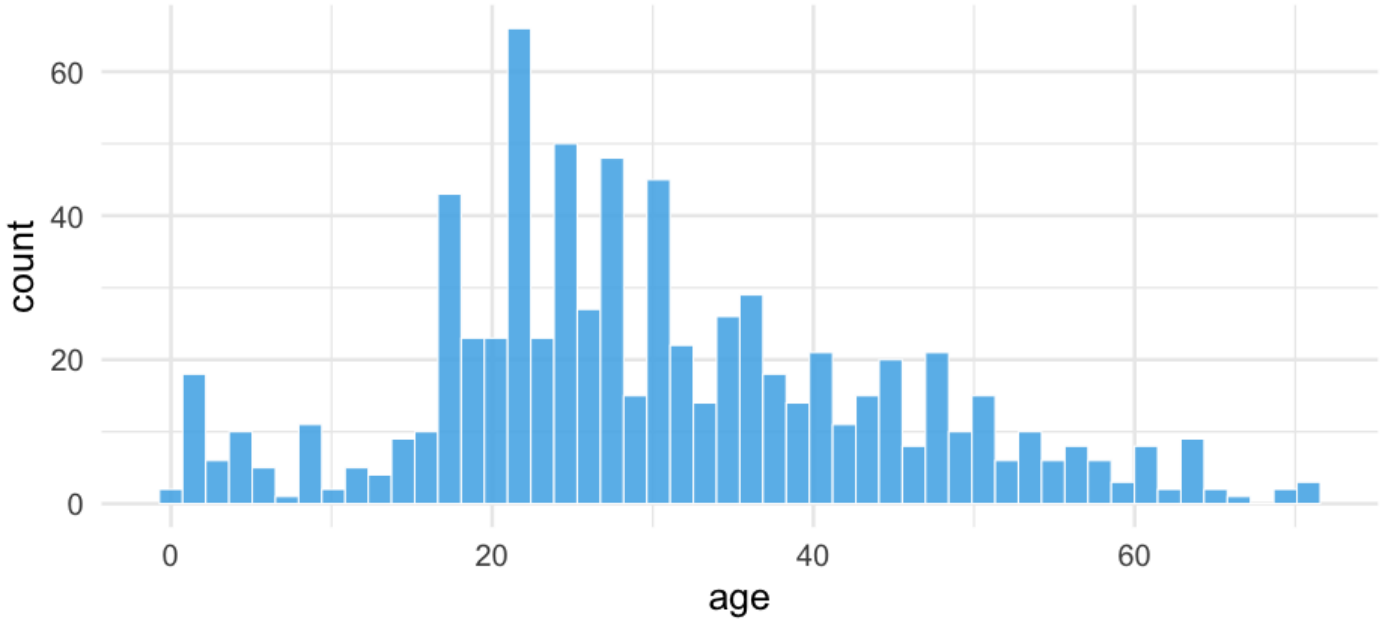

Vary the number of bins

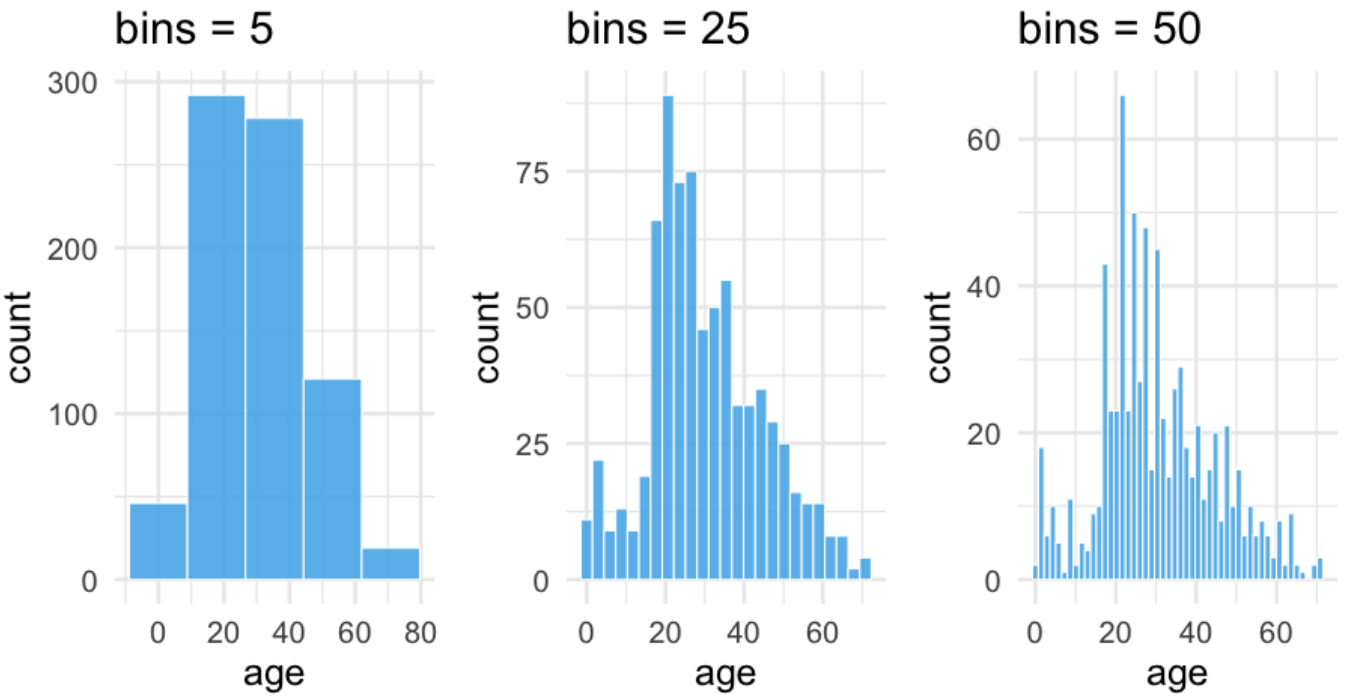

Denisty plot

ugly 😫

```
ggplot(titanic, aes(age)) +
  geom_density()
```

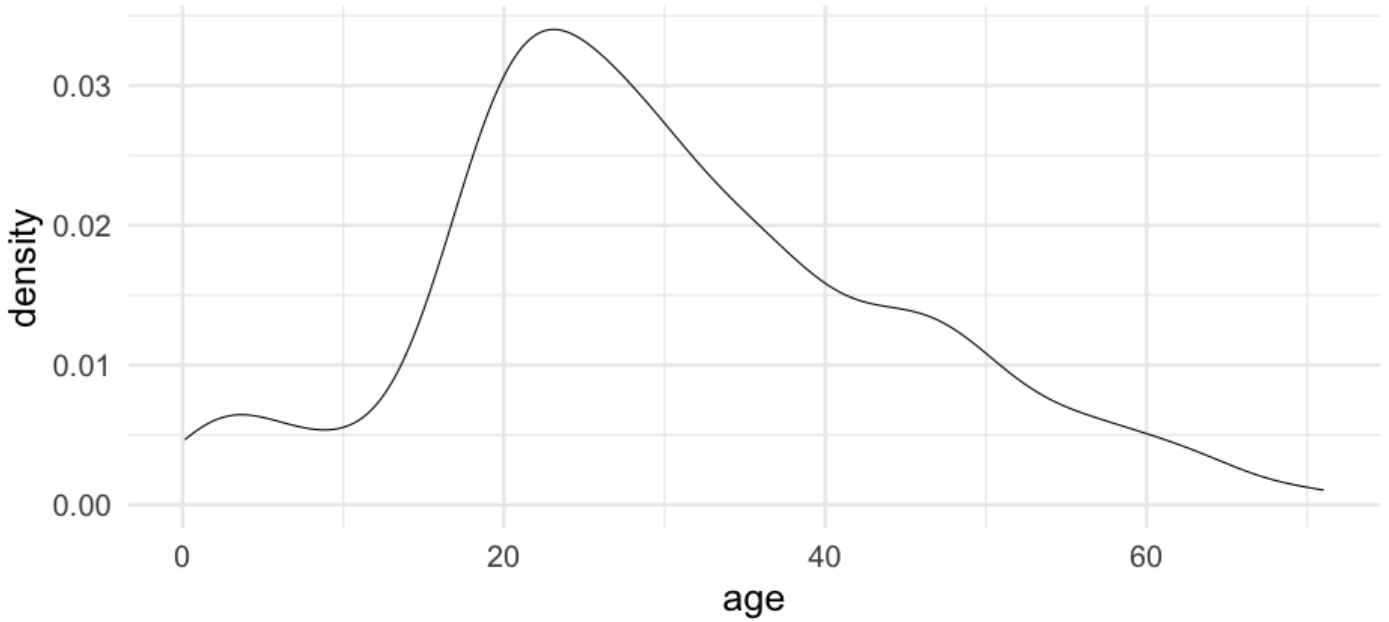

Denisty plot

Change the fill 😌

```
ggplot(titanic, aes(age)) +
  geom_density(fill = "#56B4E9")
```

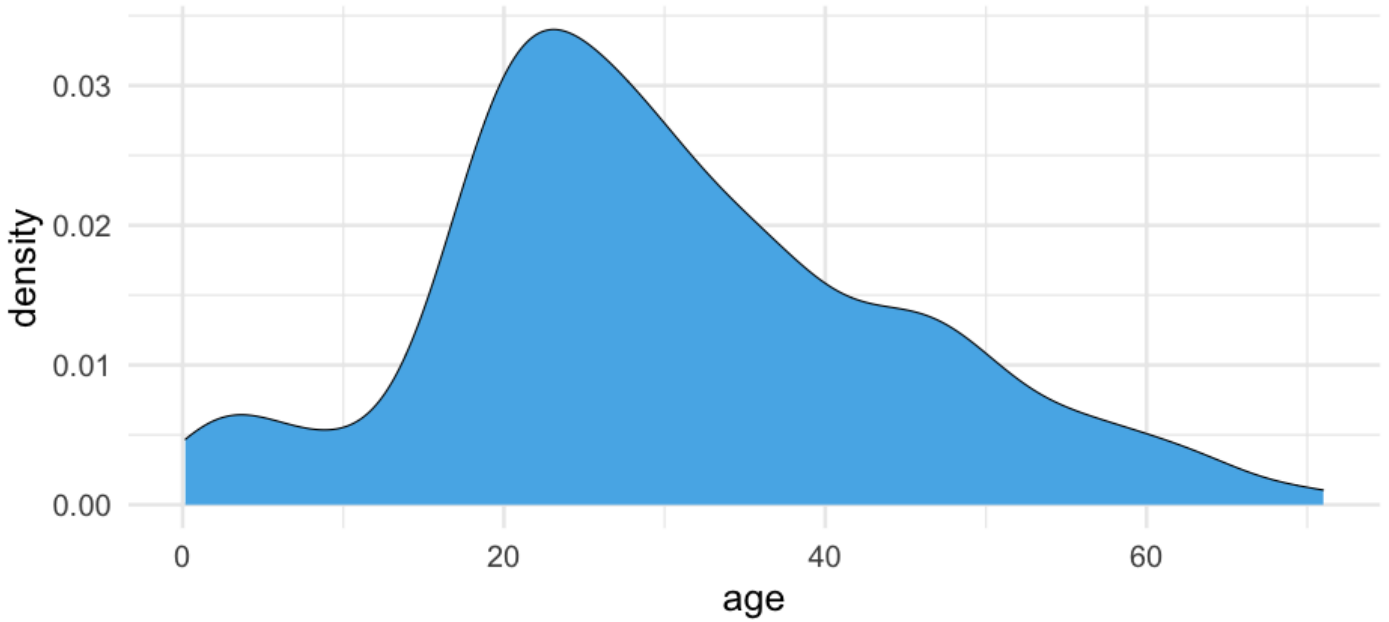

Density plot estimation

- Kernal density estimation
    - Different kernal shapes can be selected
    - Bandwidth matters most
    - Smaller bands = bend more to the data
- Approximation of the underlying continuous probability function
    - Integrates to 1.0 (y–axis is somewhat difficult to interpret)

Denisty plot

change the bandwidth

```
ggplot(titanic, aes(age)) +
  geom_density(fill = "#56B4E9",
               bw = 5)
```

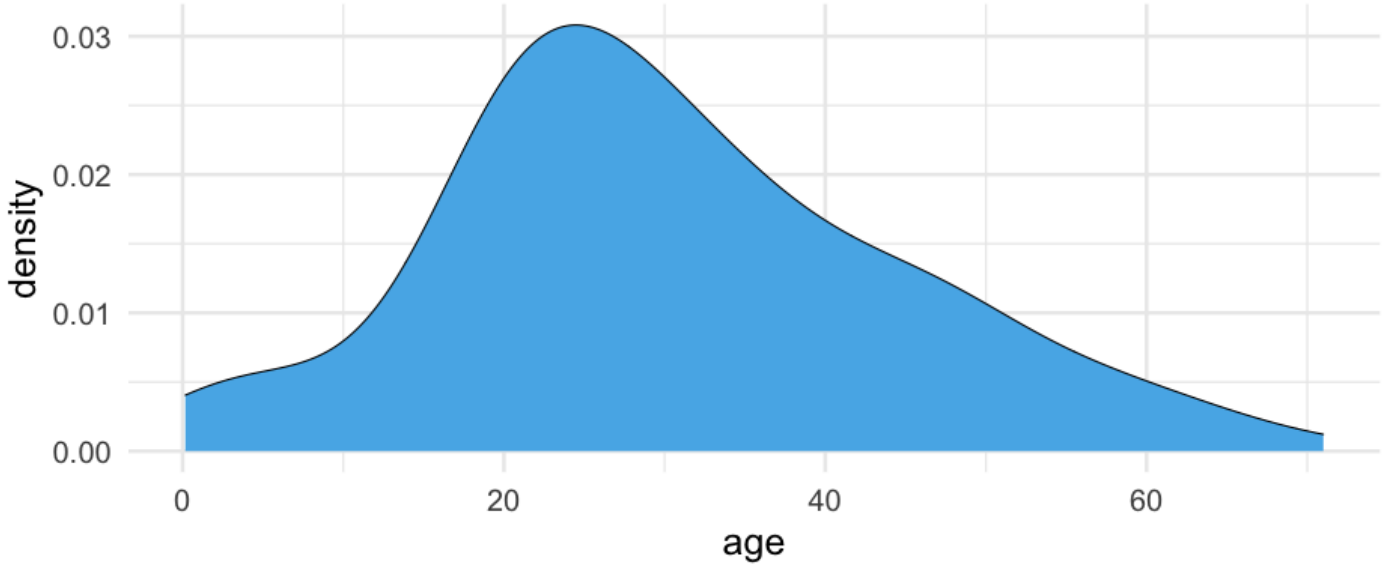

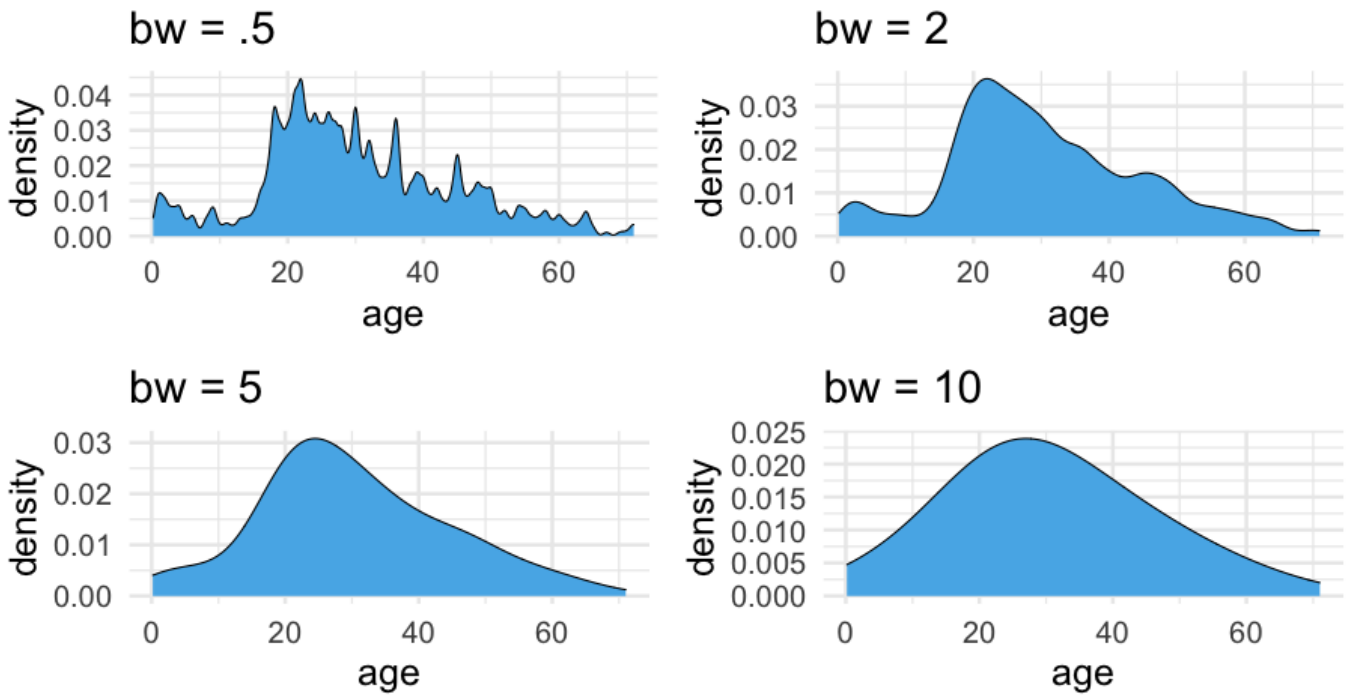
bw = .5
density
0.04
0.03
0.02
0.01
0.00
0
20
40
60
age
bw = 2
density
0.03
0.02
0.01
0.00
0
20
40
60
age
bw = 5
density
0.03
0.02
0.01
0.00
0
20
40
60
age
bw = 10
density
0.025
0.020
0.015
0.010
0.005
0.000
0
20
40
60
age

Quickly

How well does it approximate a normal distribution?

```
ggplot(titanic, aes(sample = age)) +
  stat_qq_line(color = "#56B4E9") +
  geom_qq(color = "gray40")
```

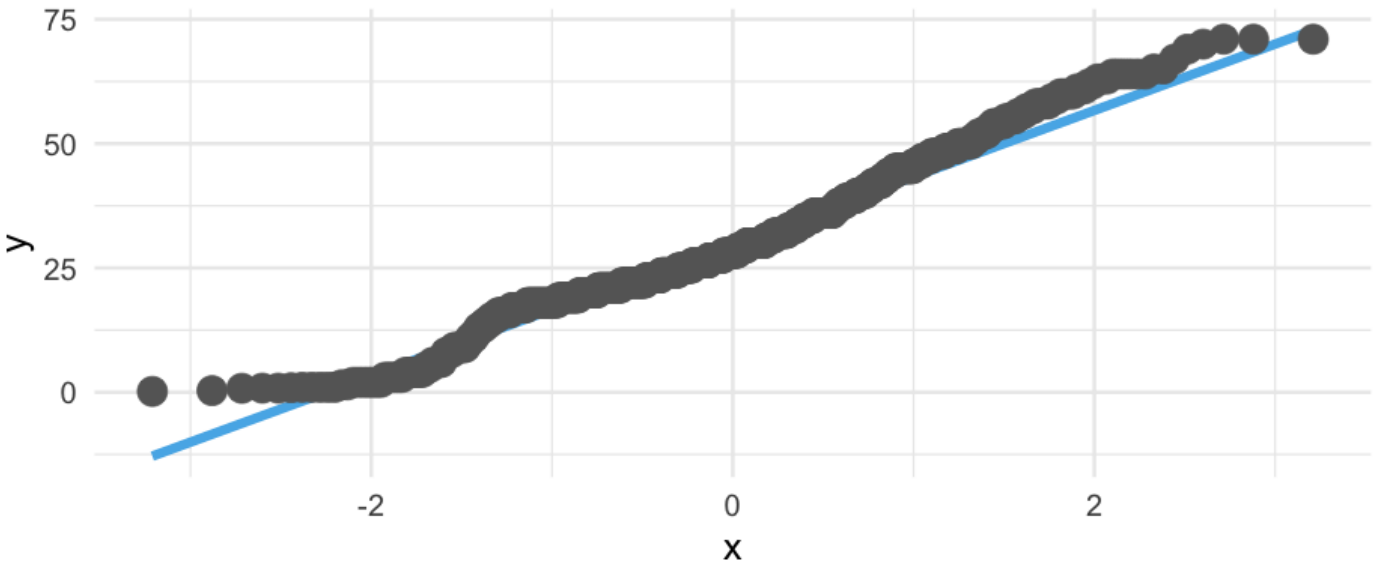

Grouped data

Distributions

How do we display more than one distribution at a time?

Boxplots

Violin plots

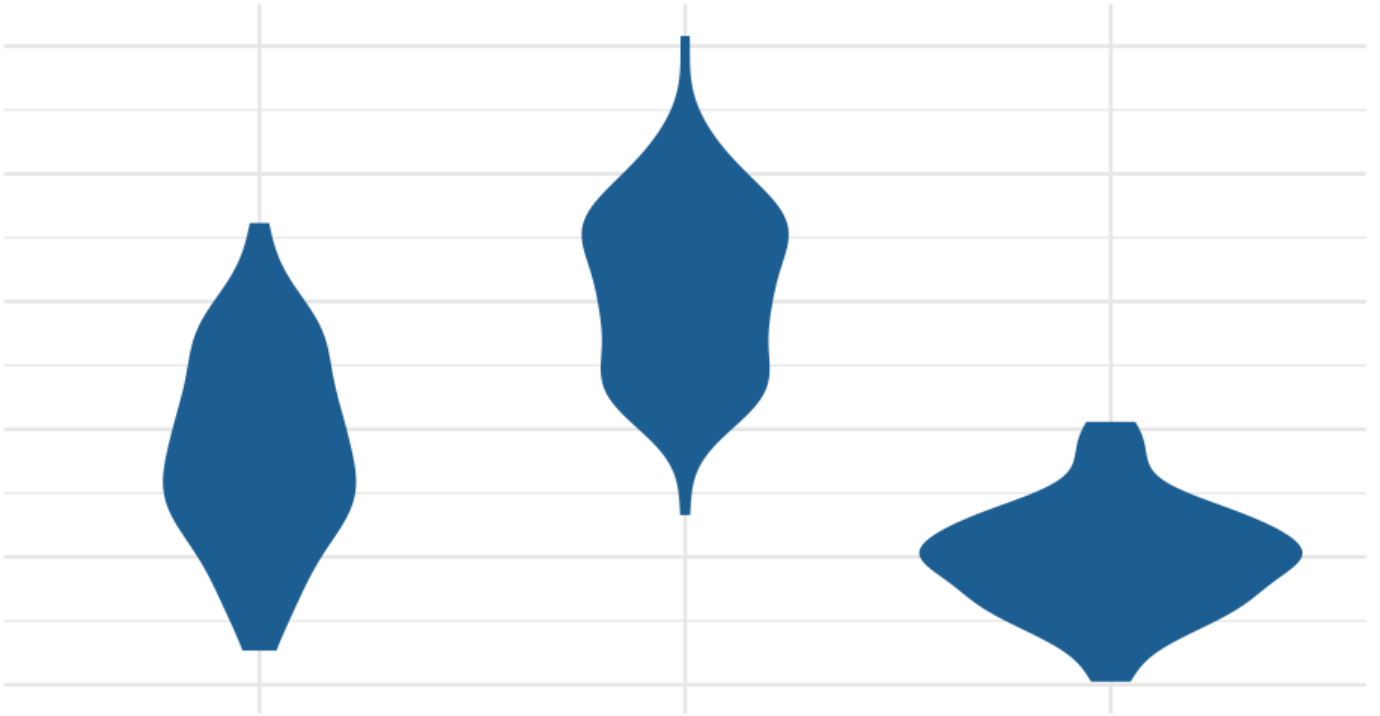

Jittered points

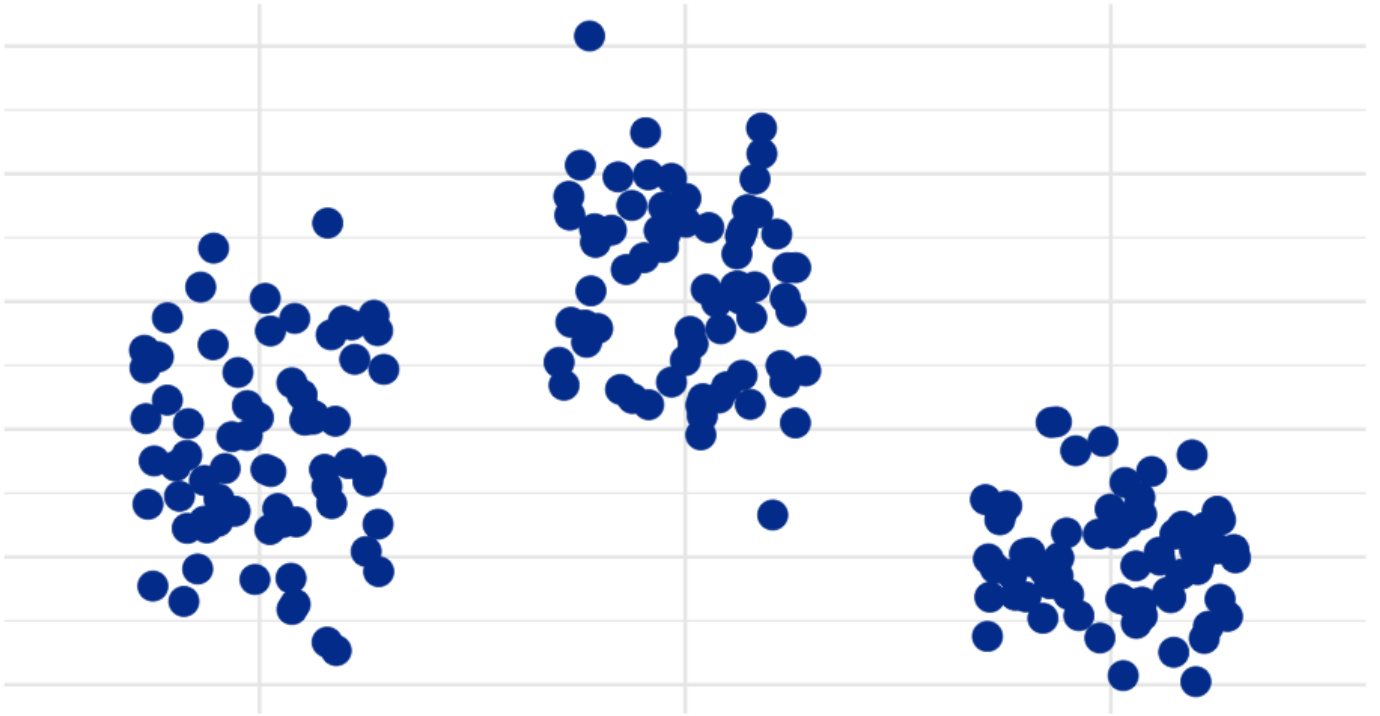

Sina plots

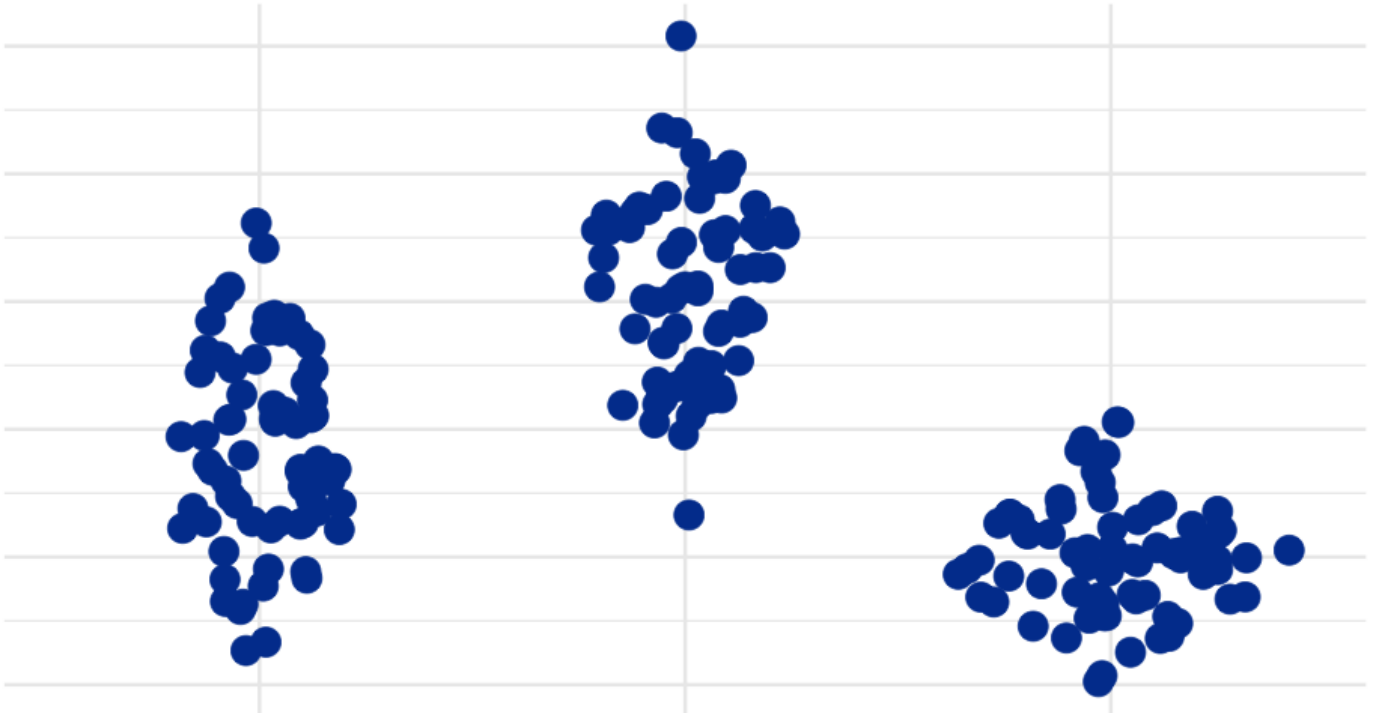

Stacked histograms

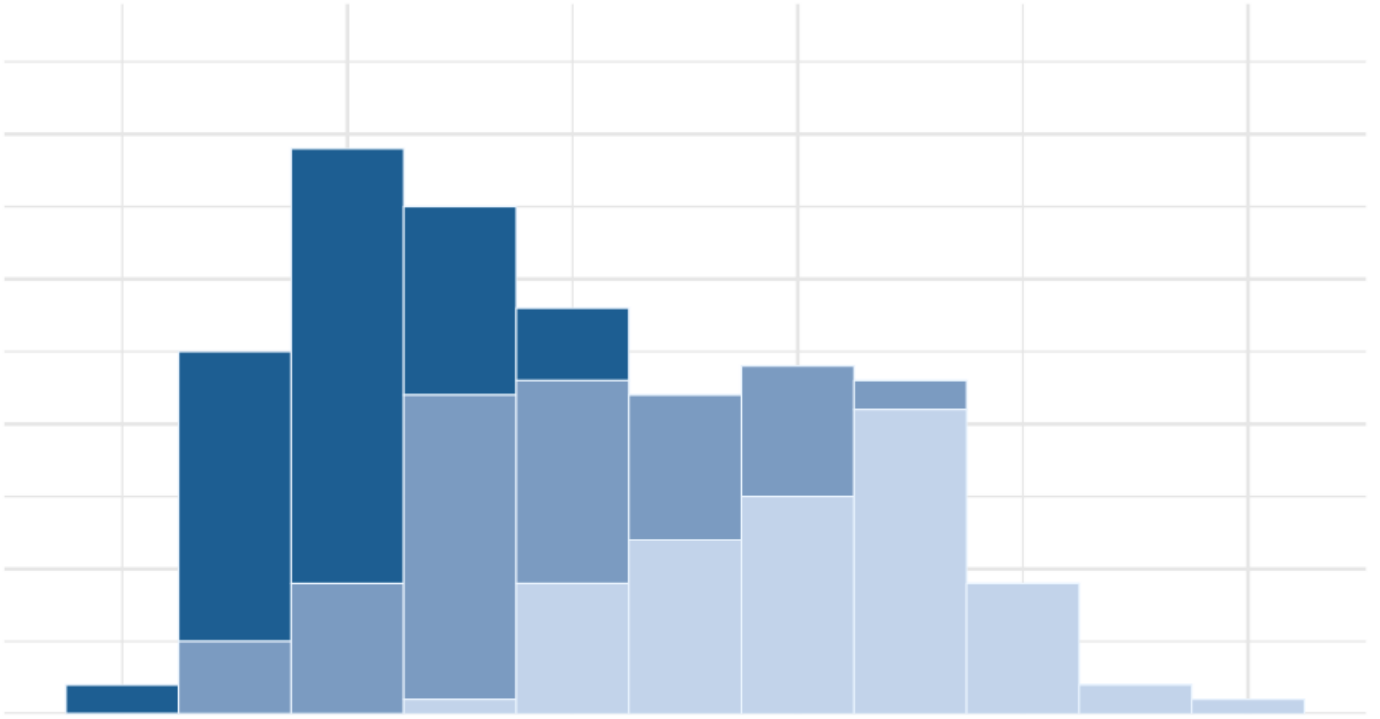

Overlapping densities

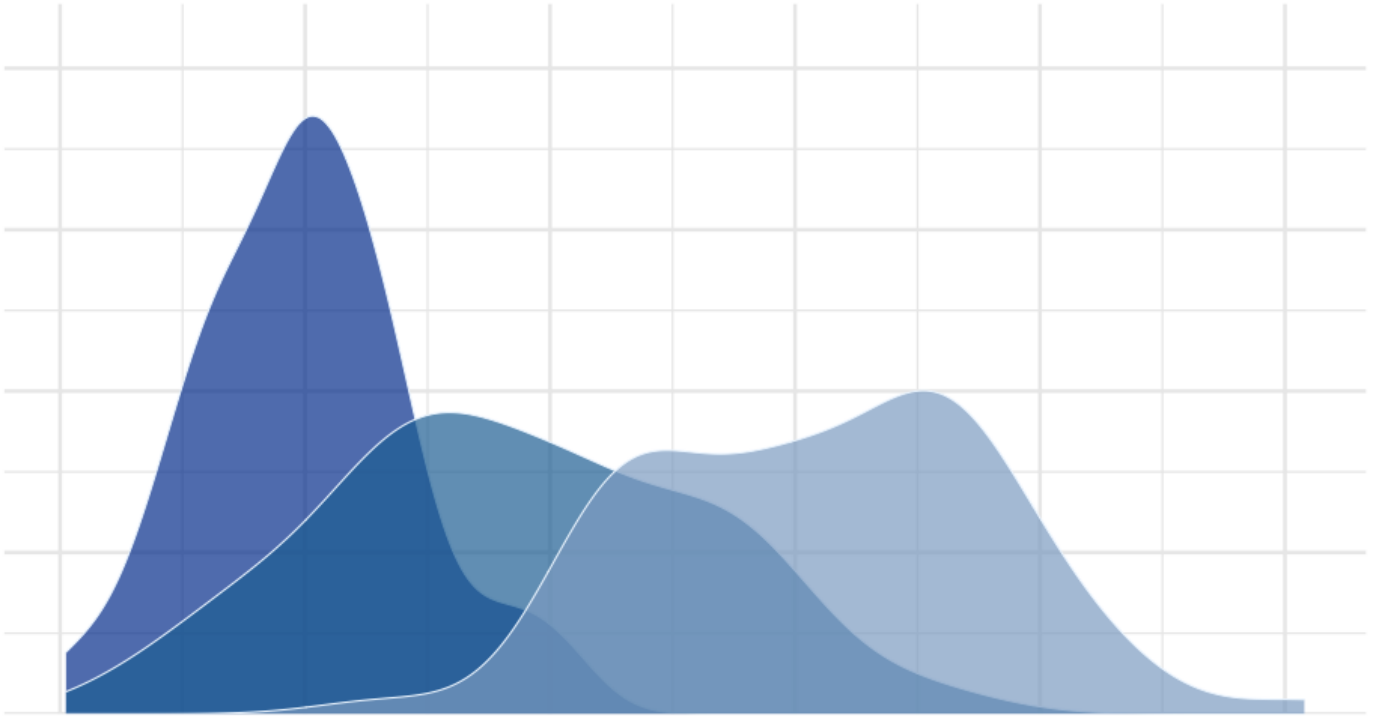

Ridgeline densities

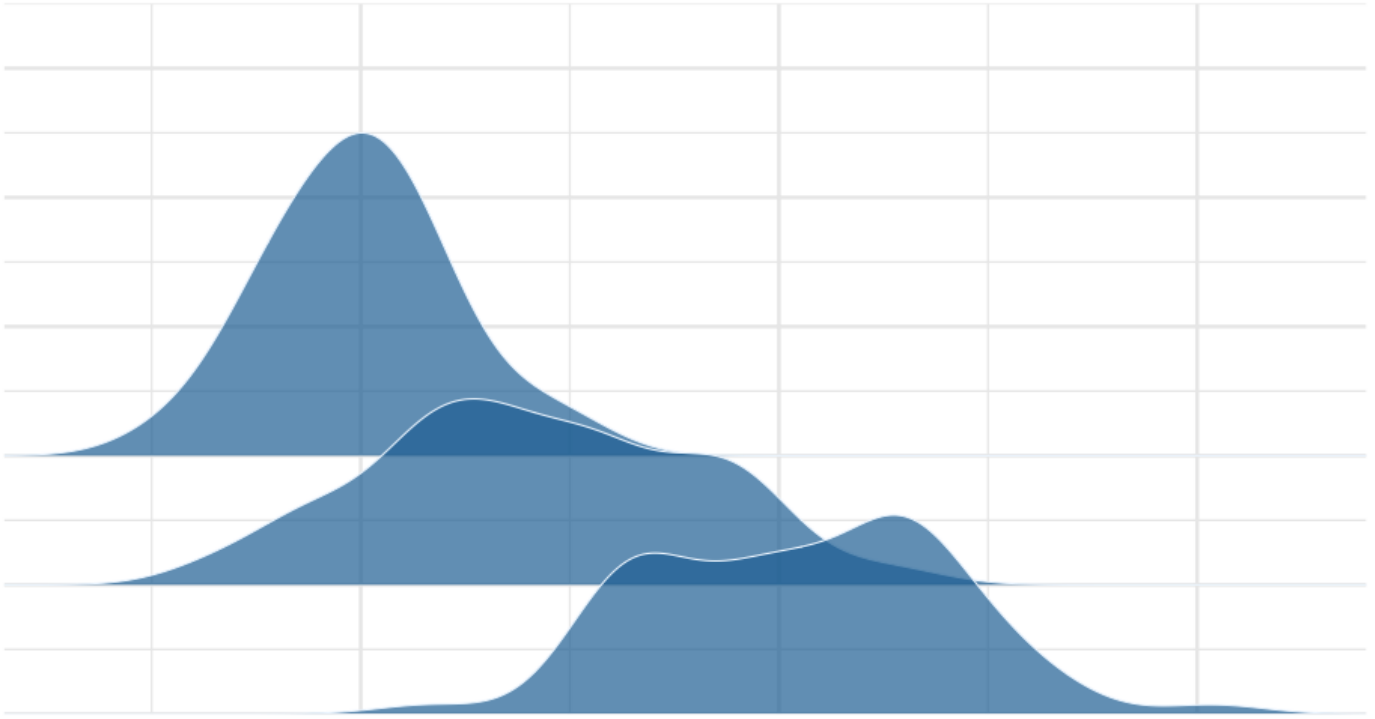

Quick empirical examples

Boxplots

```
ggplot(titanic, aes(sex, age)) +
  geom_boxplot(fill = "#A9E5C5")
```

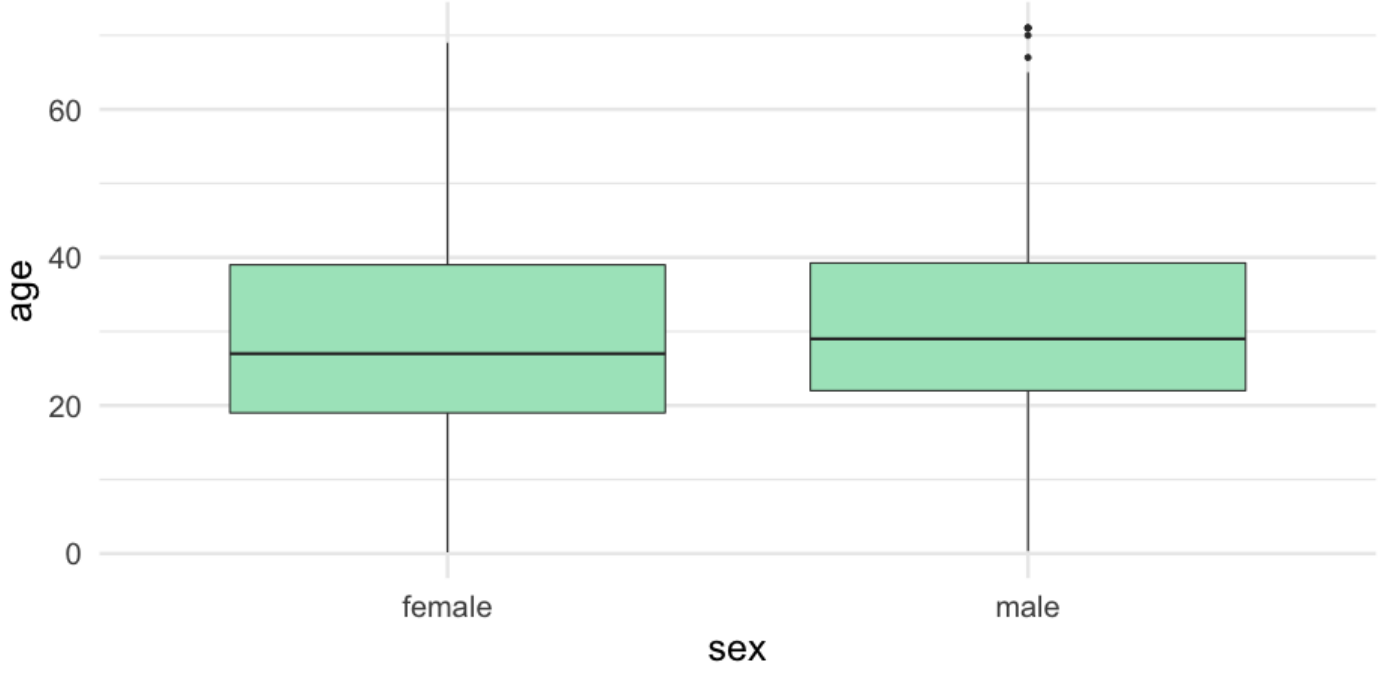

Violin plots

```
ggplot(titanic, aes(sex, age)) +
  geom_violin(fill = "#A9E5C5")
```

Jittered point plots

```
ggplot(titanic, aes(sex, age)) +
  geom_jitter(width = 0.3, height = 0)
```

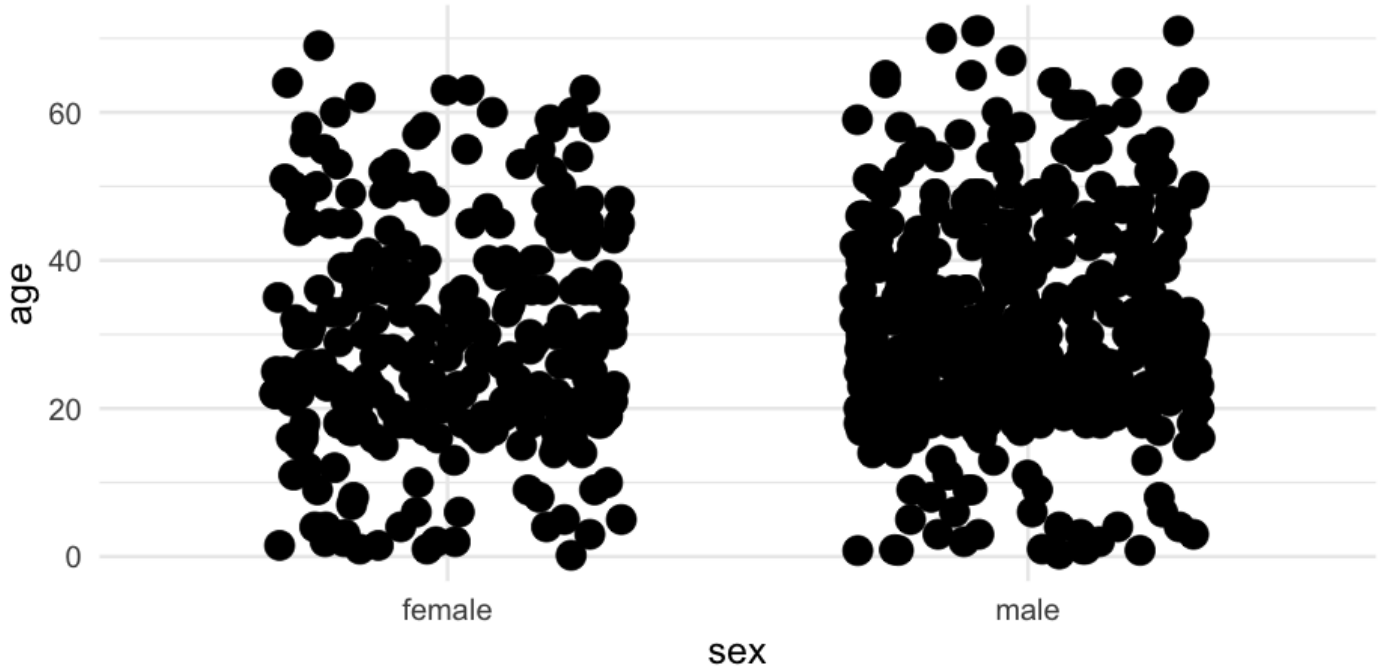

Sina plot

```
ggplot(titanic, aes(sex, age)) +
  ggforce::geom_sina()
```

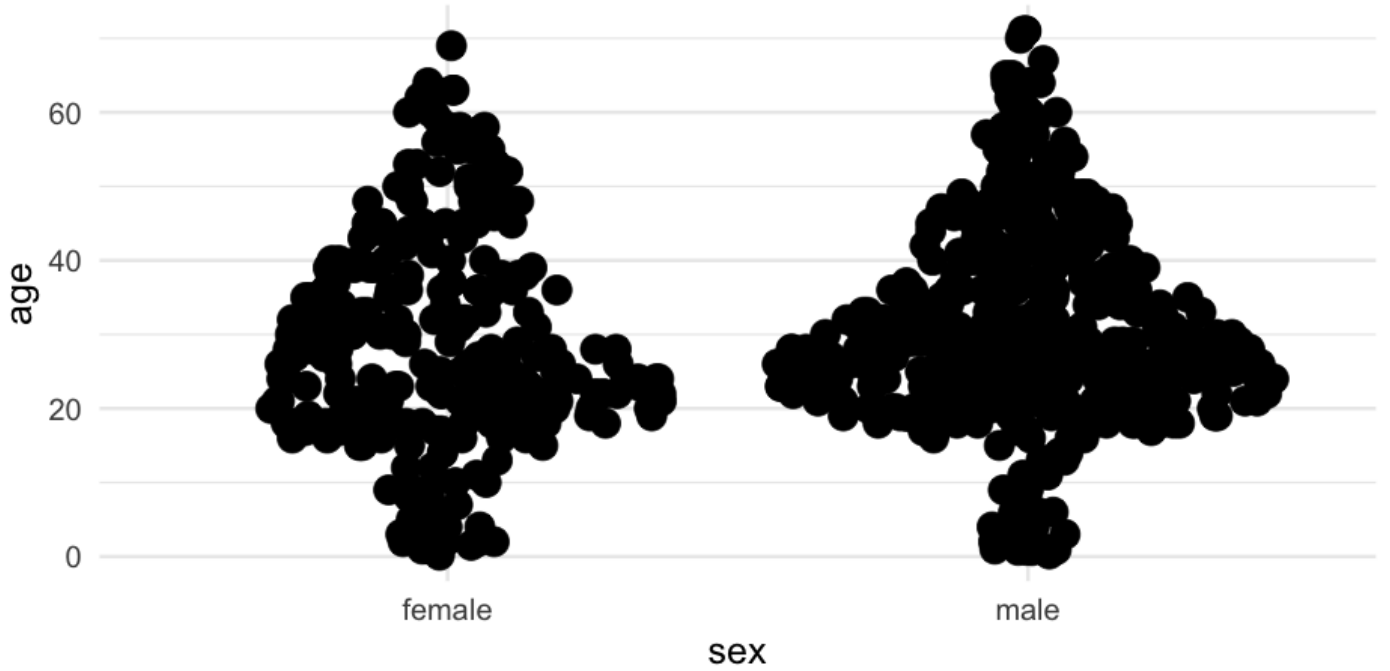

Stacked histogram

```
ggplot(titanic, aes(age)) +
  geom_histogram(aes(fill = sex))
```

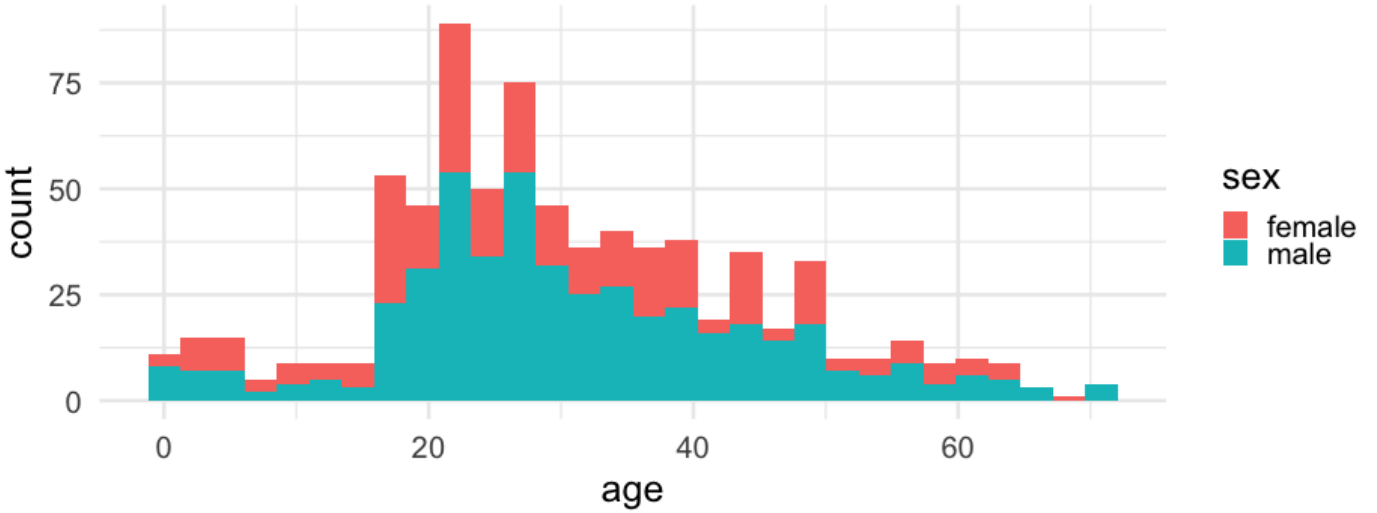

Dodged

```
ggplot(titanic, aes(age)) +
  geom_histogram(aes(fill = sex),
                 position = "dodge")
```

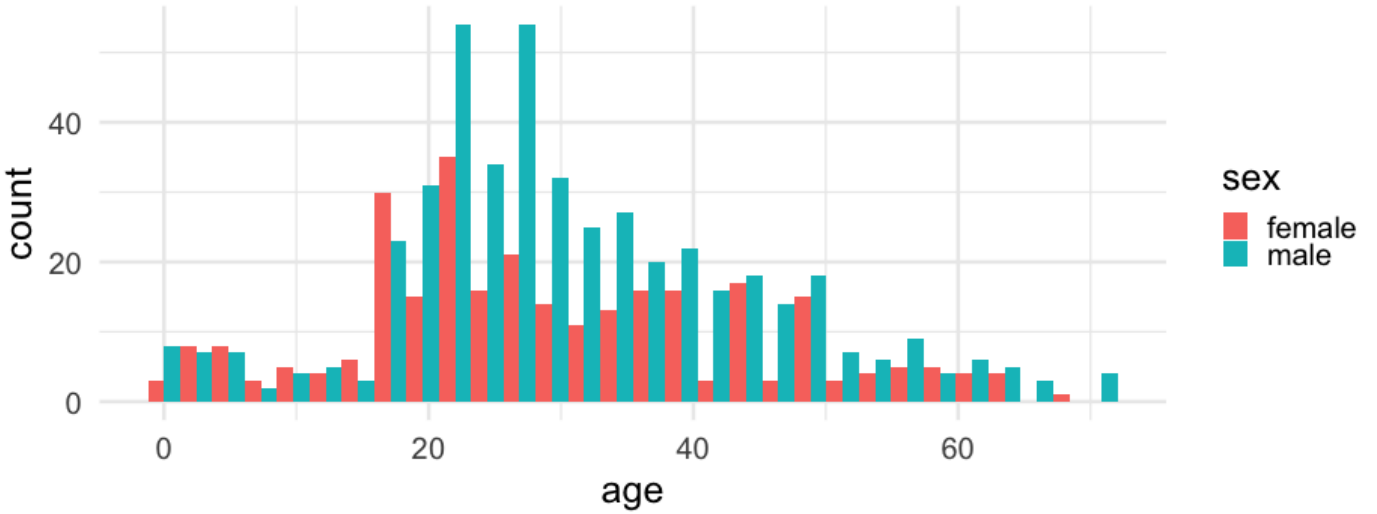

Note `position = "dodge"` does not go into `aes` (not accessing a variable in your dataset)

Better

```
ggplot(titanic, aes(age)) +
  geom_histogram(fill = "#A9E5C5",
                 color = "white",
                 alpha = 0.9,) +
  facet_wrap(~sex)
```

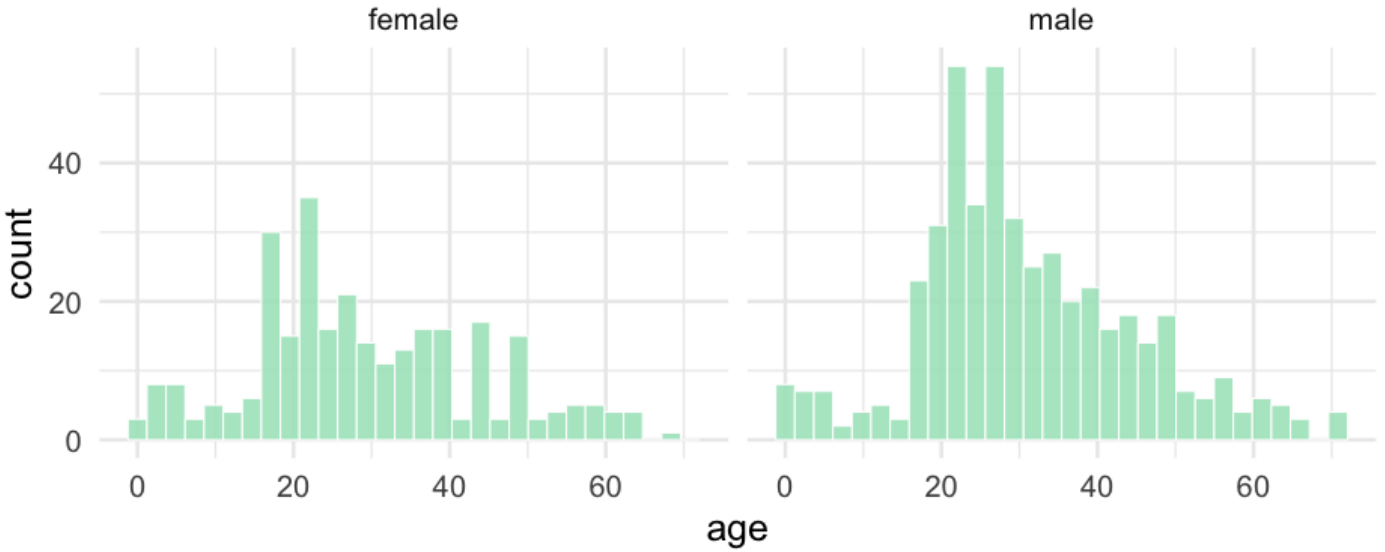

Overlapping densities

```
ggplot(titanic, aes(age)) +
  geom_density(aes(fill = sex),
               color = "white",
               alpha = 0.4)
```

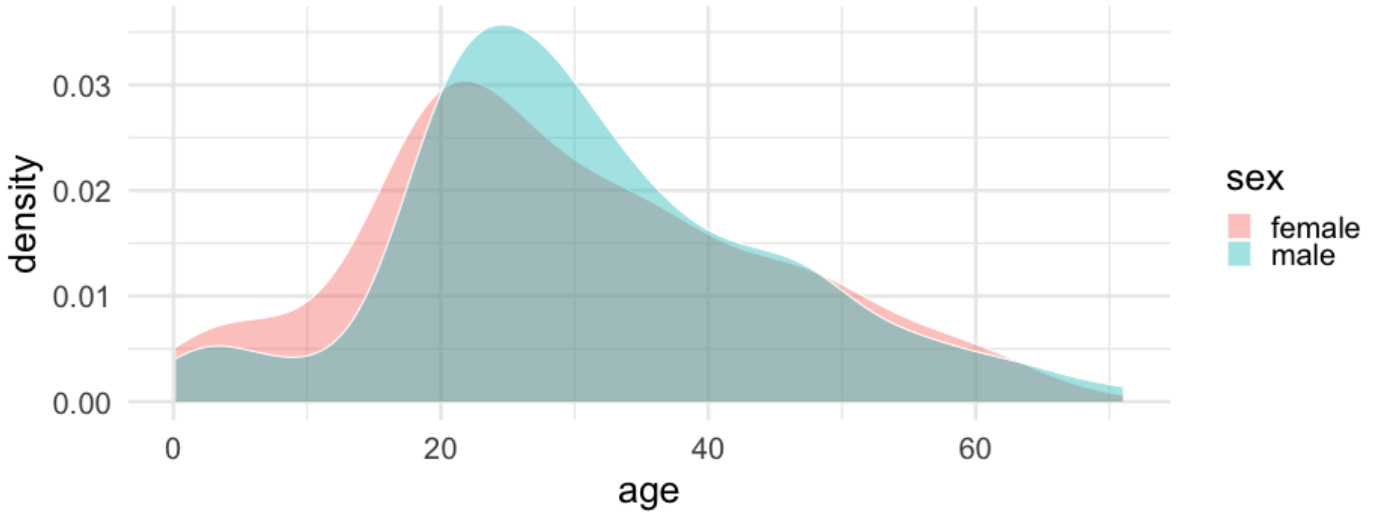

Note the default colors really don't work well in most of these

```
ggplot(titanic, aes(age)) +
  geom_density(aes(fill = sex),
               color = "white",
               alpha = 0.6) +
  scale_fill_manual(values = c("#009973", "#99ffe6"))
```

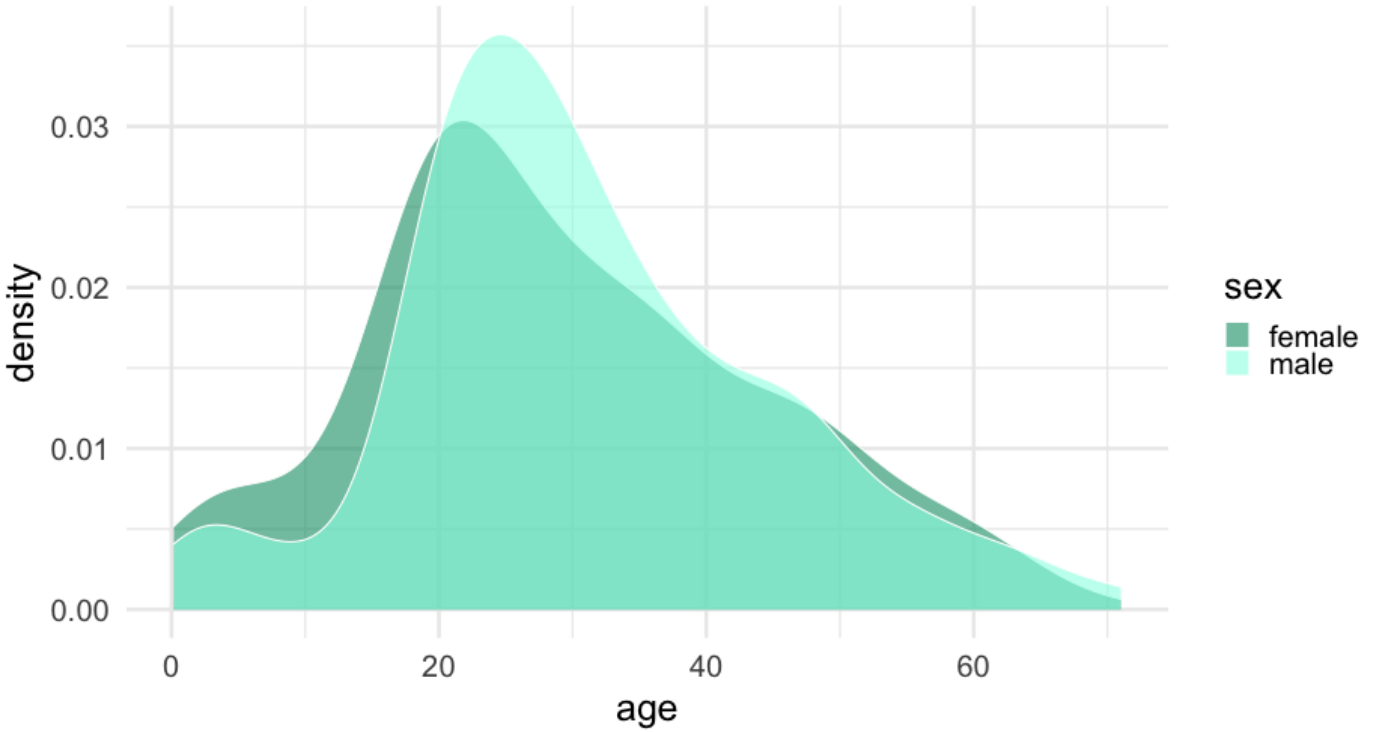

Ridgeline densities

```
ggplot(titanic, aes(age, sex)) +
  ggridges::geom_density_ridges(color = "white",
                                fill = "#A9E5C5")
```

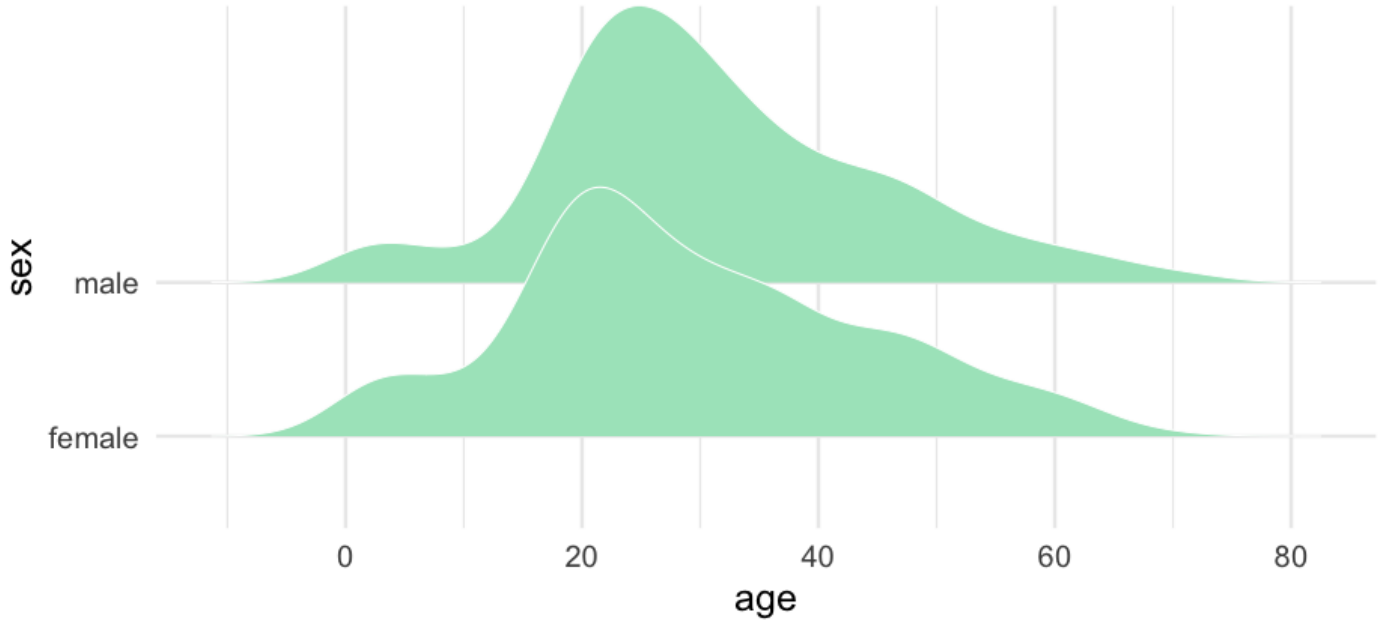

Visualizing amounts

Bar plots

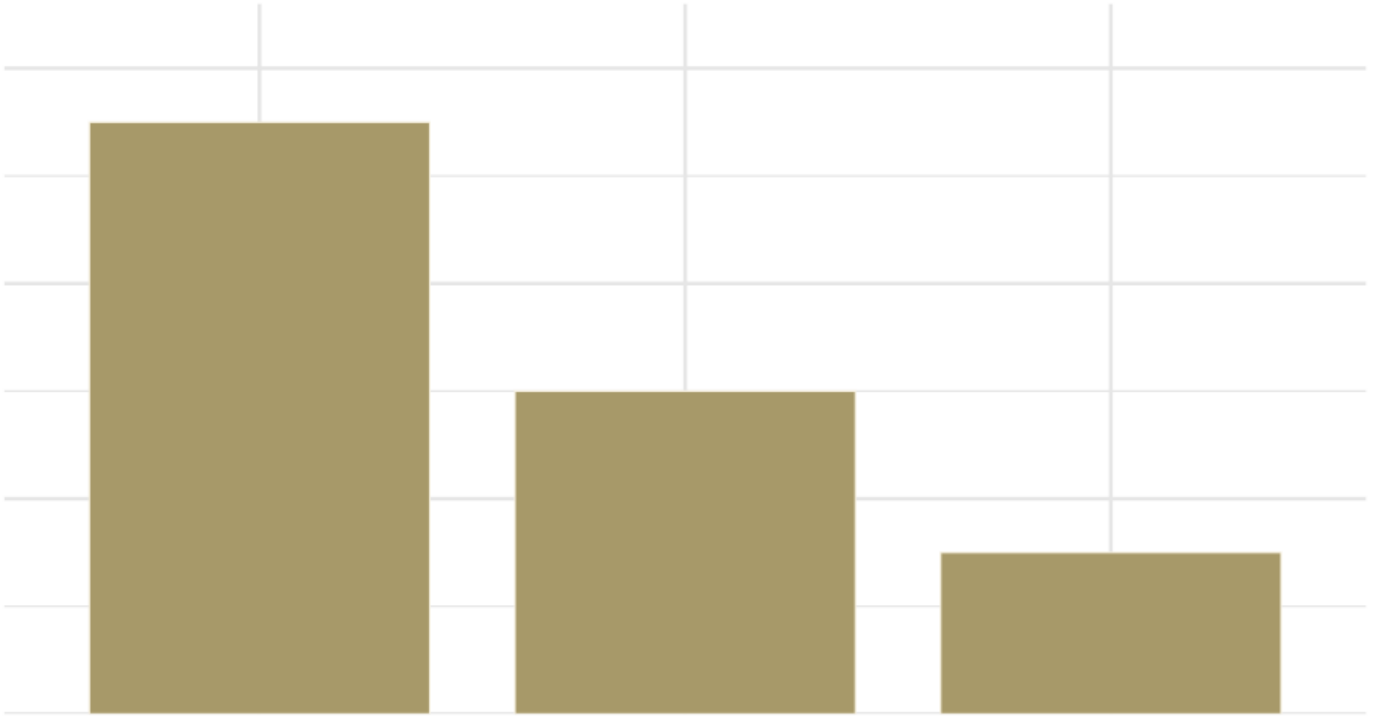

Flipped bars

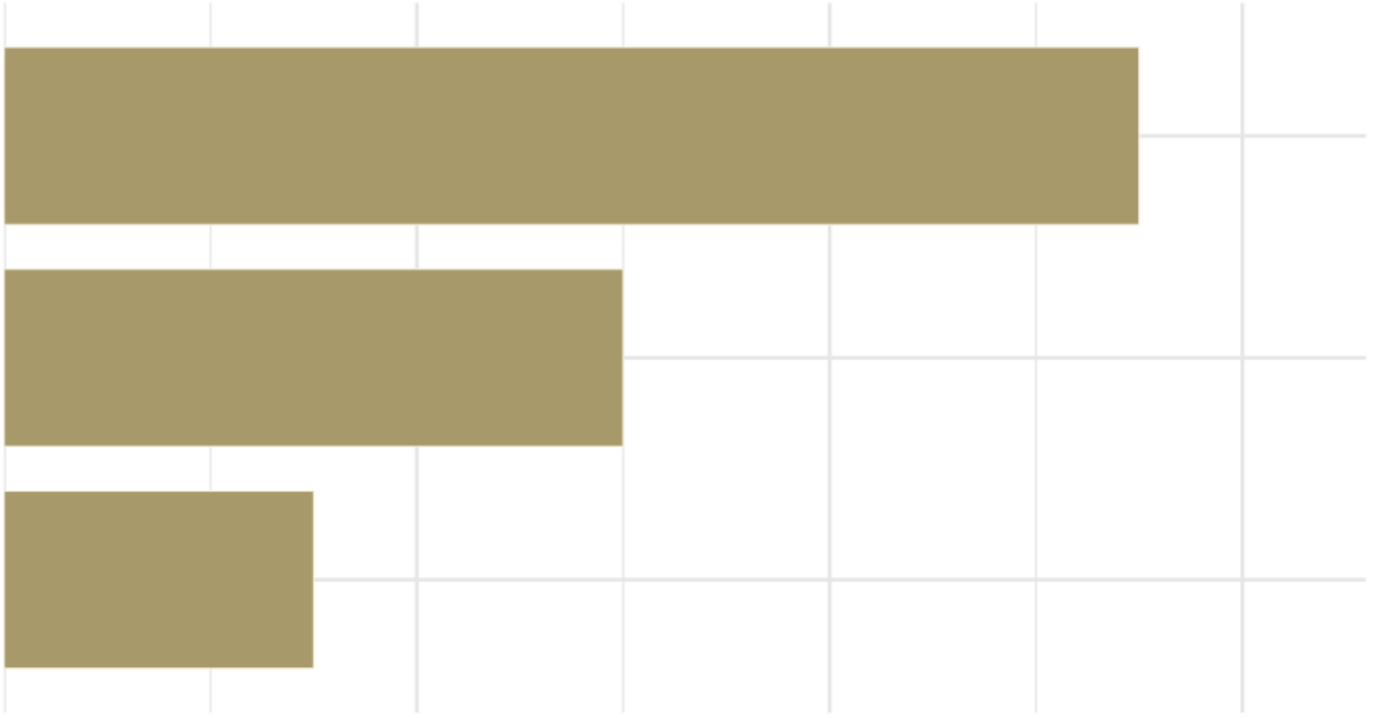

Dotplot

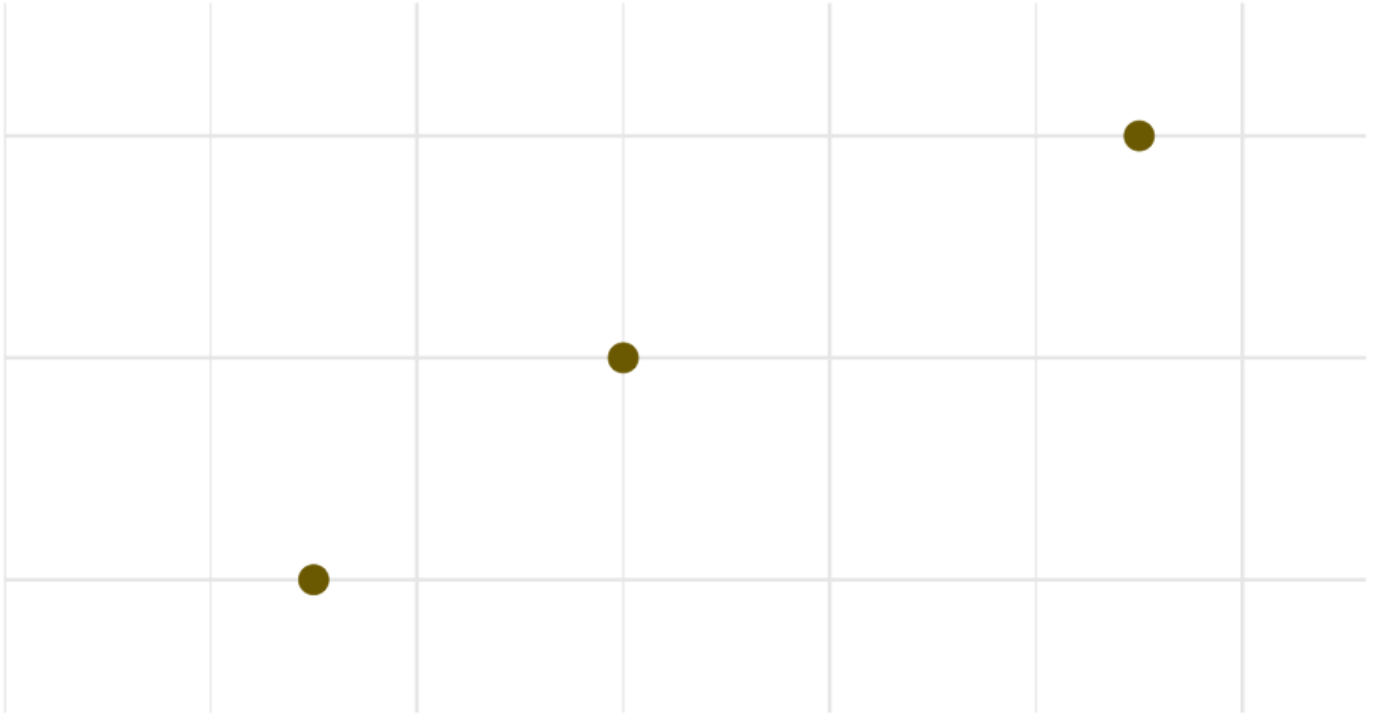

Heatmap

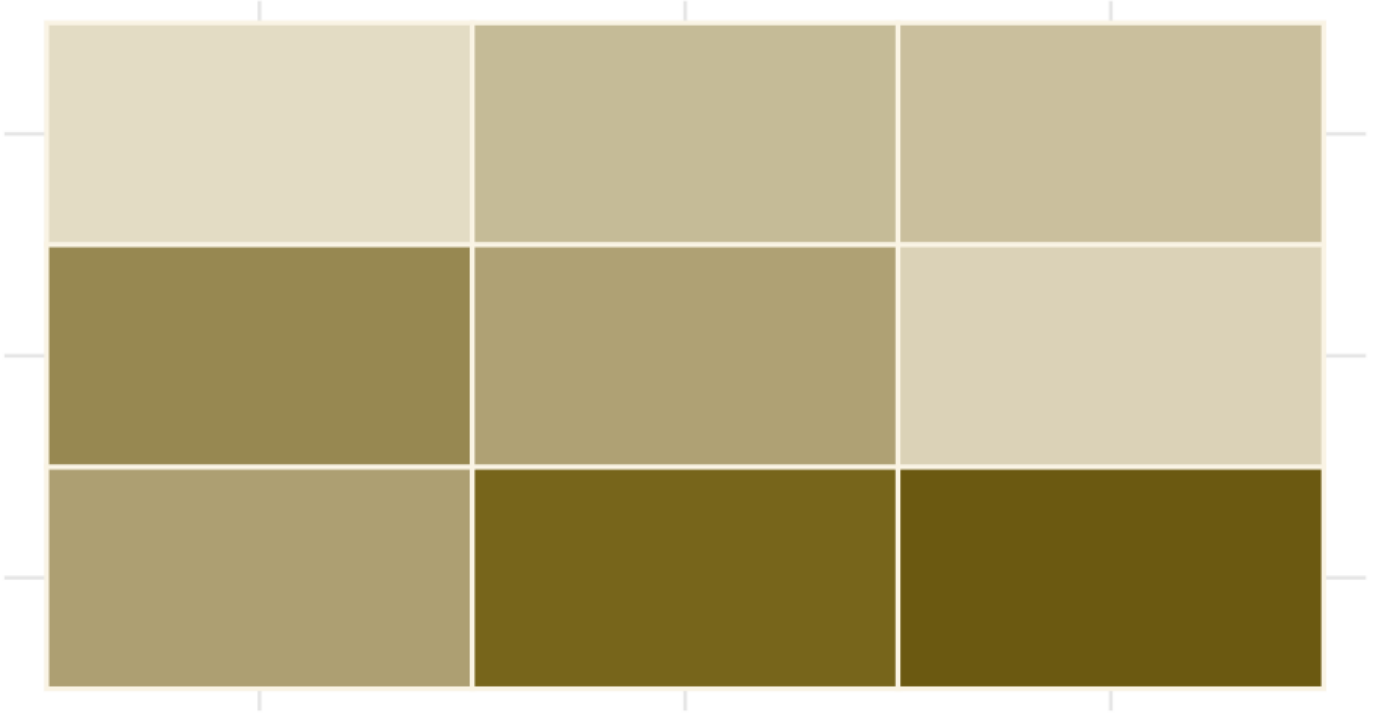

Empirical examples

How much does college cost?

```
library(here)
library(rio)
tuition <- import(here("data", "us_avg_tuition.xlsx"),
                  setclass = "tbl_df")
head(tuition)
```

```
## # A tibble: 6 x 13
##   State      `2004-05` `2005-06` `2006-07` `2007-08` `2008-09` `2009-10`
##   <chr>          <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>
## 1 Alabama     5682.838  5840.550  5753.496  6008.169  6475.092  7188.954
## 2 Alaska      4328.281  4632.623  4918.501  5069.822  5075.482  5454.607
## 3 Arizona     5138.495  5415.516  5481.419  5681.638  6058.464  7263.204
## 4 Arkansas    5772.302  6082.379  6231.977  6414.900  6416.503  6627.092
## 5 California  5285.921  5527.881  5334.826  5672.472  5897.888  7258.771
## 6 Colorado    4703.777  5406.967  5596.348  6227.002  6284.137  6948.473
## # … with 3 more variables: `2013-14` <dbl>, `2014-15` <dbl>, `2015-16` <
```

By state: 2015–16

```
ggplot(tuition, aes(State, `2015-16`)) +
  geom_col()
```

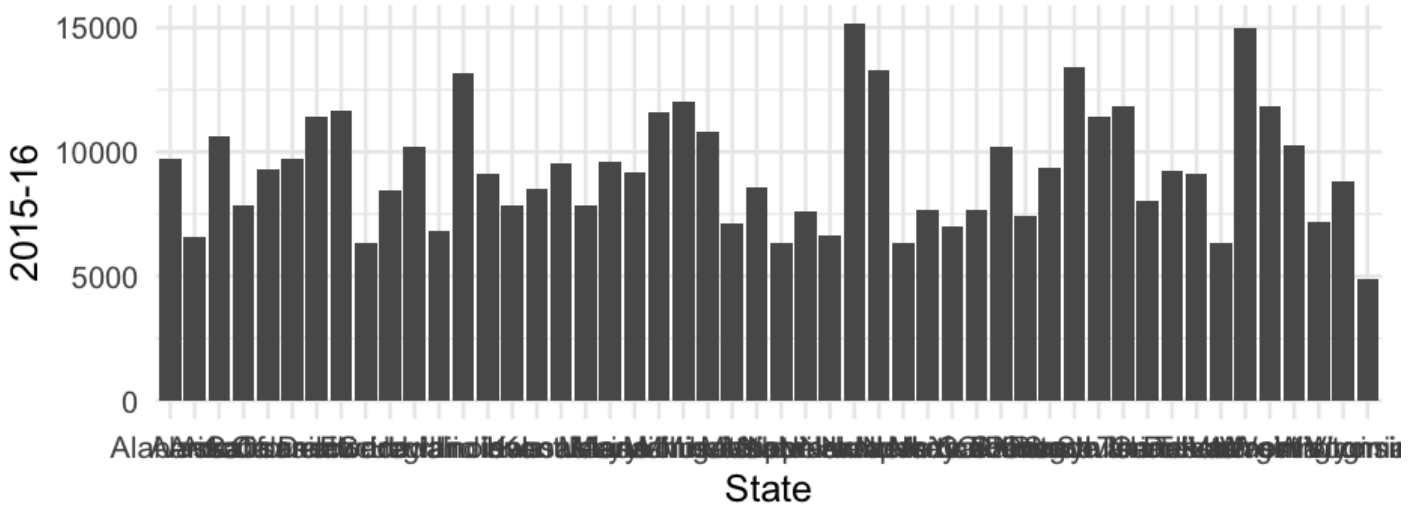

Two puke emoji version

```
ggplot(tuition, aes(State, `2015-16`)) +
  geom_col() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1, size =
```

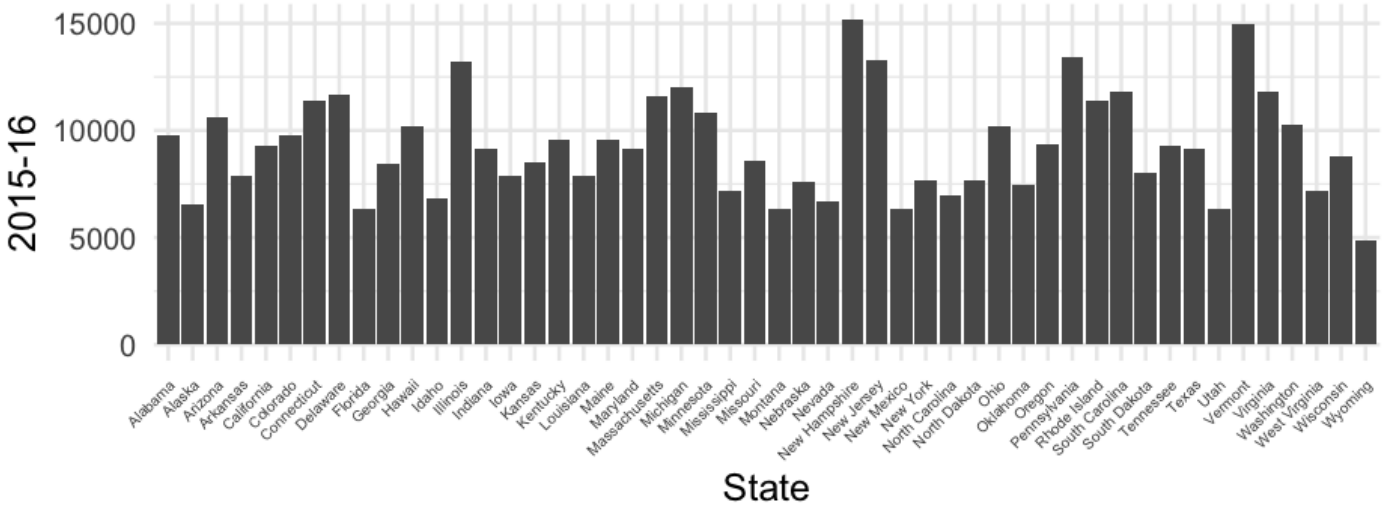

One puke emoji version

```
ggplot(tuition, aes(State, `2015-16`)) +
  geom_col() +
  coord_flip()
```

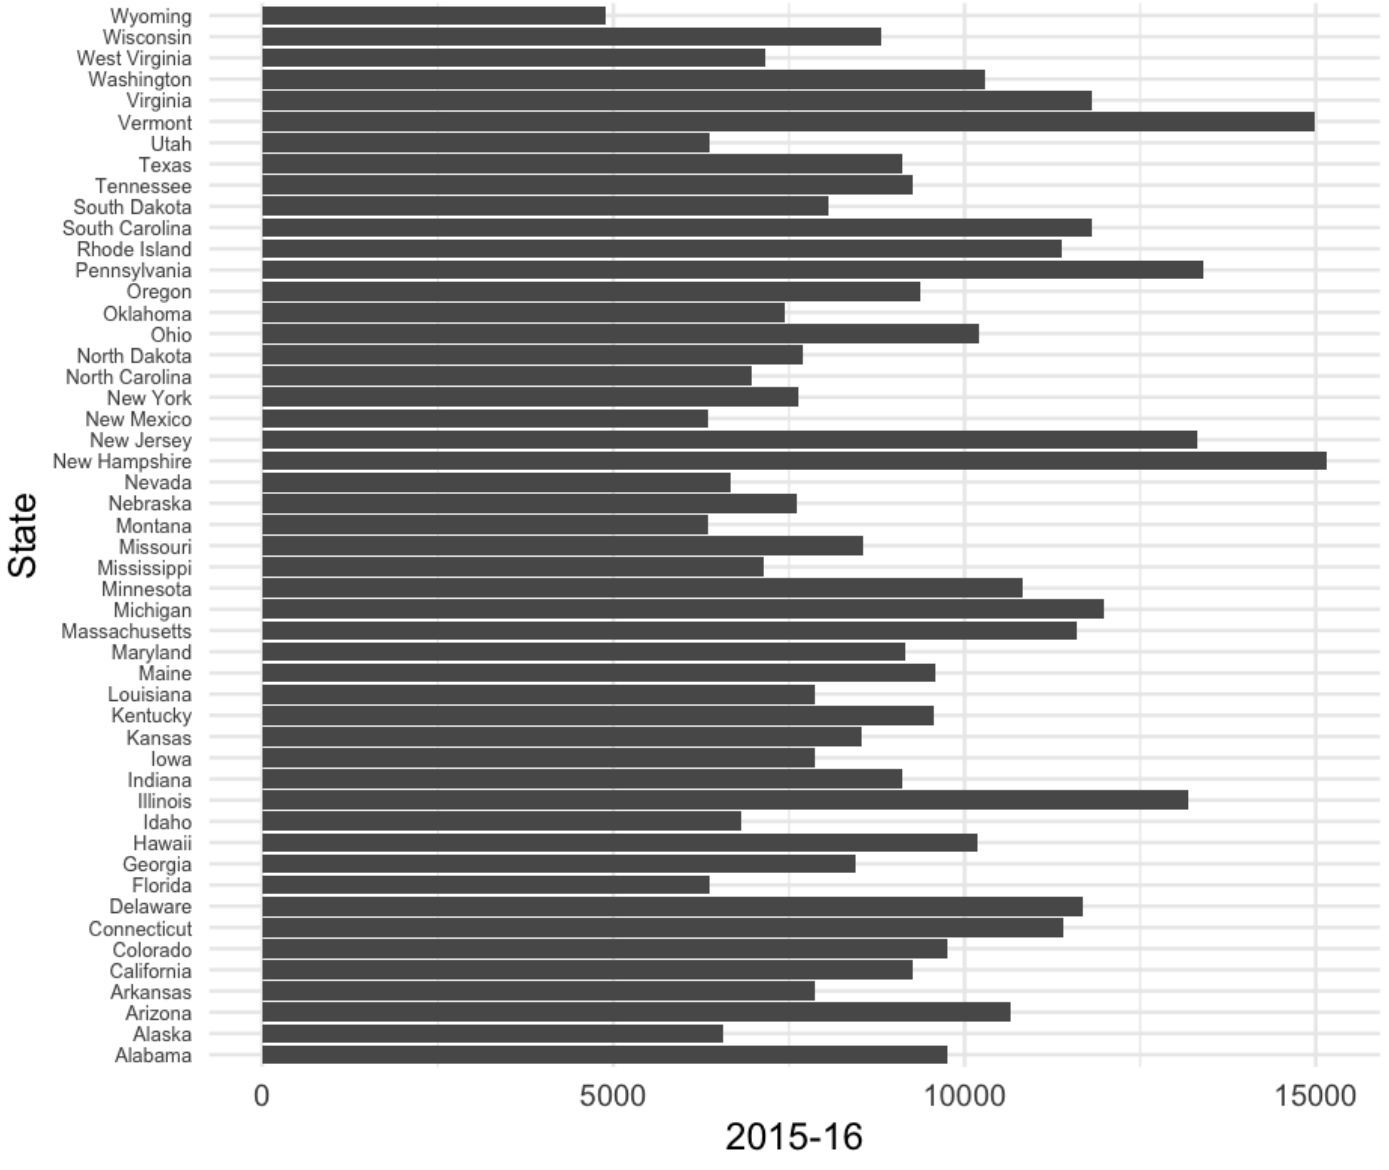
State
Wyoming
Wisconsin
West Virginia
Washington
Virginia
Vermont
Utah
Texas
Tennessee
South Dakota
South Carolina
Rhode Island
Pennsylvania
Oregon
Oklahoma
Ohio
North Dakota
North Carolina
New York
New Mexico
New Jersey
New Hampshire
Nevada
Nebraska
Montana
Missouri
Mississippi
Minnesota
Michigan
Massachusetts
Maryland
Maine
Louisiana
Kentucky
Kansas
Iowa
Indiana
Illinois
Idaho
Hawaii
Georgia
Florida
Delaware
Connecticut
Colorado
California
Arkansas
Arizona
Alaska
Alabama
0
5000
10000
15000
2015-16

Kinda smiley version

```
ggplot(tuition, aes(fct_reorder(State, `2015-16`), `2015-16`)) +
  geom_col() +
  coord_flip()
```

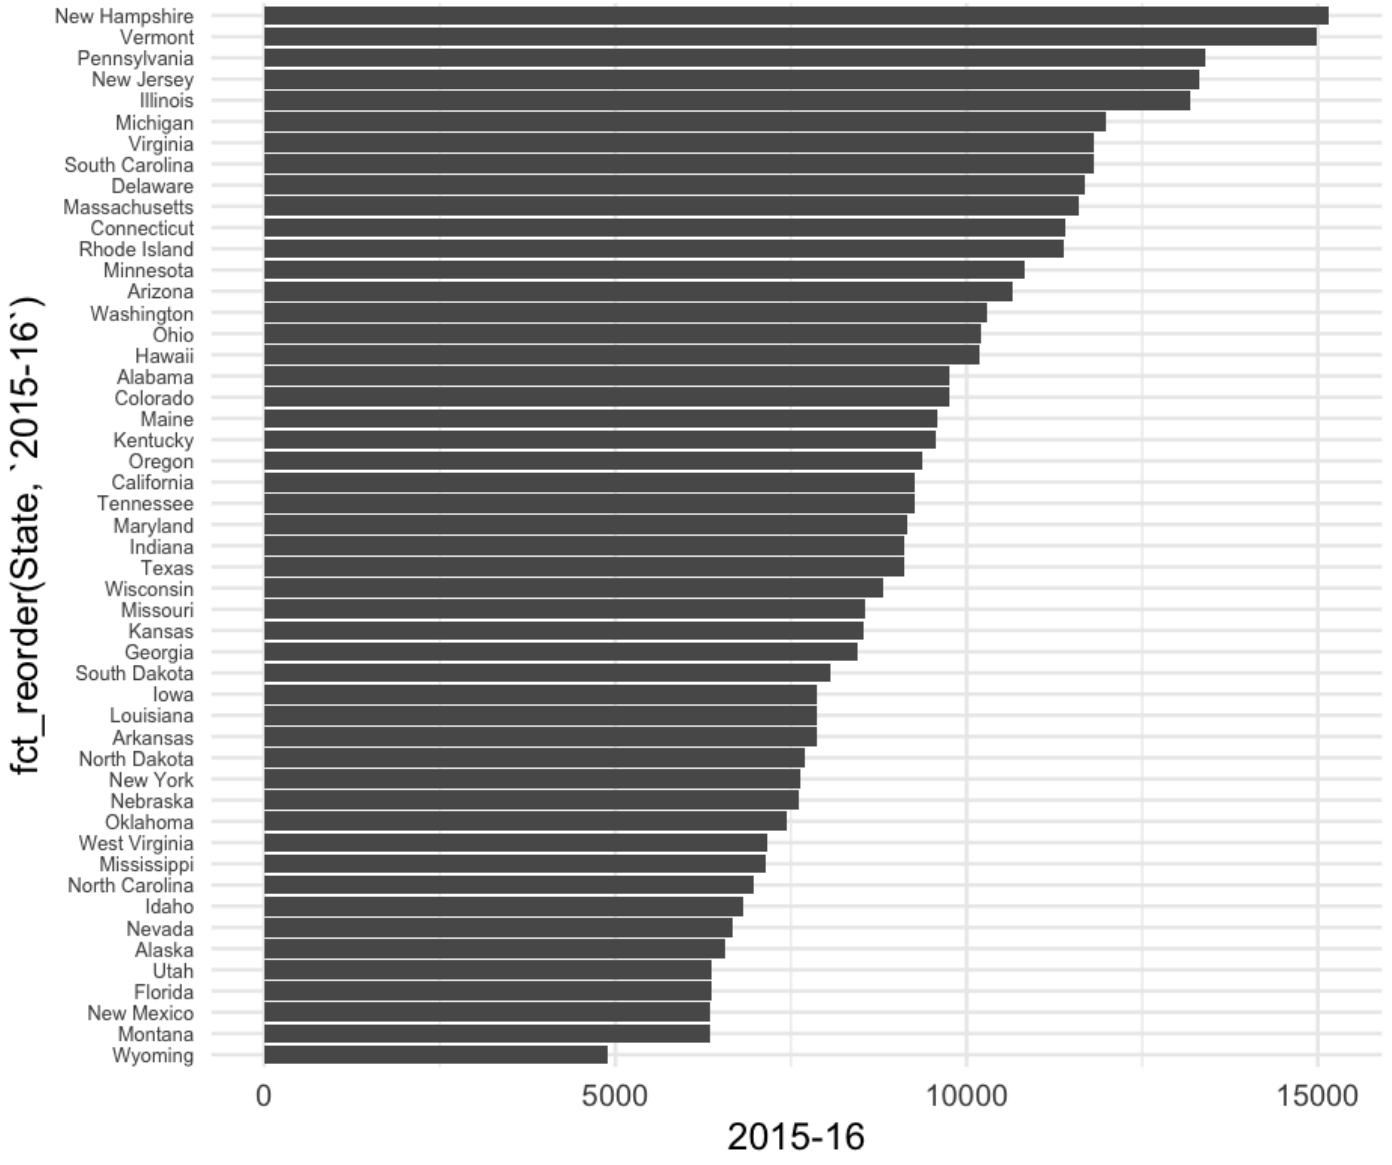
New Hampshire
Vermont
Pennsylvania
New Jersey
Illinois
Michigan
Virginia
South Carolina
Delaware
Massachusetts
Connecticut
Rhode Island
Minnesota
Arizona
Washington
Ohio
Hawaii
Alabama
Colorado
Maine
Kentucky
Oregon
California
Tennessee
Maryland
Indiana
Texas
Wisconsin
Missouri
Kansas
Georgia
South Dakota
Iowa
Louisiana
Arkansas
North Dakota
New York
Nebraska
Oklahoma
West Virginia
Mississippi
North Carolina
Idaho
Nevada
Alaska
Utah
Florida
New Mexico
Montana
Wyoming
fct_reorder(State, `2015-16`)
0
5000
10000
15000
2015-16

Highlight Oregon

```
ggplot(tuition, aes(fct_reorder(State, `2015-16`), `2015-16`)) +
  geom_col() +
  geom_col(fill = "cornflowerblue",
           data = filter(tuition, State == "Oregon")) +
  coord_flip()
```

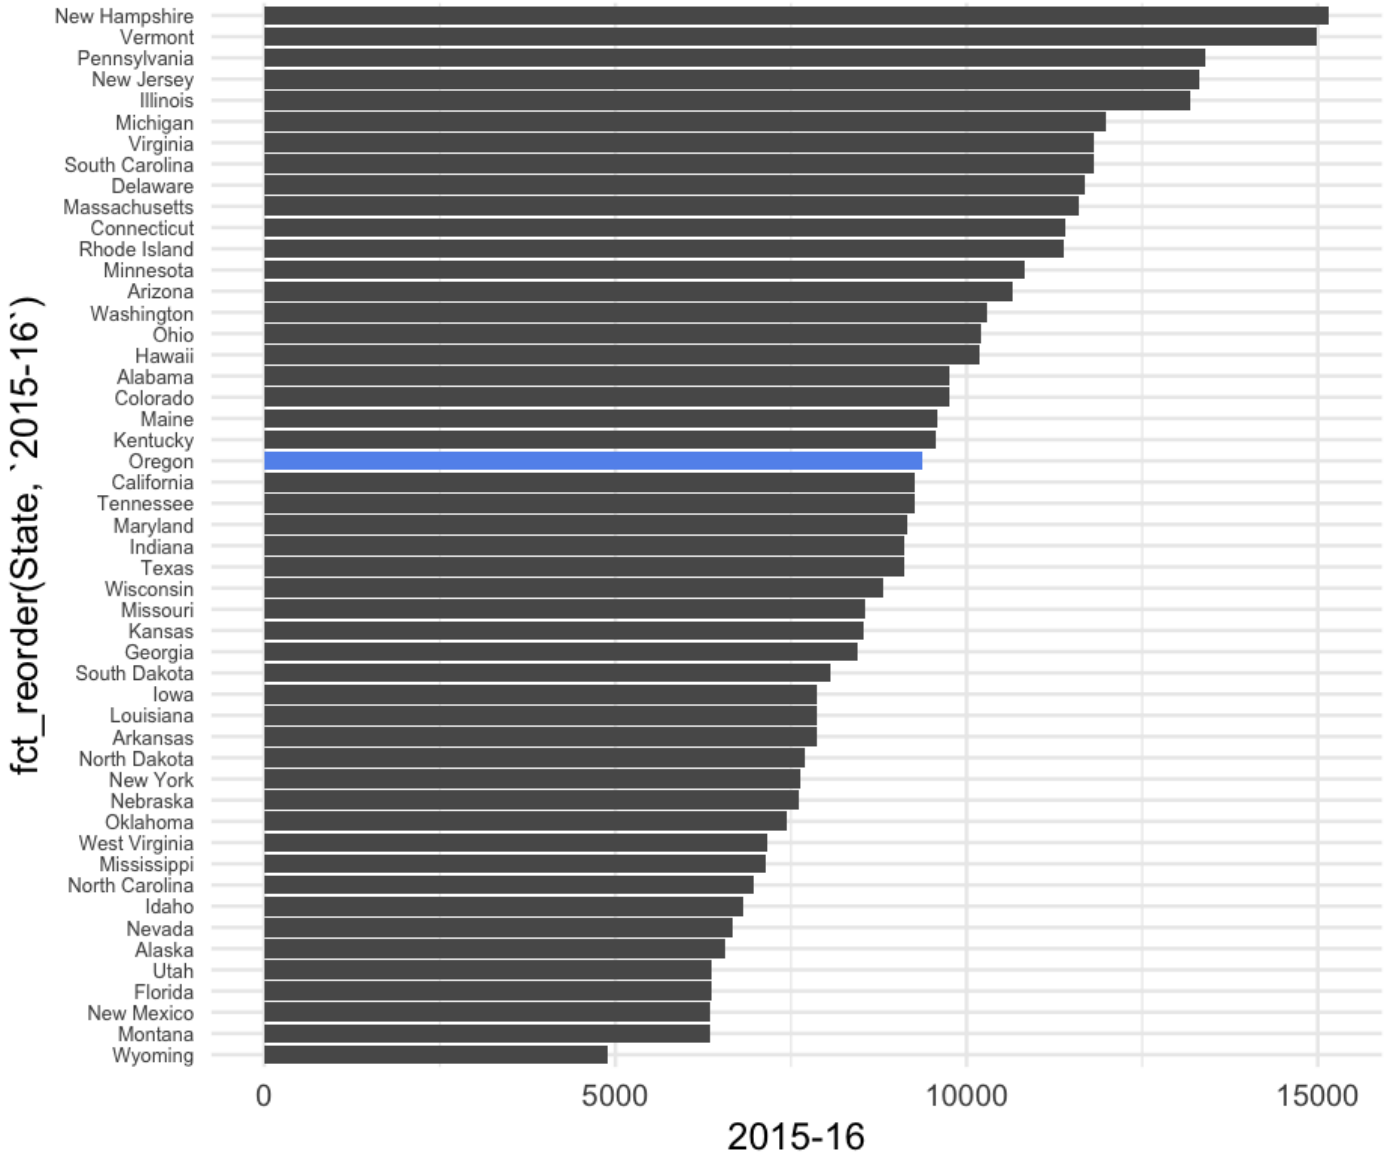
New Hampshire
Vermont
Pennsylvania
New Jersey
Illinois
Michigan
Virginia
South Carolina
Delaware
Massachusetts
Connecticut
Rhode Island
Minnesota
Arizona
Washington
Ohio
Hawaii
Alabama
Colorado
Maine
Kentucky
Oregon
California
Tennessee
Maryland
Indiana
Texas
Wisconsin
Missouri
Kansas
Georgia
South Dakota
Iowa
Louisiana
Arkansas
North Dakota
New York
Nebraska
Oklahoma
West Virginia
Mississippi
North Carolina
Idaho
Nevada
Alaska
Utah
Florida
New Mexico
Montana
Wyoming
fct_reorder(State, `2015-16`)
0
5000
10000
15000
2015-16

Not always good to sort

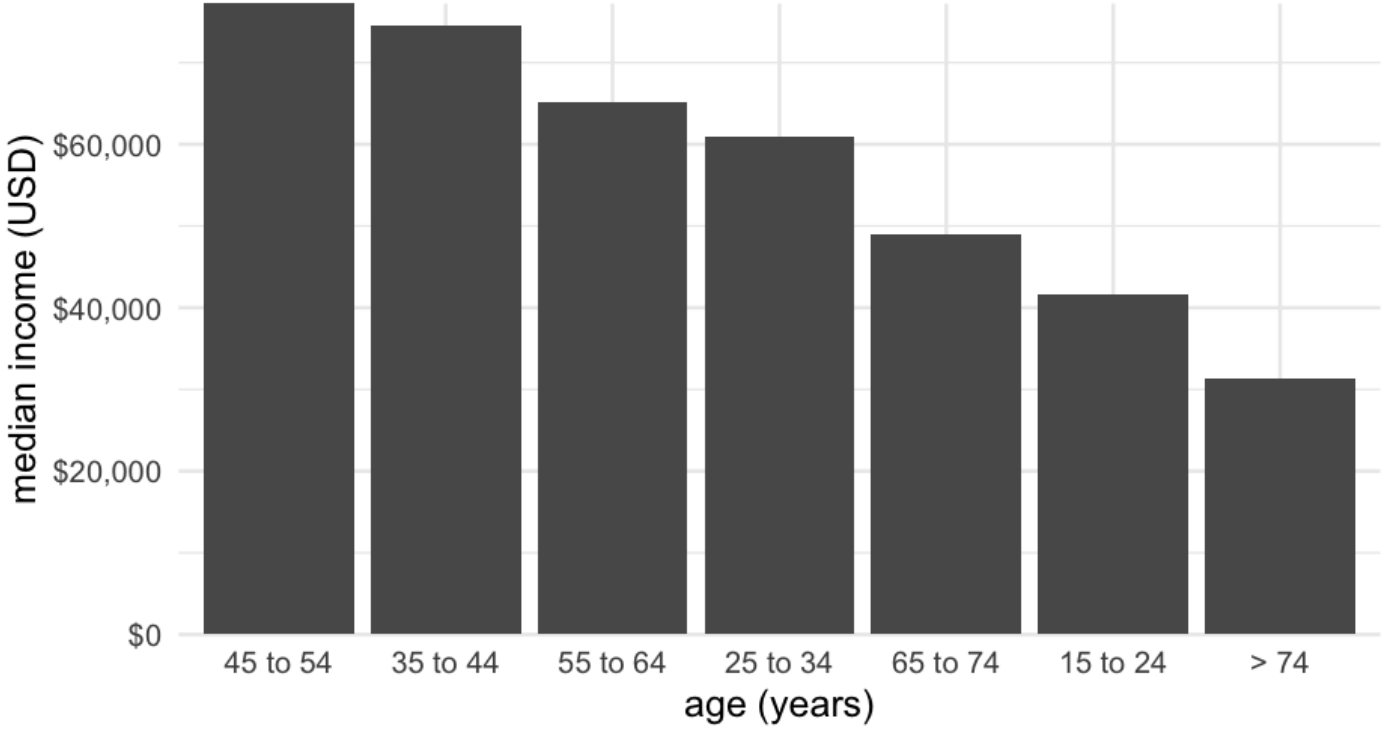

Much better

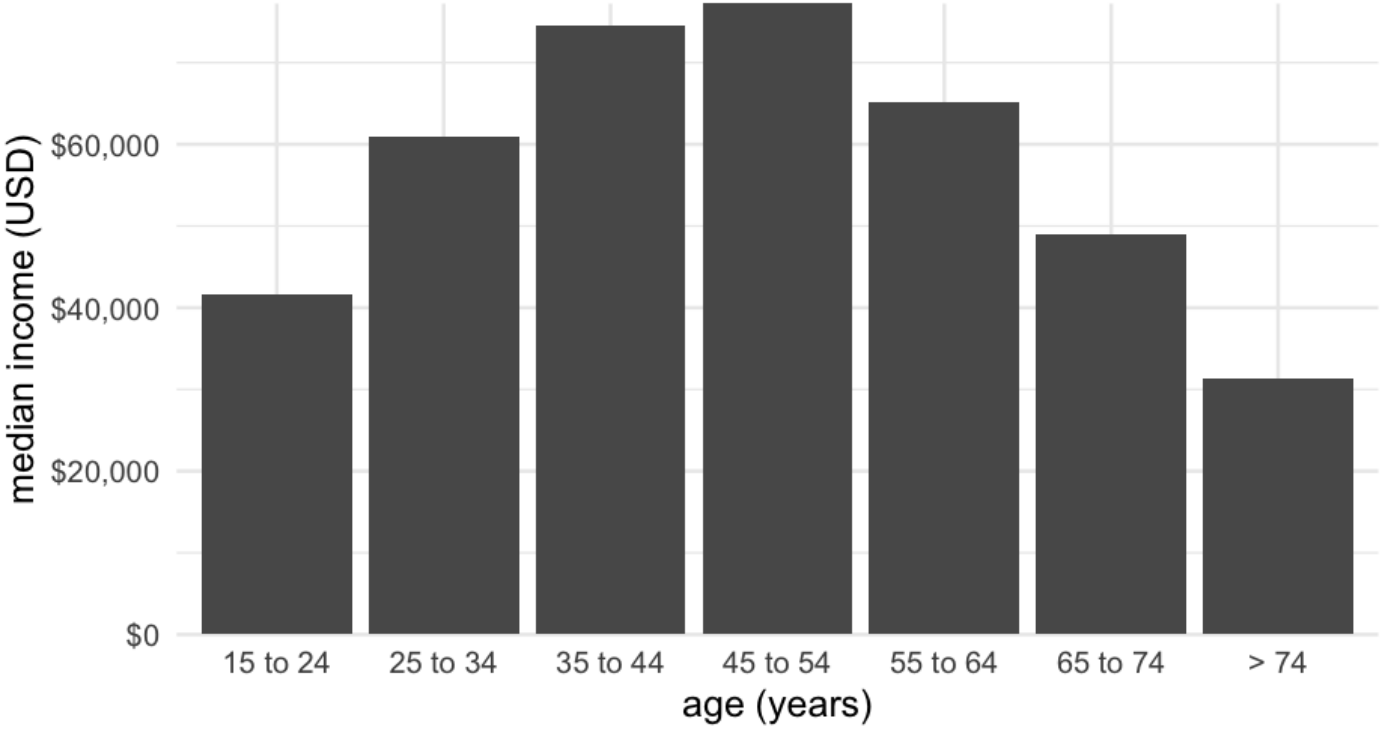
median income (USD)
$60,000
$40,000
$20,000
$0
15 to 24
25 to 34
35 to 44
45 to 54
55 to 64
65 to 74
> 74
age (years)

Averages tuition by year

How?

```
head(tuition)
```

```
## # A tibble: 6 x 13
##   State      `2004-05` `2005-06` `2006-07` `2007-08` `2008-09` `2009-10`
##   <chr>          <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>
## 1 Alabama     5682.838  5840.550  5753.496  6008.169  6475.092  7188.954
## 2 Alaska      4328.281  4632.623  4918.501  5069.822  5075.482  5454.607
## 3 Arizona     5138.495  5415.516  5481.419  5681.638  6058.464  7263.204
## 4 Arkansas    5772.302  6082.379  6231.977  6414.900  6416.503  6627.092
## 5 California  5285.921  5527.881  5334.826  5672.472  5897.888  7258.771
## 6 Colorado    4703.777  5406.967  5596.348  6227.002  6284.137  6948.473
## # … with 3 more variables: `2013-14` <dbl>, `2014-15` <dbl>, `2015-16` <
```

Rearrange

```
tuition %>%
  pivot_longer(`2004-05`:`2015-16`,
               names_to = "year",
               values_to = "avg_tuition")
```

```
## # A tibble: 600 x 3
##   State   year    avg_tuition
##   <chr>   <chr>         <dbl>
## 1 Alabama 2004-05    5682.838
## 2 Alabama 2005-06    5840.550
## 3 Alabama 2006-07    5753.496
## 4 Alabama 2007-08    6008.169
## 5 Alabama 2008-09    6475.092
## 6 Alabama 2009-10    7188.954
## # … with 594 more rows
```

Compute summaries

```
annual_means <- tuition %>%
  pivot_longer(`2004-05`:`2015-16`,
               names_to = "year",
               values_to = "avg_tuition")  %>%
  group_by(year) %>%
  summarize(mean_tuition = mean(avg_tuition))

annual_means
```

```
## # A tibble: 12 x 2
##   year    mean_tuition
## * <chr>          <dbl>
## 1 2004-05     6409.564
## 2 2005-06     6654.177
## 3 2006-07     6809.914
## 4 2007-08     7085.881
## 5 2008-09     7156.560
## 6 2009-10     7761.810
## # … with 6 more rows
```

Good

```
ggplot(annual_means, aes(year, mean_tuition)) +
  geom_col()
```

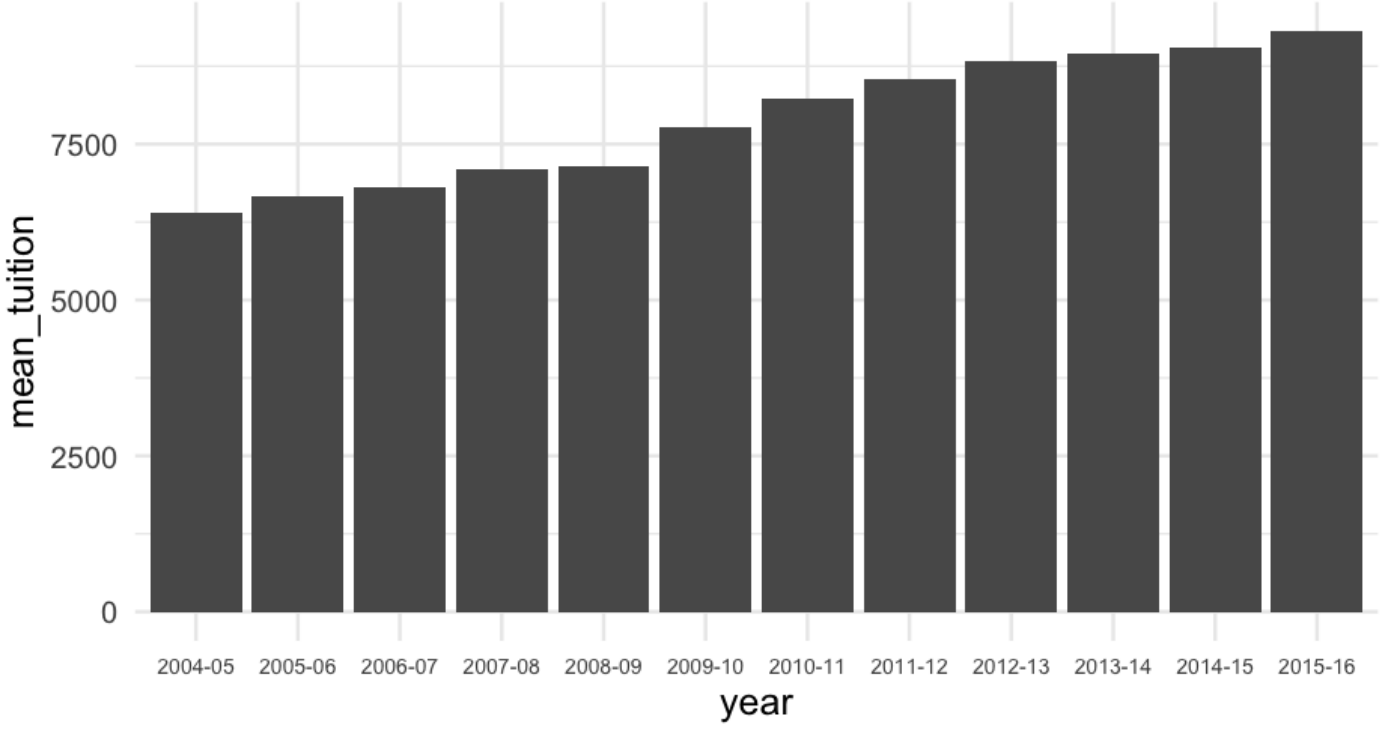

Better?

```
ggplot(annual_means, aes(year, mean_tuition)) +
  geom_col() +
  coord_flip()
```

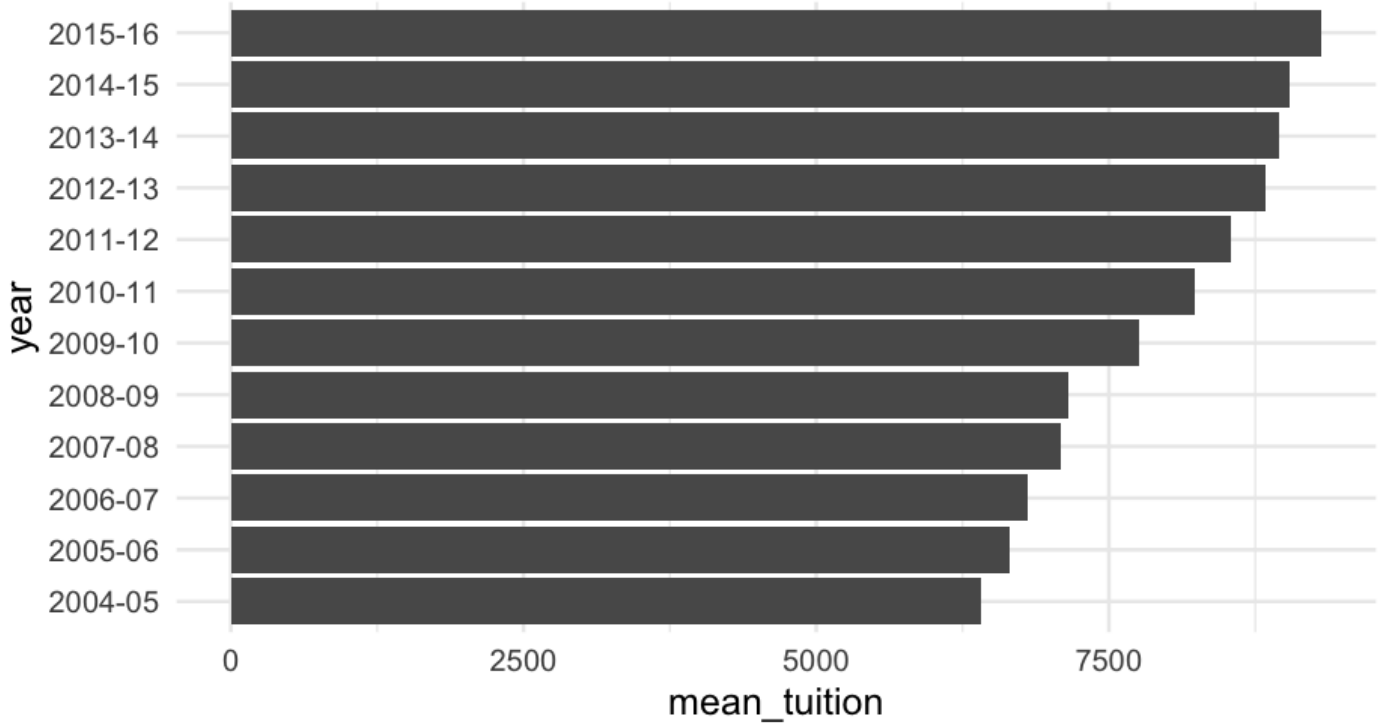

Better still?

```
ggplot(annual_means, aes(year, mean_tuition)) +
  geom_point() +
  coord_flip()
```

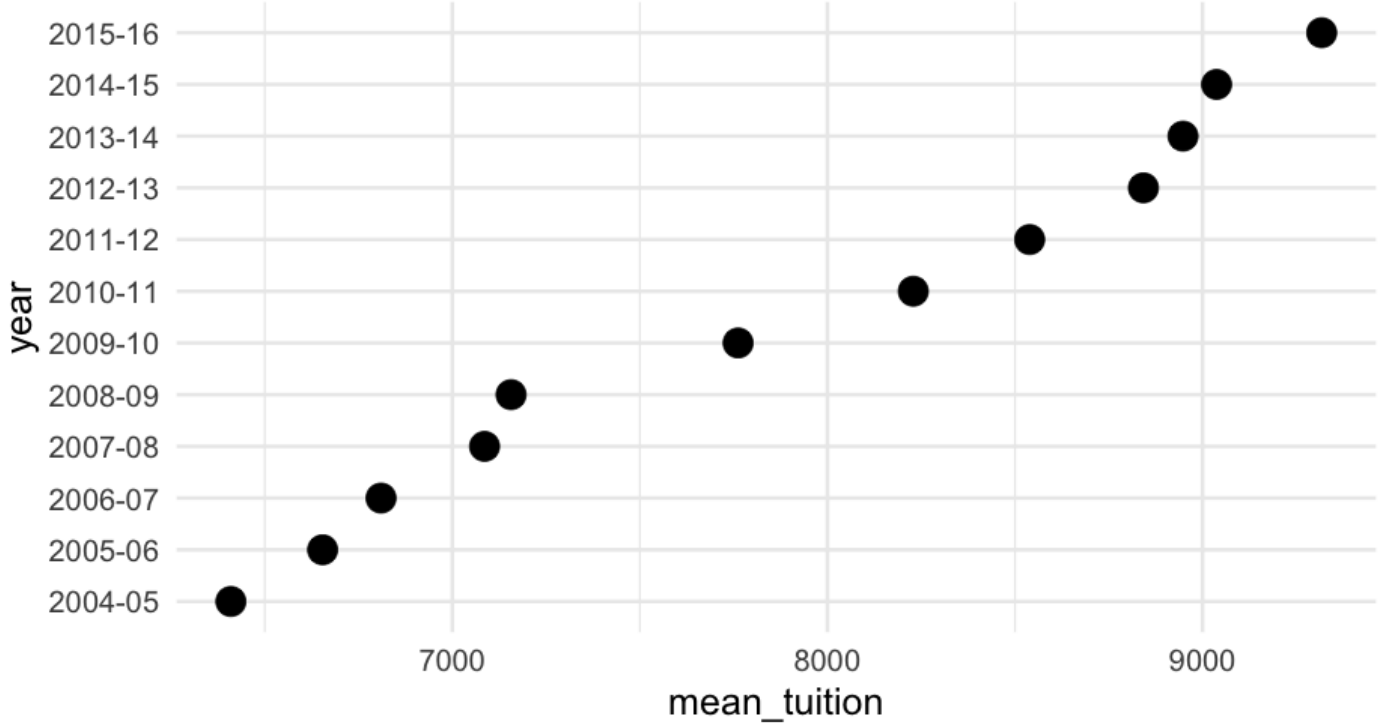

Even better

```
annual_means %>%
  mutate(year = readr::parse_number(year)) %>%
  ggplot(aes(year, mean_tuition)) +
    geom_line(color = "cornflowerblue") +
    geom_point()
```

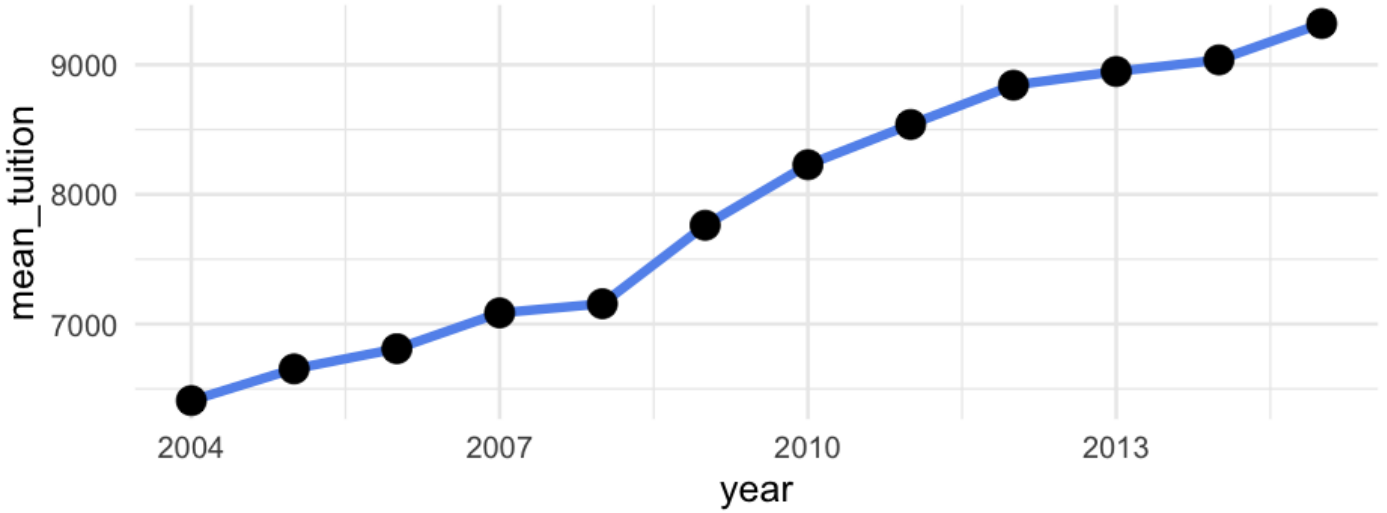

Treat time (year) as a continuous variable

Grouped points

Show change in tuition from 05–06 to 2015–16

```
tuition %>%
  select(State, `2005-06`, `2015-16`)
```

```
## # A tibble: 50 x 3
##   State      `2005-06` `2015-16`
##   <chr>          <dbl>     <dbl>
## 1 Alabama     5840.550  9751.101
## 2 Alaska      4632.623  6571.340
## 3 Arizona     5415.516 10646.28
## 4 Arkansas    6082.379  7867.297
## 5 California  5527.881  9269.844
## 6 Colorado    5406.967  9748.188
## # … with 44 more rows
```

```r
lt <- tuition %>%
  select(State, `2005-06`, `2015-16`) %>%
  pivot_longer(`2005-06`:`2015-16`,
               names_to = "Year",
               values_to = "Tuition")
lt
```

```
## # A tibble: 100 x 3
##   State   Year      Tuition
##   <chr>   <chr>       <dbl>
## 1 Alabama 2005-06  5840.550
## 2 Alabama 2015-16  9751.101
## 3 Alaska  2005-06  4632.623
## 4 Alaska  2015-16  6571.340
## 5 Arizona 2005-06  5415.516
## 6 Arizona 2015-16 10646.28
## # … with 94 more rows
```

```
ggplot(lt, aes(State, Tuition)) +
  geom_line(aes(group = State), color = "gray40") +
  geom_point(aes(color = Year)) +
  coord_flip()
```

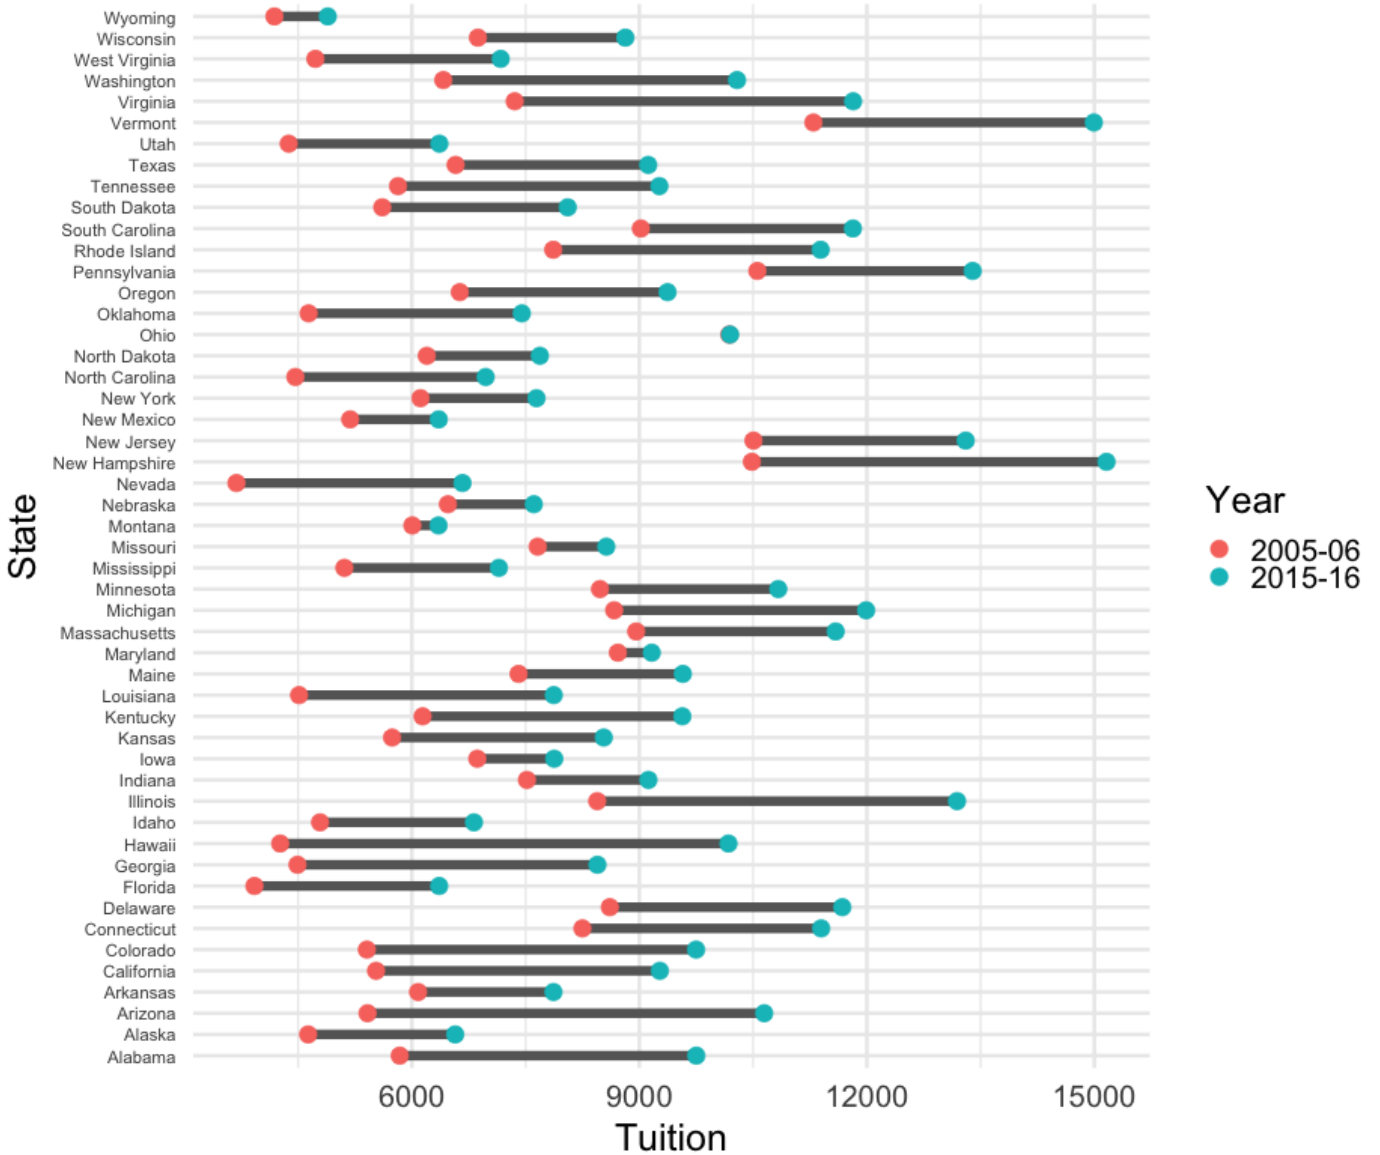
Wyoming
Wisconsin
West Virginia
Washington
Virginia
Vermont
Utah
Texas
Tennessee
South Dakota
South Carolina
Rhode Island
Pennsylvania
Oregon
Oklahoma
Ohio
North Dakota
North Carolina
New York
New Mexico
New Jersey
New Hampshire
Nevada
Nebraska
Montana
Missouri
Mississippi
Minnesota
Michigan
Massachusetts
Maryland
Maine
Louisiana
Kentucky
Kansas
Iowa
Indiana
Illinois
Idaho
Hawaii
Georgia
Florida
Delaware
Connecticut
Colorado
California
Arkansas
Arizona
Alaska
Alabama
State
6000
9000
12000
15000
Tuition
Year
2005-06
2015-16

Extensions

- I know we're probably running short on time, but we definitely would want to keep going here:
    - Order states according to something more meaningful (starting tuition, ending tuition, or difference in tuition)
    - Meaningful title, e.g., "Change in average tuition over a decade"
    - Consider better color scheme for points

Let's back up a bit

- Lets go back to our full data, but in a format that we can have a `year` variable.

```
tuition_l <- tuition %>%
    pivot_longer(-State,
                 names_to = "year",
                 values_to = "avg_tuition")

tuition_l
```

```
## # A tibble: 600 x 3
##   State   year    avg_tuition
##   <chr>   <chr>         <dbl>
## 1 Alabama 2004-05    5682.838
## 2 Alabama 2005-06    5840.550
## 3 Alabama 2006-07    5753.496
## 4 Alabama 2007-08    6008.169
## 5 Alabama 2008-09    6475.092
## 6 Alabama 2009-10    7188.954
## # … with 594 more rows
```

Heatmap

```
ggplot(tuition_l, aes(year, State)) +
  geom_tile(aes(fill = avg_tuition))
```

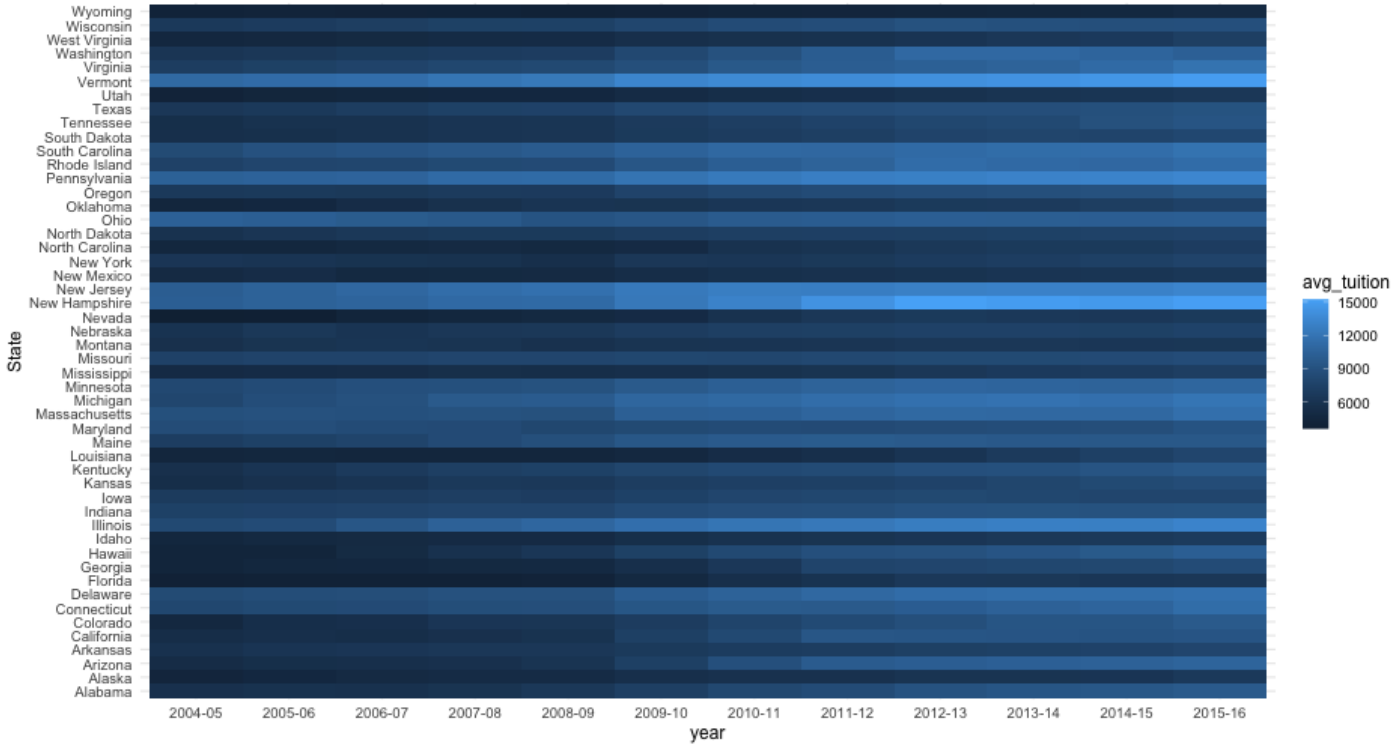

Better heatmap

```
ggplot(tuition_l, aes(year, fct_reorder(State, avg_tuition))) +
  geom_tile(aes(fill = avg_tuition))
```

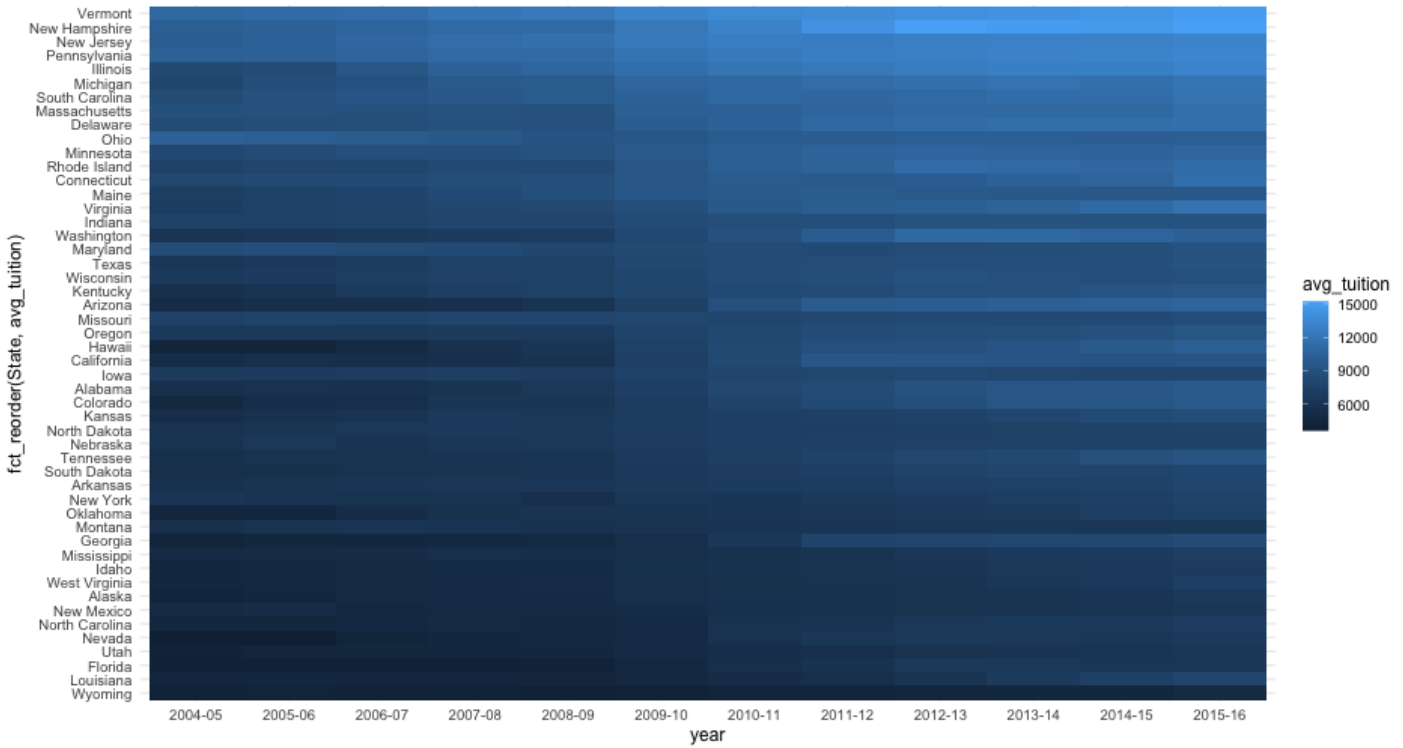

Even better heatmap

```
ggplot(tuition_l, aes(year, fct_reorder(State, avg_tuition))) +
  geom_tile(aes(fill = avg_tuition)) +
  scale_fill_viridis_c(option = "magma")
```

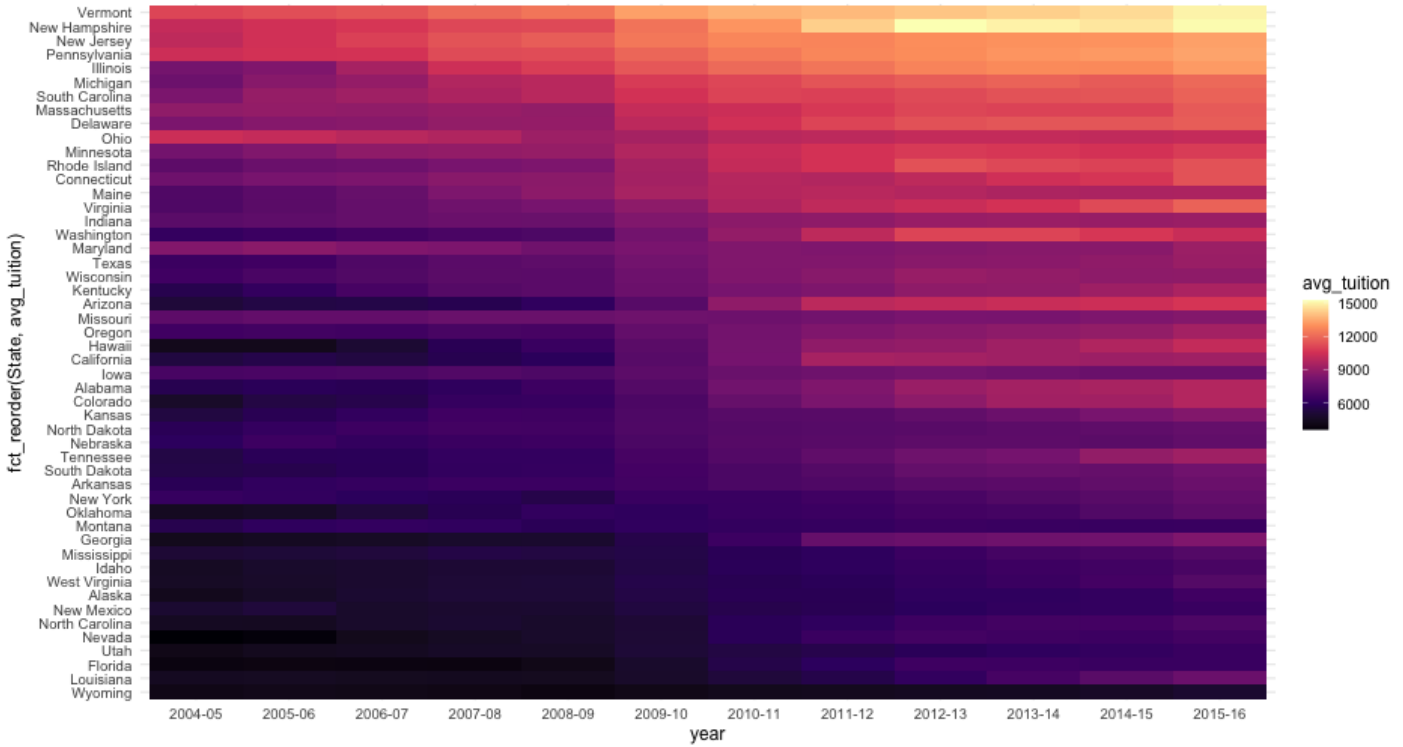

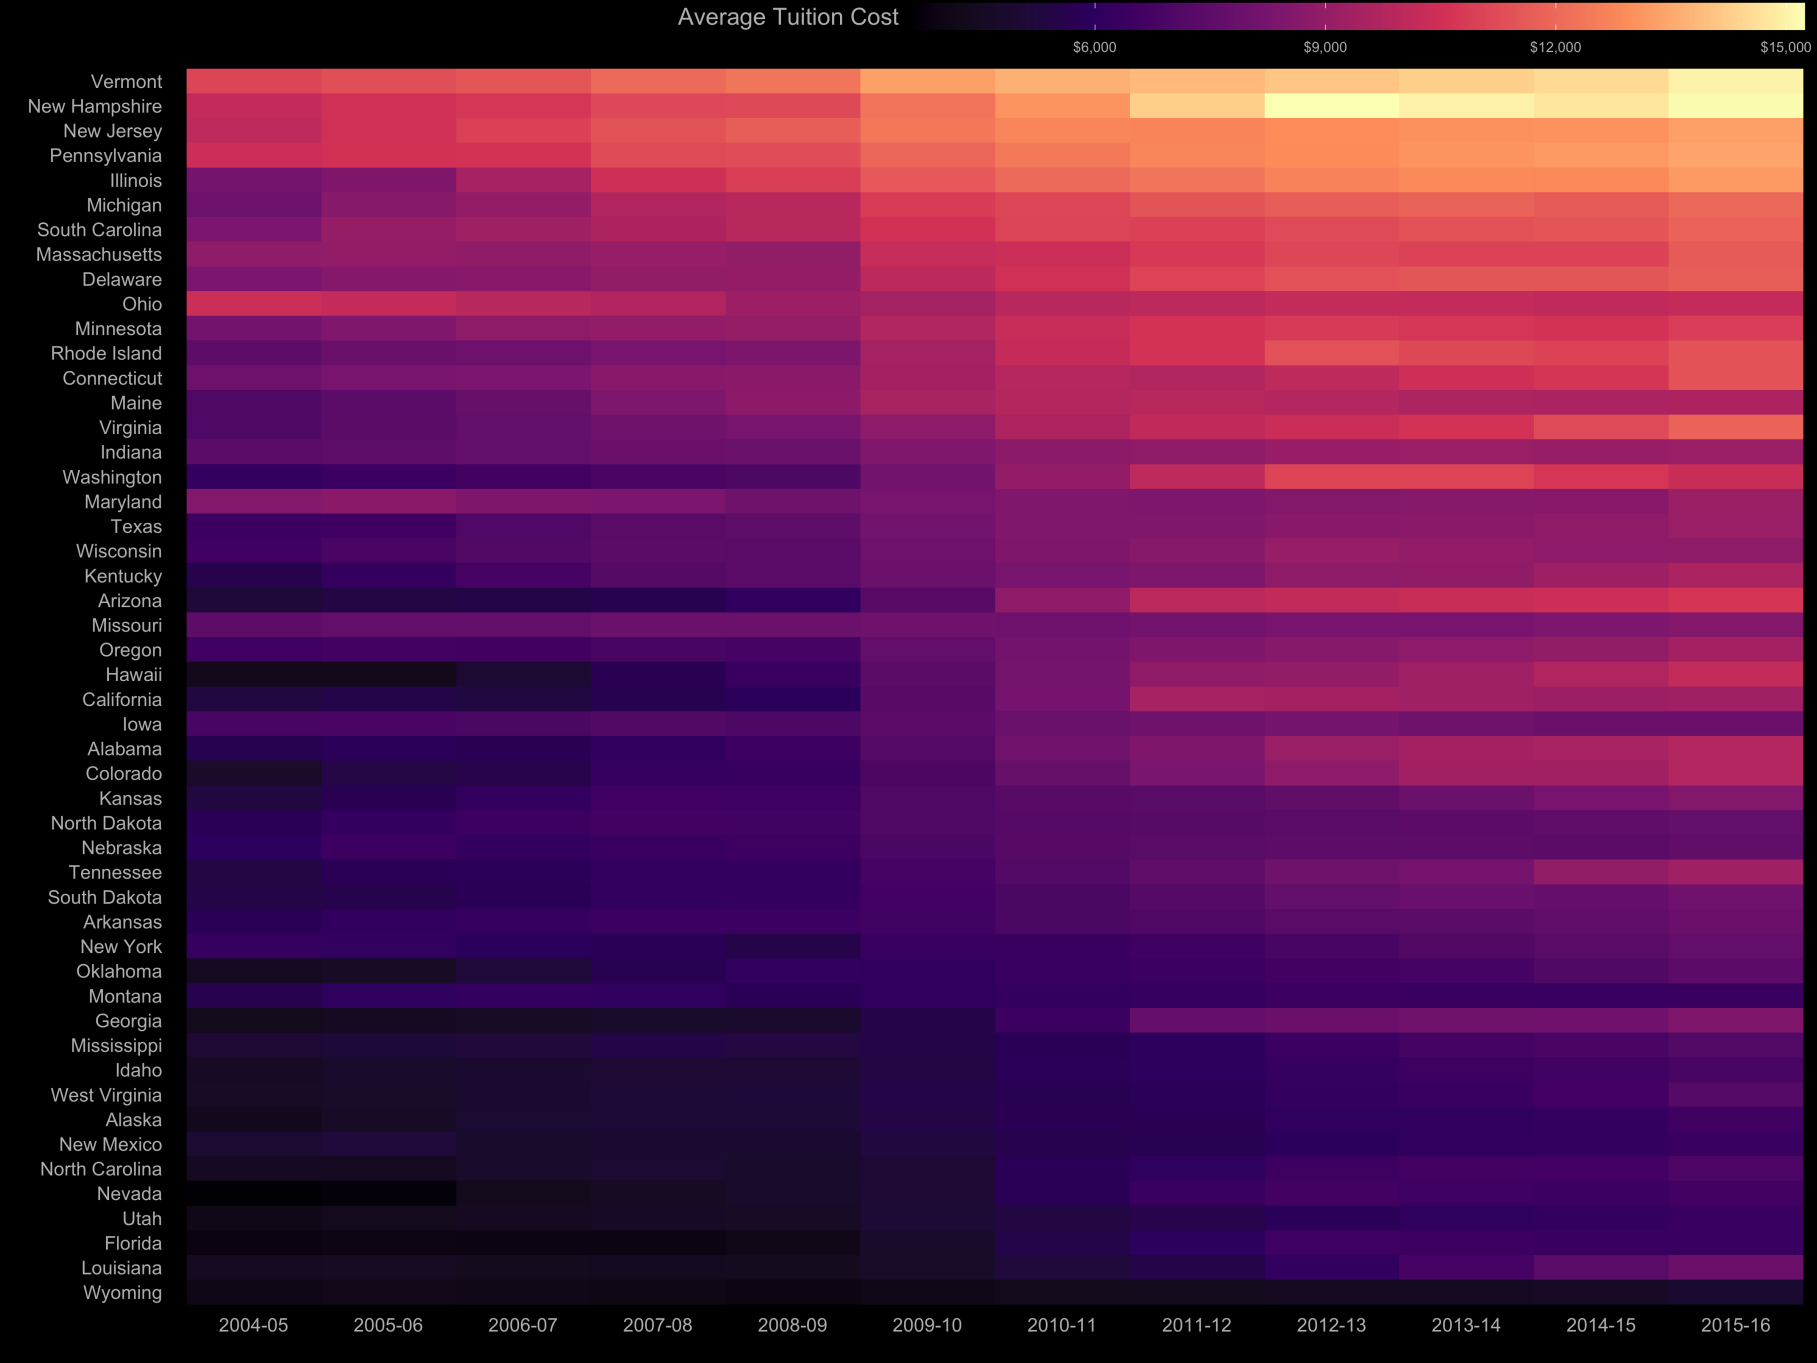
Average Tuition Cost
$6,000
$9,000
$12,000
$15,000
Vermont
New Hampshire
New Jersey
Pennsylvania
Illinois
Michigan
South Carolina
Massachusetts
Delaware
Ohio
Minnesota
Rhode Island
Connecticut
Maine
Virginia
Indiana
Washington
Maryland
Texas
Wisconsin
Kentucky
Arizona
Missouri
Oregon
Hawaii
California
Iowa
Alabama
Colorado
Kansas
North Dakota
Nebraska
Tennessee
South Dakota
Arkansas
New York
Oklahoma
Montana
Georgia
Mississippi
Idaho
West Virginia
Alaska
New Mexico
North Carolina
Nevada
Utah
Florida
Louisiana
Wyoming
2004-05
2005-06
2006-07
2007-08
2008-09
2009-10
2010-11
2011-12
2012-13
2013-14
2014-15
2015-16

Quick aside

- Think about the data you have
- Given that these are state–level data, they have a geographic component

```
#install.packages("maps")
state_data <- map_data("state") %>% # ggplot2::map_data
  rename(State = region)
```

Join it

Obviously we'll talk more about joins later

```
tuition <- tuition %>%
  mutate(State = tolower(State))
states <- left_join(state_data, tuition)
head(states)
```

```
##    long  lat group order   State subregion 2004-05 2005-06 2006-07 2007-
## 1 -87.5 30.4     1     1 alabama      <NA>    5683    5841    5753    60
## 2 -87.5 30.4     1     2 alabama      <NA>    5683    5841    5753    60
## 3 -87.5 30.4     1     3 alabama      <NA>    5683    5841    5753    60
## 4 -87.5 30.3     1     4 alabama      <NA>    5683    5841    5753    60
## 5 -87.6 30.3     1     5 alabama      <NA>    5683    5841    5753    60
## 6 -87.6 30.3     1     6 alabama      <NA>    5683    5841    5753    60
##   2011-12 2012-13 2013-14 2014-15 2015-16
## 1    8452    9098    9359    9496    9751
## 2    8452    9098    9359    9496    9751
## 3    8452    9098    9359    9496    9751
## 4    8452    9098    9359    9496    9751
## 5    8452    9098    9359    9496    9751
## 6    8452    9098    9359    9496    9751
```

Rearrange

```
states <- states %>%
  gather(year, tuition, `2004-05`:`2015-16`)
head(states)
```

```
##    long  lat group order   State subregion    year tuition
## 1 -87.5 30.4     1     1 alabama      <NA> 2004-05    5683
## 2 -87.5 30.4     1     2 alabama      <NA> 2004-05    5683
## 3 -87.5 30.4     1     3 alabama      <NA> 2004-05    5683
## 4 -87.5 30.3     1     4 alabama      <NA> 2004-05    5683
## 5 -87.6 30.3     1     5 alabama      <NA> 2004-05    5683
## 6 -87.6 30.3     1     6 alabama      <NA> 2004-05    5683
```

Plot

```
ggplot(states) +
  geom_polygon(aes(long, lat, group = group, fill = tuition)) +
  coord_fixed(1.3) +
  scale_fill_viridis_c(option = "magma") +
  facet_wrap(~year)
```

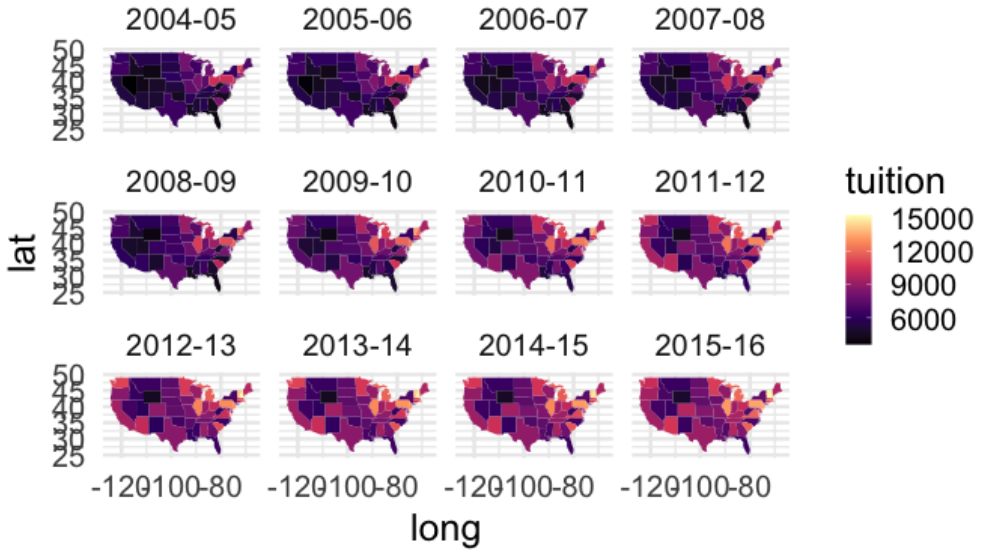

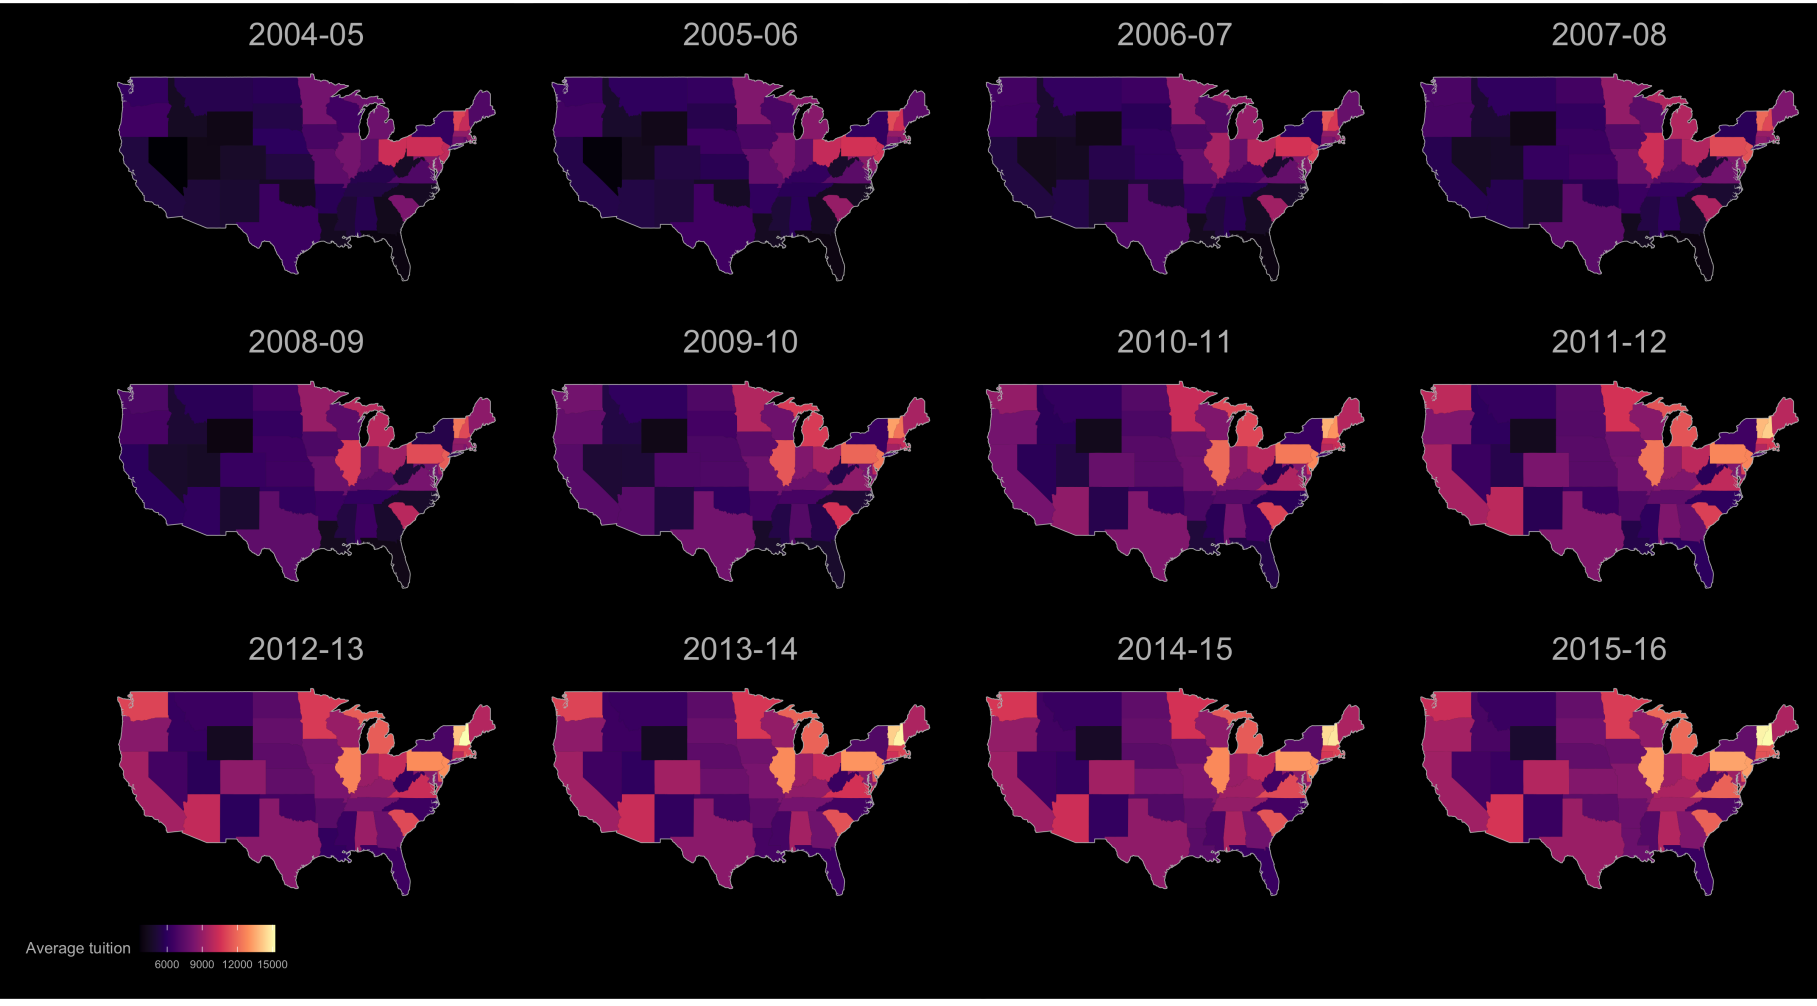
2004-05
2005-06
2006-07
2007-08
2008-09
2009-10
2010-11
2011-12
2012-13
2013-14
2014-15
2015-16
Average tuition
6000 9000 12000 15000

Average Tuition Cost

2004-05

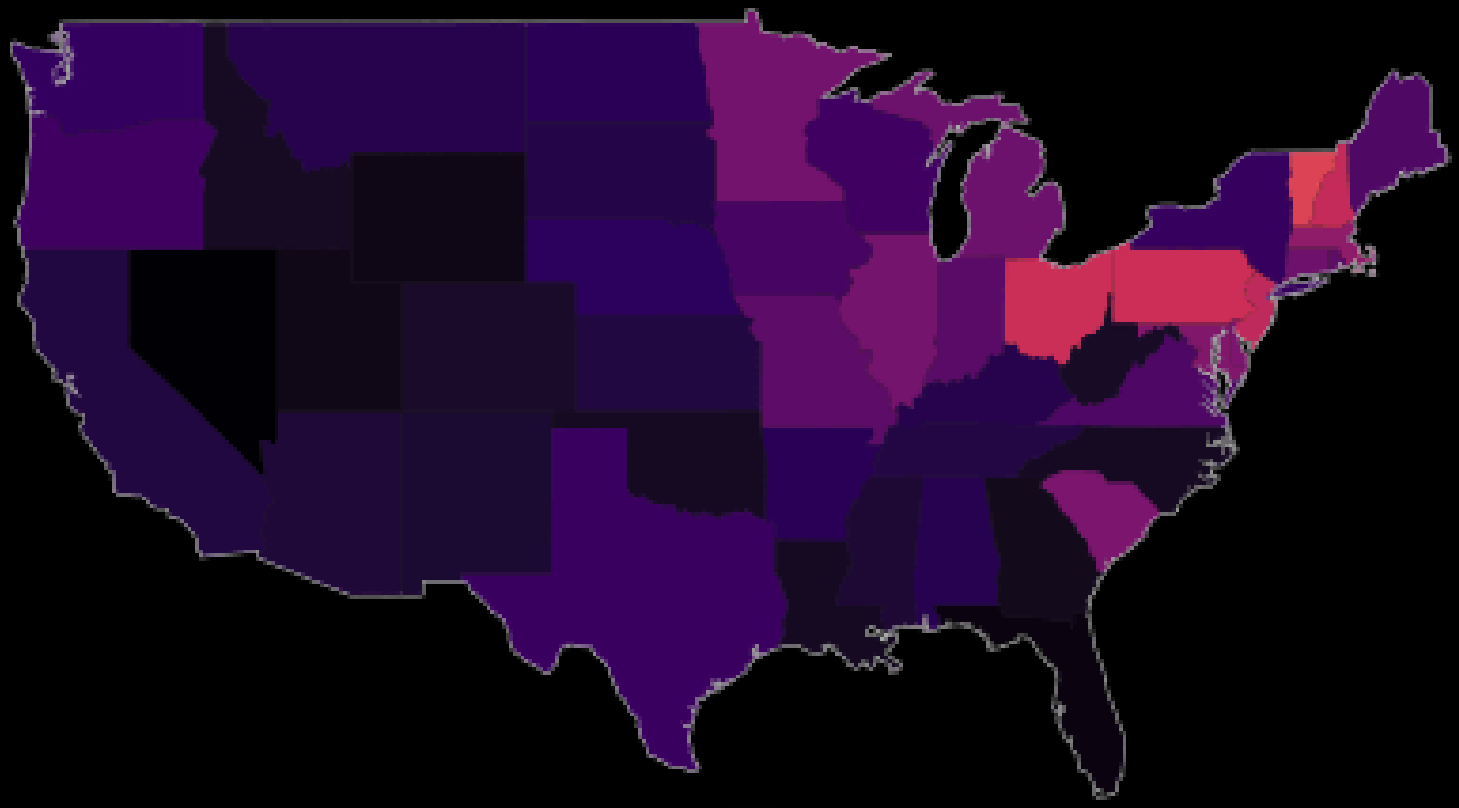

Or animated

Wrapping up

- We've got a ways to go – today was just an introduction
- The geographic part in particular was too fast, and we'll talk about better ways later (note that Alaska/Hawaii were not even included)
- We basically didn't talk about multivariate data (not even scatter plots)
- Other types of plots will be embedded within the topics later in the class

Next time

Lab 2

git/GitHub collaboration

It's already posted – feel free to start working on it whenever.

- Must be completed as a group
- Will use elements of what we talked about today, while also asking you to create branches, submit pull requests, etc.