

Learning Objectives

- Understand and be able to identify keys
- Understand different types of joins
 - left, right, inner, full
 - one-to-one, one-to-many
- Understand common ways joins fail
- Understand the difference between mutating and filtering joins

Before we get started

- Today we'll talk about both **mutating** and **filtering** joins
- Mutating joins are more common, but filtering joins can be really powerful
- Mutating joins add columns to a dataset

What if I want to add rows?

- Not technically a join (no key involved, which we'll talk about momentarily)

Quick example, binding rows

```
g3 <- tibble(sid = 1:3,  
             grade = rep(3, 3),  
             score = as.integer(rnorm(3, 200, 10)))  
  
g4 <- tibble(sid = 9:11,  
             grade = rep(4, 3),  
             score = as.integer(rnorm(3, 200, 10)))
```

g3

```
## # A tibble: 3 x 3  
##       sid grade score  
##   <int> <dbl> <int>  
## 1     1     3   215  
## 2     2     3   203  
## 3     3     3   200
```

g4

```
## # A tibble: 3 x 3  
##       sid grade score  
##   <int> <dbl> <int>  
## 1     9     4   213  
## 2    10     4   191  
## 3    11     4   209
```

bind_rows

- In examples like the previous datasets, we just want to "staple" the rows together.
- We can do so with `bind_rows`.

```
bind_rows(g3, g4)
```

```
## # A tibble: 6 x 3
##   sid grade score
##   <int> <dbl> <int>
## 1     1     3   215
## 2     2     3   203
## 3     3     3   200
## 4     9     4   213
## 5    10     4   191
## 6    11     4   209
```

Optional `.id` argument

- What if we knew the grade, but didn't have a variable in each dataset already?
- Use `.id` to add an index for each dataset

```
bind_rows(g3[, -2], g4[, -2], .id = "dataset")
```

```
## # A tibble: 6 x 3
##   dataset   sid score
##   <chr>   <int> <int>
## 1 1       1    215
## 2 1       2    203
## 3 1       3    200
## 4 2       9    213
## 5 2      10    191
## 6 2      11    209
```

```
bind_rows(g3[ , -2], g4[ , -2], .id = "dataset") %>%  
  mutate(grade = ifelse(dataset == 1, 3, 4))
```

```
## # A tibble: 6 x 4  
##   dataset    sid score grade  
##   <chr>    <int> <int> <dbl>  
## 1 1         1    215     3  
## 2 1         2    203     3  
## 3 1         3    200     3  
## 4 2         9    213     4  
## 5 2        10    191     4  
## 6 2        11    209     4
```

Even better usage

```
bind_rows(g3 = g3[, -2], g4 = g4[, -2], .id = "grade")
```

```
## # A tibble: 6 x 3
##   grade    sid score
##   <chr> <int> <int>
## 1 g3      1    215
## 2 g3      2    203
## 3 g3      3    200
## 4 g4      9    213
## 5 g4     10    191
## 6 g4     11    209
```


What if columns don't match exactly?

Pad with **NA**

```
bind_rows(g3, g4[, -2], .id = "dataset")
```

```
## # A tibble: 6 x 4
##   dataset    sid grade score
##   <chr>    <int> <dbl> <int>
## 1 1         1     3    215
## 2 1         2     3    203
## 3 1         3     3    200
## 4 2         9    NA    213
## 5 2        10    NA    191
## 6 2        11    NA    209
```

Last note – read in a bunch of files

- We'll talk about this a lot more in the next course
- `purrr::map_df` uses `bind_rows` in the background

```
dir.create("tmp")

mtcars %>%
  split(.$cyl) %>%
  walk2(c("tmp/cyl4.csv", "tmp/cyl6.csv", "tmp/cyl8.csv"),
        write_csv)

list.files("tmp")
```

```
## [1] "cyl4.csv" "cyl6.csv" "cyl8.csv"
```

Read in files

Use `purrr::map_df` with the file names Note `fs::dir_ls` is equivalent to `list.files`, but plays nicer with `purrr::map_df`

```
new_mtcars <- map_df(fs::dir_ls("tmp"), rio::import, setclass = 'mtcars',  
                    .id = "file")  
  
new_mtcars %>%  
  select(file, mpg, cyl) %>%  
  slice(1:3)
```

```
## # A tibble: 3 x 3  
##   file          mpg    cyl  
##   <chr>      <dbl> <int>  
## 1 tmp/cyl4.csv  22.8     4  
## 2 tmp/cyl4.csv  24.4     4  
## 3 tmp/cyl4.csv  22.8     4
```

```
unlink("tmp", recursive = TRUE)
```

Joins

(not to be confused with row binding)

Keys

- Uniquely identify rows in a dataset
- Variable(s) in common between two datasets to be joined
- A key can be more than one variable

Types of keys

- Small distinction that you probably won't have to worry about much, but is worth mentioning:
 - **Primary keys:** Uniquely identify observations in their dataset
 - **Foreign keys:** Uniquely identify observations in other datasets.

What's the primary key here?

```
library(rio)
library(here)
ecls <- import(here("data", "ecls-k_samp.sav"),
               setclass = "tbl_df") %>%
  characterize()
ecls
```

```
## # A tibble: 984 x 33
##   child_id teacher_id school_id k_type    school_type sex    ethnic
##   <chr>      <chr>      <chr>    <chr>    <chr>      <chr>  <chr>
## 1 0842021C 0842T02      0842    full-day public    male   BLACK OR AFF
## 2 0905002C 0905T01      0905    full-day private  male   ASIAN
## 3 0150012C 0150T01      0150    full-day private  female BLACK OR AFF
## 4 0556009C 0556T01      0556    full-day private  female HISPANIC, RA
## 5 0089013C 0089T04      0089    full-day public    male   WHITE, NON-H
## 6 1217001C 1217T13      1217    half-day public    female NATIVE HAWAI
## # ... with 978 more rows, and 26 more variables: famtype <chr>, numsibs <d
## #   SES_cat <chr>, age <dbl>, T1RSCALE <dbl>, T1MSCALE <dbl>, T1GSCALE <
## #   T2MSCALE <dbl>, T2GSCALE <dbl>, IRTreadgain <dbl>, IRTmathgain <dbl>
## #   T1ARSLIT <dbl>, T1ARSMAT <dbl>, T1ARSGEN <dbl>, T2ARSLIT <dbl>, T2AR
## #   ARSlitgain <dbl>, ARSmathgain <dbl>, ARSgkgain <dbl>, testdate1 <dat
## #   elapse <dbl>
```

Double-checking

```
ecds %>%  
  count(child_id)
```

```
## # A tibble: 984 x 2  
##   child_id      n  
## * <chr>    <int>  
## 1 0001010C      1  
## 2 0002010C      1  
## 3 0009005C      1  
## 4 0009014C      1  
## 5 0009026C      1  
## 6 0013003C      1  
## # ... with 978 more rows
```

```
ecds %>%  
  count(child_id) %>%  
  filter(n > 1)
```

```
## # A tibble: 0 x 2  
## # ... with 2 variables: child_id <chr>, n <int>
```


What about here?

```
income_ineq <- read_csv(here("data", "incomeInequality_tidy.csv")
print(income_ineq, n = 15)
```

```
## # A tibble: 726 x 6
##   Year Number.thousands realGDPperCap PopulationK percentile income
##   <dbl>         <dbl>         <dbl>         <dbl>         <dbl>    <dbl>
## 1  1947         37237         14117.32        144126         20      14243
## 2  1947         37237         14117.32        144126         40      22984
## 3  1947         37237         14117.32        144126         60      31166
## 4  1947         37237         14117.32        144126         80      44223
## 5  1947         37237         14117.32        144126         50     26764.1
## 6  1947         37237         14117.32        144126         90      41477
## 7  1947         37237         14117.32        144126         95      54172
## 8  1947         37237         14117.32        144126         99     134415
## 9  1947         37237         14117.32        144126        99.5     203001
## 10 1947         37237         14117.32        144126        99.9     479022
## 11 1947         37237         14117.32        144126        99.99  1584506
## 12 1948         38624         14451.94        146631         20      13779
## 13 1948         38624         14451.94        146631         40      22655
## 14 1948         38624         14451.94        146631         60      30248
## 15 1948         38624         14451.94        146631         80      42196
## # ... with 711 more rows
```

```
income_ineq %>%  
  count(Year, percentile) %>%  
  filter(n > 1)
```

```
## # A tibble: 0 x 3
```

```
## # ... with 3 variables: Year <dbl>, percentile <dbl>, n <int>
```

Sometimes there is no key

These tables have an *implicit* id – the row numbers. For example:

```
install.packages("nycflights13")  
library(nycflights13)
```

```
head(flights)
```

```
## # A tibble: 6 x 19  
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr  
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>  
## 1  2013     1     1     517           515           2     830  
## 2  2013     1     1     533           529           4     850  
## 3  2013     1     1     542           540           2     923  
## 4  2013     1     1     544           545          -1    1004  
## 5  2013     1     1     554           600          -6     812  
## 6  2013     1     1     554           558          -4     740  
## # ... with 8 more variables: tailnum <chr>, origin <chr>, dest <chr>, air_  
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights %>%  
  count(year, month, day, flight, tailnum) %>%  
  filter(n > 1)
```

```
## # A tibble: 11 x 6  
##   year month   day flight tailnum     n  
##   <int> <int> <int>   <int> <chr>   <int>  
## 1  2013     2     9     303 <NA>     2  
## 2  2013     2     9     655 <NA>     2  
## 3  2013     2     9    1623 <NA>     2  
## 4  2013     6     8    2269 N487WN     2  
## 5  2013     6    15    2269 N230WN     2  
## 6  2013     6    22    2269 N440LV     2  
## # ... with 5 more rows
```

Create a key

- If there is no key, it's often helpful to add one. These are called *surrogate* keys.

```
flights <- flights %>%  
  rowid_to_column()  
  
flights %>%  
  select(1:3, ncol(flights))
```

```
## # A tibble: 336,776 x 4  
##   rowid  year month time_hour  
##   <int> <int> <int> <dtm>  
## 1      1  2013     1 2013-01-01 05:00:00  
## 2      2  2013     1 2013-01-01 05:00:00  
## 3      3  2013     1 2013-01-01 05:00:00  
## 4      4  2013     1 2013-01-01 05:00:00  
## 5      5  2013     1 2013-01-01 06:00:00  
## 6      6  2013     1 2013-01-01 05:00:00  
## # ... with 336,770 more rows
```

Mutating joins

Mutating joins

- In *tidyverse*, we use `mutate()` to create new variables within a dataset.
- A mutating join works similarly, in that we're adding new variables to the existing dataset through a join.
- Two tables of data joined by a common key

Four types of joins

- **left_join**: Keep all the data in the left dataset, drop any non-matching cases from the right dataset.
- **right_join**: Keep all the data in the right dataset, drop any non-matching cases from the left dataset.
- **inner_join**: Keep only data that matches in both datasets
- **full_join**: Keep all the data in both datasets. This is also sometimes referred to as an *outer* join.

If the keys match exactly in the two tables (datasets), all of these will result in the exact same result.

Using joins to recode

Say you have a dataset like this

```
set.seed(1)
disab_codes <- c("00", "10", "20", "40", "43", "50", "60",
                 "70", "74", "80", "82", "90", "96", "98")
dis_tbl <- tibble(
  sid = 1:200,
  dis_code = sample(disab_codes, 200, replace = TRUE),
  score = as.integer(rnorm(200, 200, 10))
)
head(dis_tbl)
```

```
## # A tibble: 6 x 3
##   sid dis_code score
##   <int> <chr>   <int>
## 1     1 74      190
## 2     2 40      200
## 3     3 60      200
## 4     4 00      183
## 5     5 10      210
## 6     6 96      188
```

Codes

Code	Disability
00	'Not Applicable'
10	'Intellectual Disability'
20	'Hearing Impairment'
40	'Visual Impairment'
43	'Deaf–Blindness'
50	'Communication Disorder'
60	'Emotional Disturbance'
70	'Orthopedic Impairment'
74	'Traumatic Brain Injury'

Code	Disability
80	'Other Health Impairments'
82	'Autism Spectrum Disorder'
90	'Specific Learning Disability'
96	'Developmental Delay 0–2yr'
98	'Developmental Delay 3–4yr'

One method

```
dis_tbl %>%  
  mutate(disability = case_when(  
    dis_code == "10" ~ "Intellectual Disability",  
    dis_code == "20" ~ 'Hearing Impairment',  
    ...,  
    TRUE ~ "Not Applicable"  
  )  
)
```

Joining method

```
dis_code_tbl <- tibble(  
  dis_code = c(  
    "00", "10", "20", "40", "43", "50", "60",  
    "70", "74", "80", "82", "90", "96", "98"  
  ),  
  disability = c(  
    'Not Applicable', 'Intellectual Disability',  
    'Hearing Impairment', 'Visual Impairment',  
    'Deaf-Blindness', 'Communication Disorder',  
    'Emotional Disturbance', 'Orthopedic Impairment',  
    'Traumatic Brain Injury', 'Other Health Impairments',  
    'Autism Spectrum Disorder', 'Specific Learning Disability',  
    'Developmental Delay 0-2yr', 'Developmental Delay 3-4yr'  
  )  
)
```

dis_code_tbl

```
## # A tibble: 14 x 2
##   dis_code disability
##   <chr>      <chr>
## 1 00        Not Applicable
## 2 10        Intellectual Disability
## 3 20        Hearing Impairment
## 4 40        Visual Impairment
## 5 43        Deaf-Blindness
## 6 50        Communication Disorder
## # ... with 8 more rows
```

Join the tables

```
left_join(dis_tbl, dis_code_tbl)
```

```
## Joining, by = "dis_code"
```

```
## # A tibble: 200 x 4
```

```
##       sid dis_code score disability
```

```
##   <int> <chr>      <int> <chr>
```

```
## 1      1 74          190 Traumatic Brain Injury
```

```
## 2      2 40          200 Visual Impairment
```

```
## 3      3 60          200 Emotional Disturbance
```

```
## 4      4 00          183 Not Applicable
```

```
## 5      5 10          210 Intellectual Disability
```

```
## 6      6 96          188 Developmental Delay 0-2yr
```

```
## # ... with 194 more rows
```

Imperfect key
match?

Consider the following

```
gender <- tibble(key = 1:3, male = rbinom(3, 1, .5))  
sped <- tibble(key = c(1, 2, 4), sped = rbinom(3, 1, .5))
```

gender

```
## # A tibble: 3 x 2  
##   key male  
##   <int> <int>  
## 1     1     0  
## 2     2     1  
## 3     3     0
```

sped

```
## # A tibble: 3 x 2  
##   key sped  
##   <dbl> <int>  
## 1     1     0  
## 2     2     1  
## 3     4     0
```


left_join()?

```
left_join(gender, sped)
```

```
## # A tibble: 3 x 3
##   key  male sped
##   <dbl> <int> <int>
## 1     1     0     0
## 2     2     1     1
## 3     3     0    NA
```

right_join()?

```
right_join(gender, sped)
```

```
## # A tibble: 3 x 3
##   key  male sped
##   <dbl> <int> <int>
## 1     1     0     0
## 2     2     1     1
## 3     4    NA     0
```

inner_join()?

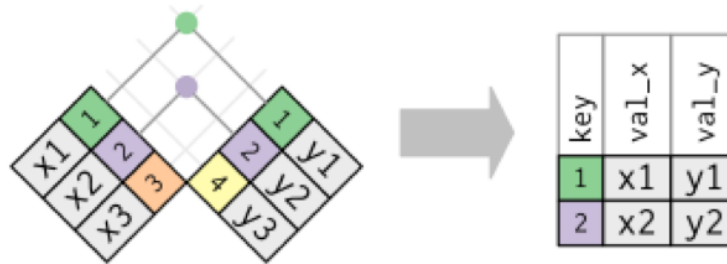
```
inner_join(gender, sped)
```

```
## # A tibble: 2 x 3
##   key  male sped
##   <dbl> <int> <int>
## 1     1     0     0
## 2     2     1     1
```

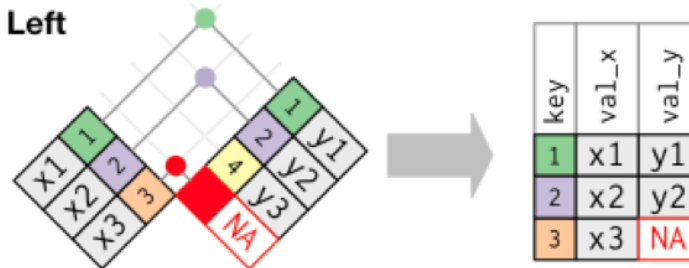
full_join()?

```
full_join(gender, sped)
```

```
## # A tibble: 4 x 3
##   key  male sped
##   <dbl> <int> <int>
## 1     1     0     0
## 2     2     1     1
## 3     3     0    NA
## 4     4    NA     0
```



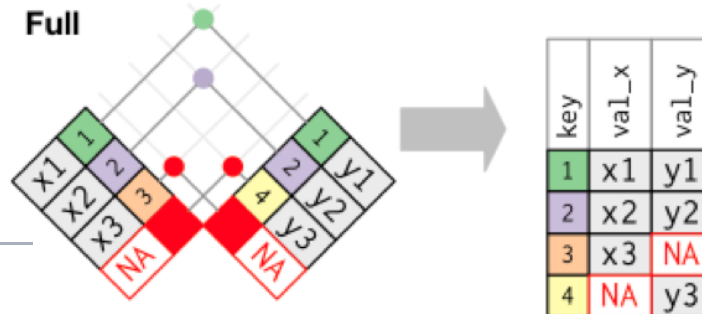
Left



Right



Full

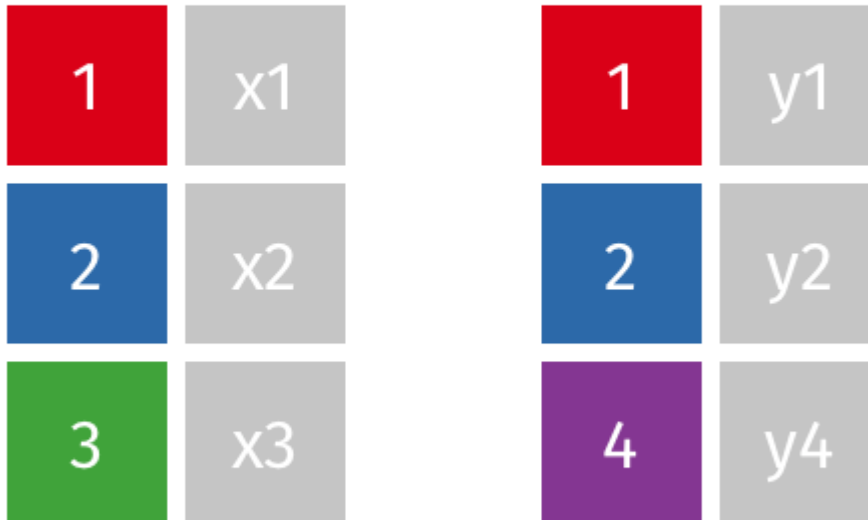


Animations

All of the following animations were created by Garrick Aden-Buie and can be found [here](#)


Animated `left_join()`

`left_join(x, y)`



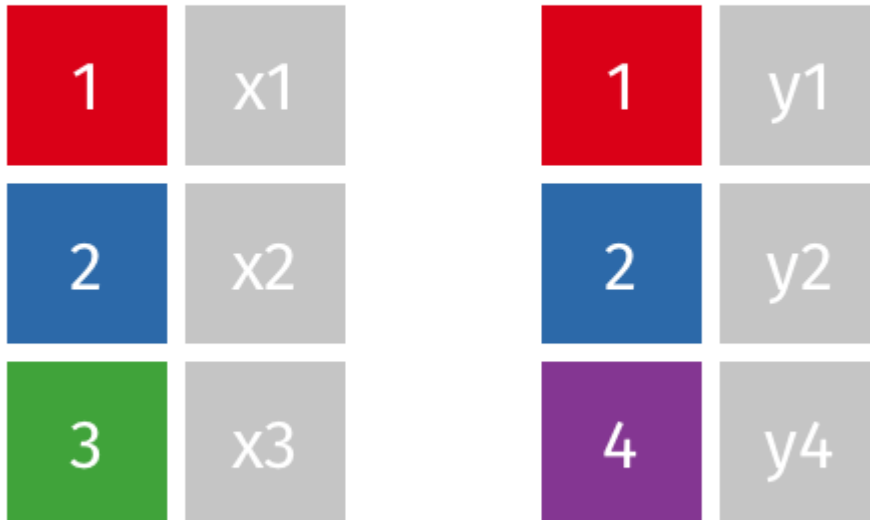
Animated `right_join`

`right_join(x, y)`



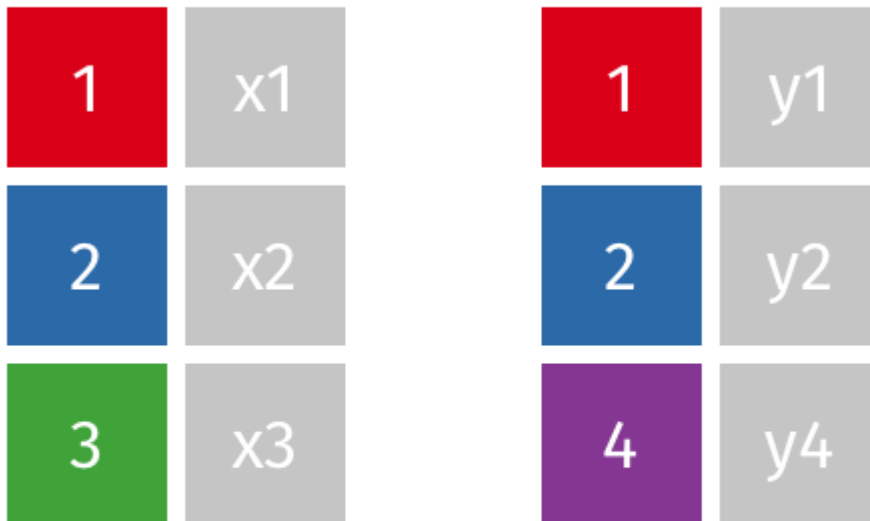
Animated `inner_join`

`inner_join(x, y)`



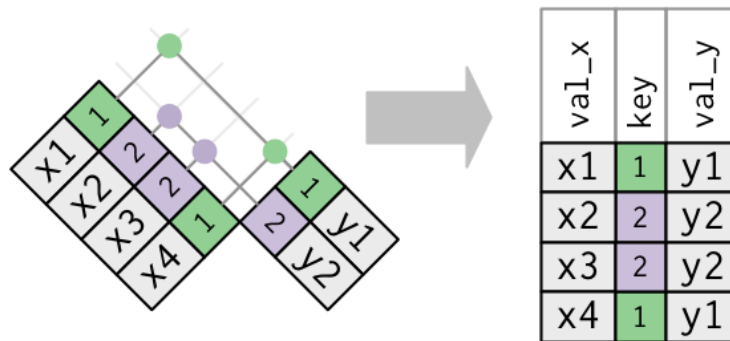
Animated `full_join`

`full_join(x, y)`



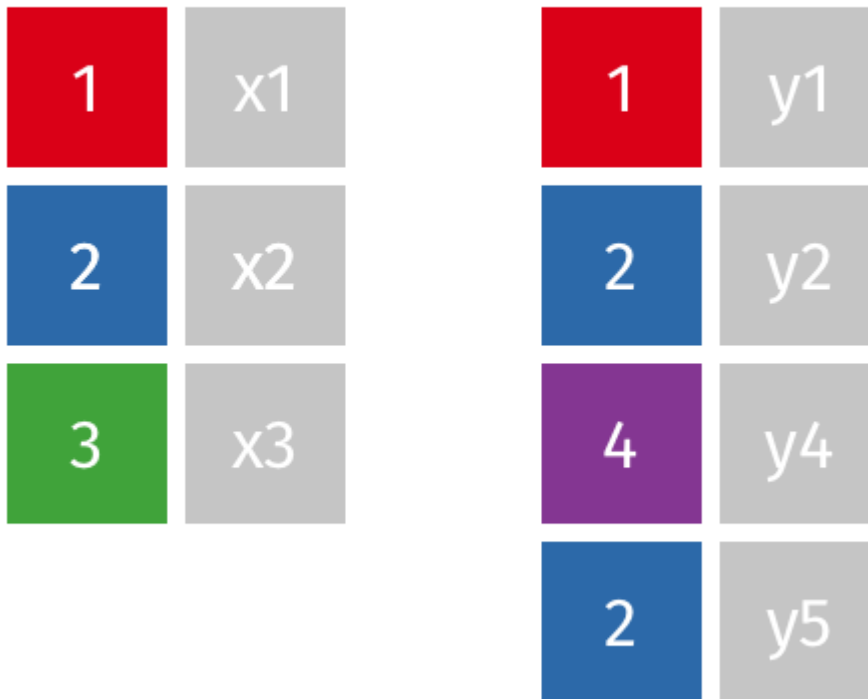
What if the key is not unique?

- Not a problem, as long as they are unique in one of the tables.
 - In this case, it's called a one-to-many join



Animated one-to-many join

`left_join(x, y)`



Example

A dataset with school IDs

```
stu <- tibble(  
  sid = 1:9,  
  scid = c(1, 1, 1, 1, 2, 2, 3, 3, 3),  
  score = c(10, 12, 15, 8, 9, 11, 12, 15, 17)  
)  
stu
```

```
## # A tibble: 9 x 3  
##   sid  scid score  
##   <int> <dbl> <dbl>  
## 1     1     1    10  
## 2     2     1    12  
## 3     3     1    15  
## 4     4     1     8  
## 5     5     2     9  
## 6     6     2    11  
## # ... with 3 more rows
```

A school-level dataset

```
schl <- tibble(  
  scid = 1:3,  
  stu_tch_ratio = c(22.05, 31.14, 24.87),  
  per_pupil_spending = c(15741.08, 11732.24, 13027.88)  
)  
schl
```

```
## # A tibble: 3 x 3  
##   scid stu_tch_ratio per_pupil_spending  
##   <int>      <dbl>          <dbl>  
## 1     1         22.05         15741.08  
## 2     2         31.14         11732.24  
## 3     3         24.87         13027.88
```

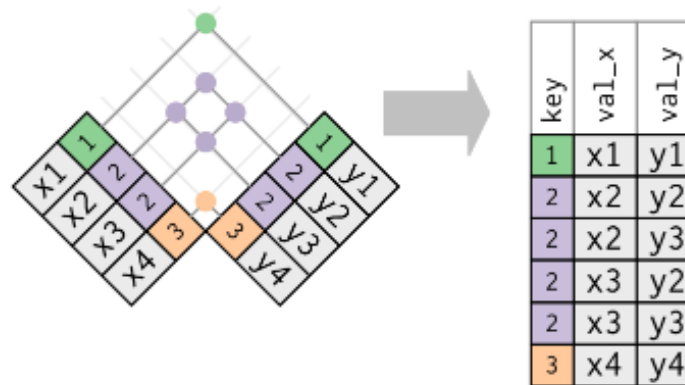
One to many

```
left_join(stu, schl)
```

```
## # A tibble: 9 x 5
##   sid  scid score stu_tch_ratio per_pupil_spending
##   <int> <dbl> <dbl>         <dbl>             <dbl>
## 1     1     1     10          22.05             15741.08
## 2     2     1     12          22.05             15741.08
## 3     3     1     15          22.05             15741.08
## 4     4     1     8           22.05             15741.08
## 5     5     2     9           31.14             11732.24
## 6     6     2    11           31.14             11732.24
## # ... with 3 more rows
```

What if key is not unique to either table?

Generally this is an error Result is probably not going to be what you want (cartesian product).



Example

```
seasonal_means <- tibble(  
  scid = rep(1:3, each = 3),  
  season = rep(c("fall", "winter", "spring"), 3),  
  mean = rnorm(3*3)  
)
```

seasonal_means

```
## # A tibble: 9 x 3  
##   scid season      mean  
##   <int> <chr>      <dbl>  
## 1     1 fall      0.3447951  
## 2     1 winter    1.539648  
## 3     1 spring   -0.3295142  
## 4     2 fall      0.9483894  
## 5     2 winter   -0.4792556  
## 6     2 spring   -1.514887  
## # ... with 3 more rows
```

```
left_join(stu, seasonal_means)
```

```
## # A tibble: 27 x 5
##   sid  scid score season      mean
##   <int> <dbl> <dbl> <chr>    <dbl>
## 1     1     1     10 fall    0.3447951
## 2     1     1     10 winter  1.539648
## 3     1     1     10 spring -0.3295142
## 4     2     1     12 fall    0.3447951
## 5     2     1     12 winter  1.539648
## 6     2     1     12 spring -0.3295142
## # ... with 21 more rows
```


How do we fix this?



In some cases, the solution is obvious. In others, it's not. But **you must have at least one unique key** to join the datasets.

In this case

Move the dataset to wide before joining

Move to wide

```
seasonal_means_wide <- seasonal_means %>%  
  pivot_wider(names_from = "season",  
              values_from = "mean")  
  
seasonal_means_wide
```

```
## # A tibble: 3 x 4  
##   scid      fall      winter      spring  
##   <int>    <dbl>    <dbl>    <dbl>  
## 1     1  0.3447951  1.539648 -0.3295142  
## 2     2  0.9483894 -0.4792556 -1.514887  
## 3     3  0.4345367 -0.5195367 -0.8345590
```

Join

One to many join

```
left_join(stu, seasonal_means_wide)
```

```
## # A tibble: 9 x 6
##   sid  scid score      fall      winter      spring
##   <int> <dbl> <dbl>      <dbl>      <dbl>      <dbl>
## 1     1     1    10  0.3447951  1.539648 -0.3295142
## 2     2     1    12  0.3447951  1.539648 -0.3295142
## 3     3     1    15  0.3447951  1.539648 -0.3295142
## 4     4     1     8  0.3447951  1.539648 -0.3295142
## 5     5     2     9  0.9483894 -0.4792556 -1.514887
## 6     6     2    11  0.9483894 -0.4792556 -1.514887
## # ... with 3 more rows
```

Move longer again?

If we did, we'd be exactly where we were with the first join.

You could make the argument it *might* make sense here

I'd still argue for *this* approach, not the cartesian product approach

More systematic, more predictable, and ultimately less error prone

Another example

- Often you want to add summary info to your dataset.
- You can do this easily with by piping arguments

ECLS-K reminder

```
ecls
```

```
## # A tibble: 984 x 33
##   child_id teacher_id school_id k_type    school_type sex    ethnic
##   <chr>      <chr>      <chr>    <chr>    <chr>      <chr>  <chr>
## 1 0842021C 0842T02      0842    full-day public    male   BLACK OR AFF
## 2 0905002C 0905T01      0905    full-day private   male   ASIAN
## 3 0150012C 0150T01      0150    full-day private   female BLACK OR AFF
## 4 0556009C 0556T01      0556    full-day private   female HISPANIC, RA
## 5 0089013C 0089T04      0089    full-day public    male   WHITE, NON-H
## 6 1217001C 1217T13      1217    half-day public    female NATIVE HAWAI
## # ... with 978 more rows, and 26 more variables: famtype <chr>, numsibs <dbl>,
## #   SES_cat <chr>, age <dbl>, T1RSCALE <dbl>, T1MSCALE <dbl>, T1GSCALE <dbl>,
## #   T2MSCALE <dbl>, T2GSCALE <dbl>, IRTreadgain <dbl>, IRTmathgain <dbl>,
## #   T1ARSLIT <dbl>, T1ARSMAT <dbl>, T1ARSGEN <dbl>, T2ARSLIT <dbl>, T2AR
## #   ARSlitgain <dbl>, ARSmathgain <dbl>, ARSgkgain <dbl>, testdate1 <date>,
## #   elapse <dbl>
```

Compute group means

```
ecls %>%  
  group_by(school_id) %>%  
  summarize(sch_pre_math = mean(T1MSCALE))
```

```
## # A tibble: 515 x 2  
##   school_id sch_pre_math  
## *   <chr>         <dbl>  
## 1 0001             20.45800  
## 2 0002             14.977  
## 3 0009             18.82  
## 4 0013             42.321  
## 5 0016             17.55100  
## 6 0022             17.8465  
## # ... with 509 more rows
```

Join right within pipeline

```
ecls %>%  
  group_by(school_id) %>%  
  summarize(sch_pre_math = mean(T1MSCALE)) %>%  
  left_join(ecls) %>%  
  select(school_id:k_type) # Just for space
```

```
## # A tibble: 984 x 5  
##   school_id sch_pre_math child_id teacher_id k_type  
##   <chr>      <dbl> <chr>    <chr>    <chr>  
## 1 0001      20.45800 0001010C 0001T01  full-day  
## 2 0002      14.977   0002010C 0002T01  half-day  
## 3 0009      18.82    0009026C 0009T01  half-day  
## 4 0009      18.82    0009014C 0009T02  half-day  
## 5 0009      18.82    0009005C 0009T01  half-day  
## 6 0013      42.321   0013003C 0013T01  full-day  
## # ... with 978 more rows
```

Default join behavior

By default, the `*_join` functions will use all columns with common names as keys.

```
flights2 <- flights %>%  
  select(year:day, hour, origin, dest, tailnum, carrier)  
flights2[1:2, ]
```

```
## # A tibble: 2 x 8  
##   year month   day hour origin dest tailnum carrier  
##   <int> <int> <int> <dbl> <chr> <chr> <chr>    <chr>  
## 1  2013     1     1     5 EWR   IAH   N14228   UA  
## 2  2013     1     1     5 LGA   IAH   N24211   UA
```

```
weather[1:2, ]
```

```
## # A tibble: 2 x 15  
##   origin year month   day hour temp dewp humid wind_dir wind_speed  
##   <chr>   <int> <int> <int> <int> <dbl> <dbl> <dbl>    <dbl>    <dbl>  
## 1 EWR    2013     1     1     1 39.02 26.06 59.37     270    10.35702  
## 2 EWR    2013     1     1     2 39.02 26.96 61.63     250     8.05546  
## # ... with 1 more variable: time_hour <dtm>
```



```
left_join(flights2, weather)
```

```
## Joining, by = c("year", "month", "day", "hour", "origin")
```

```
## # A tibble: 336,776 x 18
```

```
##   year month   day hour origin dest tailnum carrier temp dewp humid
```

```
##   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr> <dbl> <dbl> <dbl>
```

```
## 1  2013     1     1     5 EWR   IAH   N14228  UA      39.02 28.04 64.43
```

```
## 2  2013     1     1     5 LGA   IAH   N24211  UA      39.92 24.98 54.81
```

```
## 3  2013     1     1     5 JFK   MIA   N619AA  AA      39.02 26.96 61.63
```

```
## 4  2013     1     1     5 JFK   BQN   N804JB  B6      39.02 26.96 61.63
```

```
## 5  2013     1     1     6 LGA   ATL   N668DN  DL      39.92 24.98 54.81
```

```
## 6  2013     1     1     5 EWR   ORD   N39463  UA      39.02 28.04 64.43
```

```
## # ... with 336,770 more rows, and 4 more variables: precip <dbl>, pressure
```

```
## #   time_hour <dtm>
```

Use only some vars?

If we were joining *flights2* and *planes*, we would not want to use the **year** variable in the join, because **it means different things in each dataset.**

```
head(planes)
```

```
## # A tibble: 6 x 9
##   tailnum  year type      manufacturer      model      engine
##   <chr>    <int> <chr>      <chr>          <chr>      <chr>
## 1 N10156   2004 Fixed wing multi engine EMBRAER      EMB-145XR
## 2 N102UW   1998 Fixed wing multi engine AIRBUS INDUSTRIE A320-214
## 3 N103US   1999 Fixed wing multi engine AIRBUS INDUSTRIE A320-214
## 4 N104UW   1999 Fixed wing multi engine AIRBUS INDUSTRIE A320-214
## 5 N10575   2002 Fixed wing multi engine EMBRAER      EMB-145LR
## 6 N105UW   1999 Fixed wing multi engine AIRBUS INDUSTRIE A320-214
```

How?

Specify the variables with **by**

```
left_join(flights2, planes, by = "tailnum")
```

```
## # A tibble: 336,776 x 16
##   year.x month   day hour origin dest tailnum carrier year.y type
##   <int> <int> <int> <dbl> <chr> <chr> <chr>   <chr>   <int> <chr>
## 1  2013     1     1     5 EWR   IAH   N14228  UA       1999 Fixed wing
## 2  2013     1     1     5 LGA   IAH   N24211  UA       1998 Fixed wing
## 3  2013     1     1     5 JFK   MIA   N619AA  AA       1990 Fixed wing
## 4  2013     1     1     5 JFK   BQN   N804JB  B6       2012 Fixed wing
## 5  2013     1     1     6 LGA   ATL   N668DN  DL       1991 Fixed wing
## 6  2013     1     1     5 EWR   ORD   N39463  UA       2012 Fixed wing
## # ... with 336,770 more rows, and 5 more variables: model <chr>, engines <int>,
## #   speed <int>, engine <chr>
```

Mismatched names?

- What if you had data to merge like this?

```
names(schl)[1] <- "school_id"  
schl
```

```
## # A tibble: 3 x 3  
##   school_id stu_tch_ratio per_pupil_spending  
##   <int>      <dbl>      <dbl>  
## 1         1        22.05        15741.08  
## 2         2        31.14        11732.24  
## 3         3        24.87        13027.88
```

```
stu
```

```
## # A tibble: 9 x 3  
##   sid  scid score  
##   <int> <dbl> <dbl>  
## 1     1     1    10  
## 2     2     1    12  
## 3     3     1    15  
## 4     4     1     8  
## 5     5     2     9  
## 6     6     2    11  
## # ... with 3 more rows
```

Join w/mismatched names

```
left_join(stu, schl, by = c("scid" = "school_id"))
```

```
## # A tibble: 9 x 5
##   sid  scid score stu_tch_ratio per_pupil_spending
##   <int> <dbl> <dbl>         <dbl>             <dbl>
## 1     1     1    10          22.05             15741.08
## 2     2     1    12          22.05             15741.08
## 3     3     1    15          22.05             15741.08
## 4     4     1     8          22.05             15741.08
## 5     5     2     9          31.14             11732.24
## 6     6     2    11          31.14             11732.24
## # ... with 3 more rows
```

filtering joins

Filtering joins

- `semi_join()` works just like `left_join` or `inner_join` but you don't actually add the variables.
- Let's filter classrooms with extremely high math pretest average scores.

First, calculate averages

```
av_pre_mth <- ecls %>%  
  mutate(cut_high = mean(T1MSCALE) + 3*sd(T1MSCALE)) %>%  
  group_by(teacher_id, k_type) %>%  
  summarize(av_pre_mth = mean(T1MSCALE),  
            cut_high = unique(cut_high))  
av_pre_mth
```

```
## # A tibble: 707 x 4  
## # Groups:   teacher_id [707]  
##   teacher_id k_type    av_pre_mth cut_high  
##   <chr>      <chr>      <dbl>    <dbl>  
## 1 0001T01    full-day    20.45800 42.62333  
## 2 0002T01    half-day    14.977    42.62333  
## 3 0009T01    half-day    17.6475   42.62333  
## 4 0009T02    half-day    21.165    42.62333  
## 5 0013T01    full-day    42.321    42.62333  
## 6 0016T01    half-day    17.55100 42.62333  
## # ... with 701 more rows
```


Next, filter for means 3 standard deviations above the mean.

```
extr_high <- av_pre_mth %>%  
  ungroup() %>%  
  filter(av_pre_mth > cut_high)  
extr_high
```

```
## # A tibble: 3 x 4  
##   teacher_id k_type    av_pre_mth cut_high  
##   <chr>      <chr>      <dbl>      <dbl>  
## 1 0078T04    half-day    45.75      42.62333  
## 2 0663T01    full-day    42.8455    42.62333  
## 3 0944T03    half-day    45.371     42.62333
```

Finally, use `semi_join` to show the full data for these cases

```
semi_join(ecls, extr_high)
```

```
## # A tibble: 4 x 33
##   child_id teacher_id school_id k_type    school_type sex    ethnic
##   <chr>      <chr>      <chr>    <chr>    <chr>      <chr>  <chr>
## 1 0944017C 0944T03      0944    half-day private    female WHITE, NON-H
## 2 0663006C 0663T01      0663    full-day private    male   WHITE, NON-H
## 3 0663012C 0663T01      0663    full-day private    female WHITE, NON-H
## 4 0078020C 0078T04      0078    half-day public     female WHITE, NON-H
## # ... with 26 more variables: famtype <chr>, numsibs <dbl>, SES_cont <dbl>
## #   T1RSCALE <dbl>, T1MSCALE <dbl>, T1GSCALE <dbl>, T2RSCALE <dbl>, T2MS
## #   IRTreadgain <dbl>, IRTmathgain <dbl>, IRTgkgain <dbl>, T1ARSLIT <dbl>
## #   T1ARSGEN <dbl>, T2ARSLIT <dbl>, T2ARSMAT <dbl>, T2ARSGEN <dbl>, ARSL
## #   ARSmathgain <dbl>, ARSgkgain <dbl>, testdate1 <date>, testdate2 <date>
```

Filtering joins

`anti_join()` does the opposite of `semi_join`, keeping any rows that do **not** match.

```
nrow(ecls)
```

```
## [1] 984
```

```
extr_low_ecls <- anti_join(ecls, extr_high)  
nrow(extr_low_ecls)
```

```
## [1] 980
```

Why is this so beneficial?

- Sometimes the boolean logic for `filter` can be overly complicated.
- Instead, create a data frame that has only the groups you want, and `semi_join` it with your original data
- Alternatively, create a data frame that has all but the values you want.

Stop words

One more quick example

This one is probs more realistic

Jane Austen Books

```
# install.packages(c("tidytext", "janeaustenr"))  
library(tidytext)  
library(janeaustenr)  
austen_books()
```

```
## # A tibble: 73,422 x 2  
##   text                                book  
## * <chr>                            <fct>  
## 1 "SENSE AND SENSIBILITY" Sense & Sensibility  
## 2 ""                               Sense & Sensibility  
## 3 "by Jane Austen"               Sense & Sensibility  
## 4 ""                               Sense & Sensibility  
## 5 "(1811)"                       Sense & Sensibility  
## 6 ""                               Sense & Sensibility  
## # ... with 73,416 more rows
```

Get words

```
austen_books() %>%  
  unnest_tokens(word, text)
```

```
## # A tibble: 725,055 x 2  
##   book          word  
##   <fct>        <chr>  
## 1 Sense & Sensibility sense  
## 2 Sense & Sensibility and  
## 3 Sense & Sensibility sensibility  
## 4 Sense & Sensibility by  
## 5 Sense & Sensibility jane  
## 6 Sense & Sensibility austen  
## # ... with 725,049 more rows
```


Count words

```
austen_books() %>%  
  unnest_tokens(word, text) %>%  
  count(word, sort = TRUE)
```

```
## # A tibble: 14,520 x 2  
##   word      n  
##   <chr> <int>  
## 1 the    26351  
## 2 to     24044  
## 3 and    22515  
## 4 of     21178  
## 5 a      13408  
## 6 her    13055  
## # ... with 14,514 more rows
```


Plot top 15 words

```
austen_books() %>%  
  unnest_tokens(word, text) %>%  
  count(word, sort = TRUE) %>%  
  mutate(word = fct_reorder(word, n)) %>%  
  slice(1:15) %>%  
  ggplot(aes(word, n)) +  
  geom_col() +  
  coord_flip()
```



Stop words

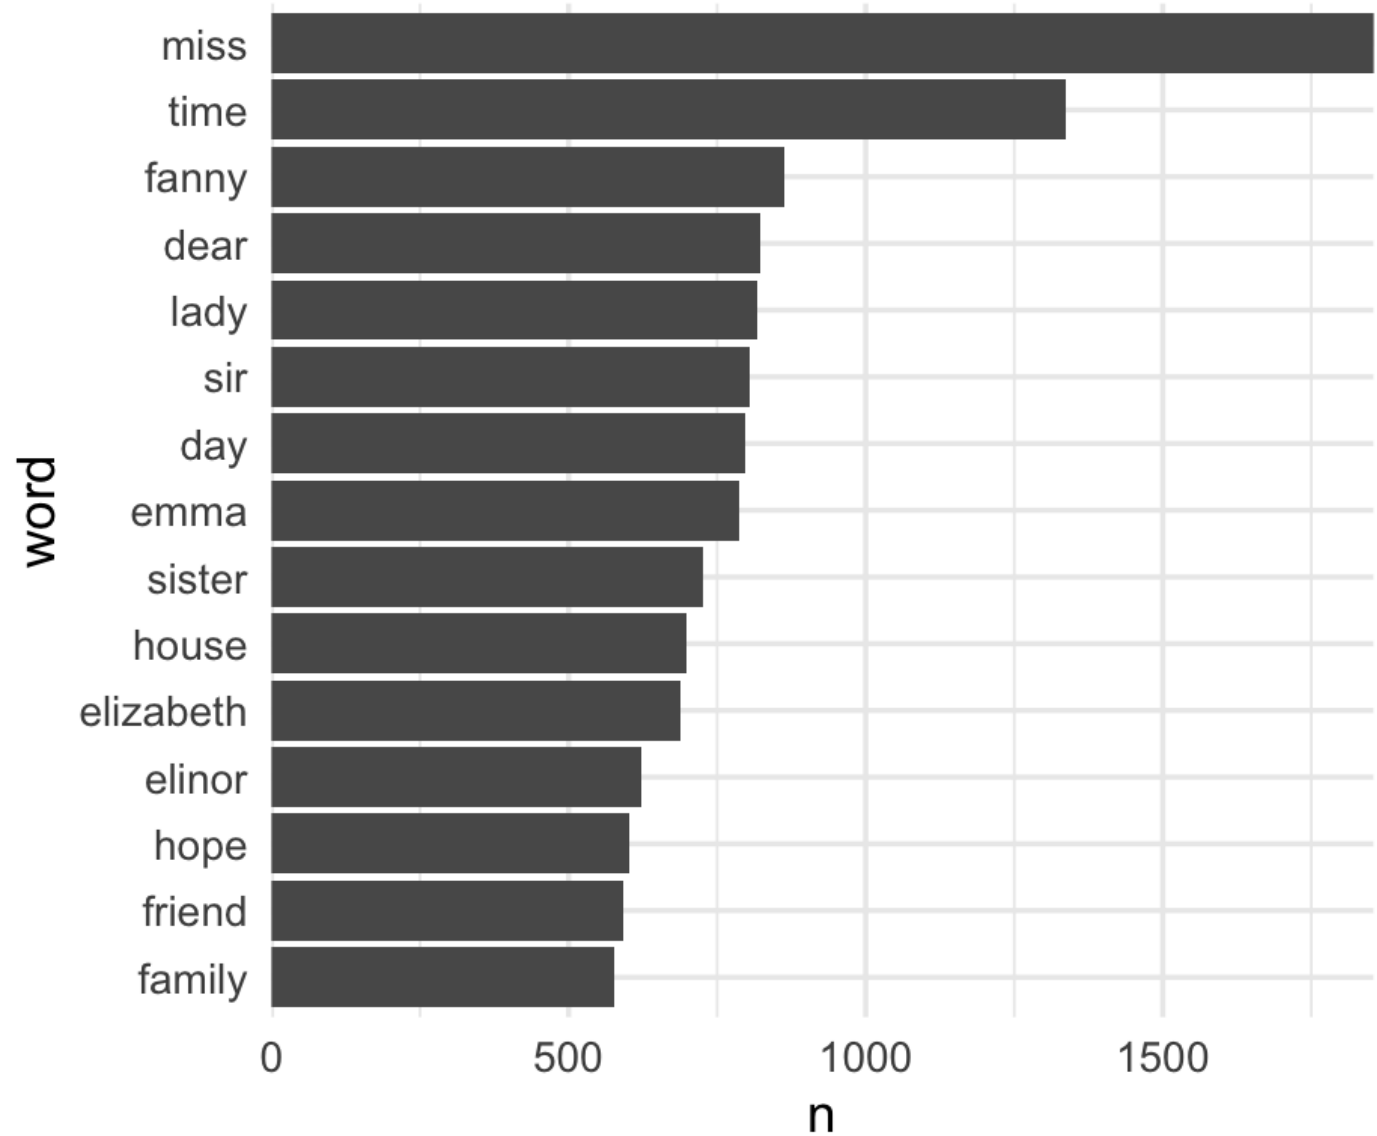
stop_words

```
## # A tibble: 1,149 x 2
##   word      lexicon
##   <chr>    <chr>
## 1 a       SMART
## 2 a's     SMART
## 3 able    SMART
## 4 about   SMART
## 5 above   SMART
## 6 according SMART
## # ... with 1,143 more rows
```

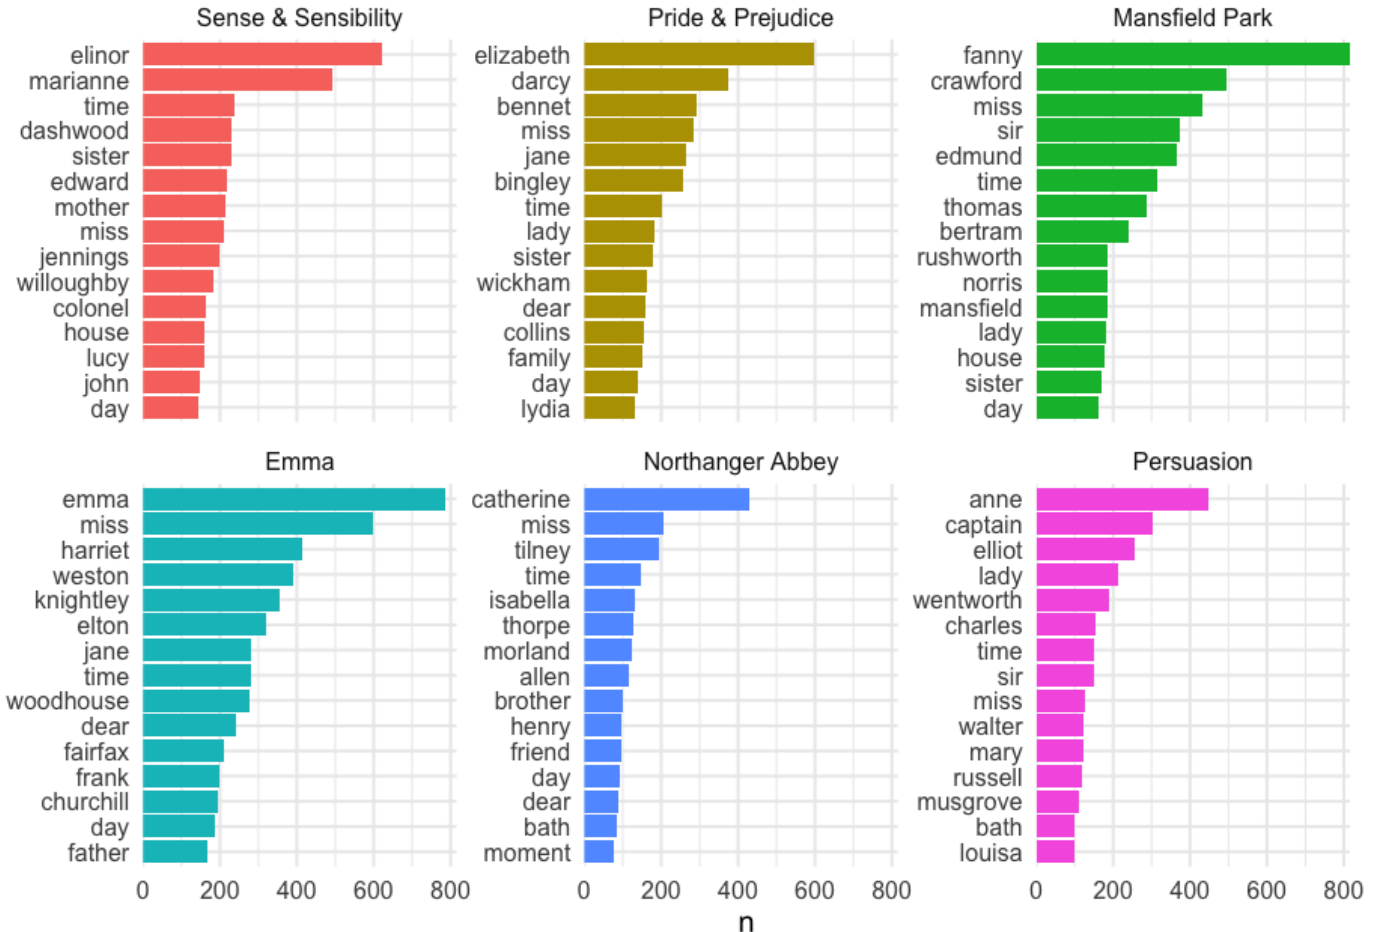
Remove stop words

```
austen_books() %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words) %>%  
  count(word, sort = TRUE)
```

```
## # A tibble: 13,914 x 2  
##   word      n  
##   <chr> <int>  
## 1 miss    1855  
## 2 time    1337  
## 3 fanny    862  
## 4 dear     822  
## 5 lady     817  
## 6 sir      806  
## # ... with 13,908 more rows
```



By book



Wrapping up

- Homework 1 assigned today
 - Be careful about keys. Likely to be rather tricky.
- Next time: Visual perception