Refining your plots

Daniel Anderson

Week 6, Class 1

Reviewing Lab 2

Data viz in Wild

Raleigh

Maggie

Ann–Marie and Murat on deck

Agenda

- Axes and aspect ratios
- Annotations
- Themes (a little bit)

What we won't get to

Each of the following are pretty fundamental to good data viz, but we won't have time to go over them today. Please make sure to read the corresponding chapters:

- Handling high data density (lots of overlapping points)
- Compound figures
  - See {patchwork} and {cowplot}
- Exporting figures

Learning Objectives

- Understand how to make a wide variety of tweaks to ggplot to essentially make it look however you want it to.
- Understand common modifications to plots to make them more clear and reduce cognitive load

Axes

- Cartesian coordinates – what we generally use

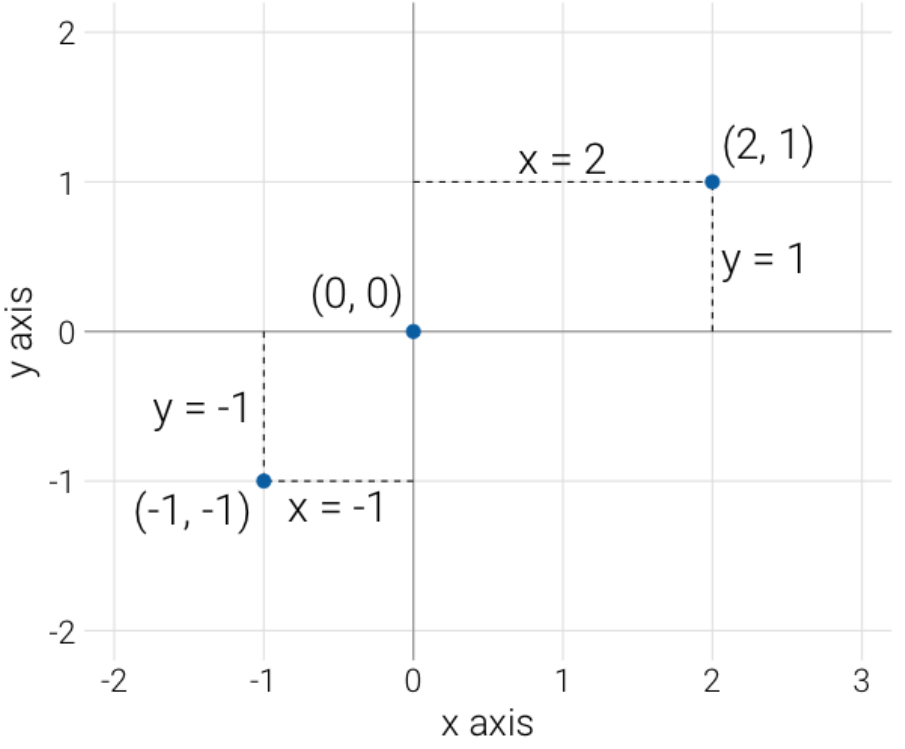

Different units

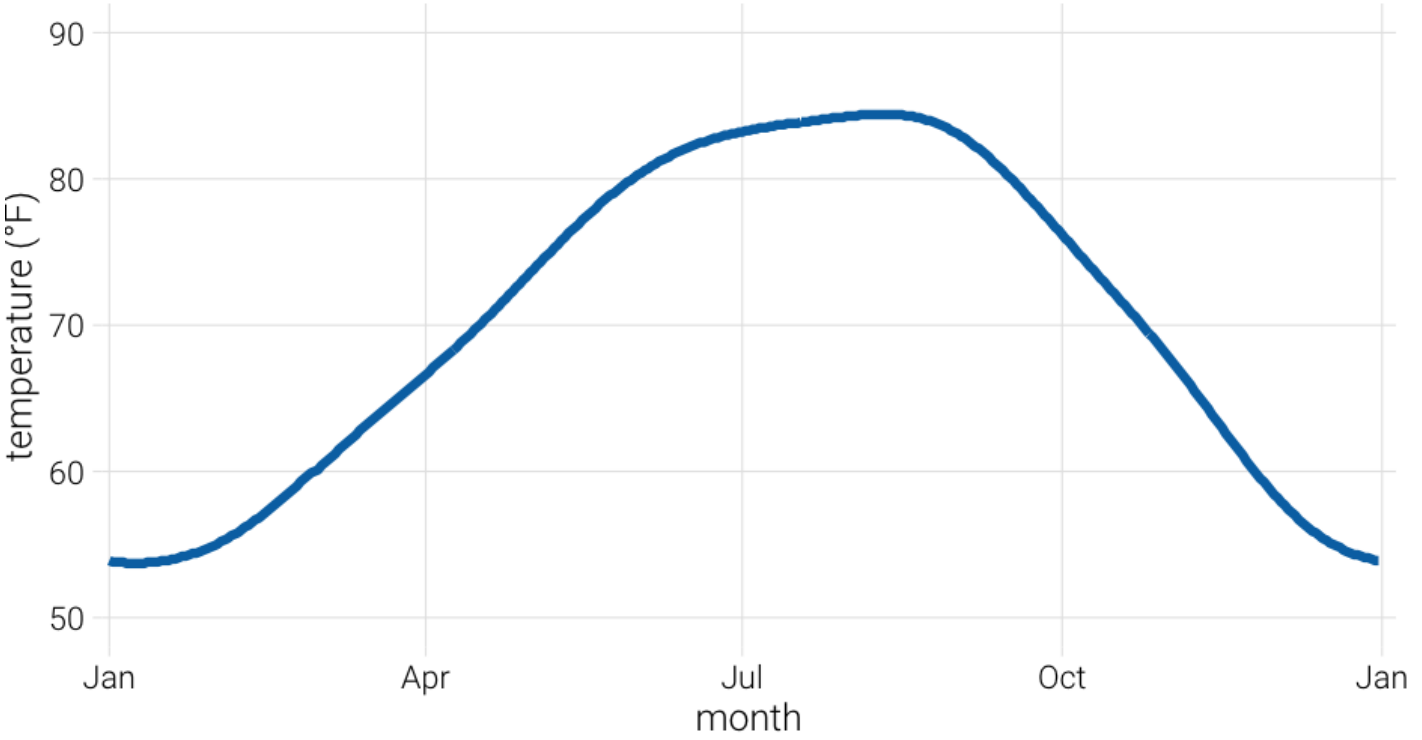

Aspect ratio

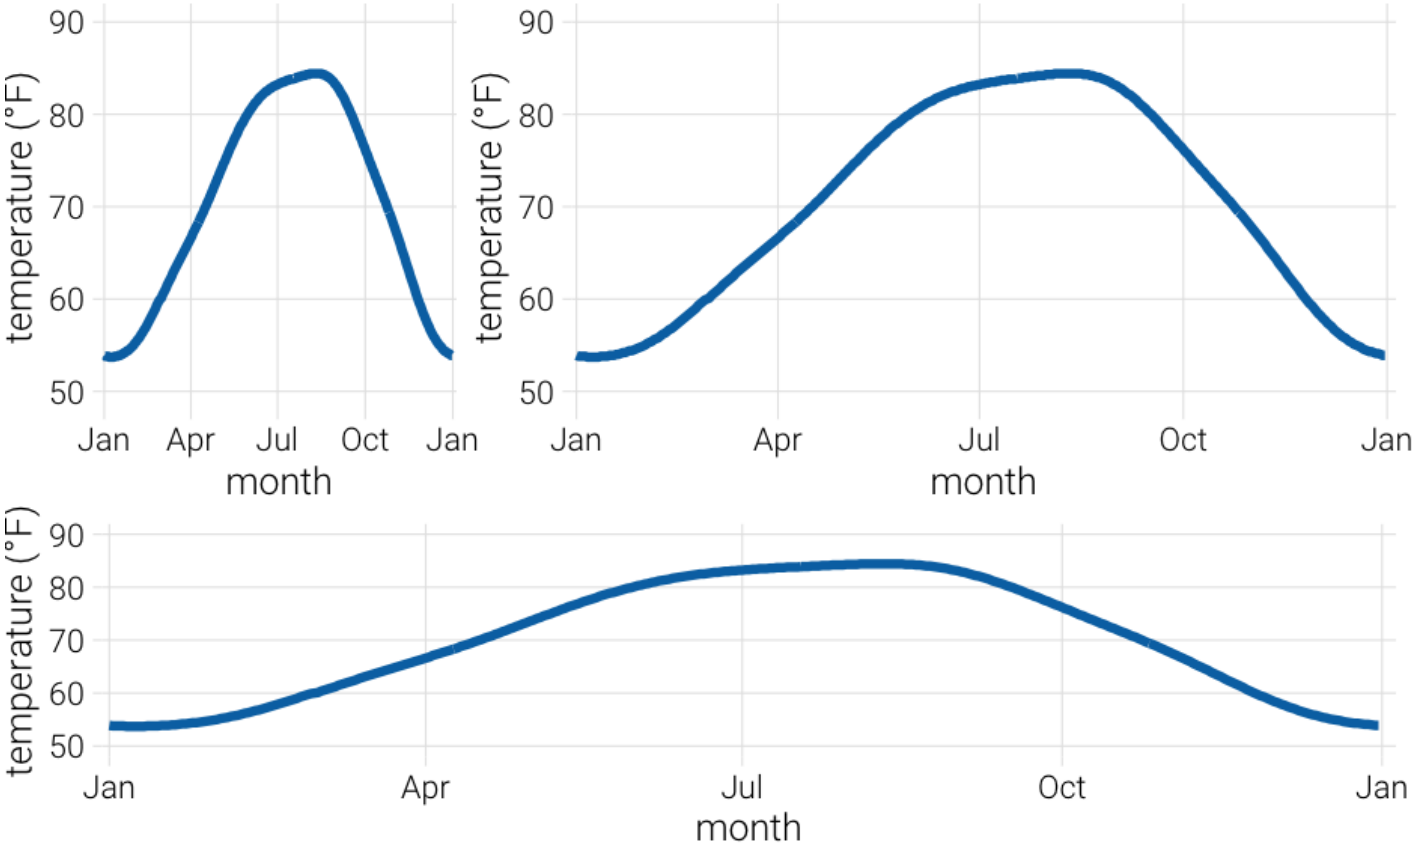

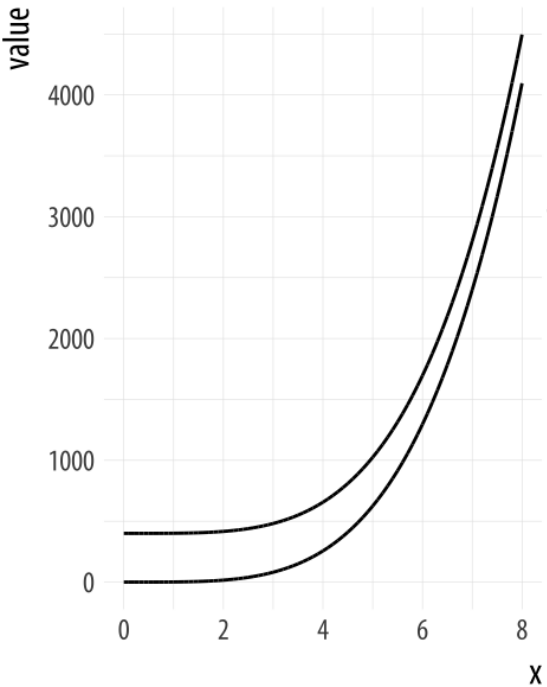
value
4000
3000
2000
1000
0
0
2
4
6
8
x

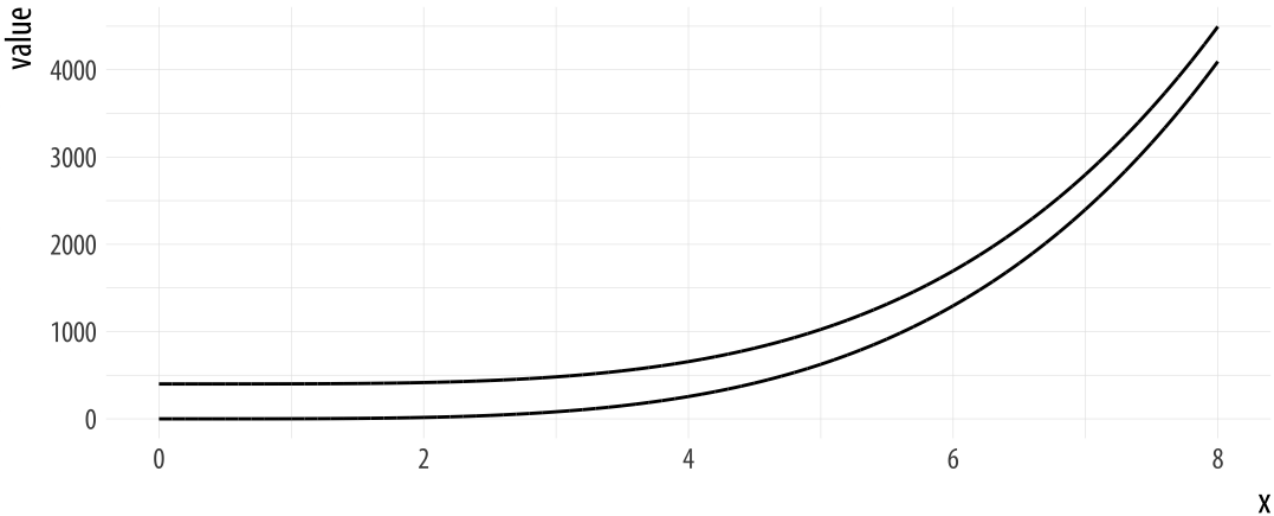
value
4000
3000
2000
1000
0
0
2
4
6
8
x

Same scales

Use `coord_fixed()`

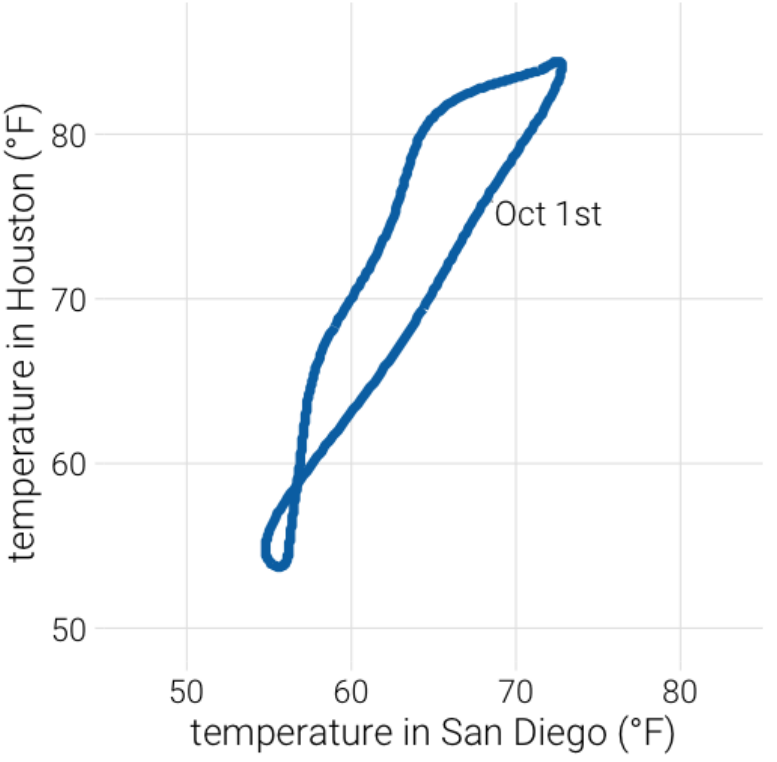

Changing aspect ratio

- Explore how your plot will look in its final size
- No hard/fast rules (if on different scales)
- Not even really rules of thumb
- Keep visual perception in mind
- Try your best to be truthful – show the trend/relation, but don't exaggerate/hide it

Handy function

(from an apparently deleted tweet from @tjmahr)

> here's my favorite helper #rstats function. preview ggsave() output
>
> ggpreview <- function (..., device = "png") {
> fname <- tempfile(fileext = paste0(".", device))
> ggplot2::ggsave(filename = fname, device = device, ...)
> system2("open", fname)
> invisible(NULL)
> }
>
> — tj mahr 🍕🍍 (@tjmahr)

Gist

(side note: gists are a good way to share things)

- See the full code/example here
- Let's take 5 minutes to play around:
  - Create a plot (could even be the example in the gist)
  - Try different aspect ratios by changing the width/length

05:00

Scale transformations

Raw scale

```
library(gapminder)
ggplot(gapminder, aes(year, gdpPercap)) +
  geom_line(aes(group = country),
            color = "gray70")
```

Log10 scale

```
ggplot(gapminder, aes(year, gdpPercap)) +
  geom_line(aes(group = country),
            color = "gray70") +
  scale_y_log10(labels = scales::dollar)
```

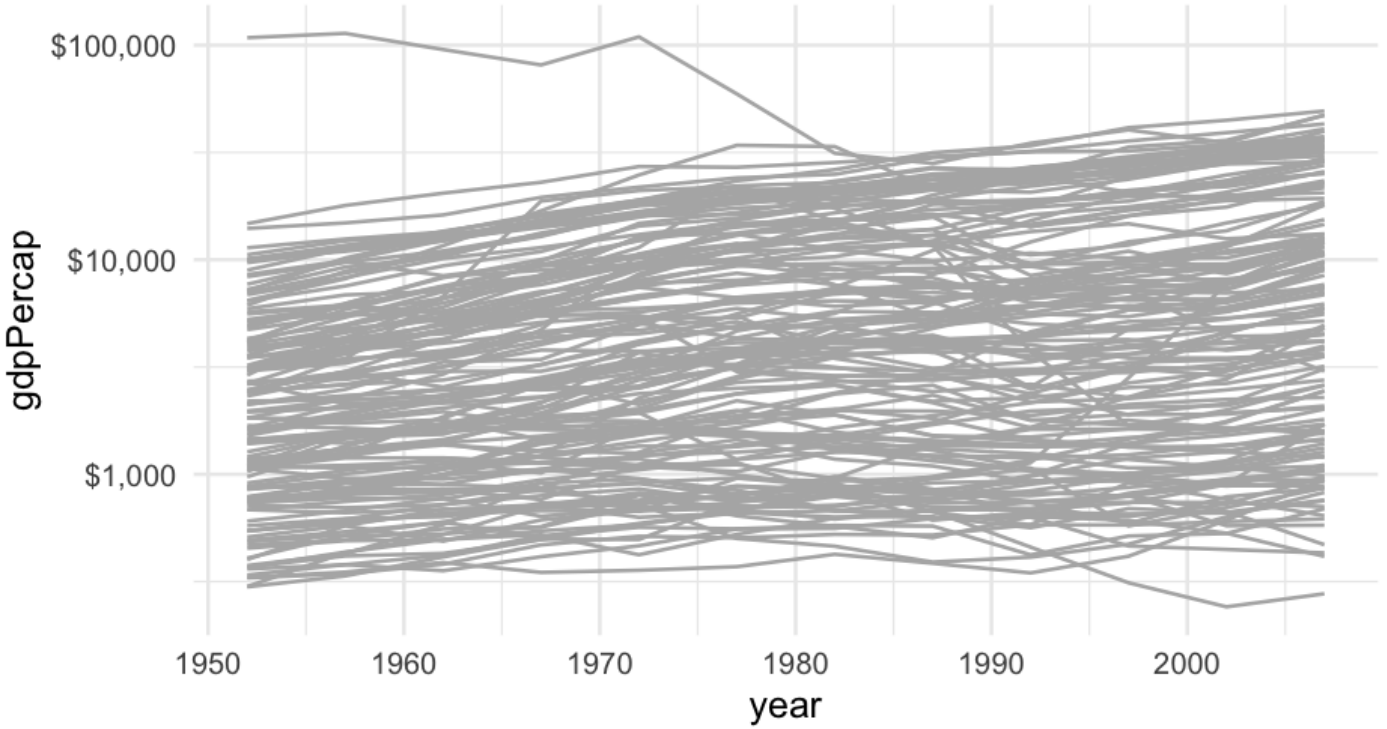

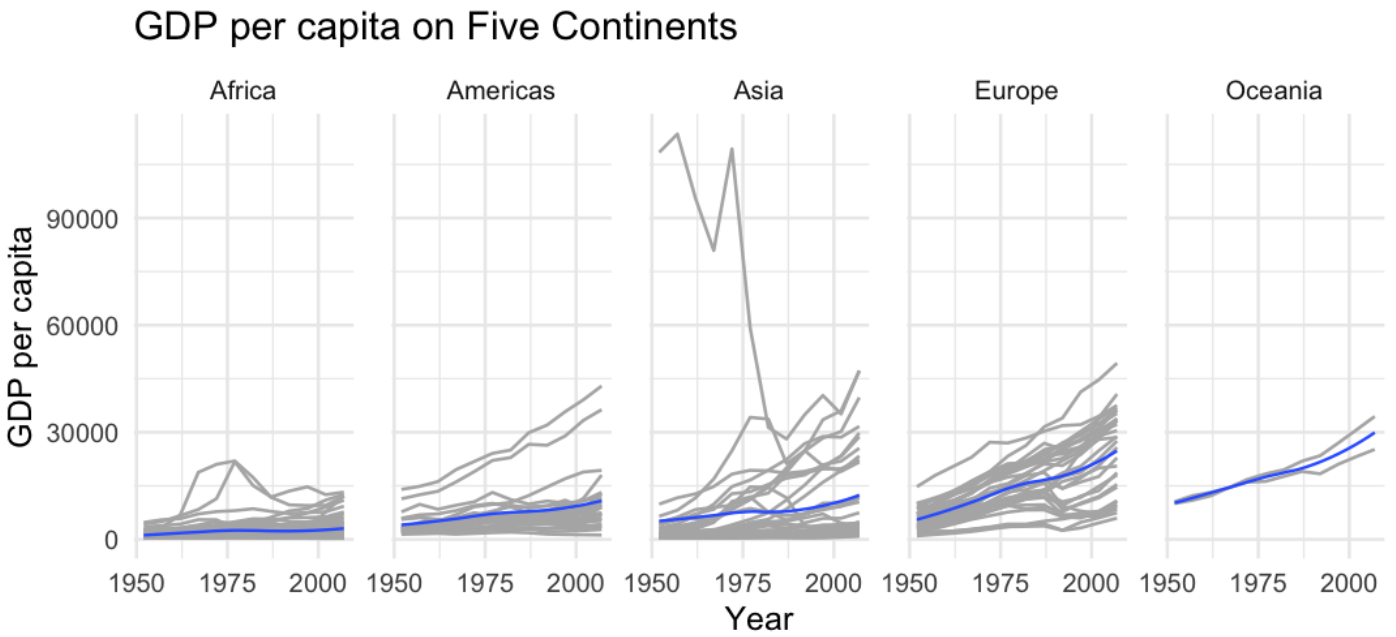
GDP per capita on Five Continents
Africa
Americas
Asia
Europe
Oceania
GDP per capita
90000
60000
30000
0
1950
1975
2000
Year

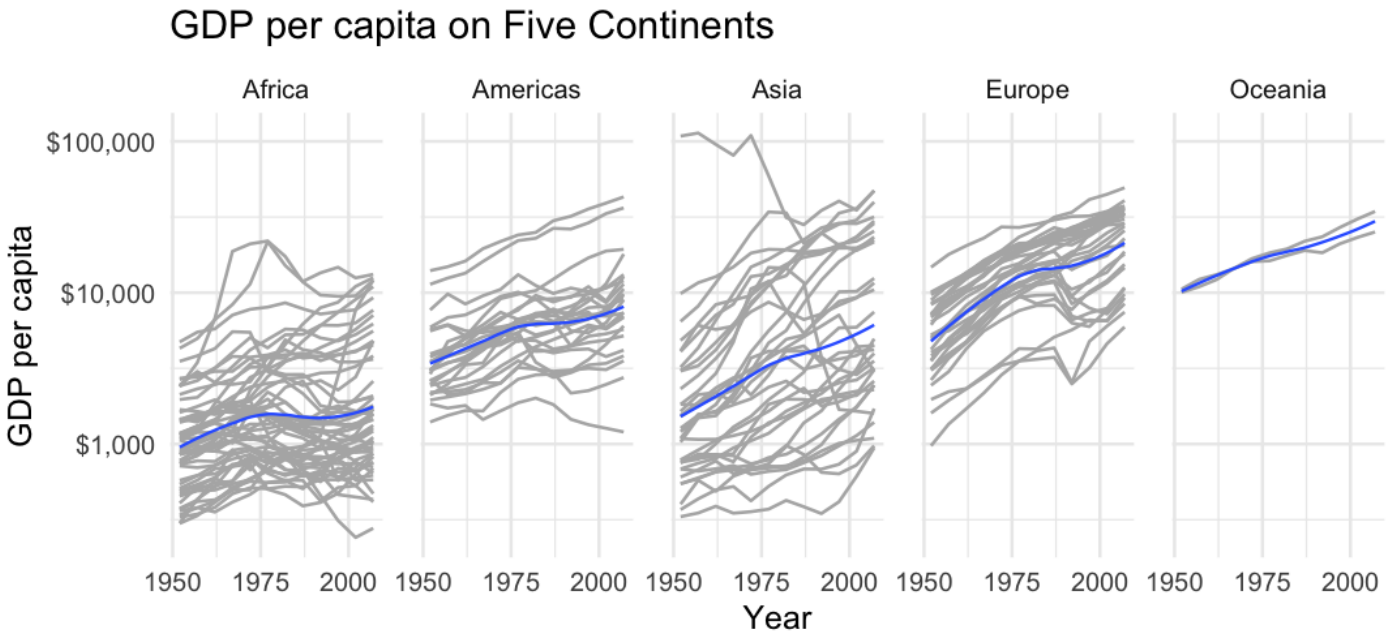
GDP per capita on Five Continents
Africa
Americas
Asia
Europe
Oceania
$100,000
$10,000
$1,000
GDP per capita
1950
1975
2000
Year

Scales

```
d <- tibble(x = c(1, 3.16, 10, 31.6, 100),
            log_x = log10(x))

ggplot(d, aes(x, 1)) +
  geom_point(color = "#0072B2")

ggplot(d, aes(x, 1)) +
  geom_point(color = "#0072B2") +
  scale_x_log10()

ggplot(d, aes(log_x, 1)) +
  geom_point(color = "#0072B2")
```

Scales

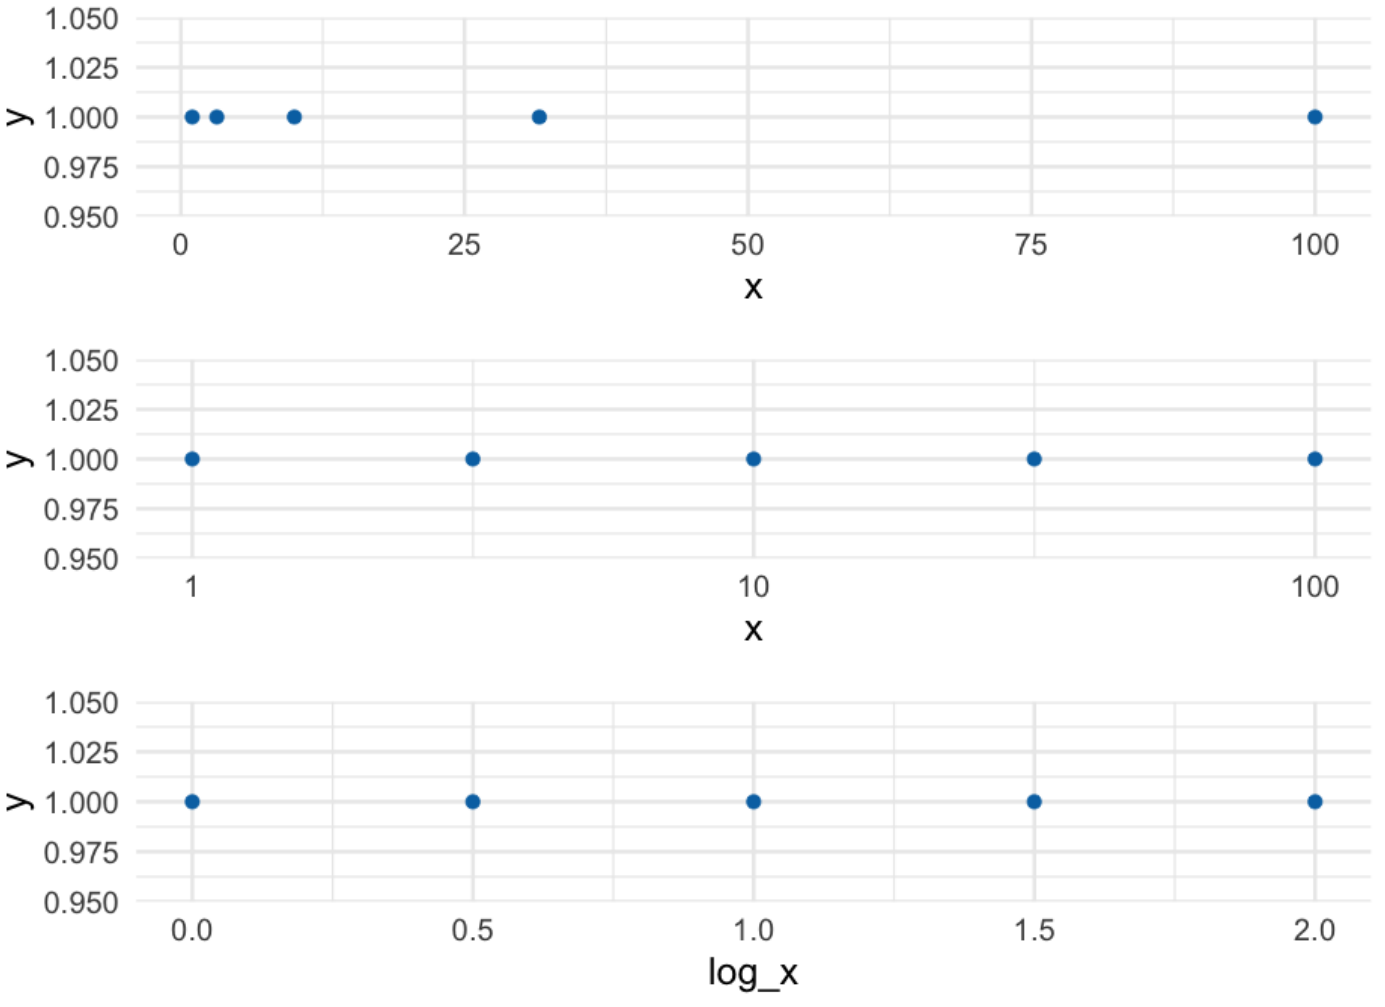
1.050
1.025
y 1.000
0.975
0.950
0
25
50
75
100
x
1.050
1.025
y 1.000
0.975
0.950
1
10
100
x
1.050
1.025
y 1.000
0.975
0.950
0.0
0.5
1.0
1.5
2.0
log_x

Don't transform twice

```
ggplot(d, aes(log_x, 1)) +
  geom_point(color = "#0072B2") +
  scale_x_log10() +
  xlim(-0.2, 2.5)
```

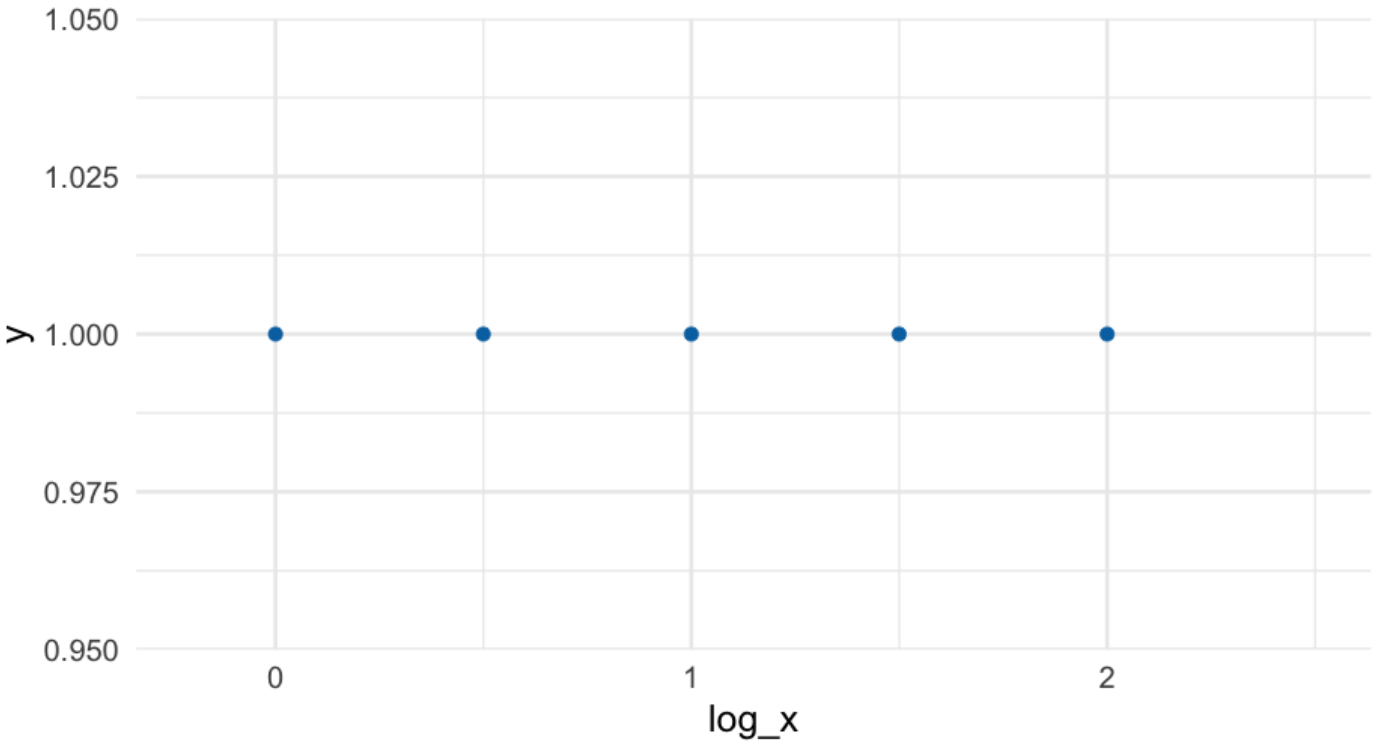

Careful with labeling

- Has the scale or the data been log transformed?
- Specify the base

```
library(ggtext)
ggplot(d, aes(log_x, 1)) +
  geom_point(color = "#0072B2") +
  labs(x = "log<sub>10</sub>(x)") +
  theme(axis.title.x = element_markdown())
```

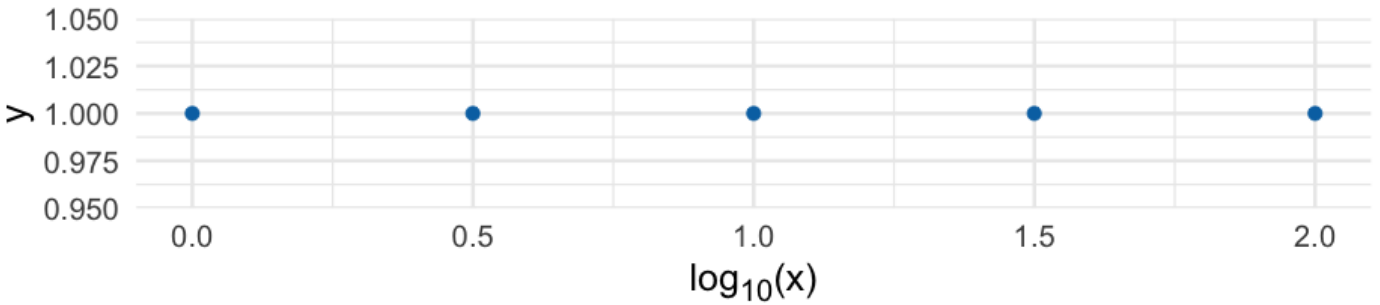

Labels should denote the data, not the scale of the axis

```
ggplot(d, aes(x, 1)) +
  geom_point(color = "#0072B2") +
  scale_x_log10()
```

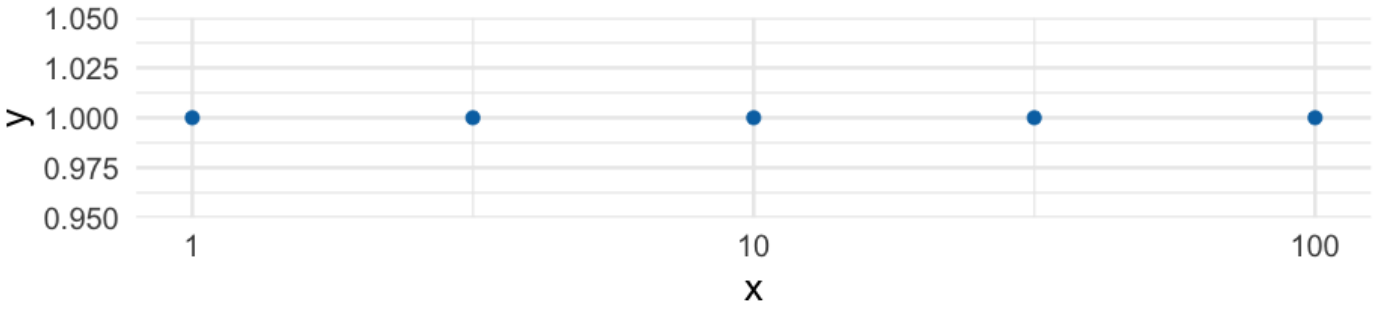

Labeling the above with $log_{10}(x)$ would be ambiguous and confusing

Labels and captions

Disclaimer

- APA style requires the labels be made in specific ways
- Much of the following discussion still applies
- Our book (Wilke) uses a similar style throughout

Title

What is the point of your figure?

What are you trying to communicate

- Figures should have only one title
- Use integrated title/subtitles for sharing with a broad audience
  - Blog posts
  - Social media
  - Reports to stakeholders
- Keep figures in subtext when there's a designated format you must adhere to
- Make sure your figure has a title
  - Should not start with "This figure displays/shows..."

Caption

Consider stating the data source

Other details relevant to the figure but not important enough for a subtitle

Axis labels

- The title for the axis
- Critical for communication
- **Never** use variable names (very common and very poor practice)
- State the measure and the unit (if quantitative)
    - e.g., "Brain Mass (grams)", "Support for Measure (millions of people)", "Dollars spent"
    - Categorical variable likely will not need to the measurement unit

Omission

- Consider omitting obvious or redundant labels
  - Use `labs(x = NULL)` or `labs(x = "")`
  - If already using `scale_x/y_*()` just supply the `name` argument

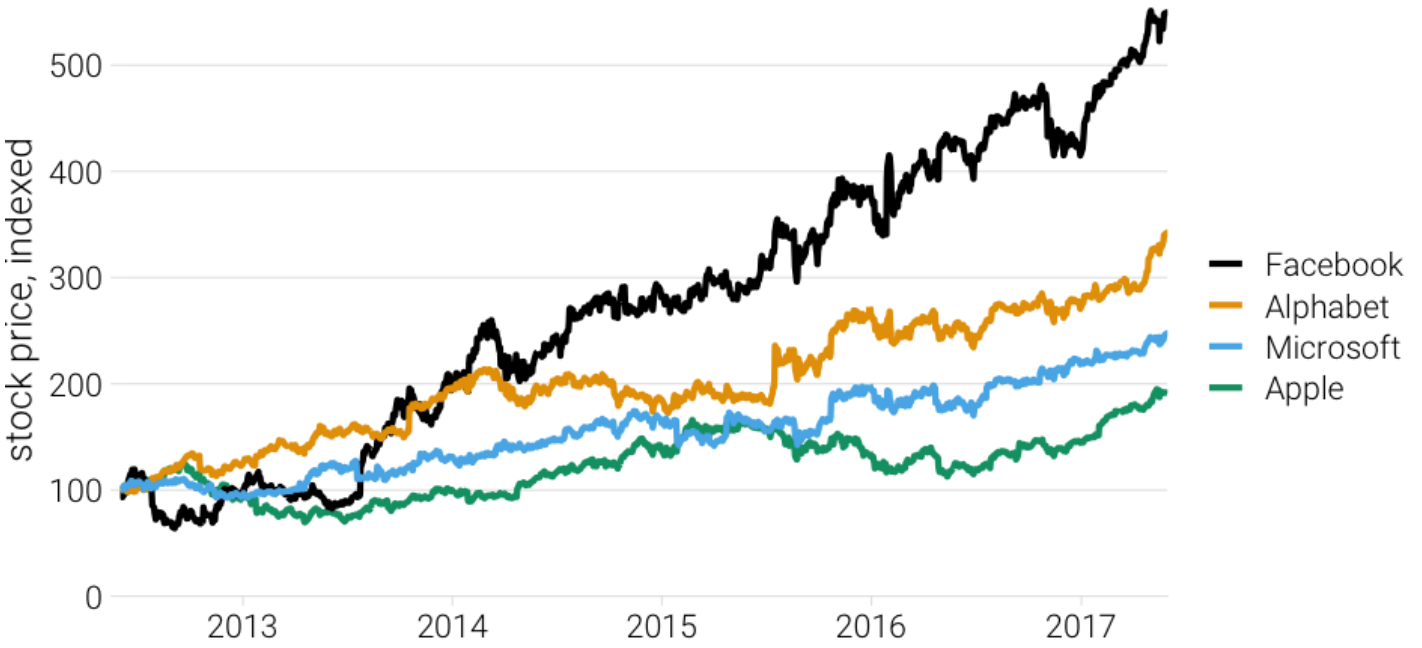

Omission

- Do not omit axis titles that are not obvious

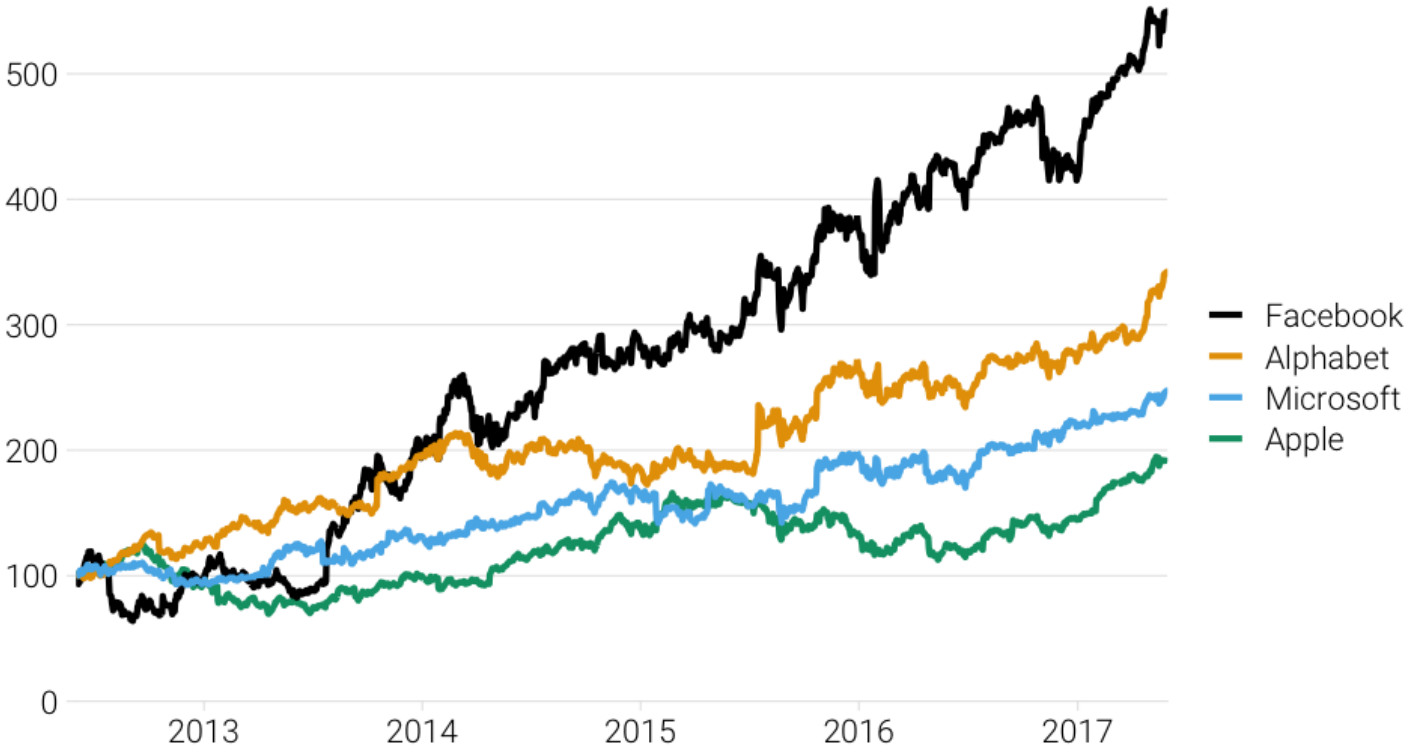

Don't overdo it

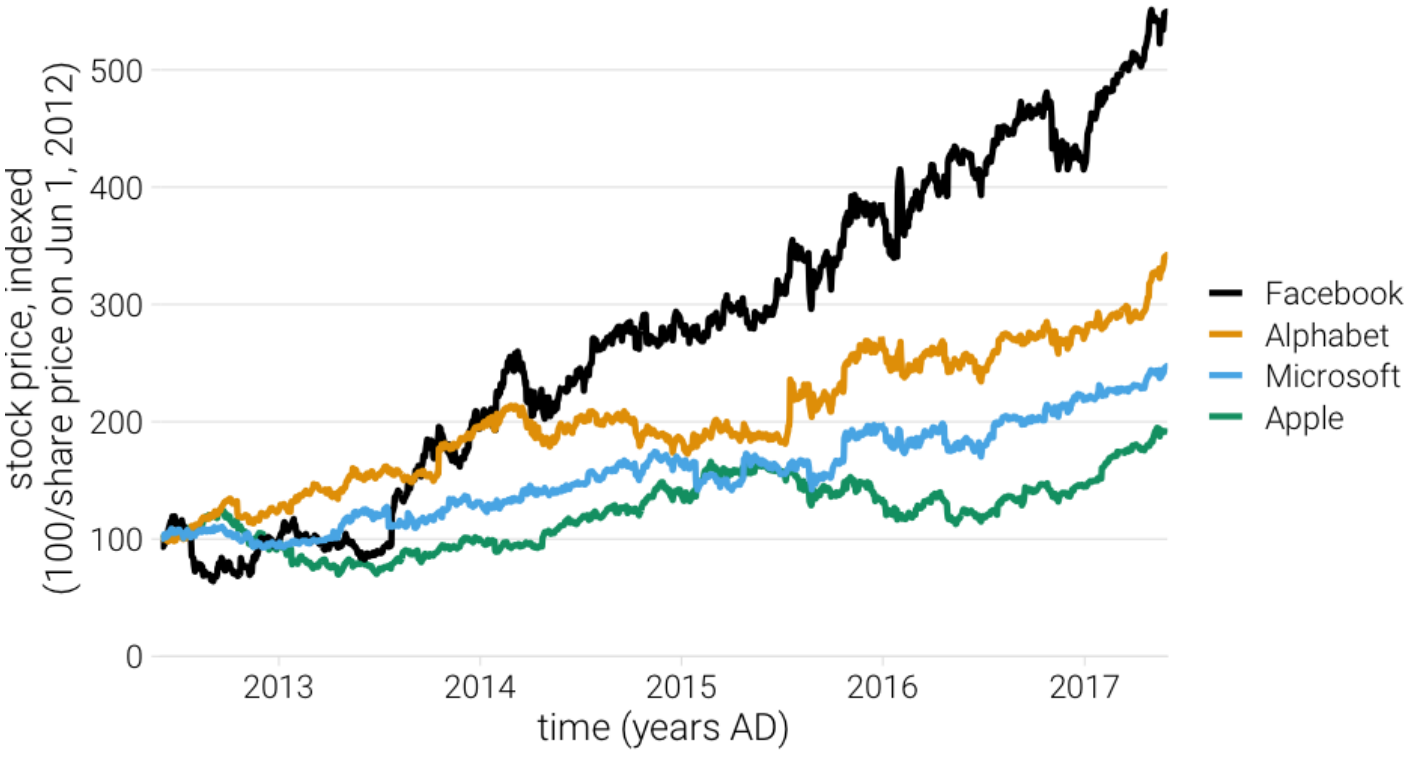

Annotations

Among the most effective

- If possible, try to remove legends, and just include annotations

Building up a plot

```
remotes::install_github("clauswilke/dviz.supp")
head(tech_stocks)
```

```
## # A tibble: 6 x 6
##   company  ticker date       price index_price price_indexed
##   <chr>    <chr>  <date>     <dbl>       <dbl>         <dbl>
## 1 Alphabet GOOG   2017-06-02 975.6        285.2      342.0757
## 2 Alphabet GOOG   2017-06-01 966.95       285.2      339.0428
## 3 Alphabet GOOG   2017-05-31 964.86       285.2      338.3100
## 4 Alphabet GOOG   2017-05-30 975.88       285.2      342.1739
## 5 Alphabet GOOG   2017-05-26 971.47       285.2      340.6276
## 6 Alphabet GOOG   2017-05-25 969.54       285.2      339.9509
```

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line()
```

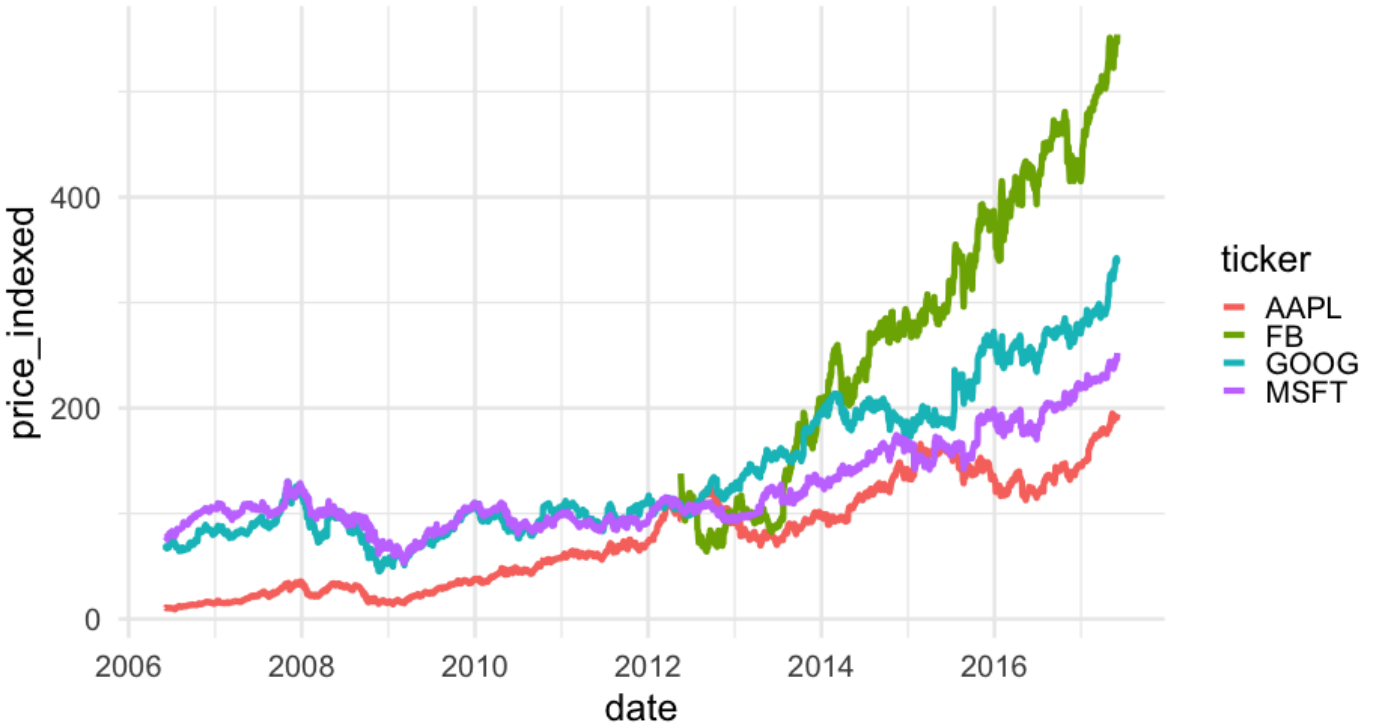

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line() +
  scale_color_OkabeIto()
```

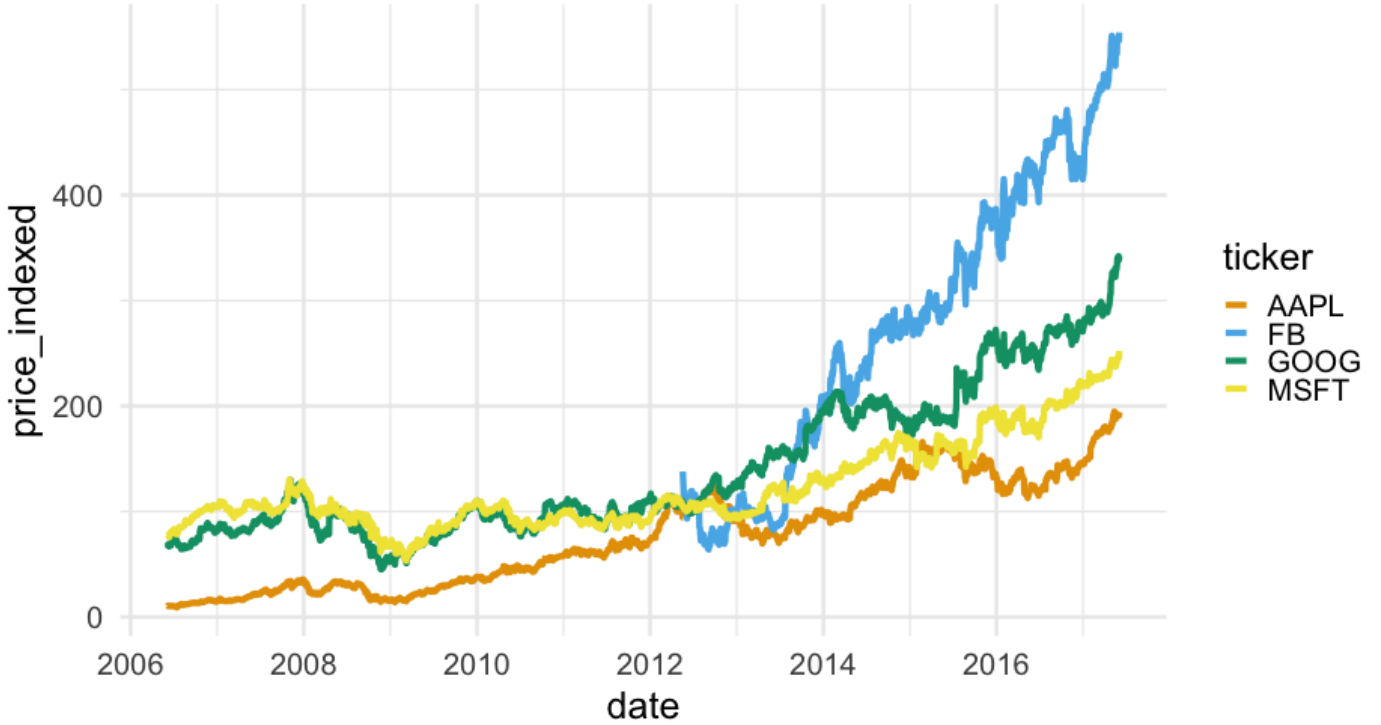

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line() +
  scale_color_OkabeIto(name = "Company",
                       breaks = c("GOOG", "AAPL", "FB", "MSFT"),
                       labels = c("Alphabet", "Apple", "Facebook
```

Bad

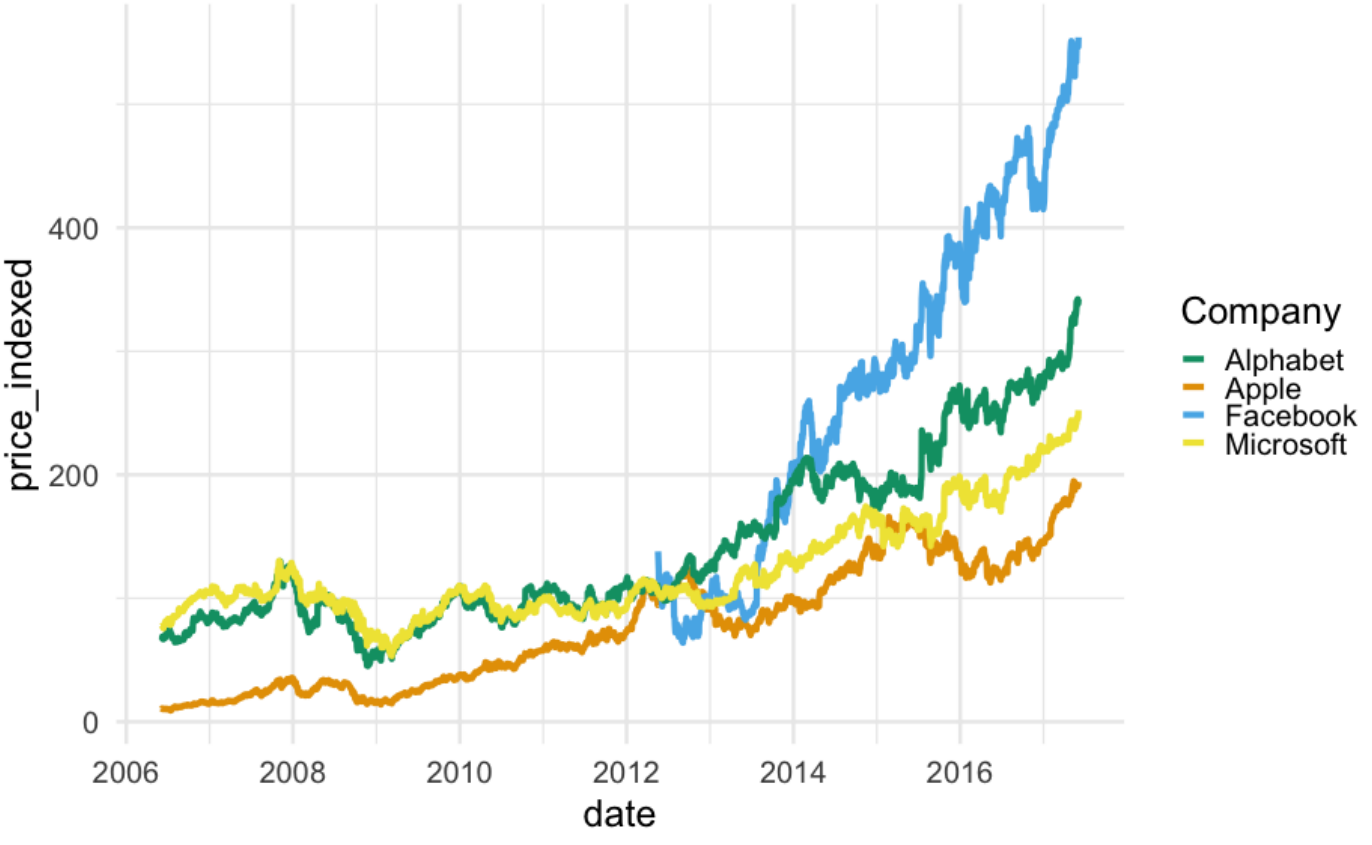
price_indexed
400
200
0
2006
2008
2010
2012
2014
2016
date
Company
Alphabet
Apple
Facebook
Microsoft

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line() +
  scale_color_OkabeIto(name = "Company",
                       breaks = c("FB", "GOOG", "MSFT", "AAPL"),
                       labels = c("Facebook", "Alphabet", "Micro
```

Good

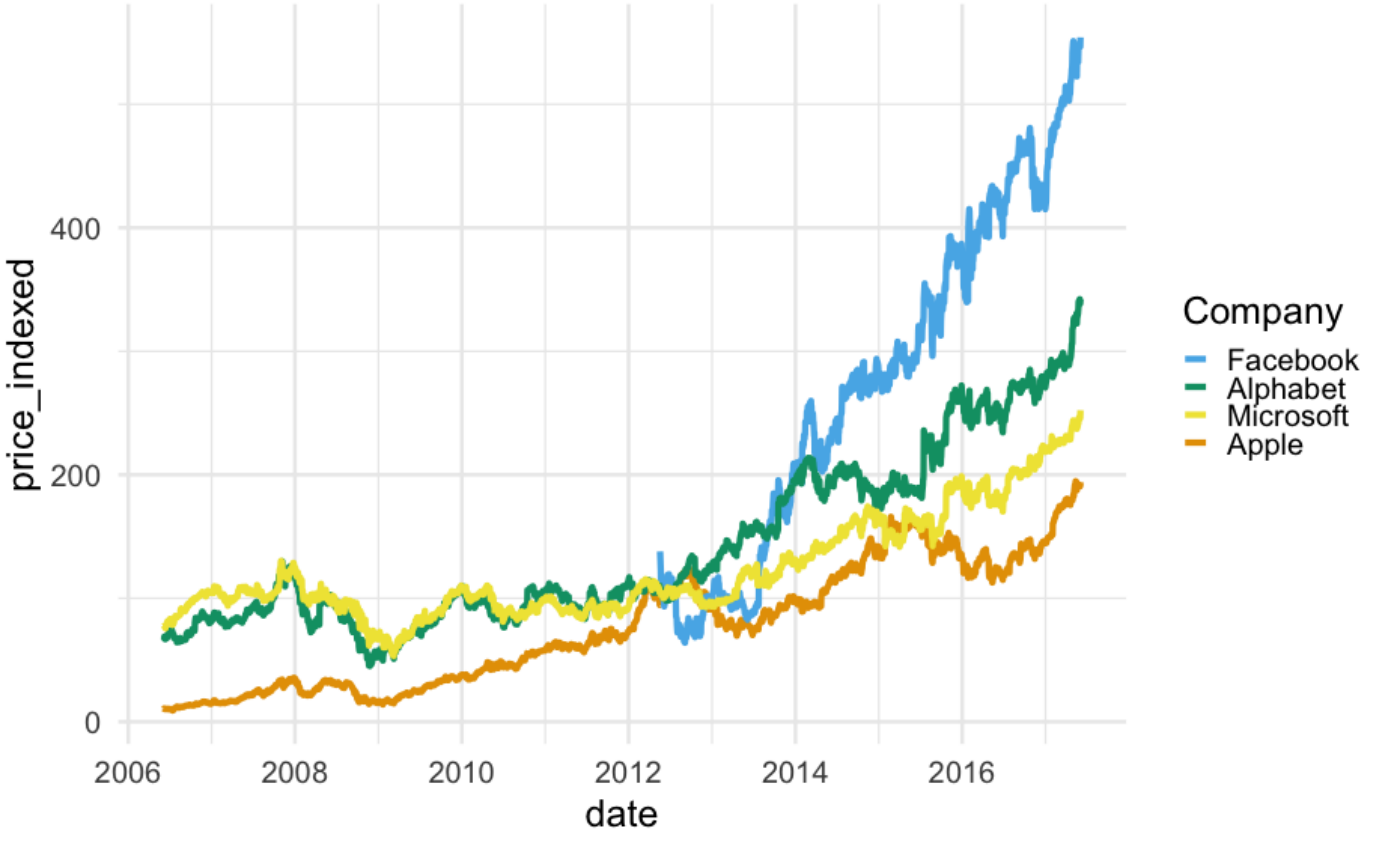
price_indexed
400
200
0
2006
2008
2010
2012
2014
2016
date
Company
Facebook
Alphabet
Microsoft
Apple

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line() +
  scale_color_OkabeIto(name = "Company",
                       breaks = c("FB", "GOOG", "MSFT", "AAPL"),
                       labels = c("Facebook", "Alphabet", "Micros
  scale_x_date(name = "year",
               limits = c(ymd("2012-06-01"), ymd("2018-12-31")),
               expand = c(0,0)) +
  geom_text(data = filter(tech_stocks, date == "2017-06-02"),
            aes(y = price_indexed, label = company),
            nudge_x = 280)
```

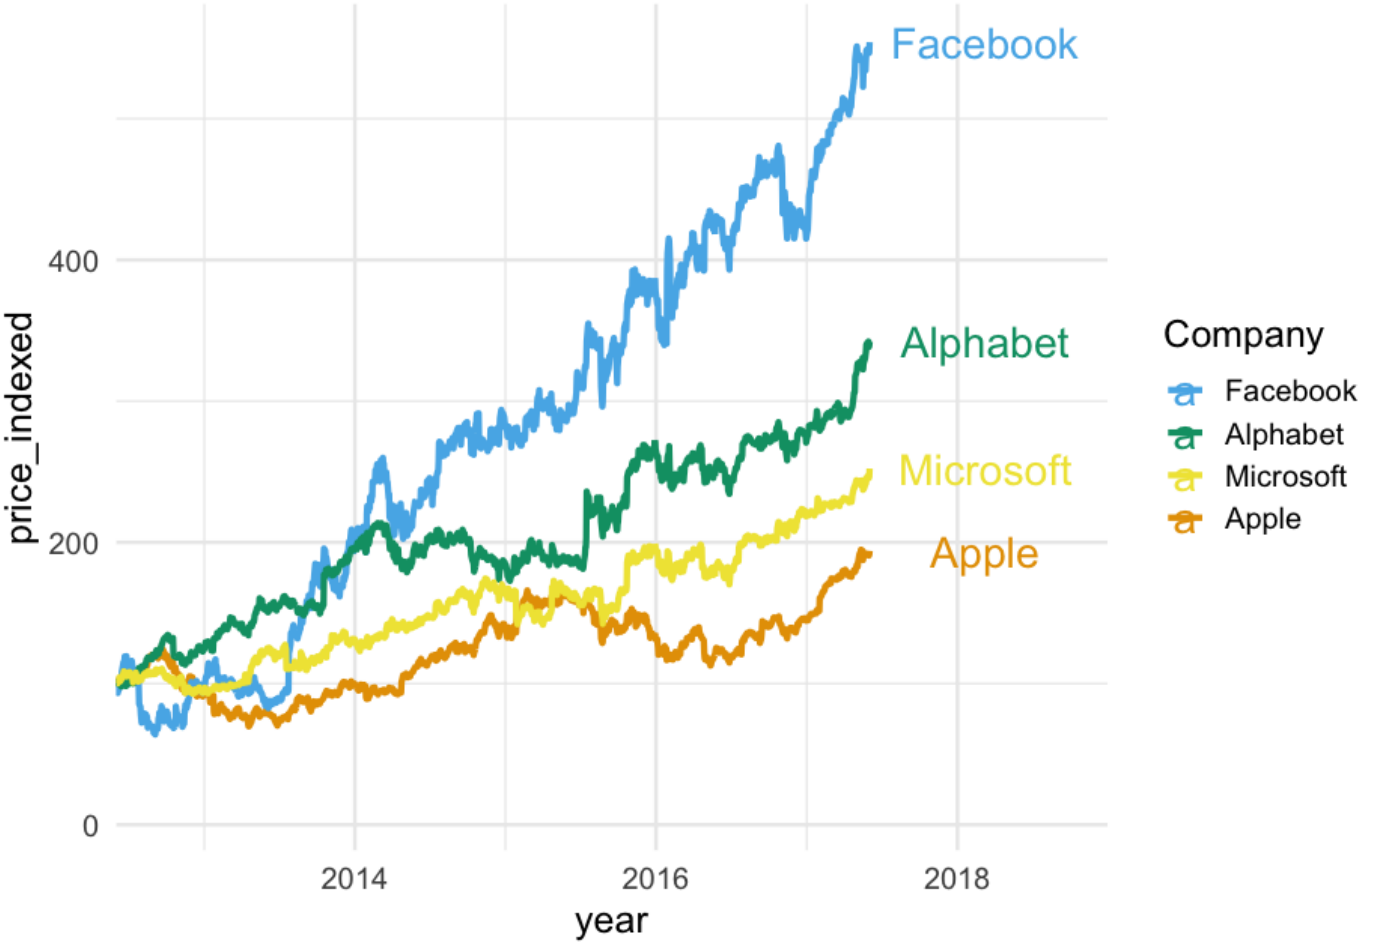
Facebook
Alphabet
Microsoft
Apple
400
200
0
price_indexed
2014
2016
2018
year
Company
Facebook
Alphabet
Microsoft
Apple

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line() +
  scale_color_OkabeIto(name = "Company",
                       breaks = c("FB", "GOOG", "MSFT", "AAPL"),
                       labels = c("Facebook", "Alphabet", "Micros
  scale_x_date(name = "year",
               limits = c(ymd("2012-06-01"), ymd("2018-12-31")),
               expand = c(0,0)) +
  geom_text(data = filter(tech_stocks, date == "2017-06-02"),
            aes(y = price_indexed, label = company),
            nudge_x = 280,
            hjust = 0)
```

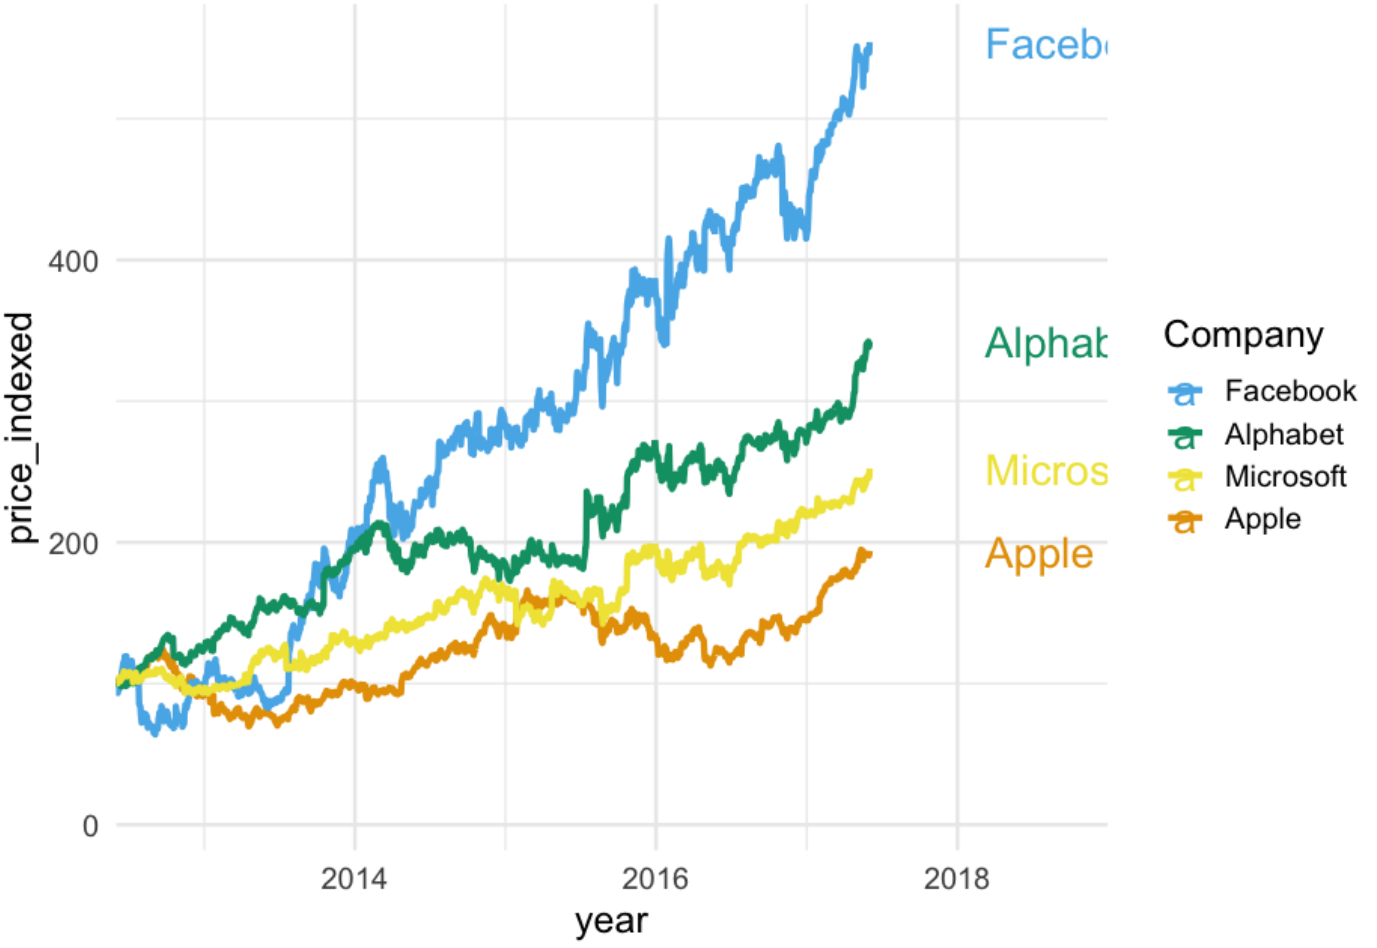
Faceb
Alphab
Micros
Apple
price_indexed
400
200
0
2014
2016
2018
year
Company
Facebook
Alphabet
Microsoft
Apple

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line() +
  scale_color_OkabeIto(name = "Company",
                       breaks = c("FB", "GOOG", "MSFT", "AAPL"),
                       labels = c("Facebook", "Alphabet", "Micros
  scale_x_date(name = "year",
               limits = c(ymd("2012-06-01"), ymd("2018-10-31")),
               expand = c(0,0)) +
  geom_text(data = filter(tech_stocks, date == "2017-06-02"),
            aes(y = price_indexed, label = company),
            color = "gray40",
            nudge_x = 20,
            hjust = 0) +
  guides(color = "none")
```

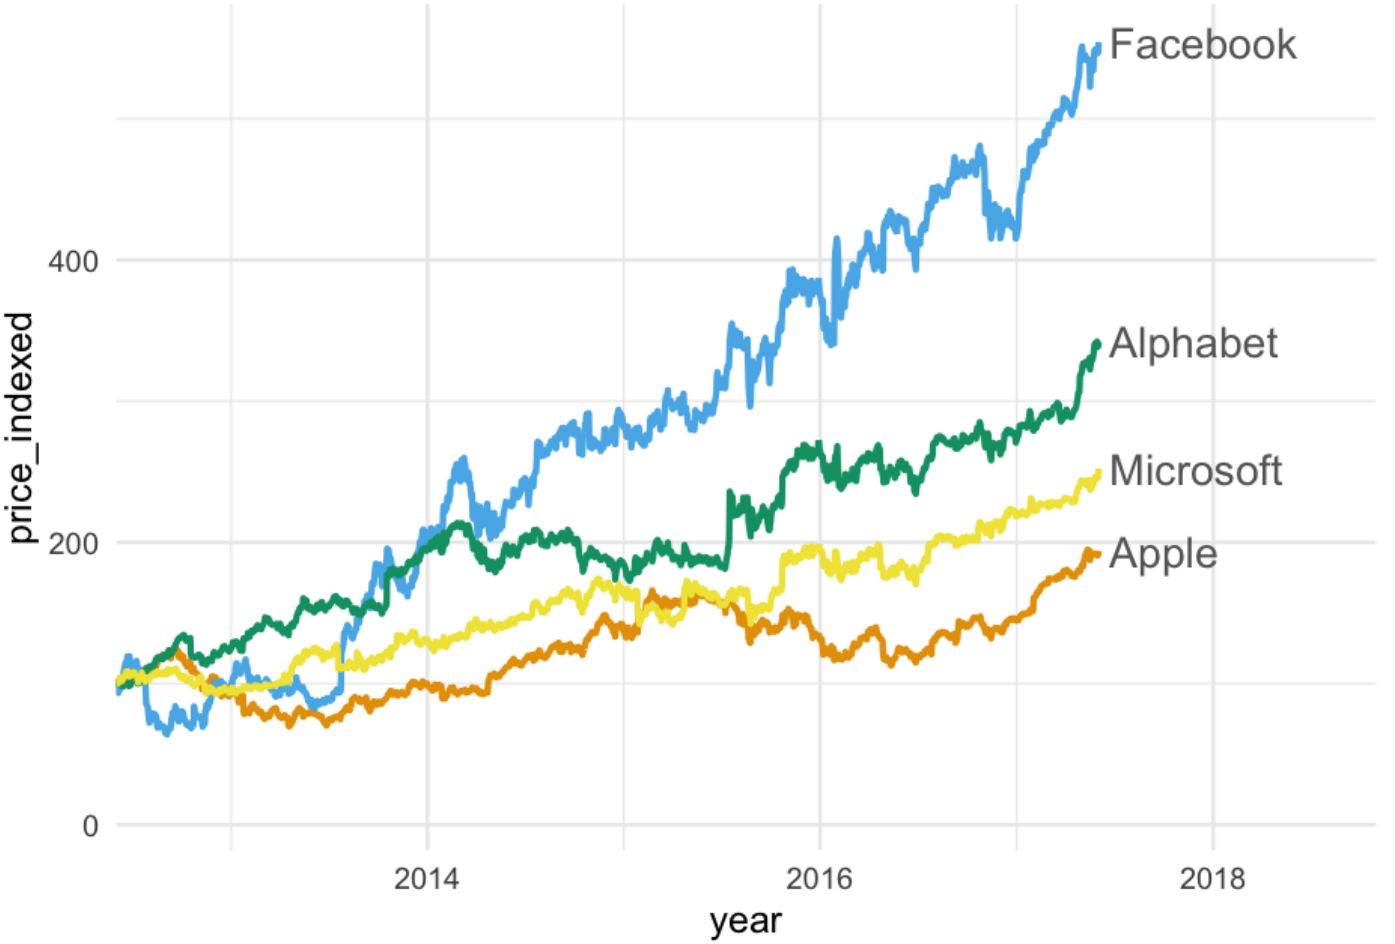
Facebook
Alphabet
Microsoft
Apple
400
200
0
price_indexed
2014
2016
2018
year

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line() +
  scale_color_OkabeIto(name = "Company",
                       breaks = c("FB", "GOOG", "MSFT", "AAPL"),
                       labels = c("Facebook", "Alphabet", "Micros
  scale_x_date(name = "",
               limits = c(ymd("2012-06-01"), ymd("2018-10-31")),
               expand = c(0,0)) +
  scale_y_continuous(name = "Stock Price, Indexed",
                     labels = scales::dollar) +
  geom_text(data = filter(tech_stocks, date == "2017-06-02"),
            aes(y = price_indexed, label = company),
            color = "gray40",
            nudge_x = 20,
            hjust = 0,
            size = 10) +
  guides(color = "none") +
  labs(title = "Tech growth over time",
       caption = "Data from Wilke (2019): Fundamentals of Data V
```

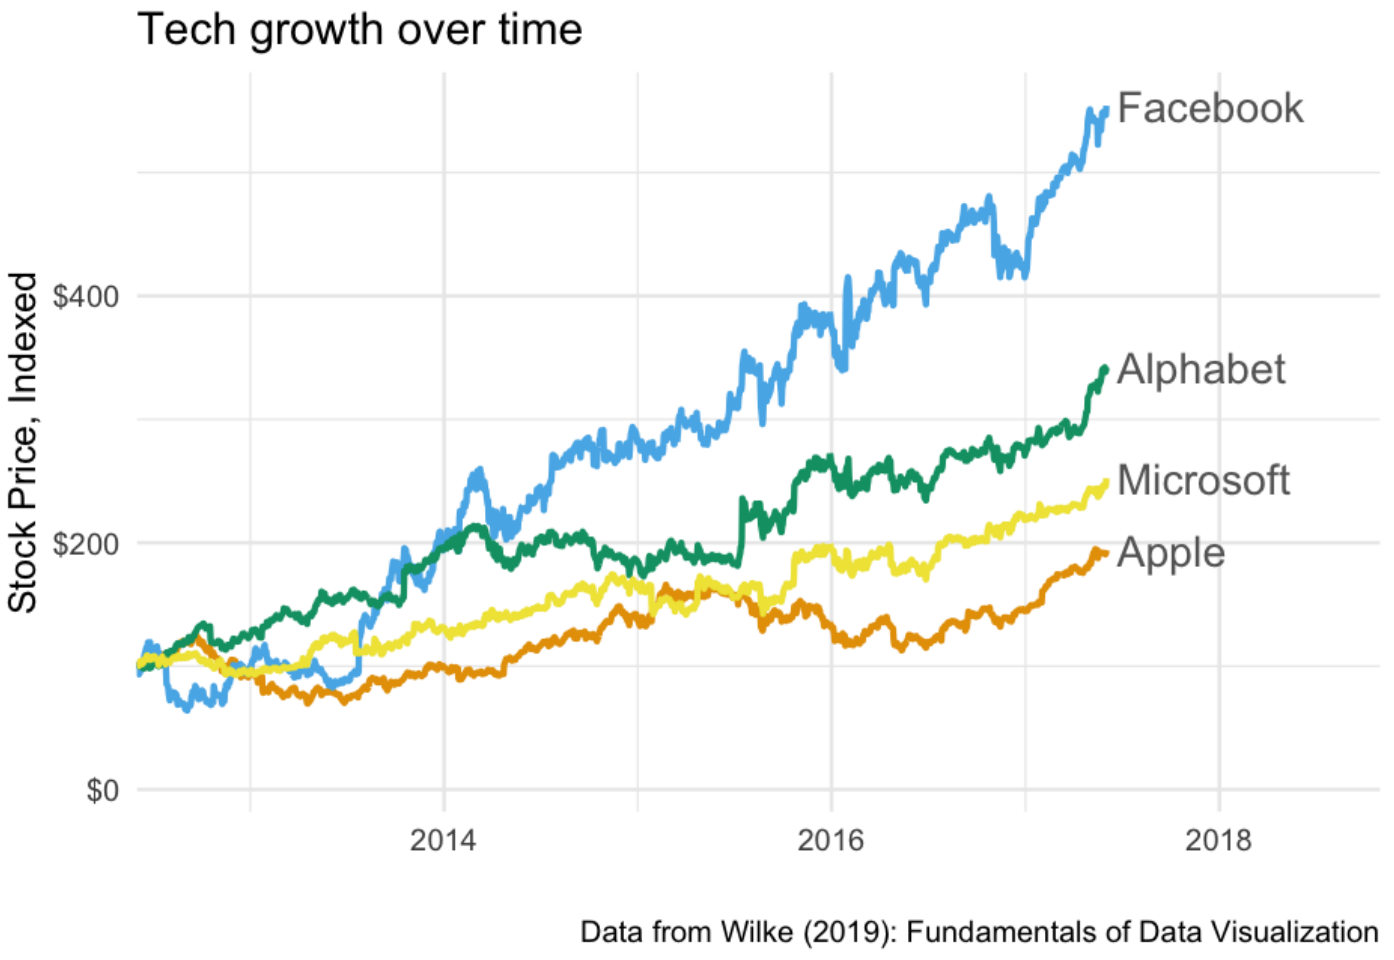
Tech growth over time
Facebook
Alphabet
Microsoft
Apple
Stock Price, Indexed
$400
$200
$0
2014
2016
2018
Data from Wilke (2019): Fundamentals of Data Visualization

Labeling bars

```
avs <- tech_stocks %>%
  group_by(company) %>%
  summarize(stock_av = mean(price_indexed)) %>%
  ungroup() %>%
  mutate(share = stock_av / sum(stock_av))
avs
```

```
## # A tibble: 4 x 3
##   company    stock_av     share
## * <chr>         <dbl>     <dbl>
## 1 Alphabet  141.0205  0.2292441
## 2 Apple      77.08241 0.1253058
## 3 Facebook  274.7427  0.4466240
## 4 Microsoft 122.3088  0.1988261
```

Bar plot

```
ggplot(avs, aes(fct_reorder(company, share), share)) +
  geom_col(fill = "#0072B2")
```

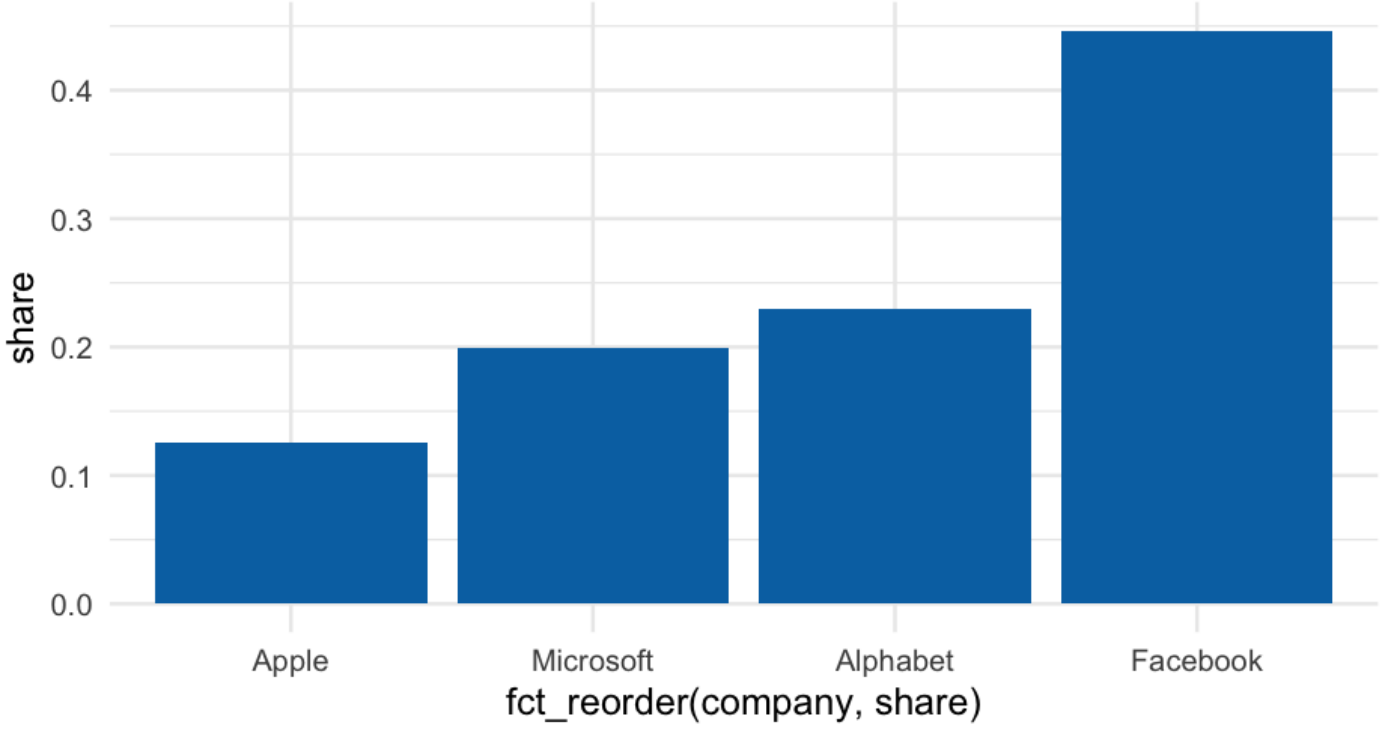

Horizontal

```
ggplot(avs, aes(share, fct_reorder(company, share))) +
  geom_col(fill = "#0072B2",
           alpha = 0.9)
```

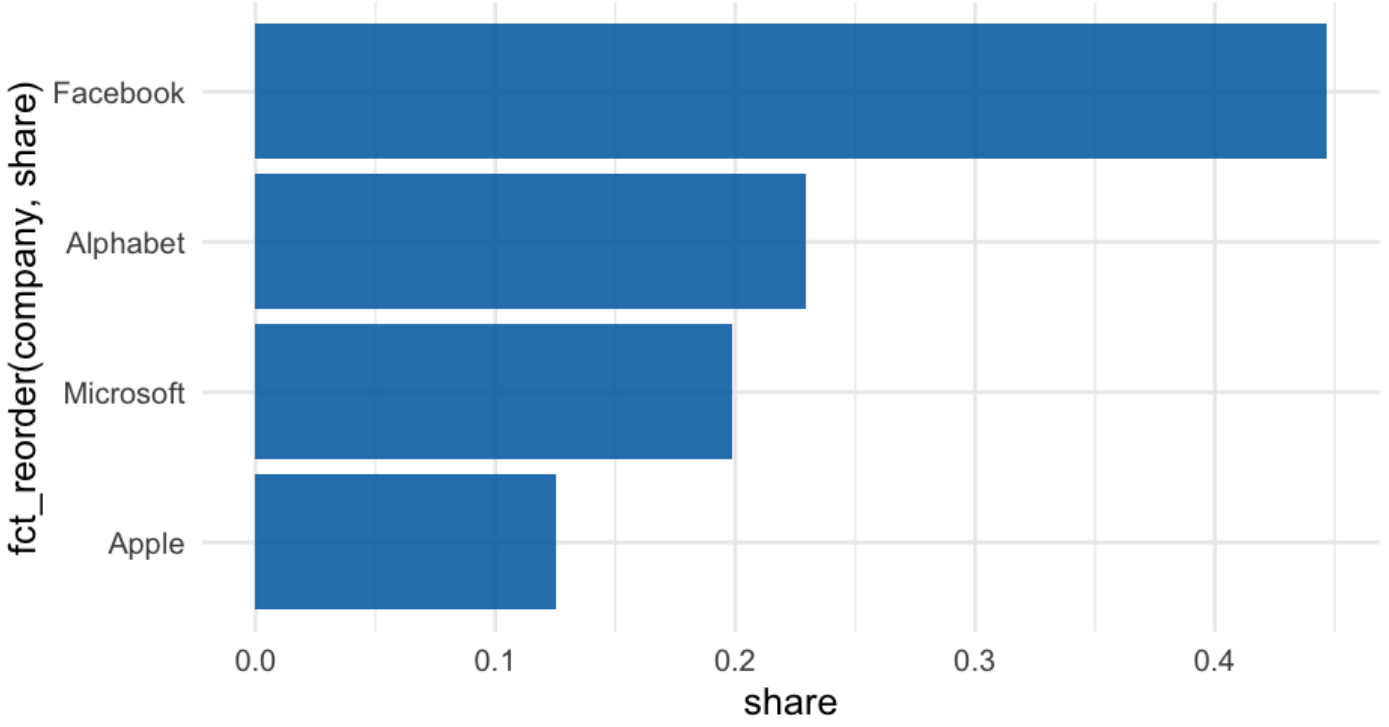

```
ggplot(avs, aes(fct_reorder(company, share), share)) +
  geom_col(fill = "#0072B2",
           alpha = 0.9) +
  coord_flip() +
  theme(panel.grid.major.y = element_blank(),
        panel.grid.minor.x = element_blank(),
        panel.grid.major.x = element_line(color = "gray80"))
```

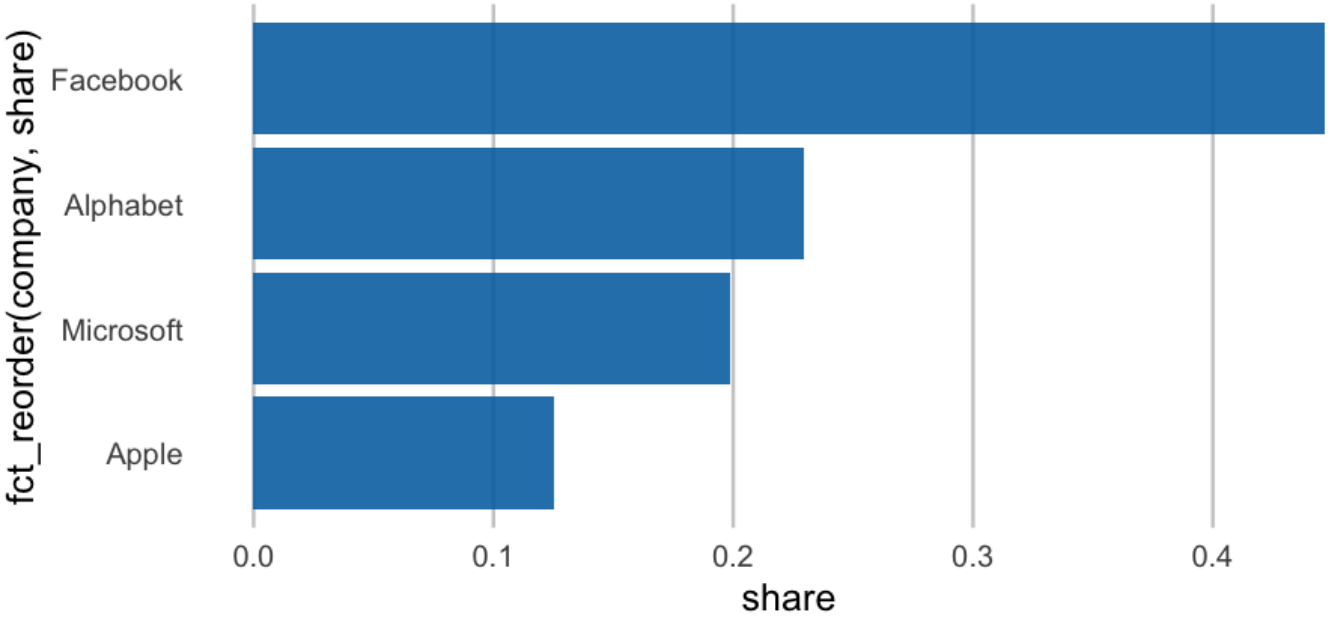

Quick aside

Let's actually make a bar plot theme

```
bp_theme <- function(...) {
  theme_minimal(...) +
    theme(panel.grid.major.y = element_blank(),
          panel.grid.minor.x = element_blank(),
          panel.grid.major.x = element_line(color = "gray80"),
          plot.title.position = "plot")
}
```

```
ggplot(avs, aes(fct_reorder(company, share), share)) +
  geom_col(fill = "#0072B2",
           alpha = 0.9) +
  geom_text(aes(company, share, label = round(share, 2)),
            nudge_y = 0.02,
            size = 8) +
  coord_flip() +
  bp_theme(base_size = 25)
```

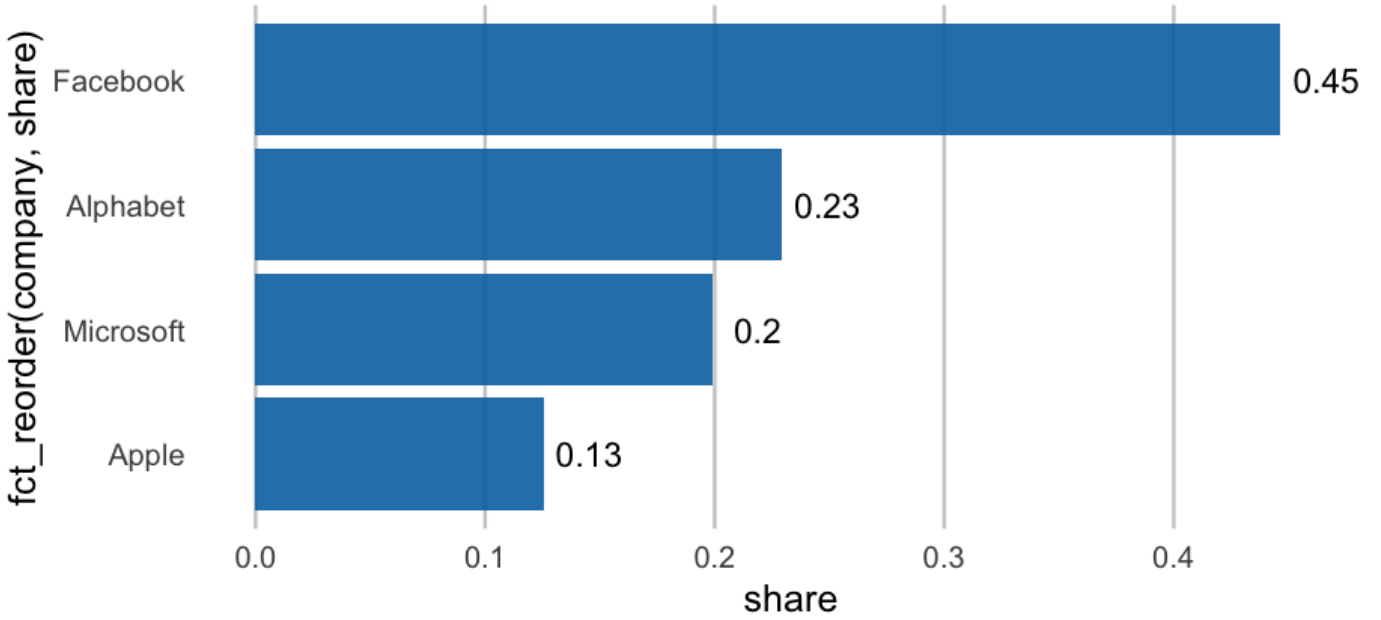

```
ggplot(avs, aes(fct_reorder(company, share), share)) +
  geom_col(fill = "#0072B2",
           alpha = 0.9) +
  geom_text(aes(company, share, label = paste0(round(share*100),
            nudge_y = 0.02,
            size = 8) +
  coord_flip() +
  scale_y_continuous("Market Share", labels = scales::percent) +
  labs(x = NULL,
       title = "Tech company market control",
       caption = "Data from Clause Wilke Book: Fundamentals of Da
  bp_theme(base_size = 25)
```

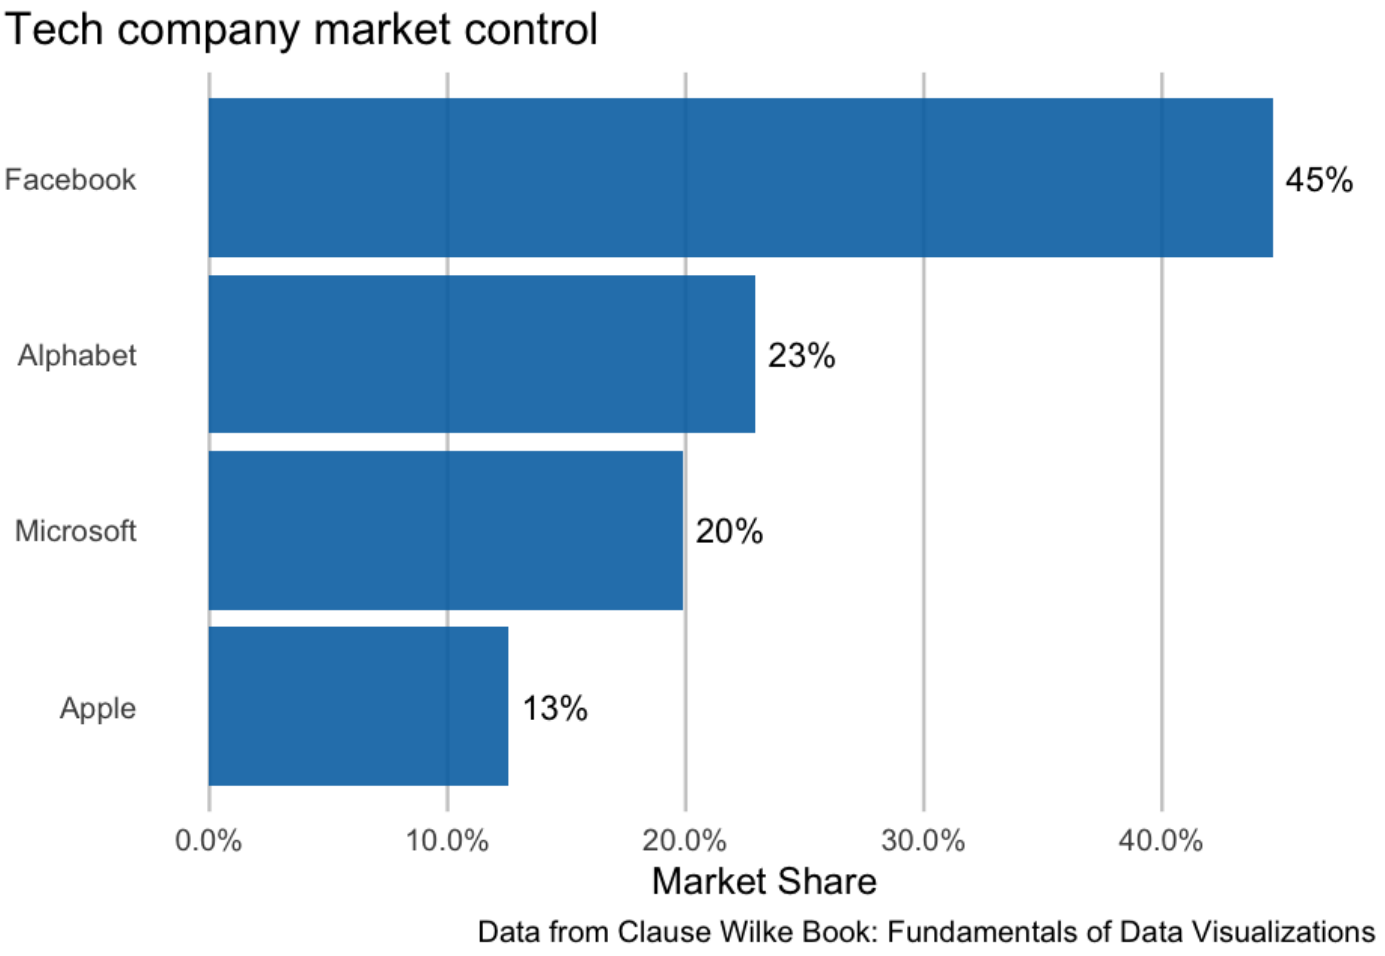
Tech company market control
Facebook
Alphabet
Microsoft
Apple
45%
23%
20%
13%
0.0%
10.0%
20.0%
30.0%
40.0%
Market Share
Data from Clause Wilke Book: Fundamentals of Data Visualizations

```
ggplot(avs, aes(fct_reorder(company, share), share)) +
  geom_col(fill = "#0072B2",
           alpha = 0.9) +
  geom_text(aes(company, share, label = paste0(round(share*100),
            nudge_y = 0.02,
            size = 8) +
  coord_flip() +
  scale_y_continuous("Market Share",
                     labels = scales::percent,
                     expand = c(0, 0, 0.05, 0)) +
  labs(x = NULL,
       title = "Tech company market control",
       caption = "Data from Clause Wilke Book: Fundamentals of Da
  bp_theme(base_size = 25)
```

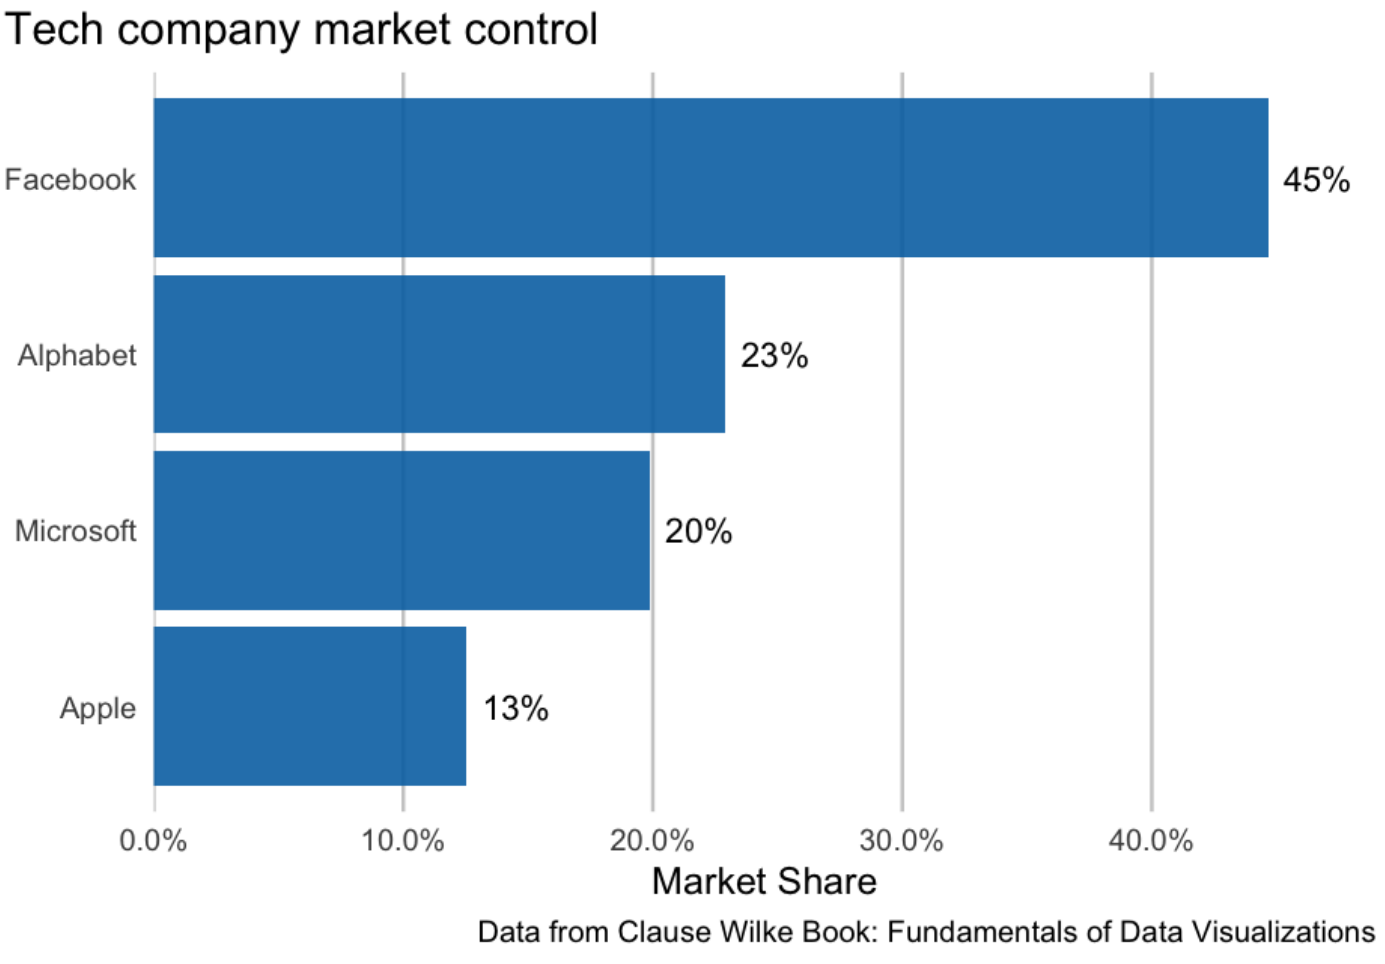
Tech company market control
Facebook
Alphabet
Microsoft
Apple
45%
23%
20%
13%
0.0%
10.0%
20.0%
30.0%
40.0%
Market Share
Data from Clause Wilke Book: Fundamentals of Data Visualizations

Last alternative

```
ggplot(avs, aes(fct_reorder(company, share), share)) +
  geom_col(fill = "#0072B2",
           alpha = 0.9) +
  geom_text(aes(company, share, label = paste0(round(share*100),
            nudge_y = -0.02,
            size = 8,
            color = "white") +
  coord_flip() +
  scale_y_continuous("Market Share",
                     labels = scales::percent,
                     expand = c(0, 0, 0.05, 0)) +
  labs(x = NULL,
       title = "Tech company market control",
       caption = "Data from Clause Wilke Book: Fundamentals of Da
  bp_theme(base_size = 25)
```

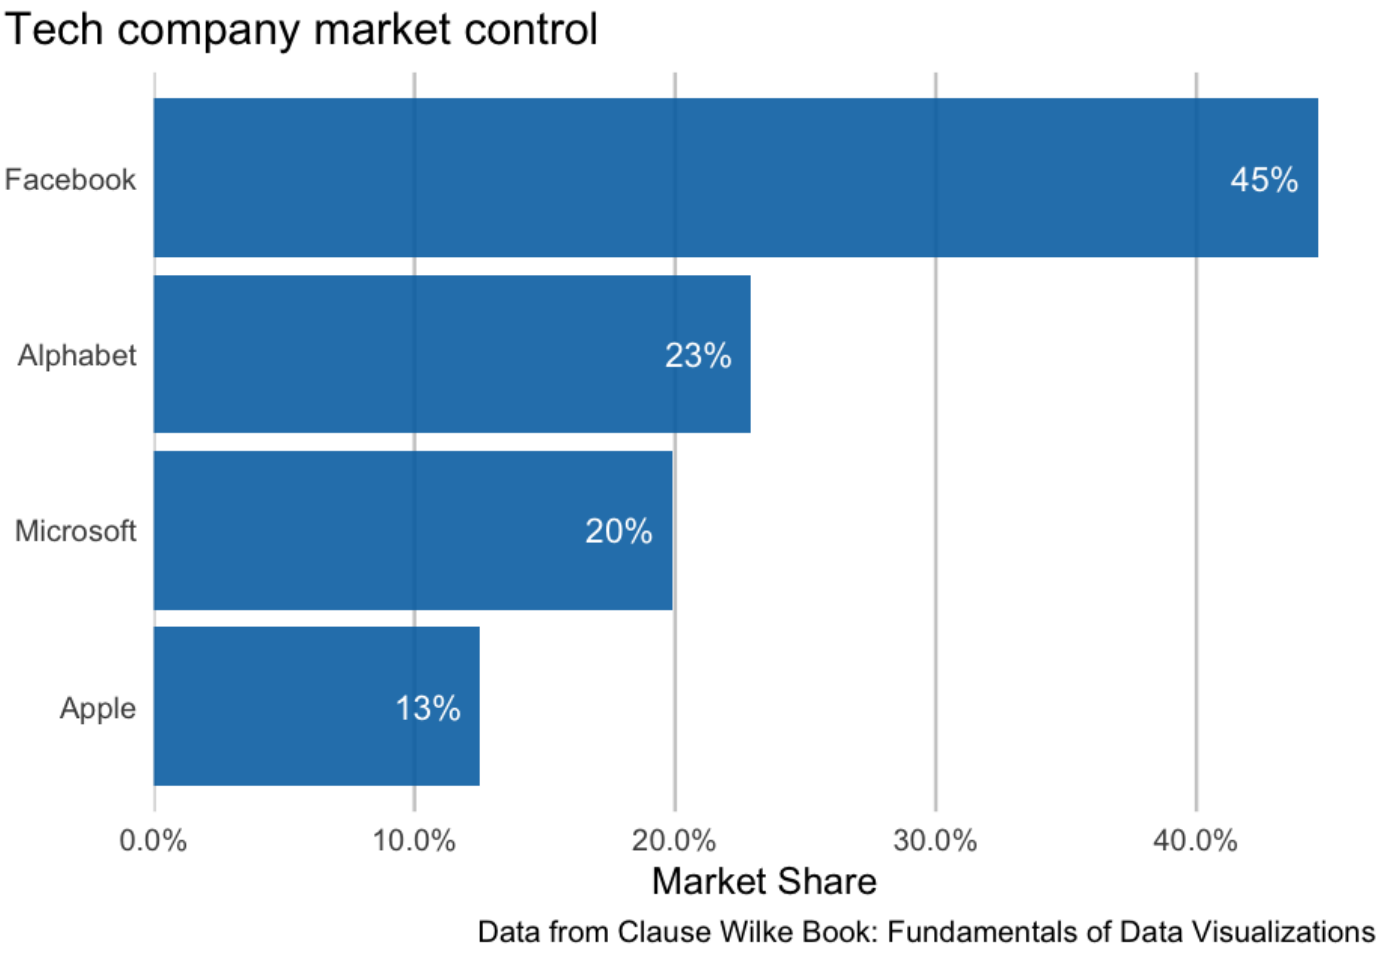
Tech company market control
Facebook
Alphabet
Microsoft
Apple
45%
23%
20%
13%
0.0%
10.0%
20.0%
30.0%
40.0%
Market Share
Data from Clause Wilke Book: Fundamentals of Data Visualizations

Distributions

```
ggplot(iris, aes(Sepal.Length, fill = Species)) +
  geom_density(alpha = 0.3,
               color = "white")
```

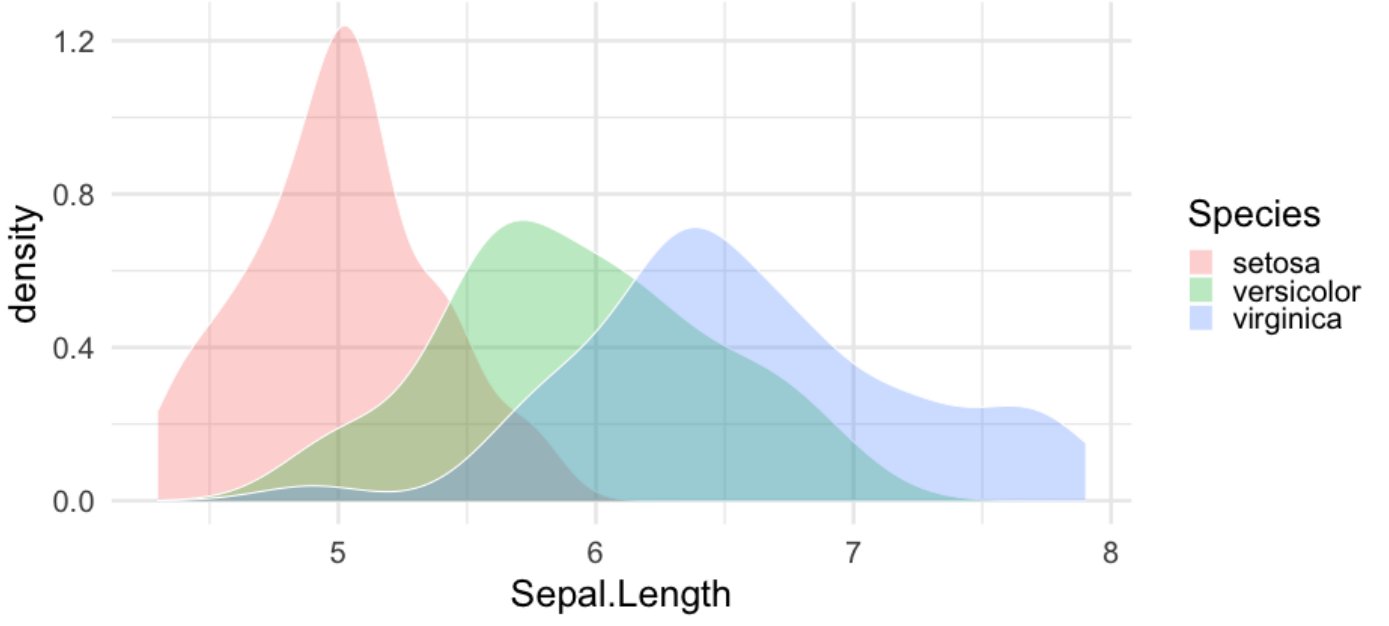

```
ggplot(iris, aes(Sepal.Length, fill = Species)) +
  geom_density(alpha = 0.3,
               color = "white") +
  scale_fill_OkabeIto()
```

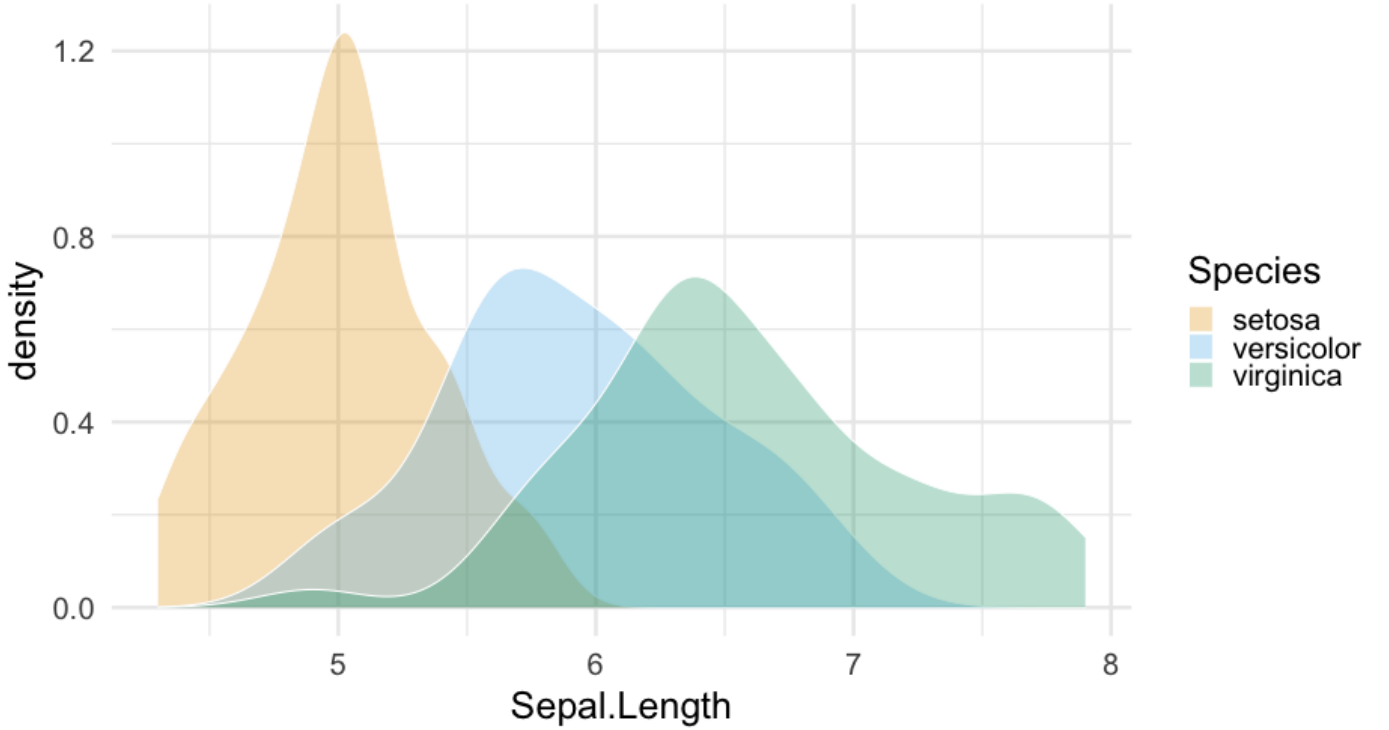

Labeling

One method

```
label_locs <- tibble(Sepal.Length = c(5.45, 6, 7),
                     density = c(1, 0.8, 0.6),
                     Species = c("setosa", "versicolor", "virgin

ggplot(iris, aes(Sepal.Length, fill = Species)) +
  geom_density(alpha = 0.3,
               color = "white") +
  scale_fill_OkabeIto() +
  geom_text(aes(label = Species, y = density, color = Species),
            data = label_locs)
```

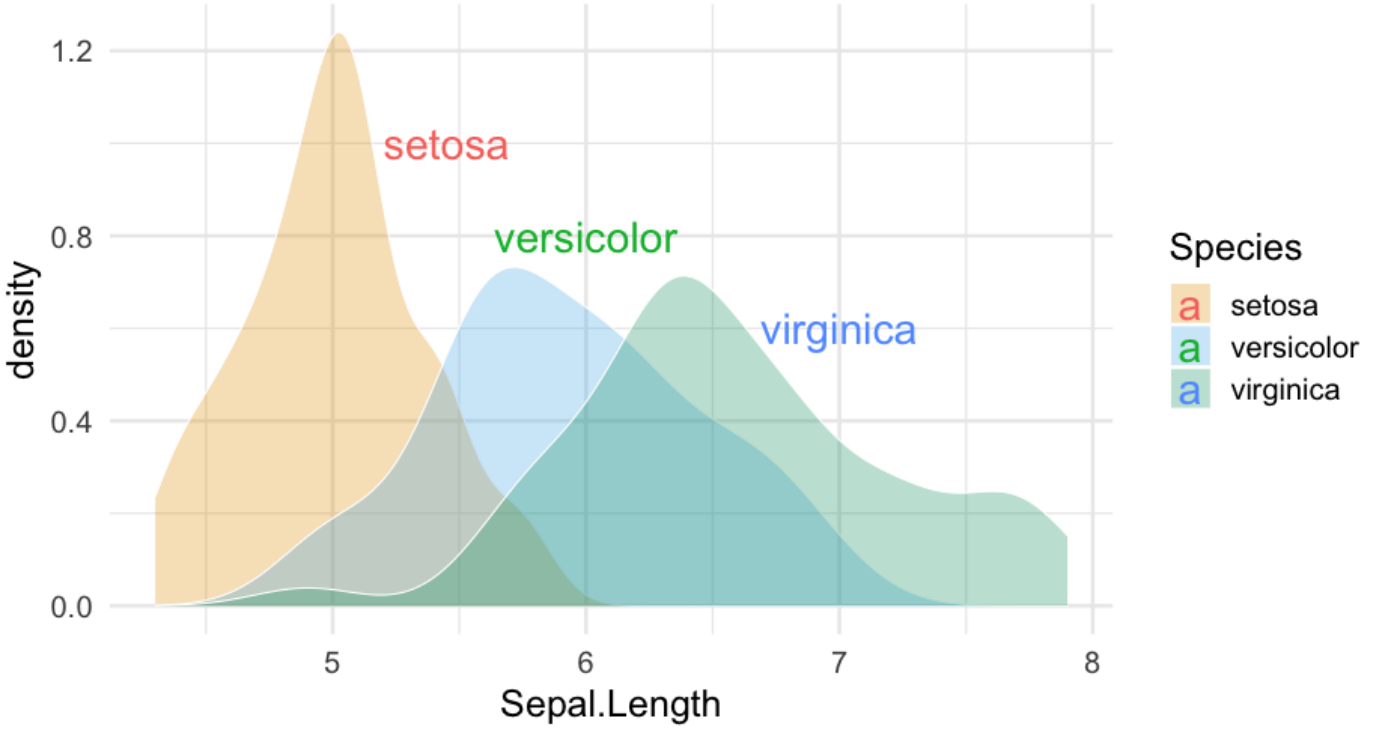
1.2
0.8
0.4
0.0
density
setosa
versicolor
virginica
5
6
7
8
Sepal.Length
Species
a setosa
a versicolor
a virginica

```
ggplot(iris, aes(Sepal.Length, fill = Species)) +
  geom_density(alpha = 0.3,
               color = "white") +
  scale_fill_OkabeIto() +
  scale_color_OkabeIto() +
  geom_text(aes(label = Species, y = density, color = Species),
            data = label_locs) +
  guides(color = "none",
         fill = "none")
```

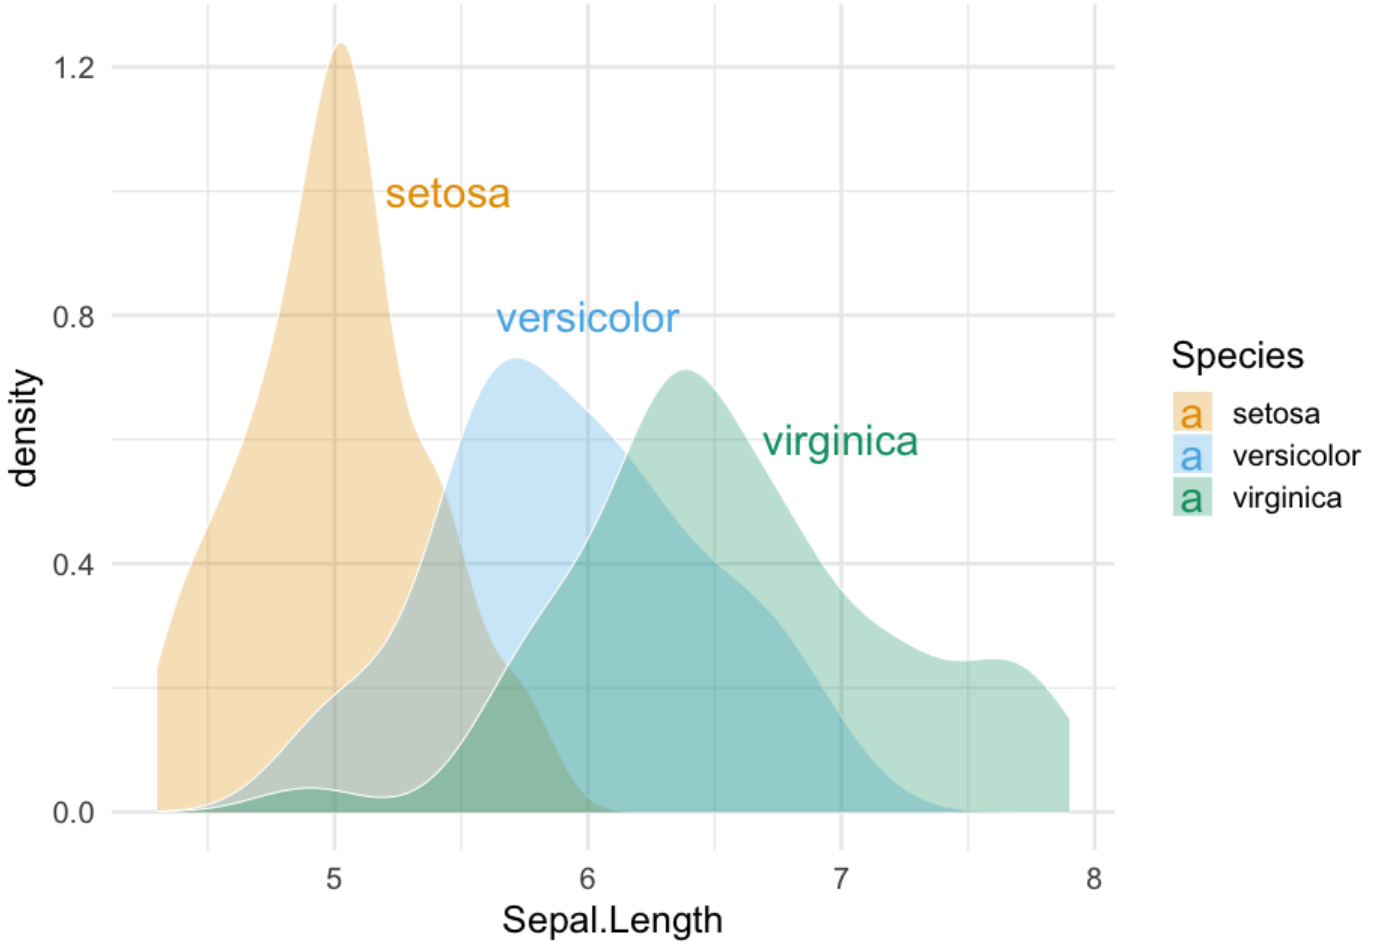
1.2
0.8
0.4
0.0
density
setosa
versicolor
virginica
5
6
7
8
Sepal.Length
Species
a setosa
a versicolor
a virginica

```r
label_locs <- tibble(Sepal.Length = c(5.4, 6, 6.9),
                     density = c(1, 0.75, 0.6),
                     Species = c("setosa", "versicolor", "virgin

ggplot(iris, aes(Sepal.Length, fill = Species)) +
  geom_density(alpha = 0.3,
               color = "white") +
  scale_fill_OkabeIto() +
  scale_color_OkabeIto() +
  geom_text(aes(label = Species, y = density),
            color = "gray40",
            data = label_locs) +
  guides(fill = "none")
```

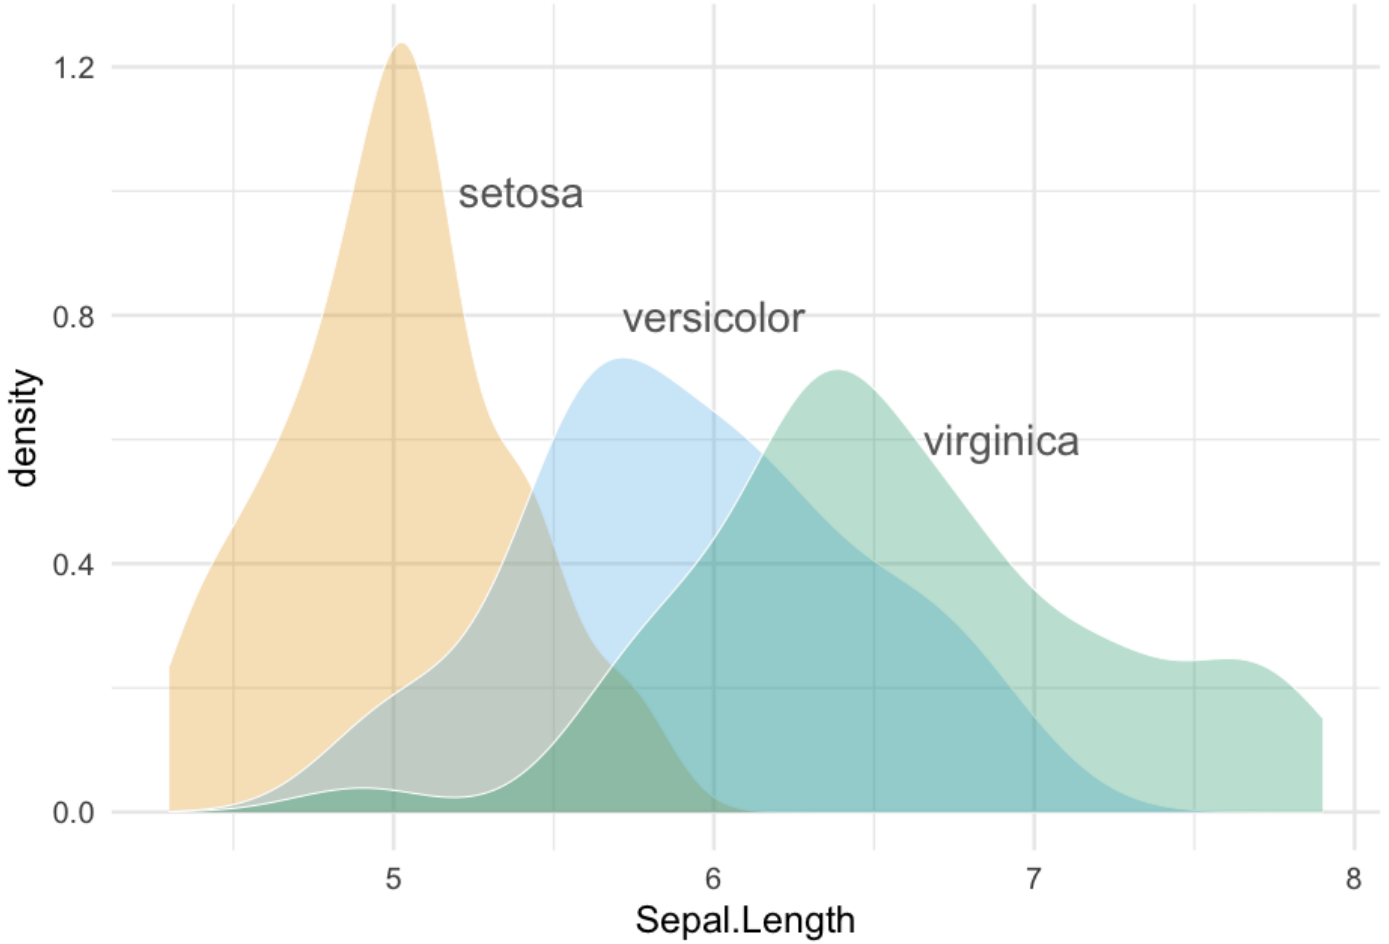
1.2
0.8
0.4
0.0
density
setosa
versicolor
virginica
5
6
7
8
Sepal.Length

Other options

- Rather than using a new data frame, you could use multiple calls to `annotate`.
- One is not necessarily better than the other, but I prefer the data frame method
- Keep in mind you can always use multiple data sources within a single plot
    - Each layer can have its own data source
    - Common in geographic data in particular

Annotate example

```
ggplot(iris, aes(Sepal.Length, fill = Species)) +
  geom_density(alpha = 0.3) +
  scale_fill_OkabeIto() +
  scale_color_OkabeIto() +
  annotate("text", label = "setosa", x = 5.45, y = 1, color = "g
  annotate("text", label = "versicolor", x = 6, y = 0.8, color =
  annotate("text", label = "virginica", x = 7, y = 0.6, color =
  guides(fill = "none")
```

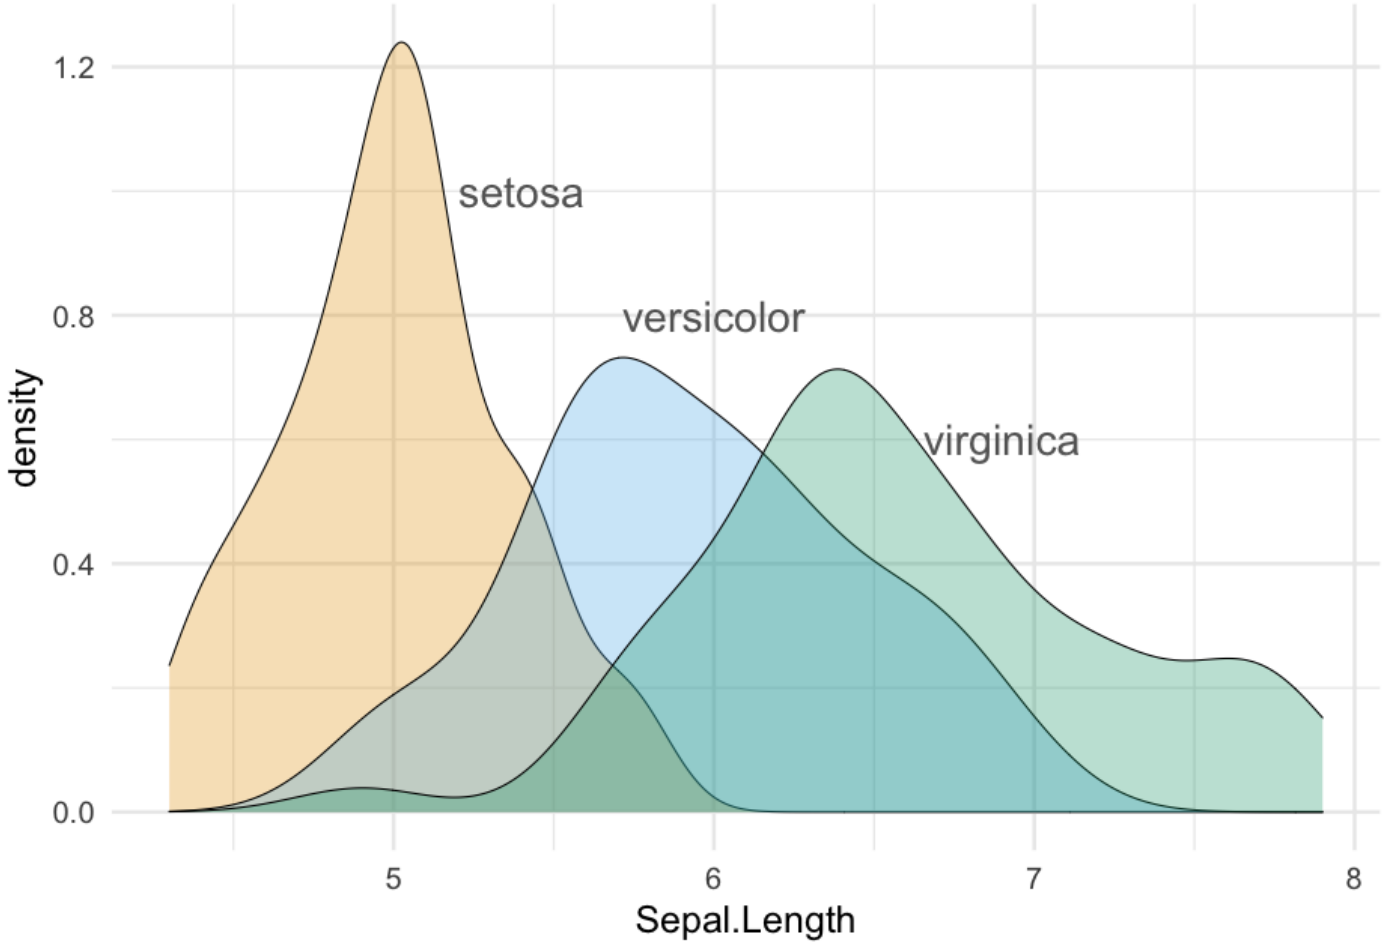
1.2
0.8
0.4
0.0
density
setosa
versicolor
virginica
5
6
7
8
Sepal.Length

ggrepel

Plot text directly

```
cars <- rownames_to_column(mtcars)

ggplot(cars, aes(hp, mpg)) +
  geom_text(aes(label = rowname))
```

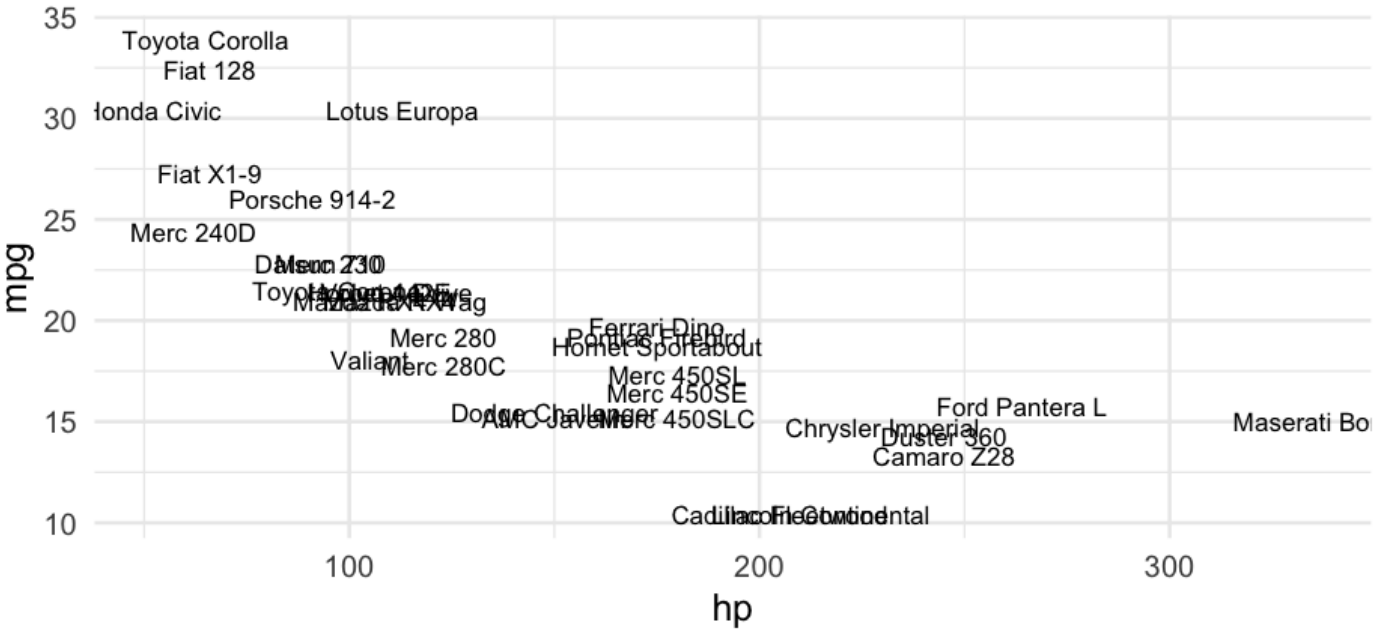

Repel text

```
library(ggrepel)
ggplot(cars, aes(hp, mpg)) +
  geom_text_repel(aes(label = rowname))
```

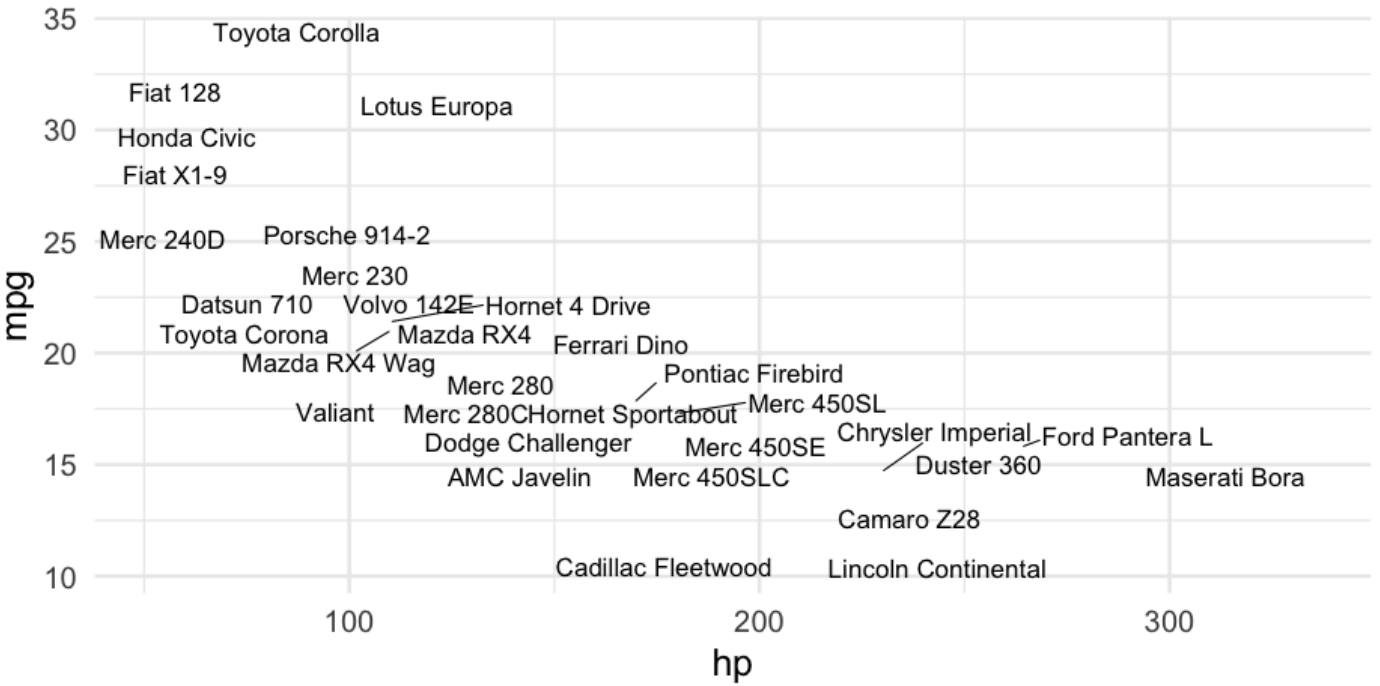

Slightly better

```
ggplot(cars, aes(hp, mpg)) +
  geom_point(color = "gray70") +
  geom_text_repel(aes(label = rowname),
                  min.segment.length = 0)
```

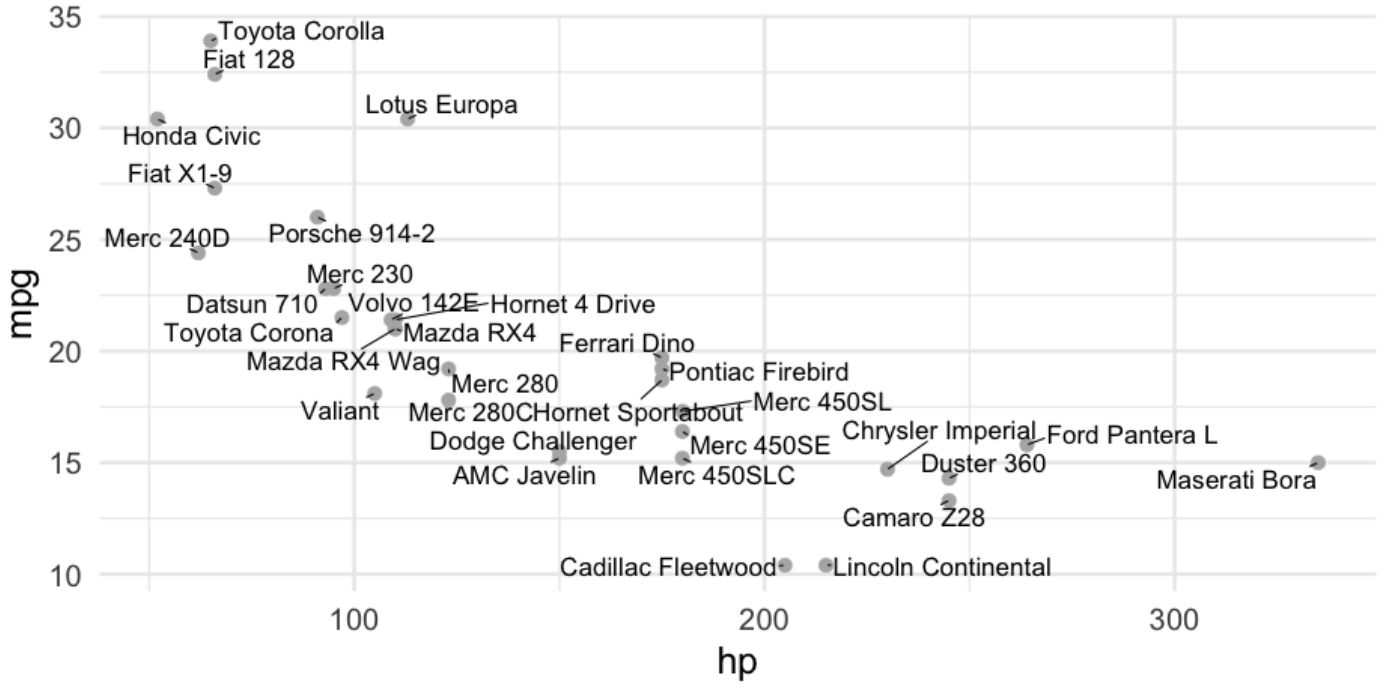

Common use cases

- Label some sample data that makes some theoretical sense (we've seen this before)
- Label outliers
- Label points from a specific group (e.g., similar to highlighting – can be used in conjunction)

Some new data

```
remotes::install_github("kjhealy/socviz")
library(socviz)
```

```
by_country <- organdata %>%
  group_by(consent_law, country) %>%
  summarize(donors_mean= mean(donors, na.rm = TRUE),
            donors_sd = sd(donors, na.rm = TRUE),
            gdp_mean = mean(gdp, na.rm = TRUE),
            health_mean = mean(health, na.rm = TRUE),
            roads_mean = mean(roads, na.rm = TRUE),
            cerebvas_mean = mean(cerebvas, na.rm = TRUE))
```

```
by_country
```

```
## # A tibble: 17 x 8
## # Groups:   consent_law [2]
##   consent_law country     donors_mean donors_sd gdp_mean health_mean roa
##   <chr>       <chr>             <dbl>     <dbl>    <dbl>       <dbl>
## 1 Informed    Australia      10.635   1.142808  22178.54    1957.5    10
## 2 Informed    Canada         13.96667 0.7511607 23711.08    2271.929  10
## 3 Informed    Denmark        13.09167 1.468121  23722.31    2054.071  10
## 4 Informed    Germany        13.04167 0.6111960 22163.23    2348.75   11
## 5 Informed    Ireland        19.79167 2.478437  20824.38    1479.929  11
## 6 Informed    Netherlands    13.65833 1.551807  23013.15    1992.786   7
## # … with 11 more rows
```

Scatterplot

```
ggplot(by_country, aes(gdp_mean, health_mean)) +
  geom_point()
```

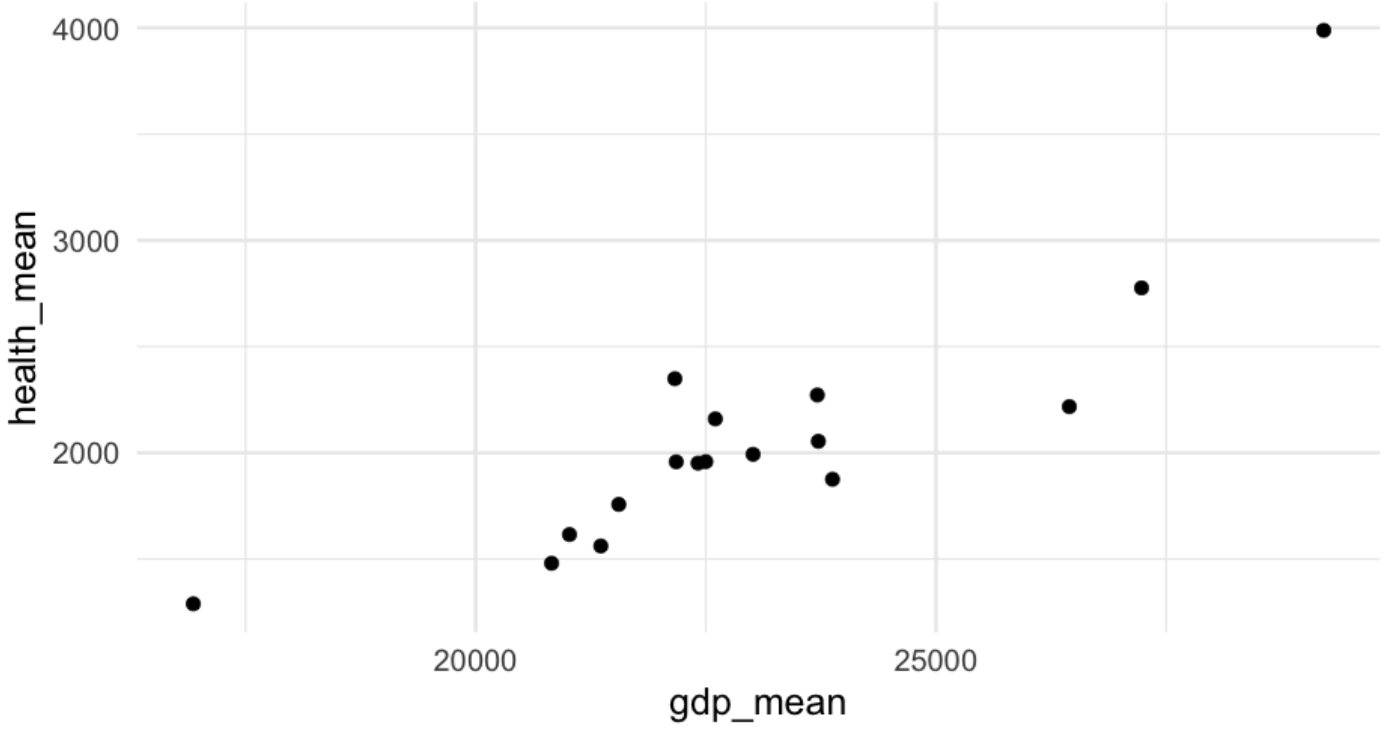

Outliers

```
ggplot(by_country, aes(gdp_mean, health_mean)) +
  geom_point() +
  geom_text_repel(data = filter(by_country,
                                gdp_mean > 25000 |
                                gdp_mean < 20000),
                  aes(label = country))
```

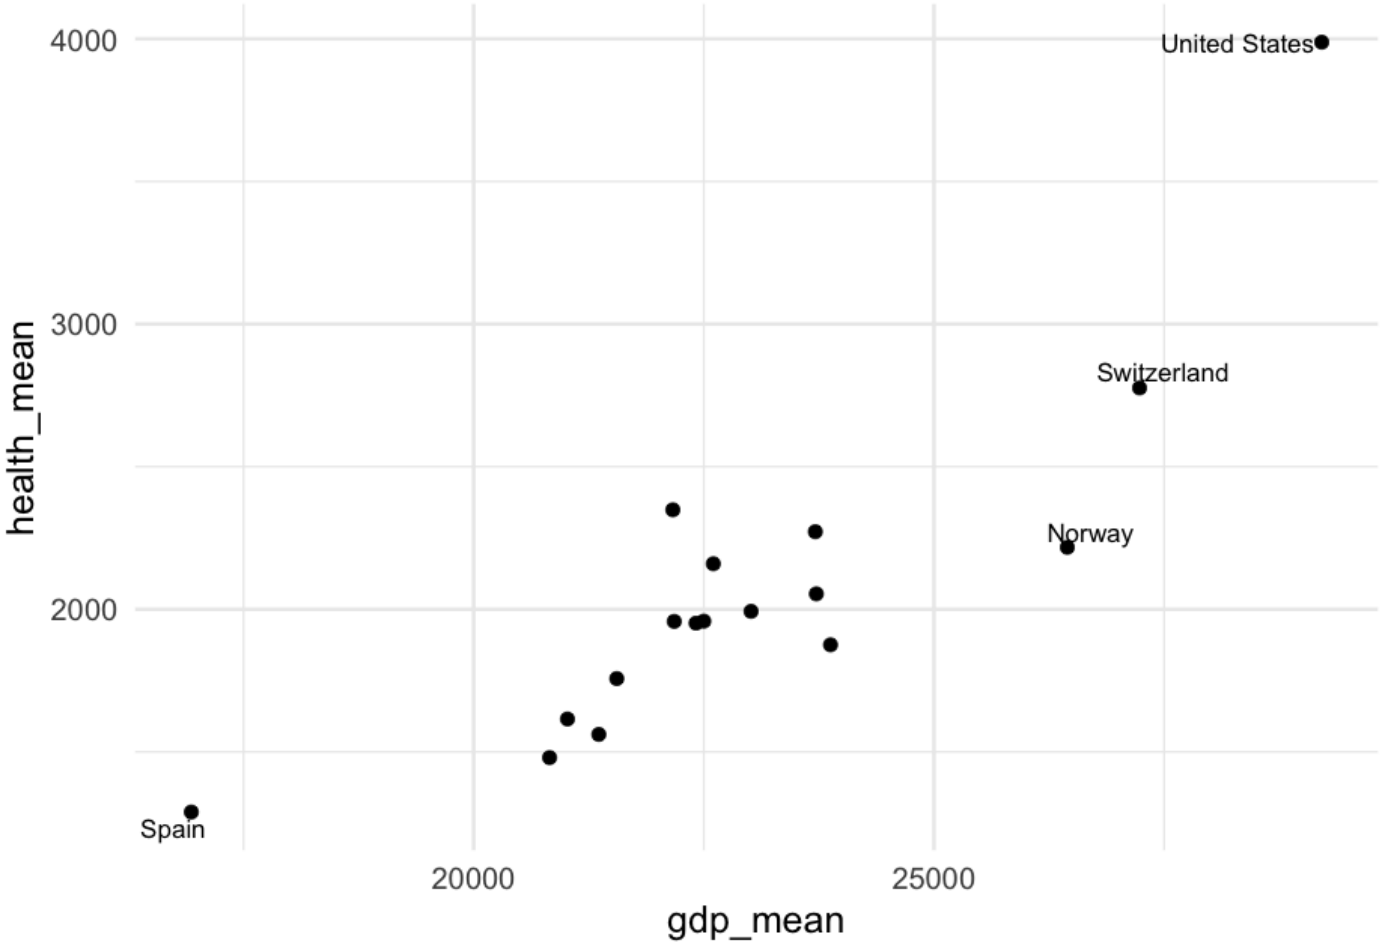
4000
3000
2000
health_mean
United States
Switzerland
Norway
Spain
20000
25000
gdp_mean

Combine with highlighting

```
library(gghighlight)
ggplot(by_country, aes(gdp_mean, health_mean)) +
  geom_point() +
  gghighlight(gdp_mean > 25000 | gdp_mean < 20000) +
  geom_text_repel(aes(label = country))
```

- Notice you only have to specify the points to highlight and `geom_text_repel` will then only label those points

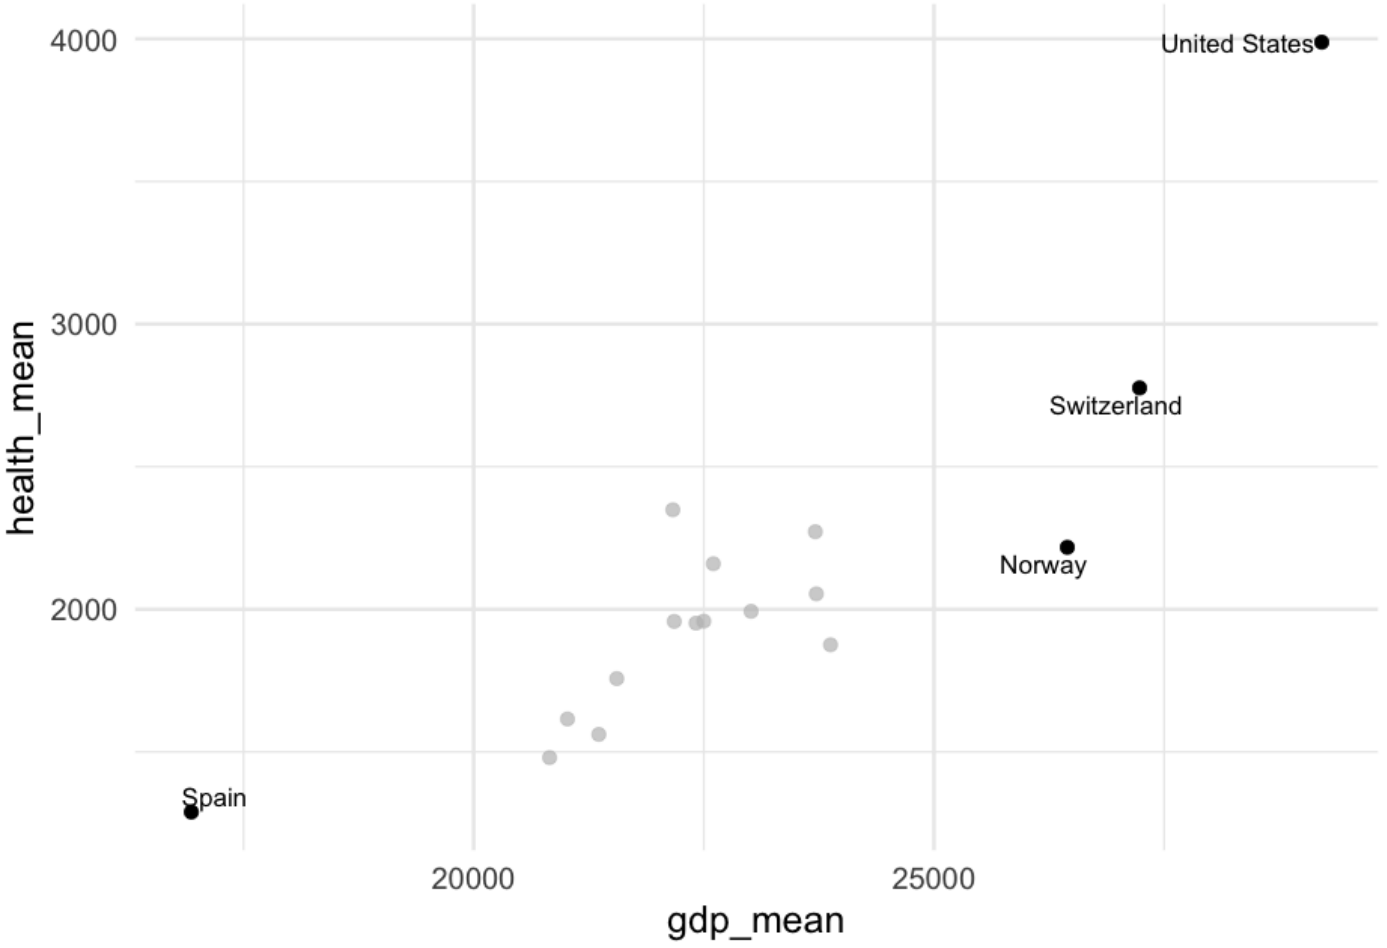
4000
3000
2000
health_mean
United States
Switzerland
Norway
Spain
20000
25000
gdp_mean

Combine with highlighting

Switch to make outliers grayed out and labeled

```
ggplot(by_country, aes(gdp_mean, health_mean)) +
  geom_point() +
  gghighlight(gdp_mean > 20000 & gdp_mean < 25000 ) +
  geom_text_repel(data = filter(by_country,
                                gdp_mean > 25000 |
                                gdp_mean < 20000),
                  aes(label = country),
                  color = "#BEBEBEB3")
```

Note I found the exact gray color by looking at the source code. Specifically, it is the output from

```
ggplot2::alpha("grey", 0.7)
```

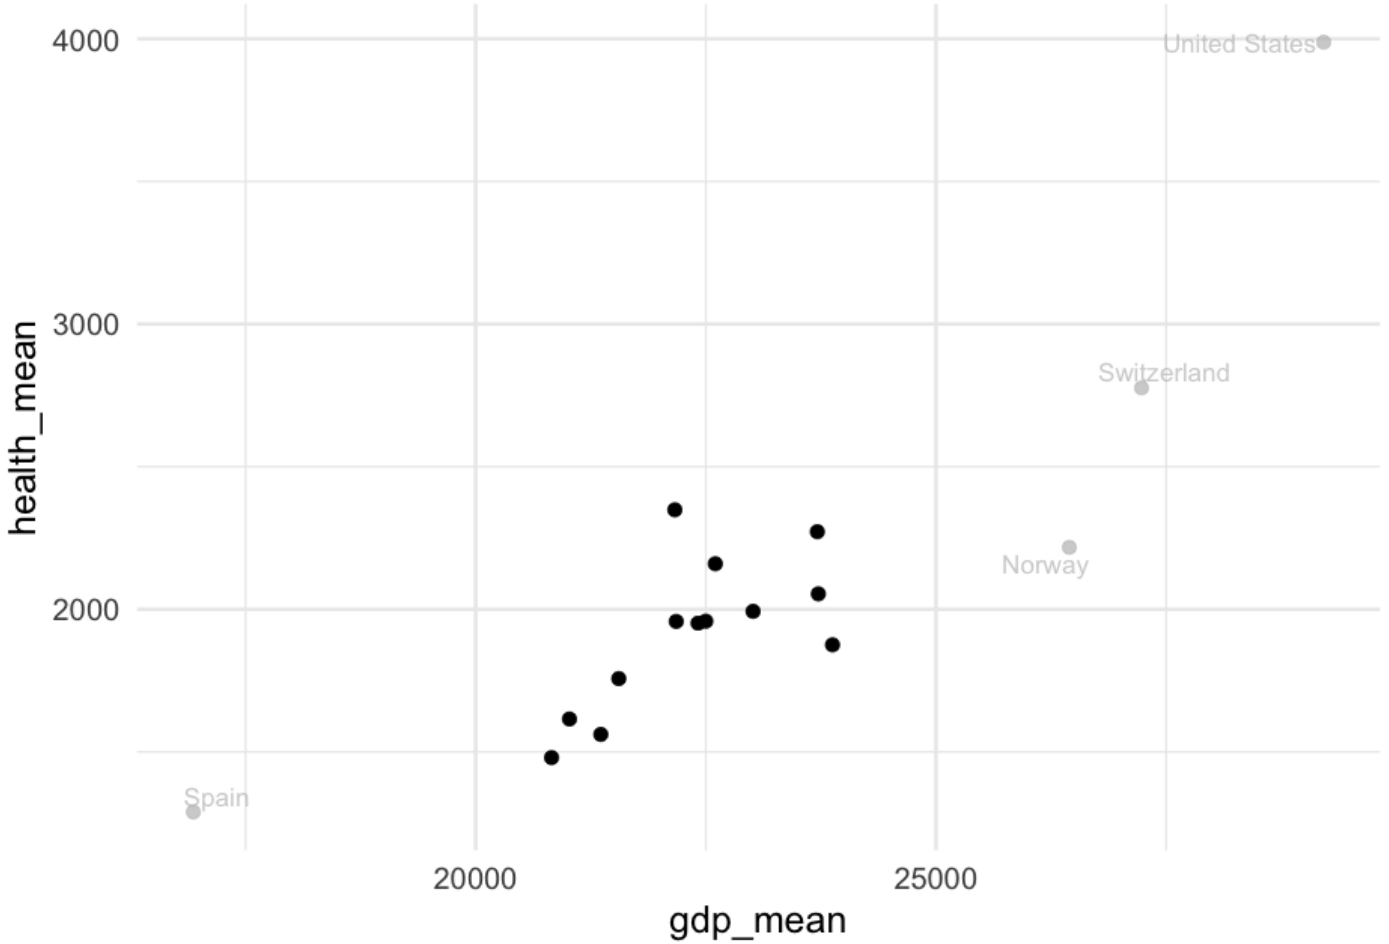
4000
3000
2000
health_mean
United States
Switzerland
Norway
Spain
20000
25000
gdp_mean

By group

```
ggplot(by_country, aes(gdp_mean, health_mean)) +
  geom_point() +
  geom_text_repel(data = filter(by_country,
                                consent_law == "Presumed"),
                  aes(label = country))
```

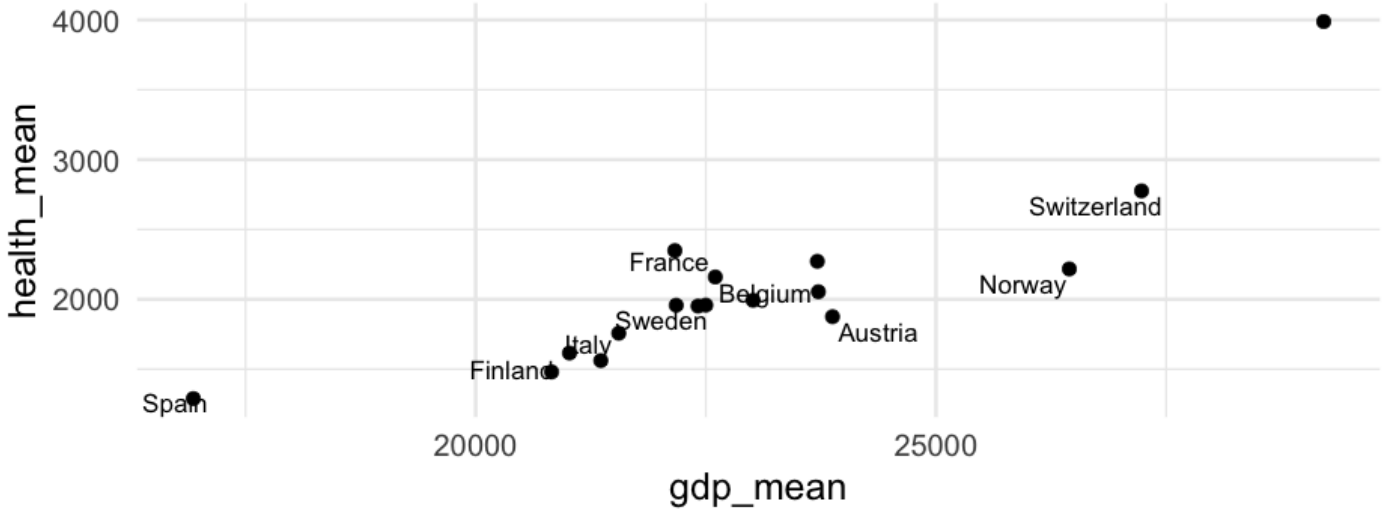

By group

```
ggplot(by_country, aes(gdp_mean, health_mean)) +
  geom_point(color = "#DC5265") +
  gghighlight(consent_law == "Presumed") +
  geom_text_repel(aes(label = country),
                  min.segment.length = 0,
                  box.padding = 0.75) +
  labs(title = "GDP and Health",
         subtitle = "Countries with a presumed organ donation co
         caption = "Data from the General Social Science Survey,
         x = "Mean GDP",
         y = "Mean Health")
```

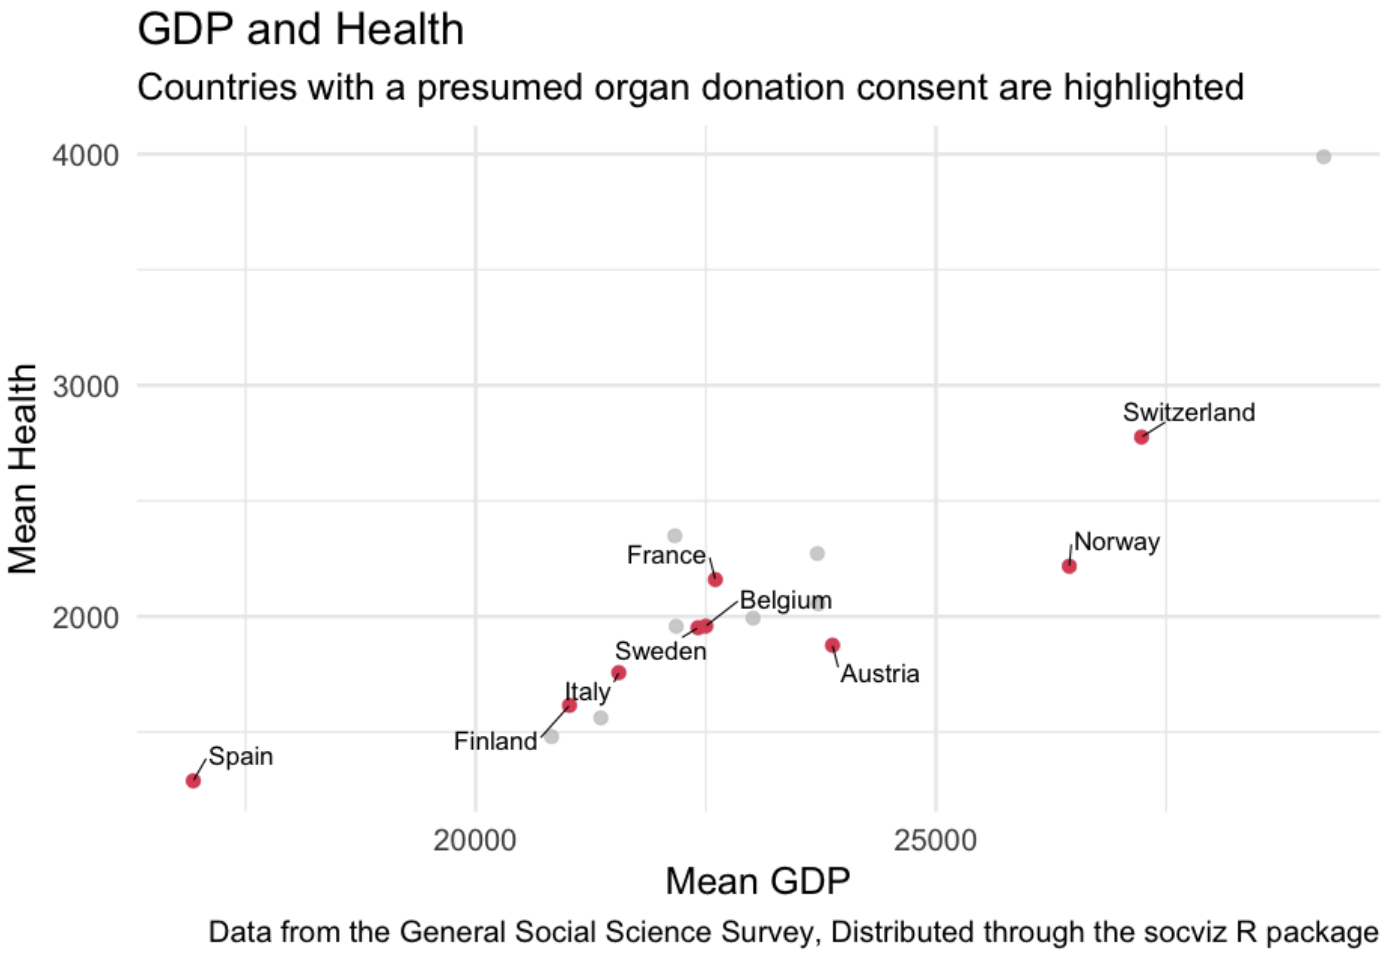
GDP and Health
Countries with a presumed organ donation consent are highlighted
Mean Health
4000
3000
2000
Switzerland
Norway
France
Belgium
Sweden
Austria
Italy
Finland
Spain
20000
25000
Mean GDP
Data from the General Social Science Survey, Distributed through the socviz R package

ggforce

Quickly

Annotating groups of points

Consider using any of the following from **ggforce** to annotate specific points

`geom_mark_rect()` `geom_mark_circle()`
`geom_mark_ellipse()` `geom_mark_hull()`

Examples

```
library(palmerpenguins)
library(ggforce)

penguins %>%
  drop_na() %>%
ggplot(aes(bill_length_mm, bill_depth_mm)) +
  geom_mark_ellipse(aes(group = species, label = species)) +
  geom_point(aes(color = species)) +
  coord_cartesian(xlim = c(25, 65), ylim = c(10, 25)) +
  guides(color = "none")
```

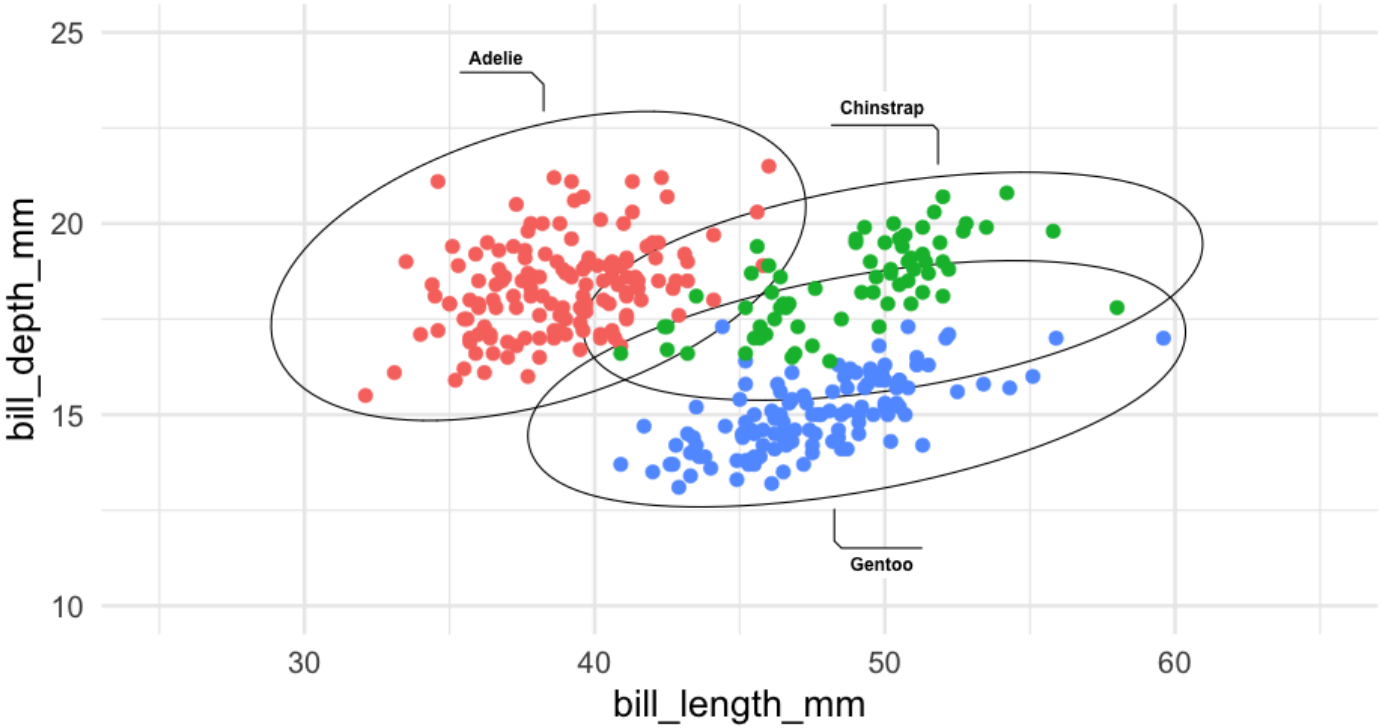
Adelie
Chinstrap
Gentoo
bill_depth_mm
bill_length_mm
25
20
15
10
30
40
50
60

Limit to a single group

```
penguins %>%
  drop_na() %>%
ggplot(aes(bill_length_mm, bill_depth_mm)) +
  geom_mark_ellipse(aes(group = species, label = species),
                    data = filter(drop_na(penguins), species ==
  geom_point(aes(color = species)) +
  coord_cartesian(xlim = c(25, 65), ylim = c(10, 25))
```

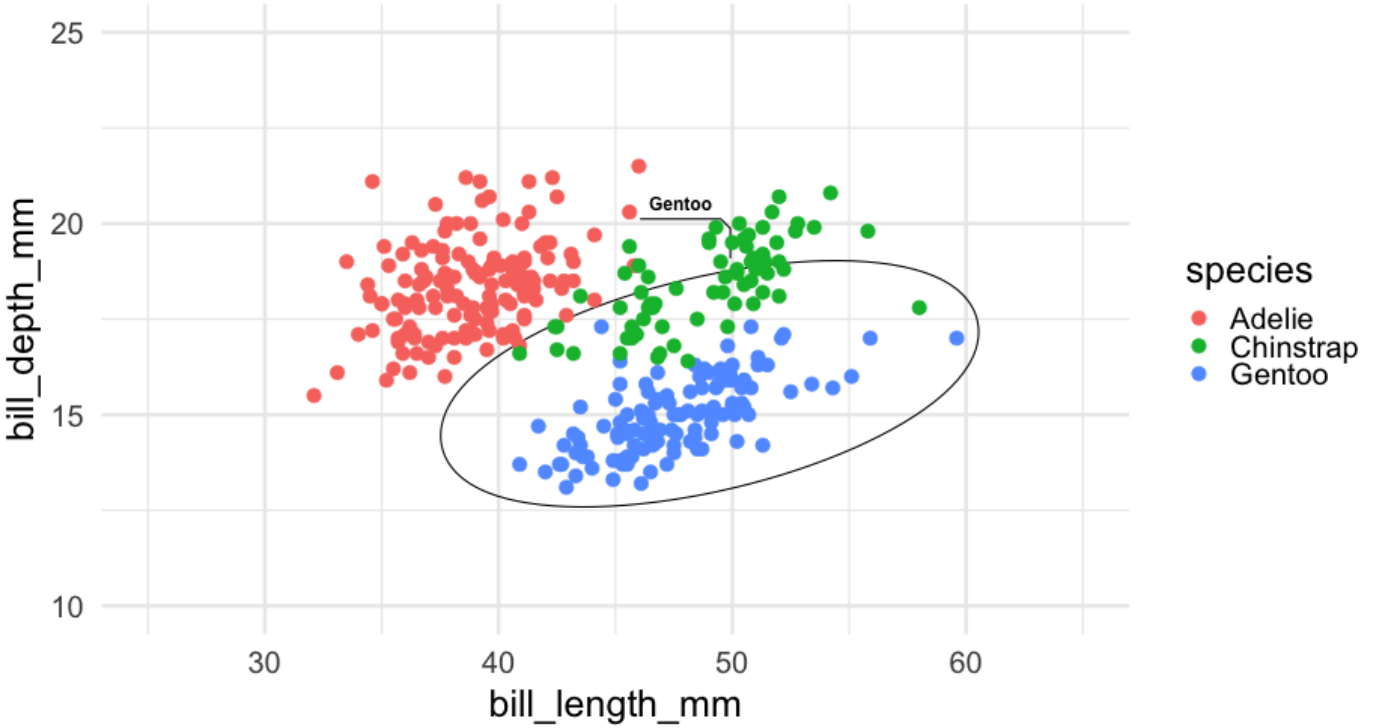
25
20
15
10
bill_depth_mm
30
40
50
60
bill_length_mm
Gentoo
species
Adelie
Chinstrap
Gentoo

Switch to hull

Note – requires the **concaveman** package be installed

```
penguins %>%
  drop_na() %>%
ggplot(aes(bill_length_mm, bill_depth_mm)) +
  geom_mark_hull(aes(group = species, label = species),
                 data = filter(drop_na(penguins), species ==
  geom_point(aes(color = species)) +
  coord_cartesian(xlim = c(25, 65), ylim = c(10, 25))
```

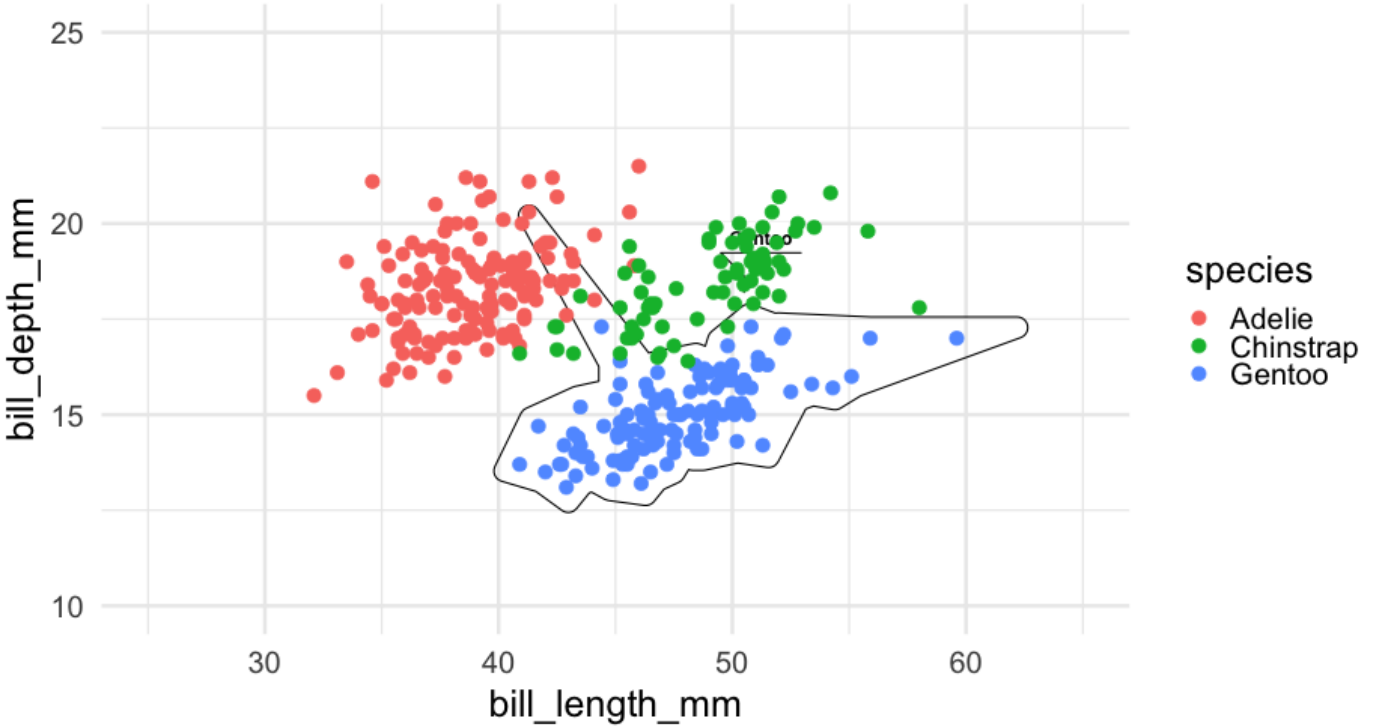
25
20
15
10
bill_depth_mm
30
40
50
60
bill_length_mm
species
Adelie
Chinstrap
Gentoo

Change expand

```
penguins %>%
  drop_na() %>%
ggplot(aes(bill_length_mm, bill_depth_mm)) +
  geom_mark_hull(aes(group = species, label = species),
                 expand = unit(1, "mm"),
                 data = filter(drop_na(penguins), species == "Ge
  geom_point(aes(color = species)) +
  coord_cartesian(xlim = c(25, 65), ylim = c(10, 25))
```

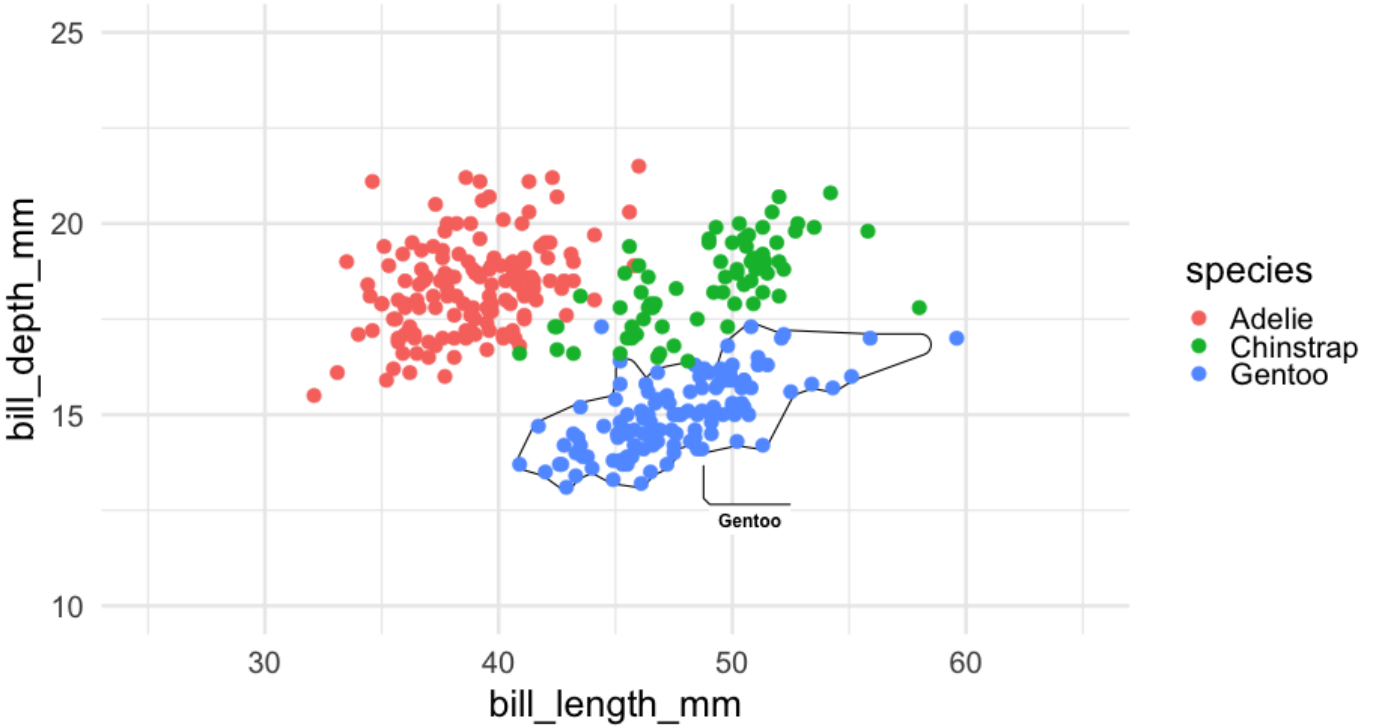
25
20
15
10
bill_depth_mm
30
40
50
60
bill_length_mm
Gentoo
species
Adelie
Chinstrap
Gentoo

Last bit

The **ggforce** package is well worth exploring more.

See here for a nice walkthrough that has good data viz and uses some of the **ggforce** functions (as well as illustrating a few other cool packages)

Themes (quickly)

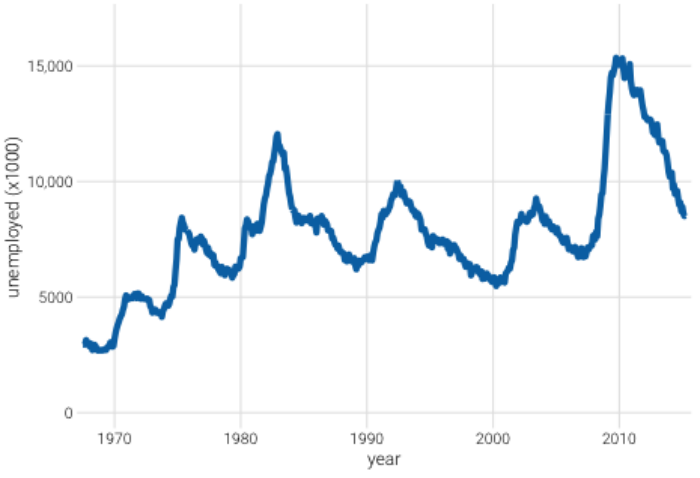
15,000
10,000
5000
0
unemployed (x1000)
1970
1980
1990
2000
2010
year

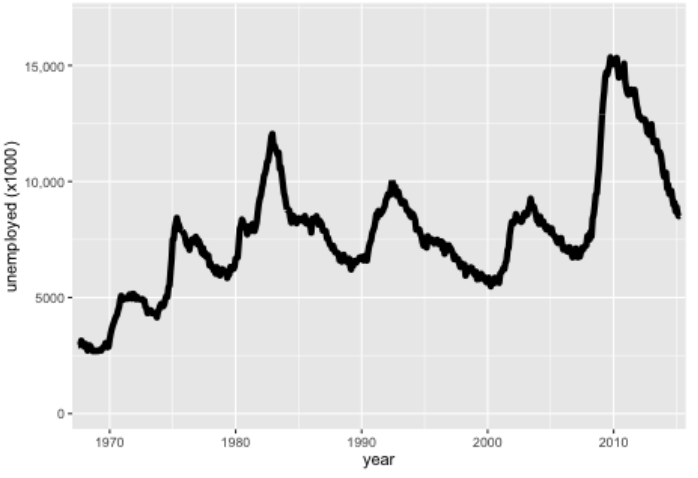
15,000
10,000
5000
0
unemployed (x1000)
1970
1980
1990
2000
2010
year

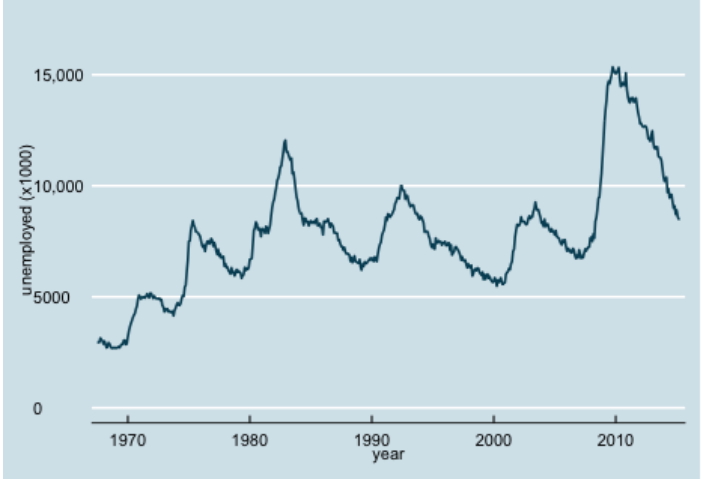
15,000
10,000
5000
0
unemployed (x1000)
1970
1980
1990
2000
2010
year

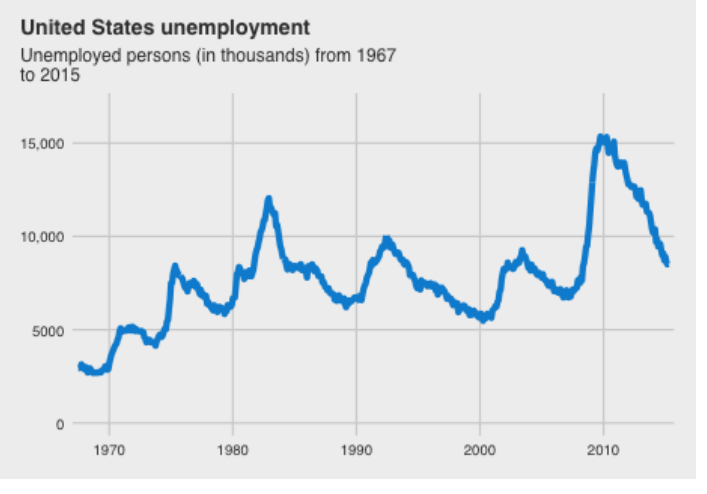
United States unemployment
Unemployed persons (in thousands) from 1967 to 2015
15,000
10,000
5000
0
1970
1980
1990
2000
2010

ggthemes

- Good place to start. All sorts of themes.
- Includes color scales, etc., that align with themes
- You can even conform with other software
  - fit into an economics conference with `theme_stata`

See the themes here

BBC

The BBC uses ggplot for most of its graphics. They've developed a package with a theme and some functions to help make it match their style more.

See the repo here

Their Journalism Cookbook is really nice too

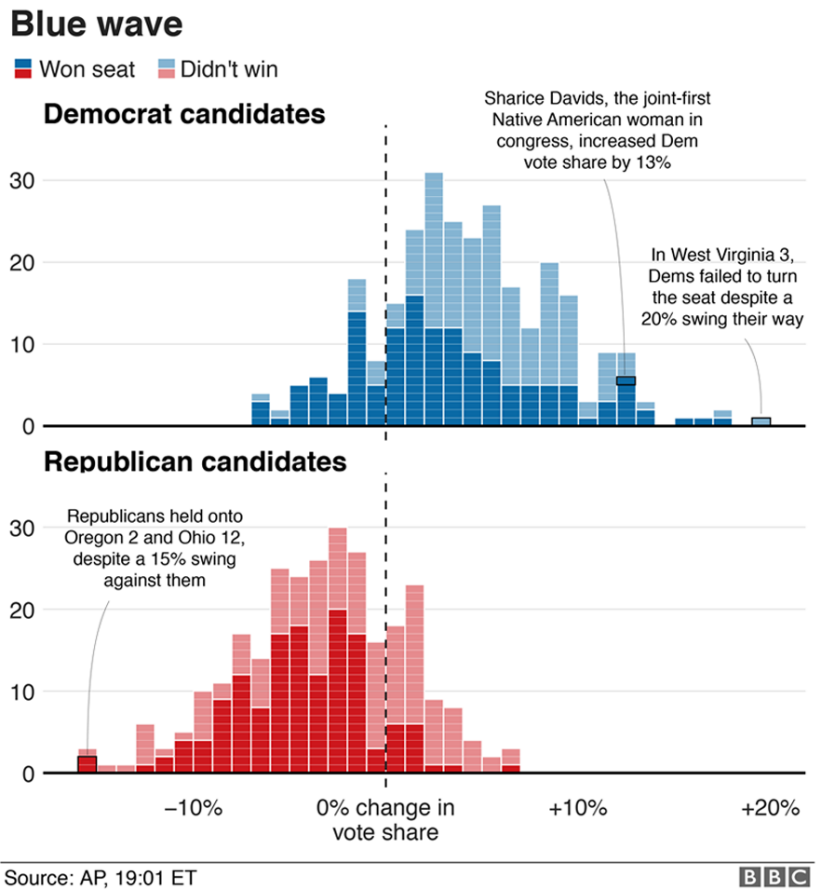
Blue wave
Won seat
Didn't win
Democrat candidates
Sharice Davids, the joint-first Native American woman in congress, increased Dem vote share by 13%
In West Virginia 3, Dems failed to turn the seat despite a 20% swing their way
30
20
10
0
Republican candidates
Republicans held onto Oregon 2 and Ohio 12, despite a 15% swing against them
30
20
10
0
−10%
0% change in vote share
+10%
+20%
Source: AP, 19:01 ET
BBC

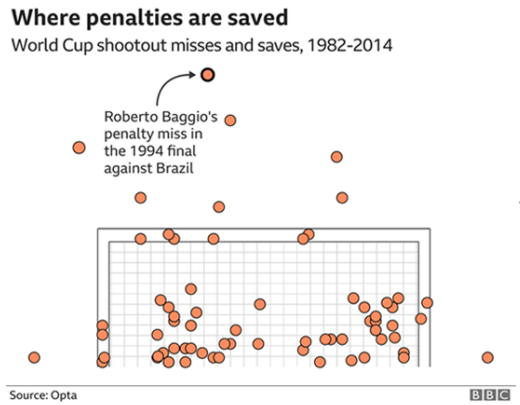
Where penalties are saved
World Cup shootout misses and saves, 1982-2014
Roberto Baggio's penalty miss in the 1994 final against Brazil
Source: Opta
BBC

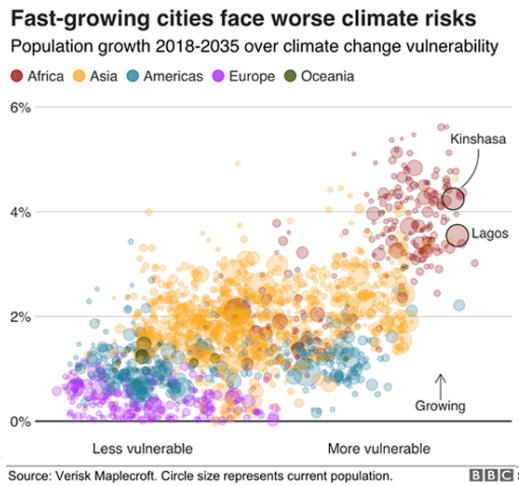
Fast-growing cities face worse climate risks
Population growth 2018-2035 over climate change vulnerability
Africa
Asia
Americas
Europe
Oceania
6%
4%
2%
0%
Kinshasa
Lagos
Growing
Less vulnerable
More vulnerable
Source: Verisk Maplecroft. Circle size represents current population.
BBC

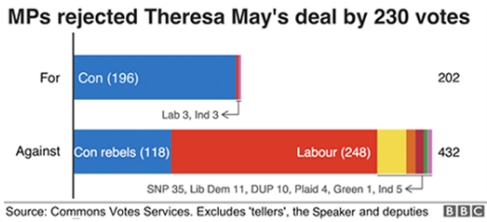
MPs rejected Theresa May's deal by 230 votes
For
Con (196)
202
Lab 3, Ind 3
Against
Con rebels (118)
Labour (248)
432
SNP 35, Lib Dem 11, DUP 10, Plaid 4, Green 1, Ind 5
Source: Commons Votes Services. Excludes 'tellers', the Speaker and deputies
BBC

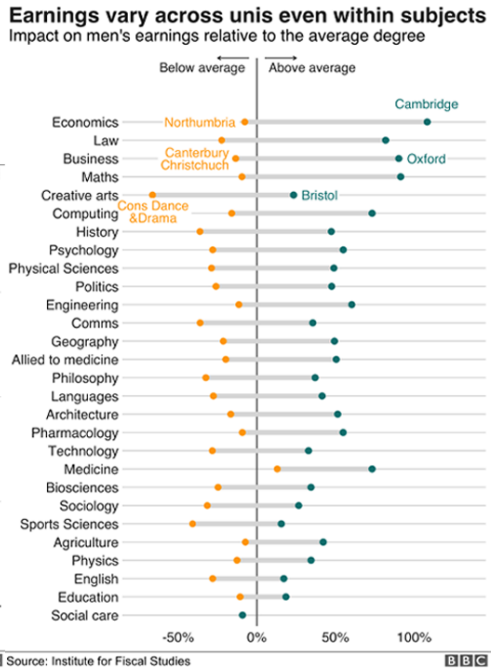
Earnings vary across unis even within subjects
Impact on men's earnings relative to the average degree
Below average
Above average
Economics
Law
Business
Maths
Creative arts
Computing
History
Psychology
Physical Sciences
Politics
Engineering
Comms
Geography
Allied to medicine
Philosophy
Languages
Architecture
Pharmacology
Technology
Medicine
Biosciences
Sociology
Sports Sciences
Agriculture
Physics
English
Education
Social care
Cambridge
Northumbria
Canterbury Christchuch
Oxford
Bristol
Cons Dance &Drama
-50%
0%
50%
100%
Source: Institute for Fiscal Studies
BBC

ggthemeassist

- Another great place to start with making major modifications/creating your own custom theme
- Can't do everything, but can do a lot
- See here

[demo]

`theme()` for everything else

- You can basically change your plot to look however you want through `theme`
- Generally a bit more complicated
- I've used ggplot for *years* and only really now gaining fluency with it

Quick example

From Lab 3

```
library(fivethirtyeight)
g <- google_trends %>%
  pivot_longer(starts_with("hurricane"),
               names_to = "hurricane",
               values_to = "interest",
               names_pattern = "_(.+)_")

landfall <- tibble(date = lubridate::mdy(c("August 25, 2017",
                                           "September 10, 2017",
                                           "September 20, 2017")),
                   hurricane = c("Harvey Landfall",
                                 "Irma Landfall",
                                 "Maria Landfall"))
```

```
p <- ggplot(g, aes(date, interest)) +
  geom_ribbon(aes(fill = hurricane, ymin = 0, ymax = interest),
              alpha = 0.6) +
  geom_vline(aes(xintercept = date), landfall,
             color = "gray80",
             lty = "dashed") +
  geom_text(aes(x = date, y = 80, label = hurricane), landfall,
            color = "gray80",
            nudge_x = 0.5,
            hjust = 0) +
  labs(x = "",
       y = "Google Trends",
       title = "Hurricane Google trends over time",
       caption = "Source: https://github.com/fivethirtyeight/data
  scale_fill_brewer("Hurricane", palette = "Set2")
```

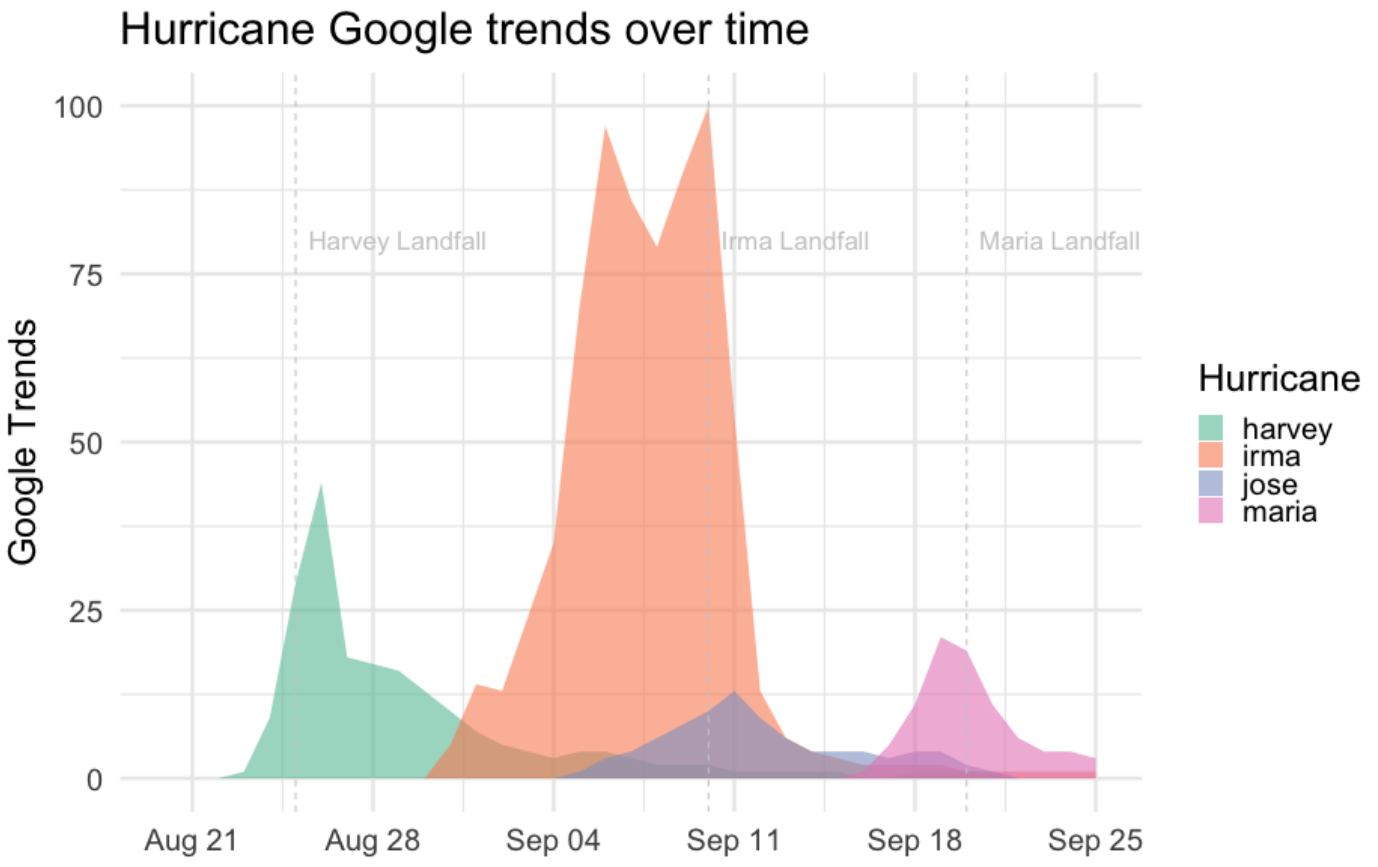

Source: https://github.com/fivethirtyeight/data/tree/master/puerto-rico-media

```
p + theme(panel.grid.major = element_line(colour = "gray30"),
          panel.grid.minor = element_line(colour = "gray30"),
          axis.text = element_text(colour = "gray80"),
          axis.text.x = element_text(colour = "gray80"),
          axis.text.y = element_text(colour = "gray80"),
          axis.title = element_text(colour = "gray80"),
          legend.text = element_text(colour = "gray80"),
          legend.title = element_text(colour = "gray80"),
          panel.background = element_rect(fill = "gray10"),
          plot.background = element_rect(fill = "gray10"),
          legend.background = element_rect(fill = NA, color = NA)
          legend.position = c(0.20, -0.1),
          legend.direction = "horizontal",
          plot.margin = margin(10, 10, b = 20, 10),
          plot.caption = element_text(colour = "gray80", vjust =
          plot.title = element_text(colour = "gray80"))
```

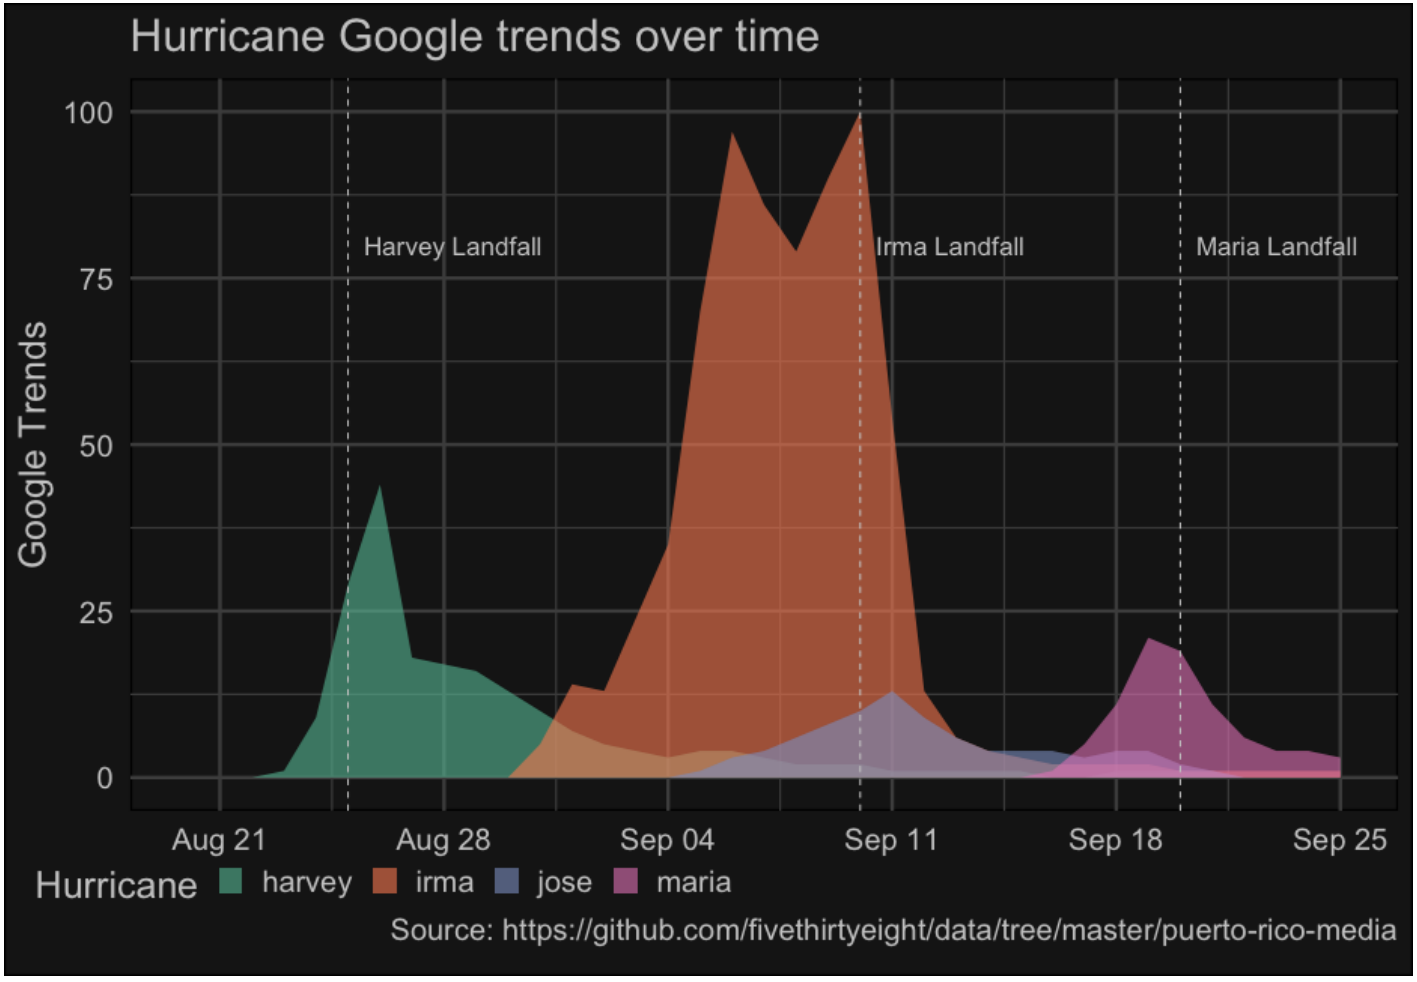
Hurricane Google trends over time
100
75
50
25
0
Google Trends
Harvey Landfall
Irma Landfall
Maria Landfall
Aug 21
Aug 28
Sep 04
Sep 11
Sep 18
Sep 25
Hurricane
harvey
irma
jose
maria
Source: https://github.com/fivethirtyeight/data/tree/master/puerto-rico-media

Next time

Visualizing uncertainty

Homework 2 is also posted currently, but is technically assigned Wednesday